

Floats (continued)


Intro to C programming

Lesson plan

- Rounding
- FP operations and caveats
- C programming: overview
- C programming: bitwise operators



IEEE Floating Point

Normalized Encoding:




$$\pm M * 2^E$$

Denormalized Encoding:



FP: Rounding



Values that are represented precisely

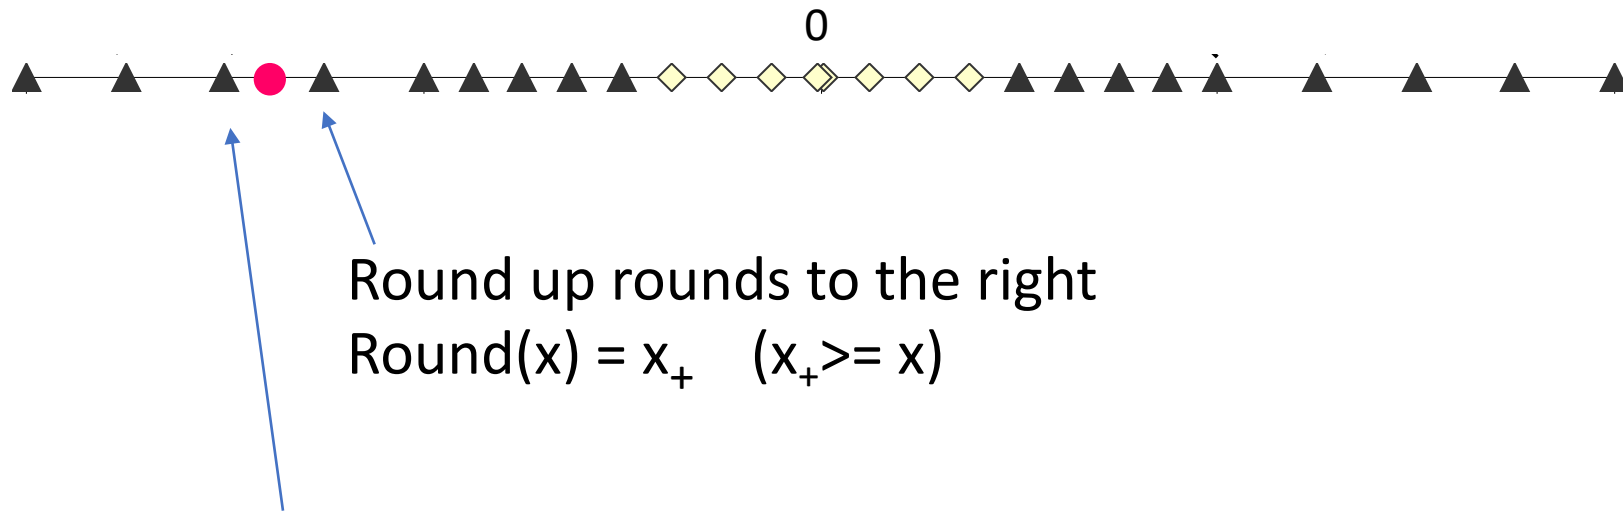
What if the result of computation is at ● ?

Rounding: Use the “closest” representable value x' for x .


4 modes:

- Round-down
- Round-up
- Round-toward-zero
- Round-to-nearest (Round-to-even in text book)

Round down vs. round up

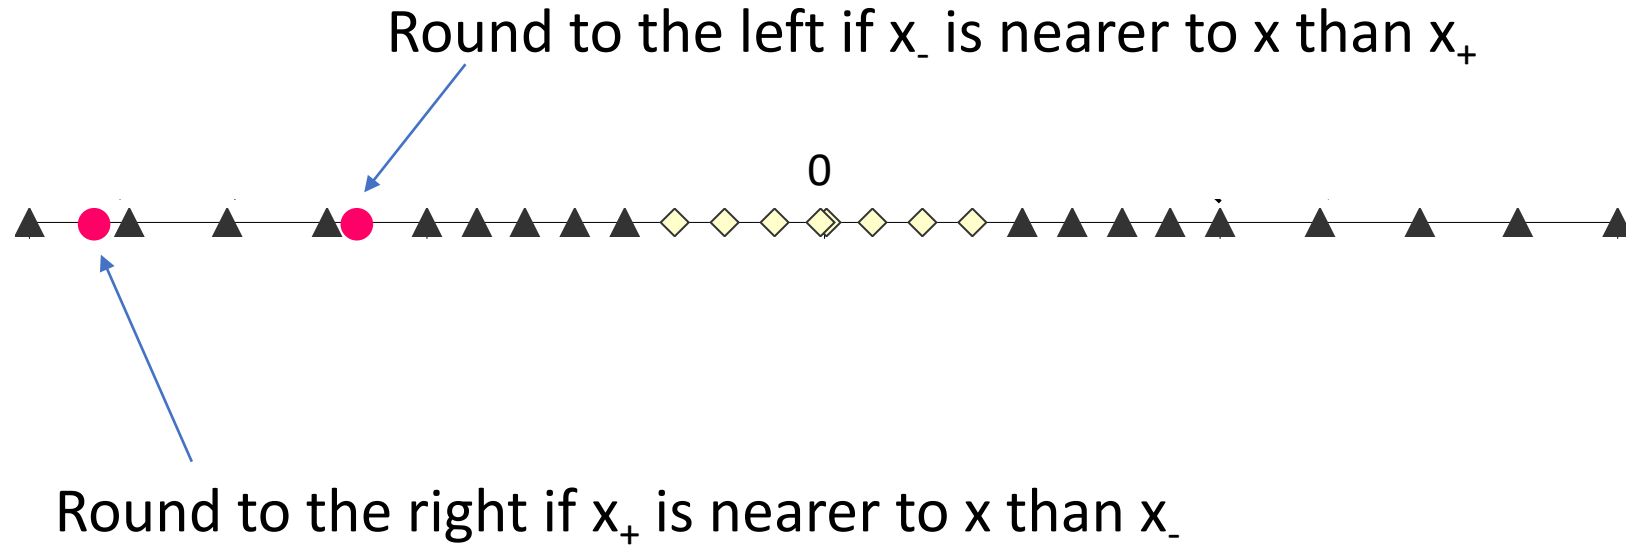


Round towards zero



Rounds to the left if $x > 0$
 $\text{Round}(x) = x_-$ if $x > 0$


Round to nearest; ties to even



In case of a tie, the one with its least significant bit equal to zero is chosen.

IEEE FP: single vs. double precision

single precision
(32 bits)



C program:


```
float f = 0.1;  
double d = 0.1;
```

double precision
(64 bits)




- What's the highest precision? (aka intervals between two denormalized numbers?)
- What's the largest positive FP?

IEEE FP: single vs. double precision



C program:

```
float f = 0.1;  
double d = 0.1;
```




- What's the highest precision? (aka intervals between two denormalized numbers?)
 - (single) 2^{-149} (double) 2^{-1045}
- What's the largest positive FP?
 - (single) $\approx 2^{128}$ (double) $\approx 2^{1024}$


How does CPU know if data is FP or integer ?

4-byte data: 0x80000001

Interpret as signed int:



Interpret as IEEE single-precision FP:



- CPU has separate registers for FPs and integers.
- CPU uses different instructions for FPs and integer operations.

Floating point operations

- Addition, subtraction, multiplication, division etc.
- Invalid operations (resulting in NaN):
 - $0/0$
 - $\text{sqrt}(-1)$
 - $\infty + \infty$
- Divide by zero: $x/0 \rightarrow \infty$
- Caveats:
 - **Overflow**: Outside the range
 - **Underflow**: $0 < \text{result} < \text{smallest denormalized value}$
 - **Inexact**: due to rounding

Floating point addition

- Commutative? $x+y == y+x$?
- Associative? $(x+y)+z = x + (y+z)$?
 - Rounding:
$$(3.14+1e10) - 1e10 = 0$$
$$3.14 + (1e10 - 1e10) = 3.14$$
 - Overflow
- Every number has an additive inverse?
 - Yes, by flipping the sign.

Floating point multiplication

- Commutative? $x * y == y * x$?
- Associative? $(x * y) * z = x * (y * z)$?
 - Overflow:
 $(1e20 * 1e20) * 1e-20 = \text{inf}, 1e20 * (1e20 * 1e-20) = 1e20$
 - Rounding
- Distributive? $(x + y) * z = x * z + y * z$?
 - $1e20 * (1e20 - 1e20) = 0.0, 1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$

FP precision decreases as value gets larger

- Storing time in computer games as a FP?
- Precision diminishes as time gets bigger

FP value (decimal)	Time value	FP precision	Time precision
1	1 sec	1.19E-07	119 nanoseconds
100	~1.5 min	7.63E-06	7.63 microseconds
10 000	~3 hours	0.000977	.976 milliseconds
1000 000	~11 days	0.0625	62.5 milliseconds

Floating point trouble

- Comparing floats for equality is a bad idea!

```
float f = 0.1;
while (f != 1.0) {
    f += 0.1;
}
```

```
f=0.2000000030
f=0.3000000119
f=0.4000000060
f=0.5000000000
f=0.6000000238
f=0.7000000477
f=0.8000000715
f=0.9000000954
f=1.0000001192
f=1.1000001431
f=1.2000001669
f=1.3000001907
f=1.4000002146
f=1.5000002384
f=1.6000002623
```



Breakout time!

Breakout exercises

- In a shooter game, the accuracy of shooting another player 1200m away is:

$$1200 = 2^{10} * (1.17)_{10} \text{ Precision: } 2^{10} * 2^{-23} = 2^{-13}$$

- Result of count?

```
int count = 0;
for (int f = 0; f <= 10; f += 1) {
    count++;
}
```

count=11



```
int count = 0;
for (float f = 0.0; f <= 1.0; f += 0.1) {
    count++;
}
```

count=10



Floating point summary

- FP format is based on normalized exponential notation
- Floating points are tricky
 - Precision diminishes as magnitude grows
 - overflow, rounding error
- Many real world disasters due to FP trickiness
 - Patriot Missile failed to intercept due to rounding error (1991)
 - Ariane 5 explosion due to overflow in converting from double to int (1996)

