


Floating point

Jinyang Li

Floating Point lesson plan

- Binary Scientific notation
- FP8 example
- IEEE FP standard
- Rounding
- FP operations caveats

Previously...



What about real numbers?

Represent real numbers: the decimal way


Real Number	Decimal Representation
$11 / 2$	$(5.5)_{10}$
$1 / 3$	$(0.333333...)_{10}$
$\sqrt{2}$	$(1.4128...)_{10}$



$$(1.4128...)_{10} = 1 * 10^0 + 4 * 10^{-1} + 1 * 10^{-2} + 2 * 10^{-3} + \dots$$


Binary Representation

$$\begin{aligned}(5.5)_{10} &= 4 + 1 + 1/2 = 2^2 + 2^0 + 2^{-1} \\ &= (101.1)_2\end{aligned}$$



Binary Representation

$$\begin{aligned}(0.1)_{10} &= 2^{-4} + 2^{-5} + 2^{-8} + 2^{-9} + 2^{-12} + 2^{-13} + \dots \\ &= (0.0001100110011\dots)_2\end{aligned}$$



Binary representation



What's the decimal value of $(10.01)_2$

Binary representation




What's the decimal value of $(10.01)_2$

Answer: 2.25

Making the representation fixed width


Strawman: fixed point



Example: $(10.011)_2$


0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

Problems of Fixed Point




Range?

Precision?



Problems of Fixed Point



- Limited range $[-2^3 + 2^{-4}, 2^3 - 2^{-4}]$
- Limited precision 2^{-4}
- Equivalent to integer with a scaling factor

Floating Point is based on scientific notation

Scientific notation in decimal:

$$365.25 = 3.6525 * 10^2$$

$$0.0123 = 1.23 * 10^{-2}$$

$\pm M * 10^E$, where $1 \leq M < 10$

The diagram illustrates the components of scientific notation. A blue arrow labeled "M: mantissa" points to the variable M in the expression $\pm M * 10^E$. Another blue arrow labeled "E: exponent" points to the variable E . A third blue arrow labeled "Normalize mantissa" points to the value 3.6525 in the first example, indicating the process of adjusting the mantissa to be between 1 and 10.

Floating Point: binary scientific notation

Binary scientific notation

$\pm M * 2^E$, where $1 \leq M < 2$, aka $M = (1.b_1b_2b_3\dots b_n)_2$

$$(5.5)_{10} = (101.1)_2 = (1.011)_2 * 2^2$$



(Binary) normalized representation of $(10.25)_{10}$?



(Binary) normalized representation of $(10.25)_{10}$?

Answer: $(10.25)_{10} = (1010.01)_2 = (1.01001)_2 * 2^3$

Floating Point lesson plan


- Scientific notation
- FP8 example
- IEEE FP standard
- Rounding
- FP operations caveats

An example FP8 representation: E4M3

$\pm M * 2^E$, where $1 \leq M < 2$, aka $M = (1.b_1b_2b_3)_2$

Example: $(0.1875)_{10} = (0.0011)_2 = (1.1)_2 * 2^{-3}$

exponent (E) fraction(M)



An example FP8 representation: E4M3


$\pm M * 2^E$, where $1 \leq M < 2$, aka $M = (1.b_1b_2b_3)_2$

Example: $(0.1875)_{10} = (0.0011)_2 = (1.1)_2 * 2^{-3}$


exponent = $E + \text{bias}$, bias = $2^{(e-1)} - 1 = 7$



exponent (E) fraction (M)




FP8 (E4M3) on the number line




Larger the bit pattern, larger FP magnitude



FP8 (E4M3) on the number line



Larger the bit pattern, larger FP magnitude



Having normalized representation only leaves a disproportionately large gap around zero.


FP8 (E4M3): denormalized (subnormal) number

exponent (E)					fraction (M)		
s	0	0	0	0	b ₁	b ₂	b ₃

$\pm M * 2^{1-\text{bias}}$, where $M = (0.b_1b_2b_3)_2$

Denormalize $(1.01)_2 \times 2^{-7}$?

$= (0.101)_2 \times 2^{-6}$?



IEEE Floating Point Standard

- Lots of FP implementations in 60s/70s
 - Code was not portable across processors
- IEEE formed a committee (IEEE.754) to standardize FP format and specification.
 - IEEE FP standard published in 1985
 - Led by William Kahan



Prof. William Kahan
University of California at Berkeley
Turing Award (1989)

IEEE Floating Point Standard

- This class only covers basic FP materials
- A deep understanding of FP is crucial for numerical/scientific computing
 - More FP is covered in undergrad/grad classes on numerical methods



Numerical Computing with IEEE Floating Point Arithmetic

Including One Theorem, One Rule of Thumb,
and One Hundred and One Exercises

Michael L. Overton

Courant Institute of Mathematical Sciences
New York University
New York, New York

Goals of IEEE Standard

- Consistent representation of floating point numbers at various widths
- Correctly rounded floating point operations, using several rounding modes.
- Consistent treatment of exceptional situations such as division by zero


IEEE FP32 normalized + denormalized



If (exp!=0 && exp!=255) n = (1.F)₂ * 2^{exp-127} (normalized)



$$n = (1.01001)_2 * 2^{130-127}$$



If (exp == 0) n = (0.F)₂ * 2⁻¹²⁶ (denormalized)



$$n = (0.01001)_2 * 2^{-126}$$


IEEE FP32: special values

Special Value's Encoding:




values	sign	frac
$+\infty$	0	all zeros
$-\infty$	1	all zeros
NaN	any	non-zero

IEEE FP32: single vs. double precision




```
float f = 0.1;  
double d = 0.1;
```



Floating Point lesson plan

- Scientific notation
- FP8 example
- IEEE FP standard
- Rounding
- FP operations caveats

FP: Rounding



Values that are represented precisely


What if the result of computation is at ● ?

Rounding: Use the “closest” representable value x' for x .

4 modes:


- Round-down
- Round-up
- Round-toward-zero
- Round-to-nearest (Round-to-even in text book)

Round up vs. round down



Round down rounds to the left
 $\text{Round}(x) = x_- \quad (x_- \leq x)$

Round towards zero




Rounds to the right if $x < 0$

$$\text{Round}(x) = x_+ \quad \text{if } x < 0$$

Rounds to the left if $x > 0$

$$\text{Round}(x) = x_- \quad \text{if } x < 0$$

Round to nearest; ties to even



In case of a tie, the one with its least significant bit equal to zero is chosen.

How does CPU know if some 4-byte value should be interpreted as IEEE FP or integers?

CPU uses separate registers for floating point and ints.

CPU uses different instructions for floating points and int operations.

Floating Point lesson plan

- Scientific notation
- FP8 example
- IEEE FP standard
- Rounding
- FP operations caveats

Floating point operations

- FP Caveats:
 - Invalid operation: $0/0$, $\text{sqrt}(-1)$, $\infty+\infty$
 - Divide by zero: $x/0 \rightarrow \infty$
 - Overflows: result too big to fit
 - Underflows: $0 < \text{result} < \text{smallest denormalized value}$
 - Inexact: round it!
- FP addition: commutative but not always associative
- FP multiplication: commutative but not always associative and distributive

Floating point trouble

- Comparing floats for equality is a bad idea!

```
float f = 0.1;  
while (f != 1.0) {  
    f += 0.1;  
}
```

```
f=0.2000000030  
f=0.3000000119  
f=0.4000000060  
f=0.5000000000  
f=0.6000000238  
f=0.7000000477  
f=0.8000000715  
f=0.9000000954  
f=1.0000001192  
f=1.1000001431  
f=1.2000001669  
f=1.3000001907  
f=1.4000002146  
f=1.5000002384  
f=1.6000002623
```

Floating point trouble

- FP is not associative: the order of operations affects results
 - $(a + b) + c \neq a + (b + c)$

```
0.1+1e20 - 1e20
>>> 0
0.1 + (1e20-1e20)
>>> 0.1
```

```
import random

vals = [1e-10, 1e-5, 1e-2, 1]
vals = vals + [-v for v in vals]

results = []
random.seed(42)
for _ in range(10000):
    random.shuffle(vals)
    results.append(sum(vals))

results = sorted(set(results))
print(f"There are {len(results)} unique results: {results}")

# Output:
# There are 102 unique results: [-8.326672684688674e-17, -7.45931094670027e-17, ...]
```

FP point trouble

- Many real world disasters are due to FP trickiness
 - Patriot Missile failed to intercept due to rounding error (1991)
 - Ariane 5 explosion due to overflow in converting from double to int (1996)



Floating point summary

- FP format is based on binary scientific notation
- IEEE FP format
 - Normalized, denormalized, special values
- Floating points are tricky
 - Precision diminishes as magnitude grows
 - overflow, rounding error