

Large C Program organization, I/O

Jinyang Li

Organization of large C programs

- Breaking a large program into multiple files
 - *.h and *.c files
- C pre-processing

Linked list: one big file

```
typedef struct {
    int val;
    struct node *next;
}node;

node* insert(node *head, int val) {
    node *new_head = (node *)malloc(sizeof(node));
    new_head->next = head;
    new_head->val = val;
}

int main() {
    node *head = NULL;
    for (int i = 0; i < 3; i++)
        head = insert(head, i);
}
```

list.c

What if another program also wants to use this linked list implementation?

linked list: multiple files

```
typedef struct {  
    int val;  
    struct node *next;  
}node;  
node *insert(node *head, int val)
```

header file includes
type definitions and
exported function
signatures

list.h

If not included, gcc
does have info on the
node type to compile
list.c

```
#include "list.h"  
node* insert(node *head, int val) {  
    node *new_head = (node *)malloc(sizeof(node));  
    new_head->next = head;  
    new_head->val = val;  
}
```

list.c

```
$ gcc -c list.c
```

← generate object file list.o

```
$ gcc list.c
```

← will not work since main() is not defined

linked list: multiple files

```
#include "list.h"
int main() {
    node *head = NULL;
    for (int i = 0; i < 3; i++)
        head = insert(head, i);
}
```

test1.c

```
#include "list.h"
int main() {
    node *head;
    for (int i = 0; i < 3; i++)
        head=insert(head, i);
}
```

test2.c

```
$ gcc -c test1.c
```

← generate object file test1.o

```
$ gcc test1.o list.o
```

← link test1.o and list.o to form executable a.out

```
$ ./a.out
```

Exporting global variables

```
typedef struct {  
    int val;  
    struct node *next;  
}node;  
node *insert(node *head, int val);
```

list.h

```
#include "list.h"  
int debug;  
node* insert(node *head, int val) {  
    ...  
    if (debug > 0)  
        printf("inserted val %d\n", val);  
}
```

list.c

```
#include "list.h"  
int main() {  
    debug = 1;  
    ...  
}
```

test1.c

Exporting global variables

```
typedef struct {  
    int val;  
    struct node *next;  
}node;  
extern int debug;  
node *insert(node *head, int val);
```

Declares debug variable but does not allocate space

list.h

```
#include "list.h"  
int debug;  
node* insert(node *head, int val) {  
    ...  
    if (debug > 0)  
        printf("inserted val %d\n", val);  
}
```

list.c

```
#include "list.h"  
int main() {  
    debug = 1;  
    ...  
}
```

test1.c

C does not have explicit namespace

- Scope of a global variable / function by default is across all files (linked together)
- To restrict the scope of a global variable / function to this file only, prefix with “static” keyword

```
#include "list.h"
static int debug;
static node* insert(node *head, int val) {
    ...
    if (debug > 0)
        printf("inserted val %d\n", val);
}
```

No other files can use the debug variable and insert function

list.c

static prefixing local variables means different things

- Normal local variables are de-allocated upon function exit
- Static local variables are not de-allocated
 - offers private, persistent storage across function invocation

```
node* insert(node *head, int val) {  
    static int n_inserts = 0;  
    ...  
    n_inserts++;  
    printf("number of inserts %d\n", n_inserts);  
}
```

initialized once,
never deallocated
(like a global
variable, except
with local scope)

C standard library

<assert.h> assert

<ctype.h> isdigit(c), isupper(c), isspace(c), tolower(c), toupper(c) ..

<math.h> log(f) log10(f) pow(f, f), sqrt(f), ...

<stdio.h> fopen, fclose, fread, fwrite, printf, ...

<stdlib.h> malloc, free, atoi, rand

<string.h> strlen, strcpy, strcat, strcmp

Section 3 of
manpage is
dedicated to
C std library

To read manual, type
man 3 strlen

The C pre-processor

- All the hashtag directives are processed by C pre-processor **before** compilation
- `#include <stdio.h>`
 - insert text of included file in the current file
 - with `<...>` , preprocessor searches system path for specified file
 - with `"..."` , preprocessor searches local directory as well as system path

C Macros

- #define name replacement_text

```
#define NITER 10000
```

It's better to write:
`static const int niter = 10000;`

```
int main()  
    for (int i = 0; i < NITER; i++) {  
        ....  
    }  
}
```

C Macros

- Macro can have arguments
- Macro is NOT a function call

```
#define SQUARE(X) X*X
```

```
a = SQUARE(2);
```

```
a = 2*2;
```

```
b = SQUARE(i+1);
```

```
b = i+1*i+1;
```

```
c = SQUARE(i++);
```

C Macros

- Macros can have arguments
- Macro is NOT a function call

```
#define SQUARE(X) (X)*(X)
```

```
a = SQUARE(2);
```

```
a = (2)*(2);
```

```
b = SQUARE(i+1);
```

```
b = (i+1)*(i+1);
```

```
c = SQUARE(i++);
```

```
c = (i++)*(i++);
```



Macro is hard to debug, avoid it if you can


Doing I/O in C

I/O in C

- I/O facilities are not part of core C language
 - provided by library using OS facilities.
- Two interfaces
 - (high level) Buffered I/O:
 - implemented by stdio library
 - uses low level interface internally
 - (low level) UNIX(Unbuffered) I/O:
 - an API provided by OS to invoke its I/O functionalities.

Buffered I/O

- each I/O stream is represented by a file pointer of type **FILE***
- Obtain the file pointer using **fopen**
 - file should be closed upon finish: **fclose**
- Access the file using file pointer with functions
 - fread, fwrite, fgetc, fgets



Type
man stdio

Buffered I/O

- each I/O stream is represented by a file pointer of type **FILE***
- Special streams: no need to explicitly open them
 - stdin
 - stdout
 - stderr

Buffered I/O example

- Count # of lines in a file

```
// open file using (fopen)

// while not end of file stream
    read file line by line (fgets)
    increment counter

// close file (fclose)
// print out counter value
```

Buffered I/O example

```
#include <stdio.h>

int main(int argc, char **argv)
{
    //open file based on argument

    int n = countlines(fp);

    //close file

    printf("# of lines %d\n", n);
}
```

Type "man fopen"

**FILE *fopen(const char *path,
const char *mode);**

fopen opens the file whose name is the string pointed to by **path** and associates a stream with it.

The argument **mode** points to a string beginning with one of the following sequences

- r** Open file for reading.
- r+** Open for reading and writing.
- w** Truncate file to zero length or create file for writing.

....

Buffered I/O example

```
int main(int argc, char **argv)
{
    //open file based on argument
    FILE *fp = fopen(argv[1], "r");

    int n = countlines(fp);

    //close file
    fclose(fp);

    printf("# of lines %d\n", n);
}
```

Buffered I/O example

```
int countlines(FILE *fp)
{
    int count = 0;

    while (!feof(fp)) {
        fgets(...)
        count++;
    }

    return count;
}
```

char *fgets(char *s, int size, FILE *stream);

fgets() reads in at most one less than size characters from stream and stores them into the buffer pointed to by s.

Reading stops after an **EOF** or a newline. If a newline is read, it is stored into the buffer. A terminating null byte ('\0') is stored after the last character in the buffer.

Buffered I/O example

```
#define BUFSZ 1000
int countlines(FILE *fp)
{
    int count = 0;
    char buf[BUFSZ];

    while (!feof(fp)) {
        fgets(buf, BUFSZ, fp);
        count++;
    }

    return count;
}
```

Buffered I/O example

```
int countlines(FILE *fp)
{
    int count = 0;
    char *buf;

    while (!feof(fp)) {
        buf = (char *)malloc(BUFSZ);
        fgets(buf, 1000, fp);
        count++;
    }

    return count;
}
```


Buffered I/O example

```
int countlines(FILE *fp)
{
    int count = 0;
    char buf[BUFSZ];

    while (!feof(fp)) {
        fgets(buf, BUFSZ, fp);
        count++;
    }

    return count;
}
```

char *fgets(char *s, int size, FILE *stream);

fgets() reads in at most one less than **size** characters from **stream** and stores them into the buffer pointed to by **s**.

Reading stops after an **EOF** or a newline. If a newline is read, it is stored into the buffer. A terminating null byte ('\0') is stored after the last character in the buffer.

fgets() returns **s** on success, and **NULL** on error or when end of file occurs while no characters have been read.

Buffered I/O example

```
int countlines(FILE *fp)
{
    int count = 0;
    char buf[BUFSZ];

    while (!feof(fp)) {
        if (!fgets(buf, BUFSZ, fp))
            break;
        count++;
    }
    return count;
}
```

Buffered I/O example

```
int countlines(FILE *fp)
{
    int count = 0;
    char buf[BUFSZ];

    while (!feof(fp)) {
        fgets(buf, BUFSZ, fp);
        count++;
    }

    return count;
}
```

char *fgets(char *s, int size, FILE *stream);

fgets() reads in at most one less than size characters from stream and stores them into the buffer pointed to by s.

Reading stops after an **EOF** or a newline. If a newline is read, it is stored into the buffer. A terminating null byte ('\0') is stored after the last character in the buffer.

fgets() returns s on success, and NULL on error or when end of file occurs while no characters have been read.

Buffered I/O example

```
int countlines(FILE *fp)
{
    int count = 0;
    char buf[BUFSZ];

    while (!feof(fp)) {
        if (!fgets(buf, BUFSZ, fp))
            break;
        if (buf[strlen(buf)-1]=='\n') {
            count++;
        }
    }

    return count;
}
```

Buffered I/O example

```
int countlines(FILE *fp)
{
    int count = 0;
    char buf[1000];
    while (!feof(fp)) {
        if(!fgets(buf, 1000, fp)) ← buffer allocated by caller
            break;
        if(buf[strlen(buf)-1]=='\n'){
            count++;
        }
    }
    return count;
}
```

```
BufferedReader br = new BufferedReader(new FileReader(file));
String line;
int count = 0;
while ((line = br.readLine()) != null) {
    count++;
}
```

(Low-level) UNIX I/O

- Used by stdio library to implement buffer I/O
- A thin wrapper to interface with OS kernel
- Each I/O stream is represented by an integer (called file descriptor).
- Special file descriptors:
 - 0: standard input
 - 1: standard output
 - 2: standard error

system call interface



UNIX I/O example: Count lines

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int main(int argc, char **argv)
{
    //open file based on argument
    int fd = open(argv[1], O_RDONLY);

    int n = countlines(fd);

    //close file
    close(fd);

    printf("# of lines %d\n", n);
}
```



type **"man 2 open"**

UNIX I/O example: count lines

```
#include <unistd.h>
int countlines(int fd)
{
    int count = 0;
    char buf[BUFSZ];
    ssize_t n;

    while ((n = read(fd, buf, BUFSZ)) > 0)
        for (ssize_t i = 0; i < n; i++) {
            if (buf[i] == '\n') {
                count++;
            }
        }

    return count;
}
```

typedef long ssize_t

Type “man 2 read”

ssize_t read(int fd, void
*buf, size_t count);

read() attempts to read up to **count** bytes from file descriptor **fd** into the buffer starting at **buf**.

On success, the number of bytes read is returned (zero indicates end of file), On error, -1 is returned...

What is **FILE***

```
typedef struct {  
    int cnt;    // characters left in buffer  
    char *ptr;  // next character in the buffer  
    char *base; // location of buffer  
    int mode;   // mode of file access  
    int fileno; // file descriptor  
} FILE;
```

Can you implement fopen, fclose, fgets using open, close, and read?
see page 176-177 of K&R