

# Machine-Level Programming V: Memory layout

Jinyang Li

Some slides adapted from Bryant and O'Hallaron

# x86 Procedure Recap

## ■ **call**

- push return address on stack, jump to label

## ■ **ret**

- pop 8 bytes from stack into PC

## ■ **Argument passing from caller to callee**

- First 6 arguments passed in registers (%rdi, %rsi, %rdx, %rcx, %r8, %r9)
- Rest on stack

## ■ **Return value passing from callee to caller**

- %rax

## ■ **Local variable**

- either registers, or allocated on stack (deallocated before ret)

## ■ **Caller vs. callee-save registers**

- Caller-save: all “argument” registers, %rax, %r10, %r11
- Callee-save: %rbx, %r12, %r13, %r14, %rbp

# Recap: Procedure call example

```
int add2(int a, int b)
{
    return a + b;
}
```

```
int add3(int a, int b, int c)
{
    int r = add2(a, b);
    r = r + c;
    return r;
}
```

```
add2:
    leal    (%rdi,%rsi), %eax
    ret
```

```
add3:
    pushq   %rbx
    movl    %edx, %ebx
    movl    $0, %eax
    call    add2
    addl    %ebx, %eax
    popq    %rbx
    ret
```

a: %edi  
b: %esi  
c: %edx

%edx (containing c)  
is needed after call,  
so save in %ebx

## Registers

First 6 Arguments: %rdi, %rsi, %rdx, %rcx, %r8, %9

Return value: %rax

# Today

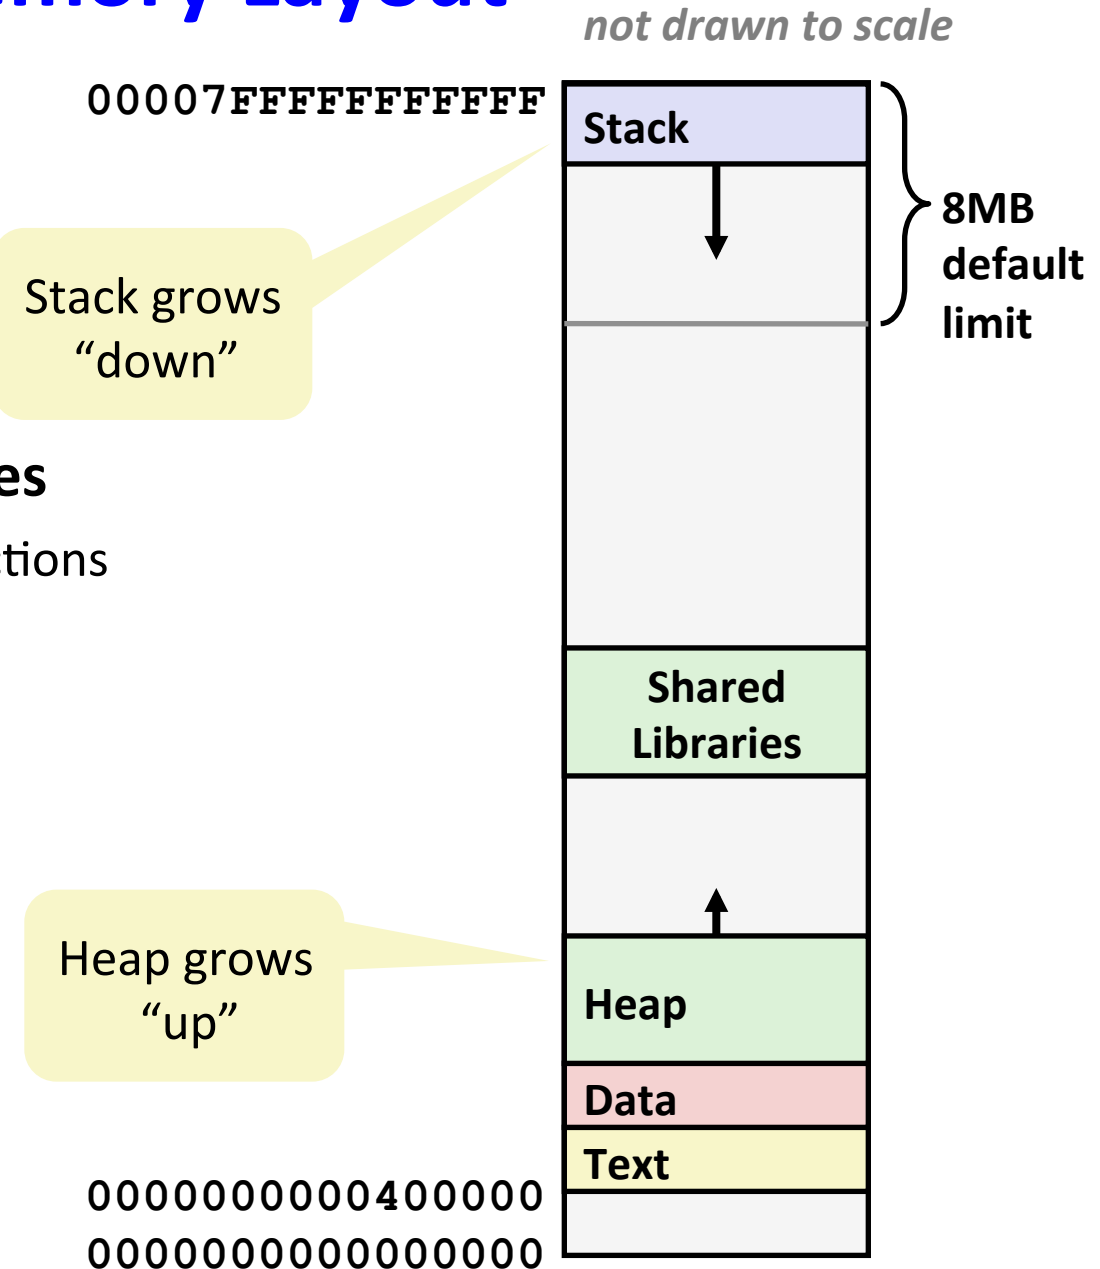
- Memory layout
- Demo: Using gdb for binary forensics

# OS loads a program to memory

- OS loads different parts of a program into different memory regions
- Parts of a running program:
  - Stack
    - e.g. local variables
  - Heap
    - e.g. malloc(), new
  - (statically allocated) Data
    - global variables, string constants
  - Executable instructions
- Why different regions?
  - need to grow independently
  - need different permissions

# x86-64 Linux Memory Layout

- **Stack**
- **Heap**
- **Data**
- **Text / Shared Libraries**
  - aka executable instructions



*not drawn to scale*


# Memory Allocation Example

```
char big_array[1<<24]; /* 16 MB */
char huge_array[1<<31]; /* 2 GB */

int global = 0;

int useless() { return 0; }

int main ()
{
    void *p1, *p2, *p3, *p4;
    int local = 0;
    p1 = malloc(1 << 28); /* 256 MB */
    p2 = malloc(1 << 8); /* 256 B */
    p3 = malloc(1 << 32); /* 4 GB */
    p4 = malloc(1 << 8); /* 256 B */
    ...
}
```




*Where does everything go?*

# x86-64 Example Addresses

address range  $\sim 2^{47}$

```
local
p1
p3
p4
p2
big_array
huge_array
main()
useless()
```

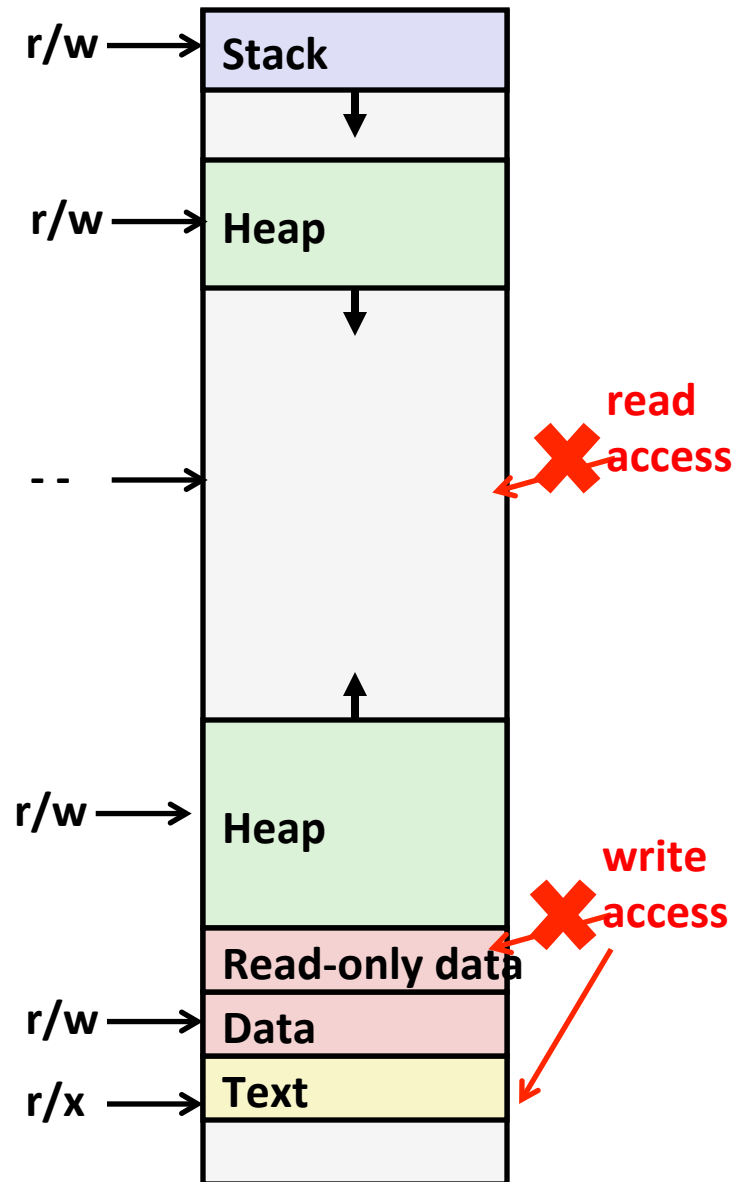
```
0x00007ffe4d3be87c
0x00007f7262a1e010
0x00007f7162a1d010
0x000000008359d120
0x000000008359d010
0x0000000080601060
0x0000000000601060
0x000000000040060c
0x0000000000400590
```






# Segmentation Fault

- Each memory segment can be readable (r), executable (x), writable(w), or none at all (-)
- Segmentation fault occurs when program tries to access “illegal” memory
  - Read from segment with no permission
  - Write to read-only segments



# Segmentation fault example

```
int main() {  
    char s1[100] = "hello world1";  
    char *s2 = "hello world2";  
    printf("str1 %p str2 %p\n", s1, s2);  
    s1[0] = 'H';  
    s2[0] = 'H';  
    ...  
}
```



# Not all Bad Memory Access lead to immediate segmentation

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

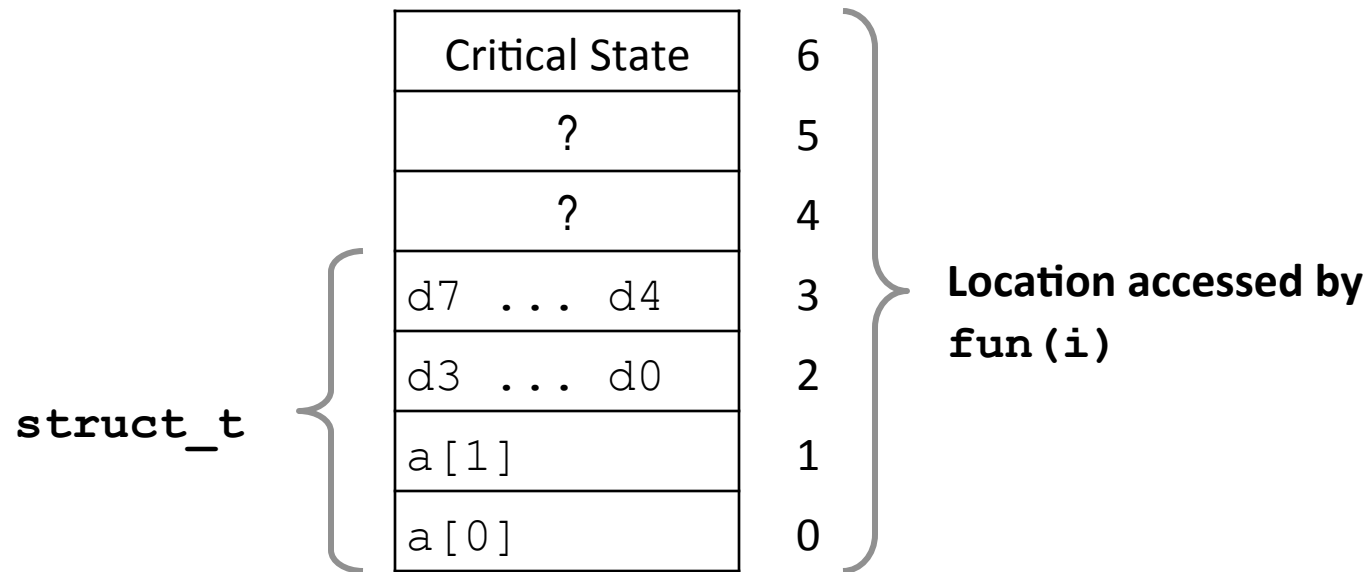
<b>fun(0)</b>	<b>→</b>	<b>3.14</b>
<b>fun(1)</b>	<b>→</b>	<b>3.14</b>
<b>fun(2)</b>	<b>→</b>	<b>3.1399998664856</b>
<b>fun(3)</b>	<b>→</b>	<b>2.00000061035156</b>
<b>fun(4)</b>	<b>→</b>	<b>3.14</b>
<b>fun(6)</b>	<b>→</b>	<b>Segmentation fault</b>

- Result is system specific

# Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault



# Such problems are a BIG deal


- **Generally called a “buffer overflow”**
  - when exceeding the memory size allocated for an array
- **Why a big deal?**
  - It's the #1 technical cause of security vulnerabilities
    - #1 overall cause is social engineering / user ignorance
- **Most common form**
  - Unchecked lengths on string inputs
  - Particularly for bounded character arrays on the stack
    - sometimes referred to as stack smashing

# Today

- Memory layout
- **Demo: Using gdb for binary forensics**

# **gdb cheat sheet**

- **info registers**
- **info proc mappings**
- **b <function>**
- **nexti**
- **continue**
- **bt: print backtrace**
- **disass <function>**
- **x/4xb <address> : print 4 bytes starting at address in hex**
- **x/4i <address>: print 4 instructions starting at address**
- **p/x \$rax**



**(gdb) help x**