

Dynamic Memory Allocation

Jinyang Li

Some are based on Tiger Wang's slides

Why dynamic memory allocation?

```
typedef struct node {
    int val;
    struct node *next;
} node;

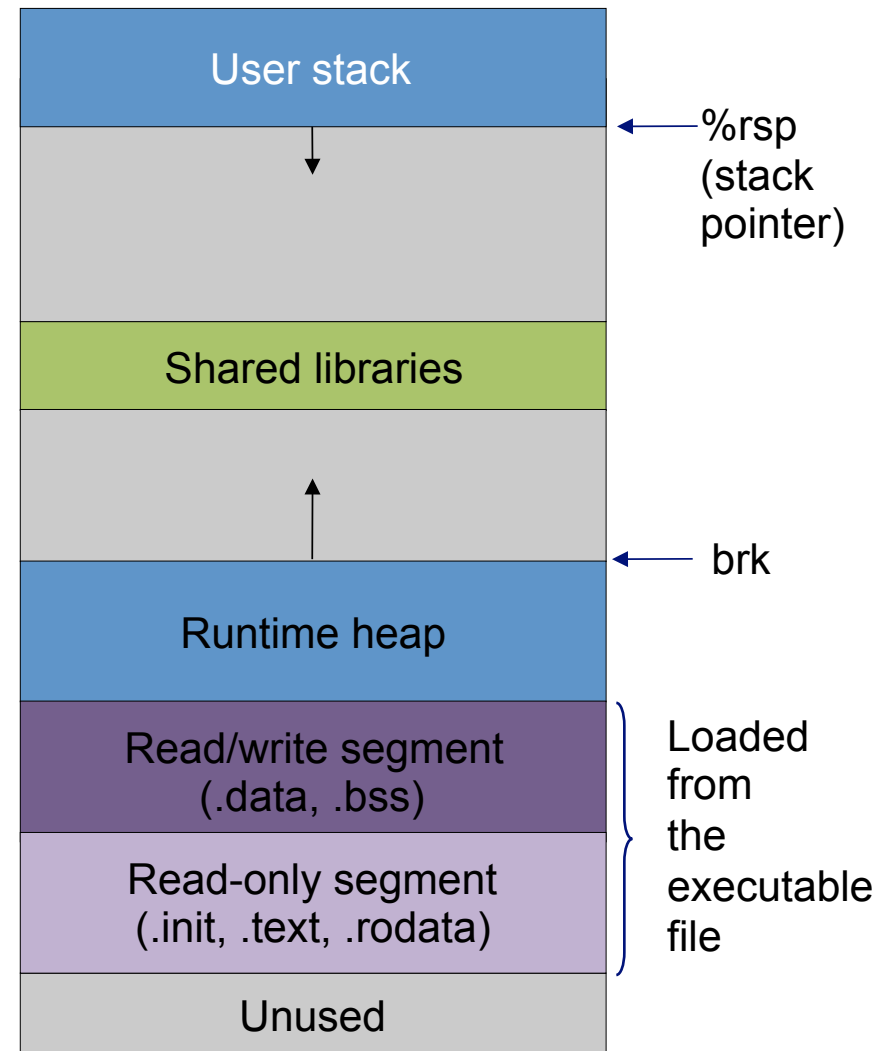
void
list_insert(node **head, int v) {
    node *np = malloc(sizeof(node));
    np->next = *head;
    np->val = v;
    *head = np;
}

int
main(void) {
    node *head = NULL;
    int n = atoi(argv[1]);
    for (int i = 0; i < n; i++) {
        list_insert(&head, i);
    }
}
```

How many nodes to allocate is only known at runtime (when the program executes)

Dynamic allocation on heap

Question: can one dynamically allocate memory on stack?




Dynamic allocation on heap

Question: Is it possible to dynamically allocate memory on stack?

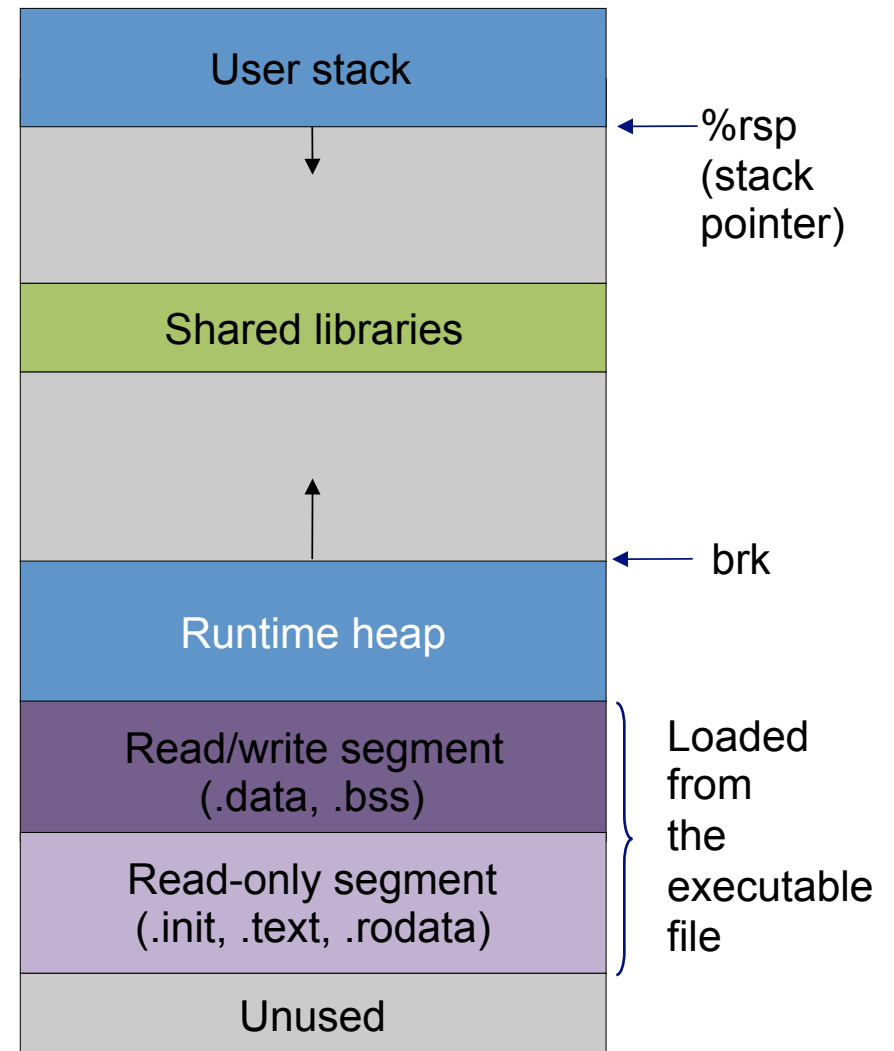
Answer: Yes, but space is freed upon function return

```
void  
list_insert(node *head, int v) {  
    node n;  
    node *np = &n;  
    np->next = head;  
    np->val = v;  
    *head = np;  
}
```



Dynamic allocation on heap

Question: How to allocate memory on heap?



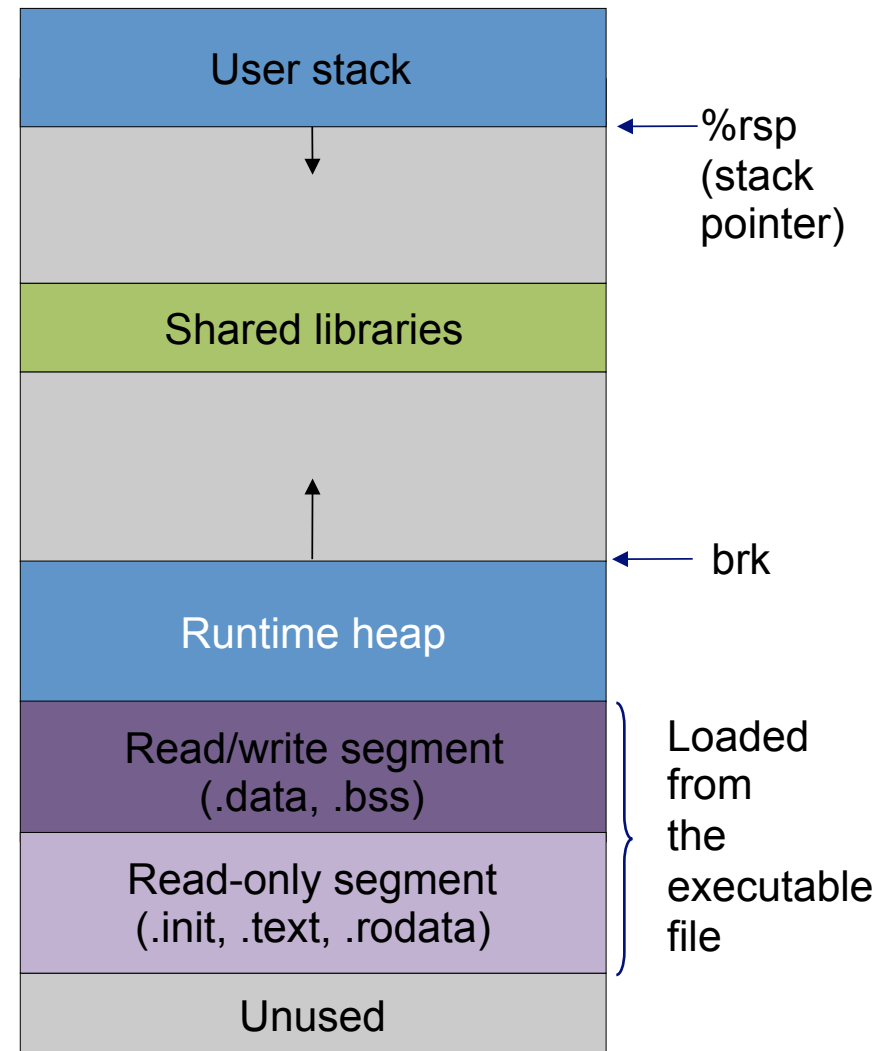
Dynamic allocation on heap

Question: How to allocate memory on heap?

Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

It increases the top of heap by “size” and returns a pointer to the base of new storage. The “size” can be a negative number.



Dynamic allocation on heap

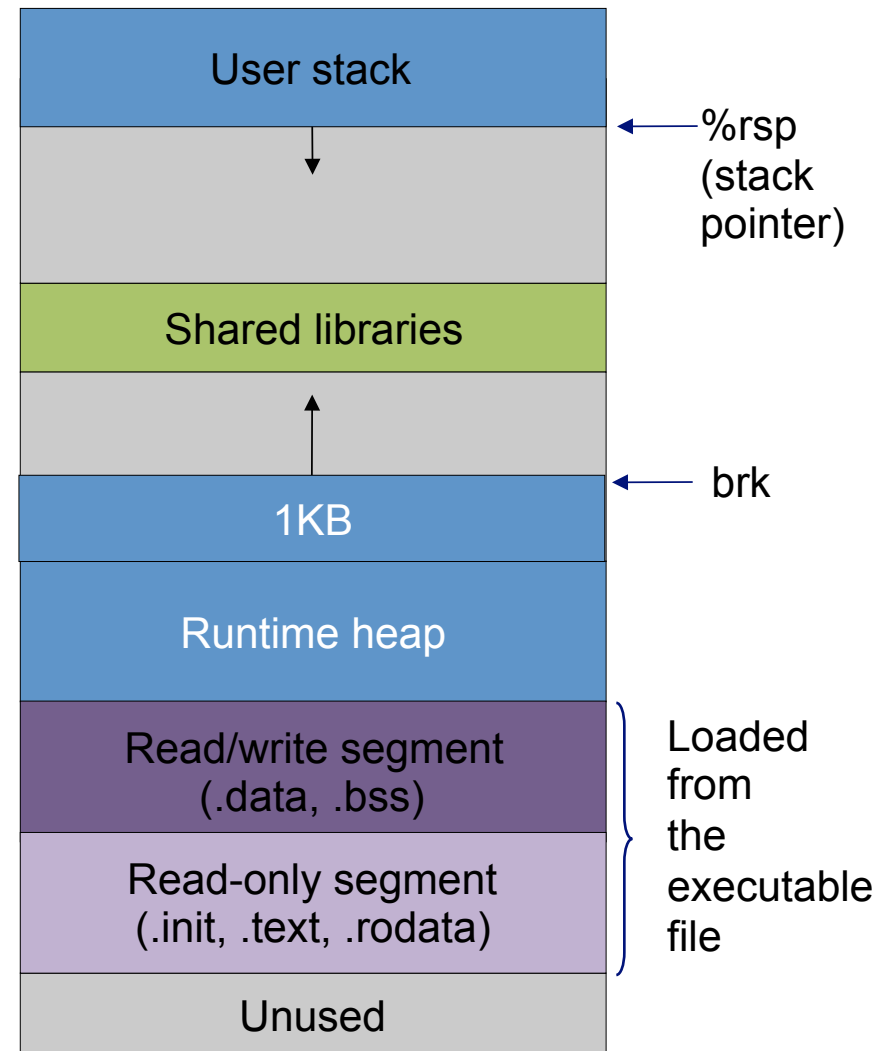
Question: How to allocate memory on heap?

Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

It increases the top of heap by “size” and returns a pointer to the base of new storage. The “size” can be a negative number.

```
p = sbrk(1024) //allocate 1KB
```



Dynamic allocation on heap

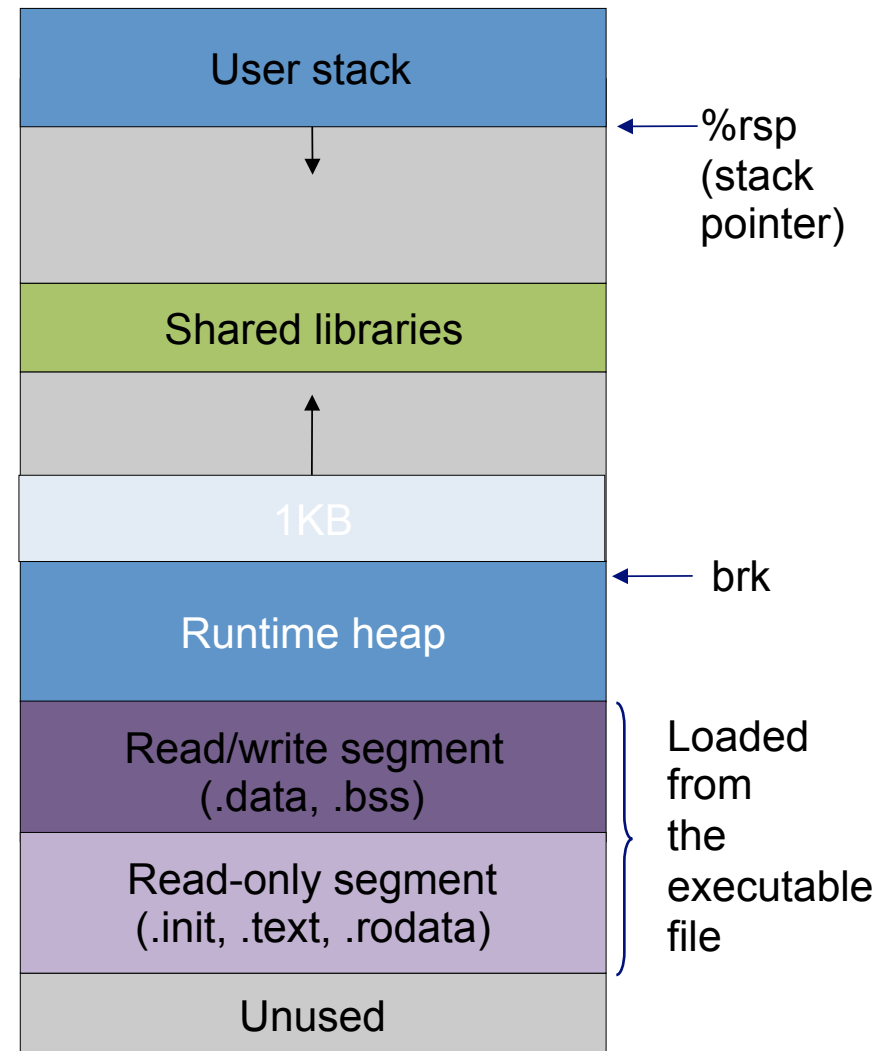
Question: How to allocate memory on heap?

Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

It increases the top of heap by “size” and returns a pointer to the base of new storage. The “size” can be a negative number.

```
p = sbrk(1024) //allocate 1KB  
sbrk(-1024) //free p
```



Dynamic allocation on heap

Question: How to allocate memory on heap?

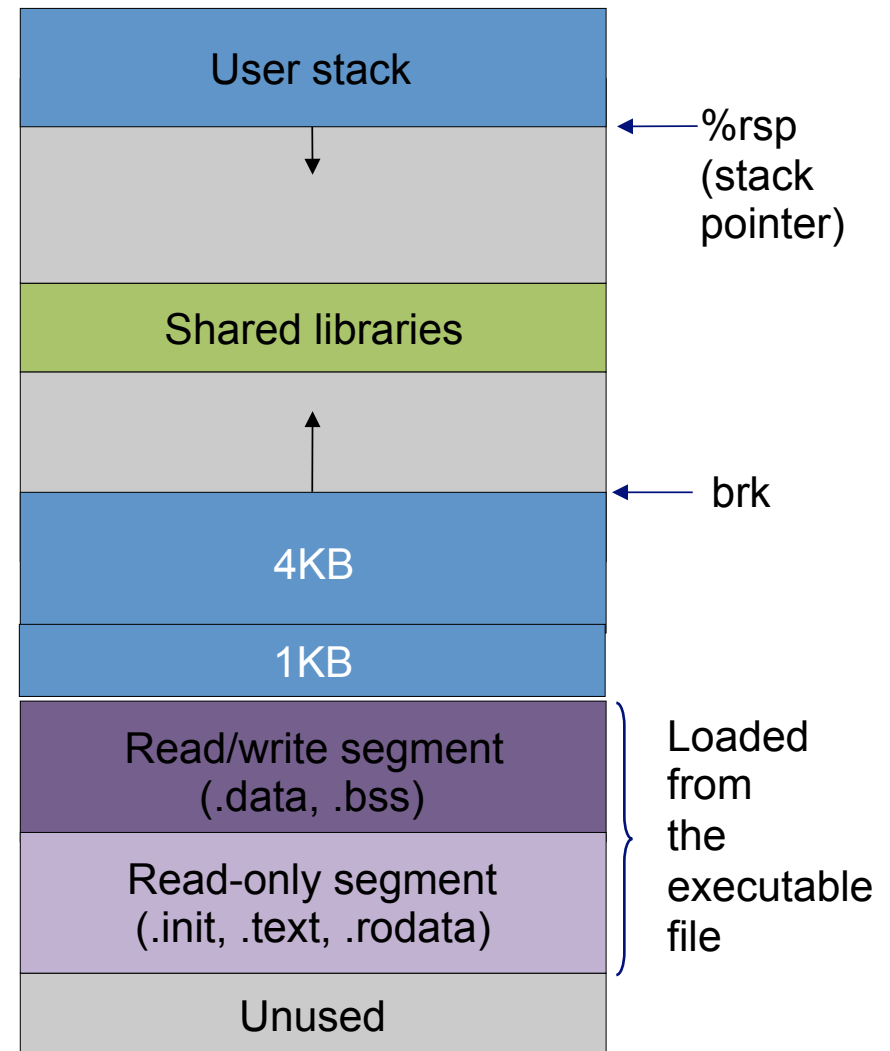
Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

Issue 1 – can only free the memory on the top of heap

```
p1 = sbrk(1024) //allocate 1KB  
p2 = sbrk(4096) //allocate 4KB
```

How to free p1?



Dynamic allocation on heap

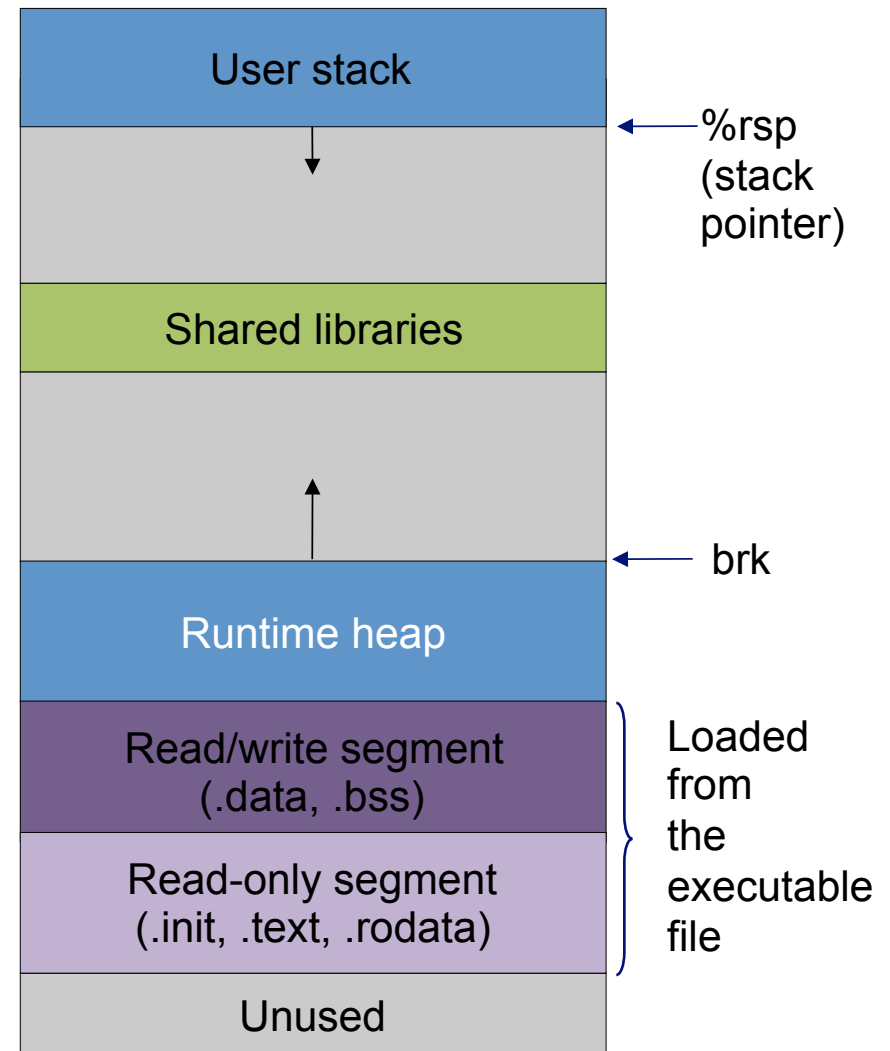
Question: How to allocate memory on heap?

Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

Issue I – can only free the memory on the top of heap

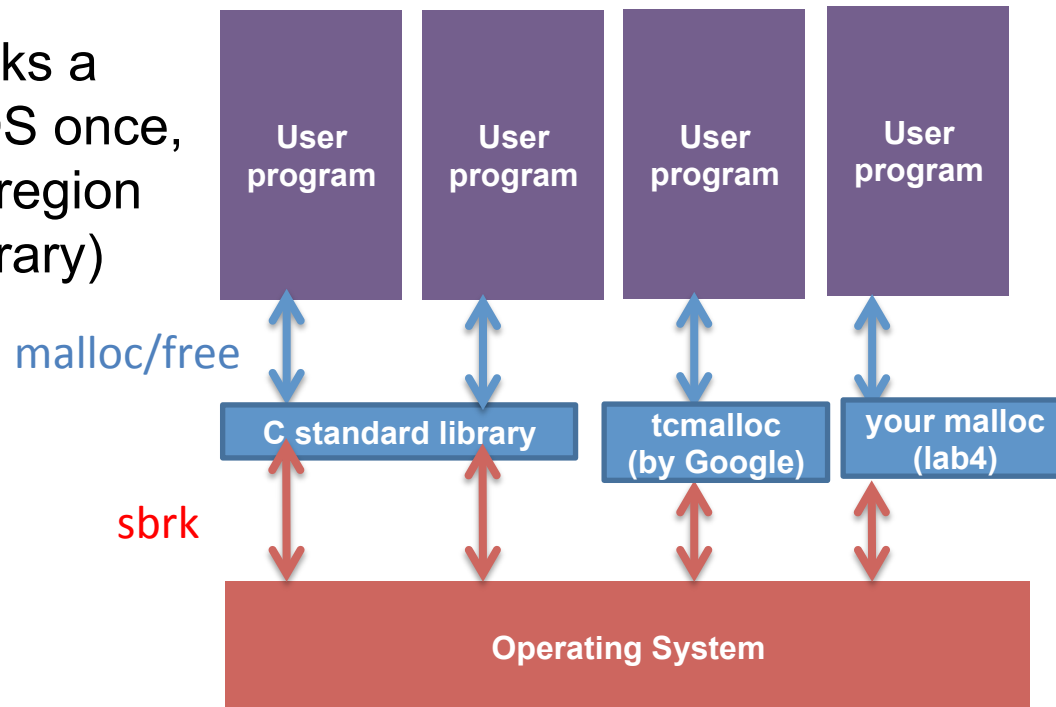
Issue II – system call has high performance cost > 10X



Dynamic allocation on heap

Question: How to efficiently allocate memory on heap?

Basic idea: user program asks a large memory region from OS once, then manages this memory region by itself (using a “malloc” library)



How to implement a memory allocator?

API:

- `void* malloc(size_t size);`
- `void free(void *ptr);`

Goal:

- Efficiently utilize acquired memory with high throughput
 - high throughput – how many mallocs / frees can be done per second
 - high utilization – fraction of allocated size / total heap size

How to implement a memory allocator?

Assumed behavior of applications:

- Issue an arbitrary sequence of malloc/free
- Argument of free must be the return value of a previous malloc
- No double free

Restrictions on the allocator:

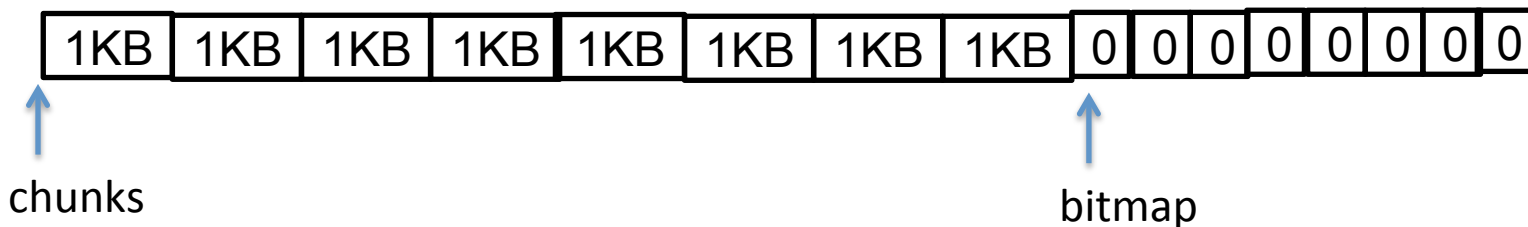
- Once allocated, space cannot be moved around

Malloc design challenges

1. (Basic book-keeping) How to keep track which bytes are free and which are not?
2. (Allocation decision) Which free chunk to allocate?
3. (API restriction) free is only given a pointer, how to find out the allocated chunk size?

How to bookkeep? Strawman #1

- Structure heap as n 1KB chunks + n metadata

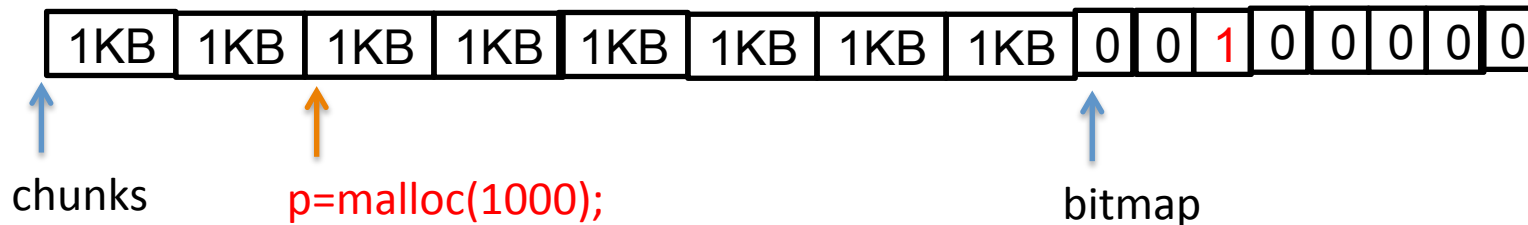


```
#define CHUNKSIZE 1<<10;
typedef char[CHUNKSIZE] chunk;
char *bitmap;
chunk *chunks;
size_t n_chunks;
```

```
void init() {
    n_chunks = 128;
    sbrk(n_chunks*sizeof(chunk)+ n_chunks/8);
    chunks = (chunk *)heap_lo();
    bitmap = heap_lo() + n_chunks *CHUNKSIZE;
}
```

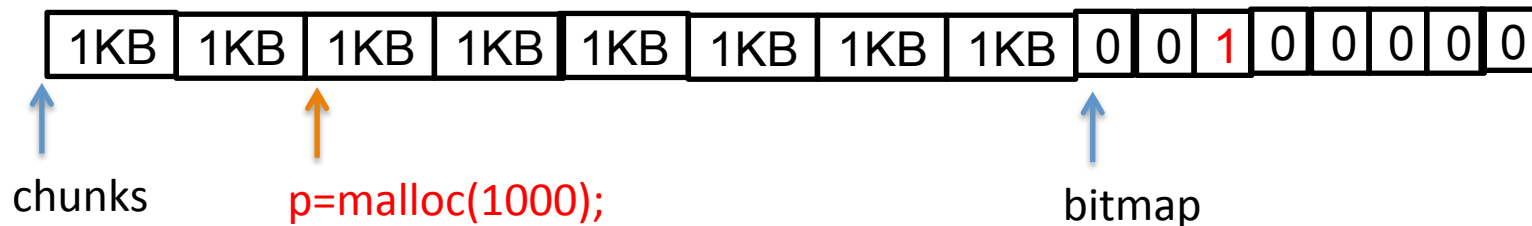
Assume allocator asks for
enough memory from OS
in the beginning

How to bookkeep? Strawman #1



```
void* malloc(size_t sz) {  
    // find out # of chunks needed to fit sz bytes  
    CSZ = ...  
  
    //find csz consecutive free chunks according to bitmap  
    int i = find_consecutive_chunks(bitmap);  
  
    // return NULL if did not find csz free consecutive chunks  
    if (i < 0)  
        return NULL;  
  
    // set bitmap at positions i, i+1, ... i+csz-1  
    bitmap_set_pos(bitmap, i, csz);  
    return (void *)&chunks[i];  
}
```


How to bookkeep? Strawman #1



```
void free(void *p) {  
    i = ((char *)p - (char *)chunks)/sizeof(chunk);  
    bitmap_clear_pos(bitmap, i); //how many bits to clear??  
}
```

- Problem with strawman?
 - free does not know how many chunks allocated
 - wasted space within a chunk (internal fragmentation)
 - wasted space for non-consecutive chunks (external fragmentation)

How to bookkeep? Other Strawmans

- How to support a variable number of variable-sized chunks?
 - Idea #1: use a hash table to map address → [chunk size, status]
 - Idea #2: use a linked list in which each node stores [address, chunk size, status] information.

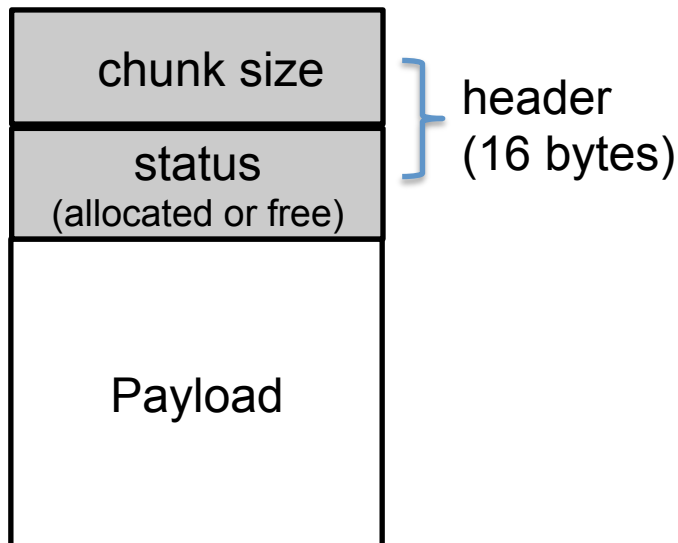
Problems of strawmans?

Implementing a hash table and linked list requires use of a dynamic memory allocator!


How to implement a “list” without use of malloc

Implicit list

- Each chunk contains metadata + payload
 - Metadata (chunk header) stores chunk size and status
- Alignment requirement: 16-byte (aka the starting address of payload must be multiples of 16).
 - Make header 16-byte in size
 - Make chunk multiples of 16 in size.

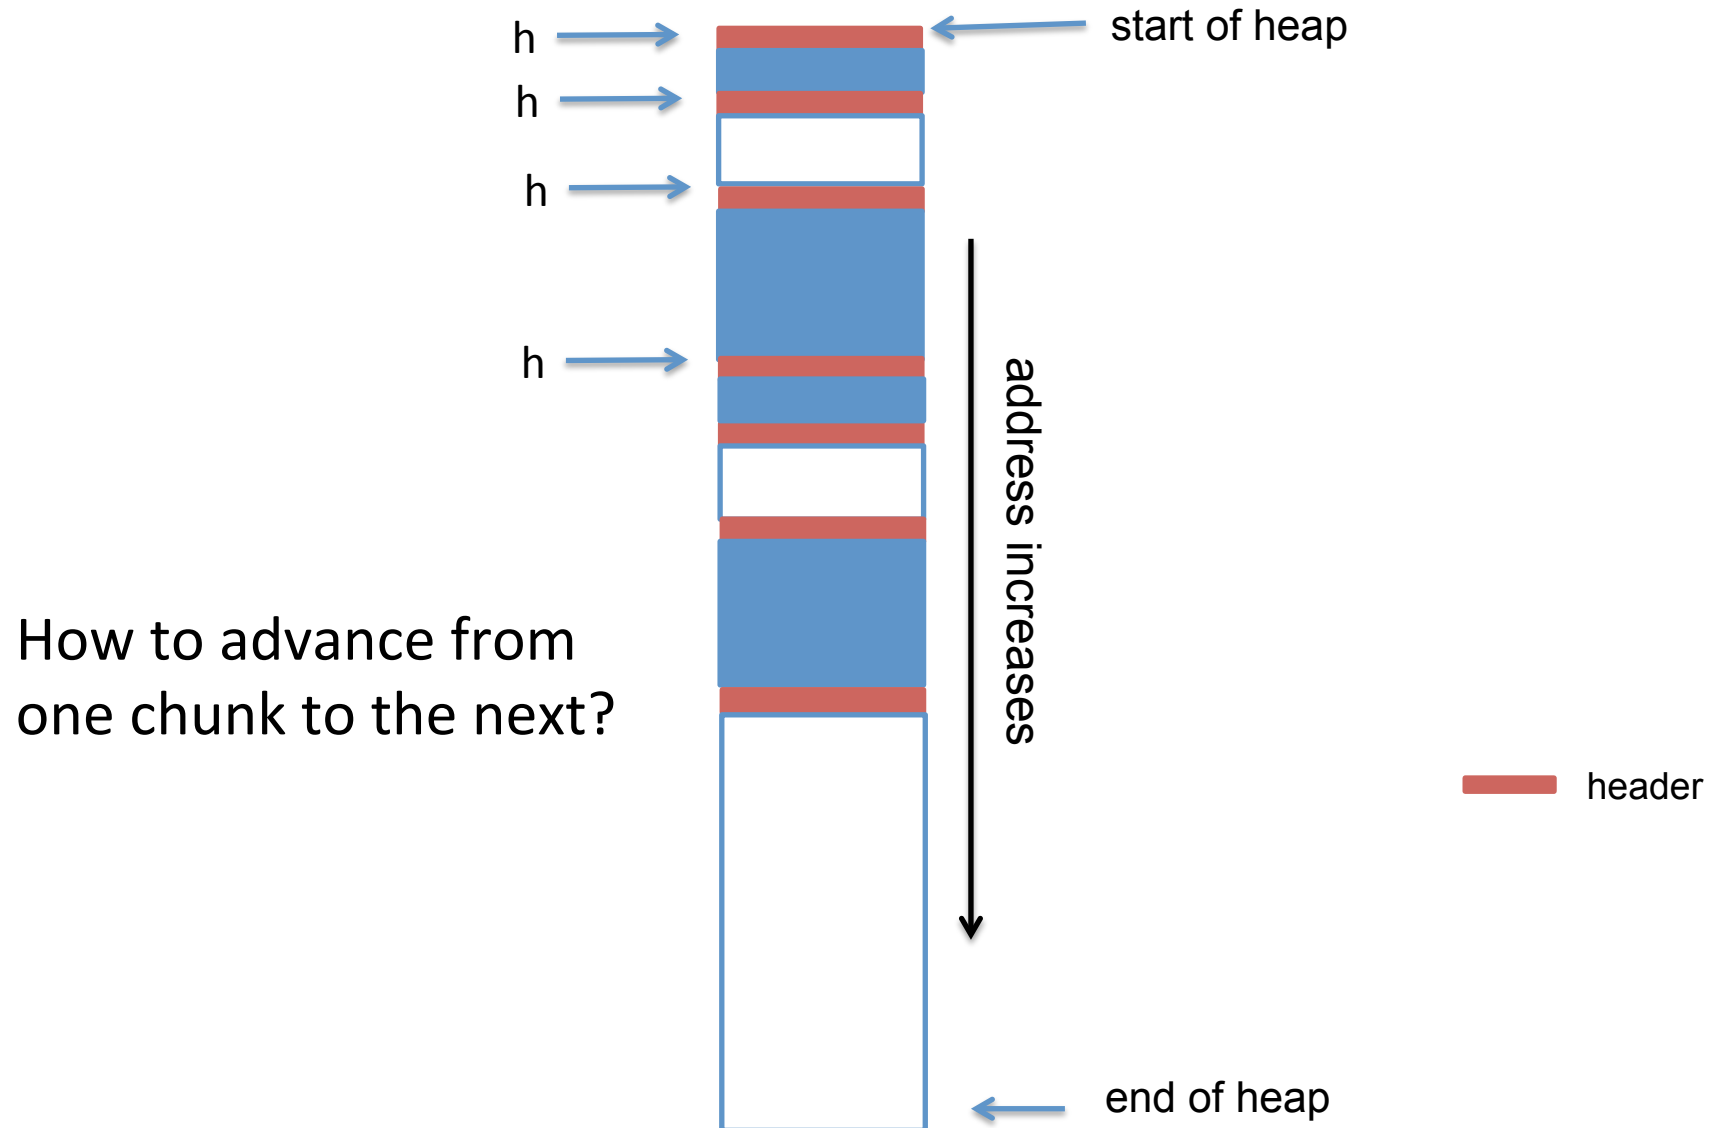


Implicit list: chunk format



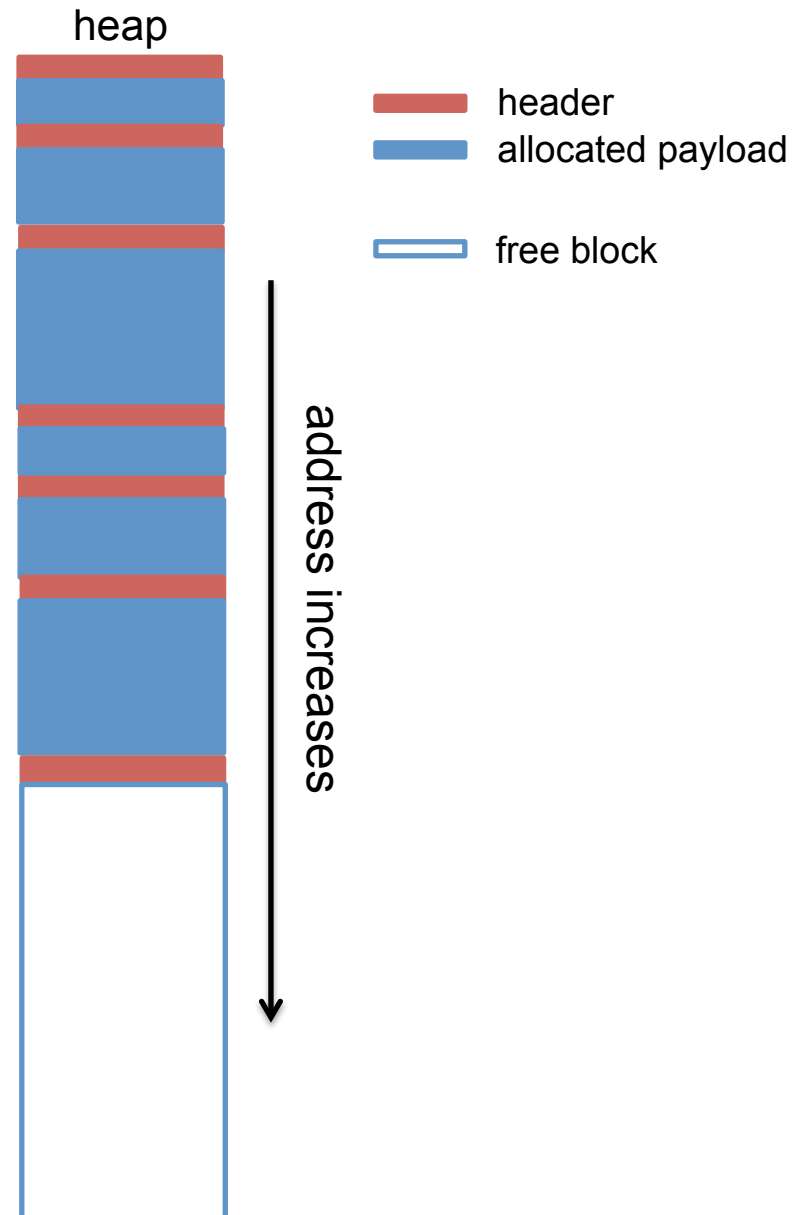
`p = malloc(20)`

Implicit list: heap traversal



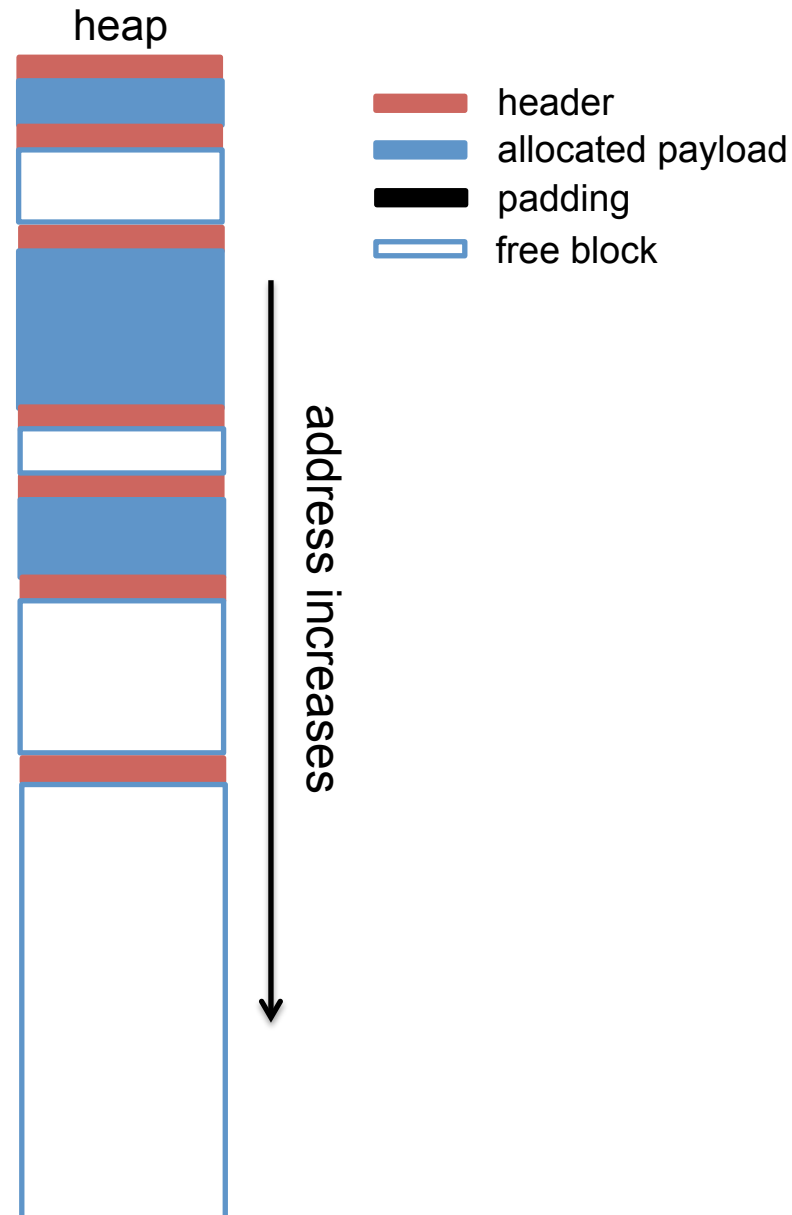
Which chunk to allocate?

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
```



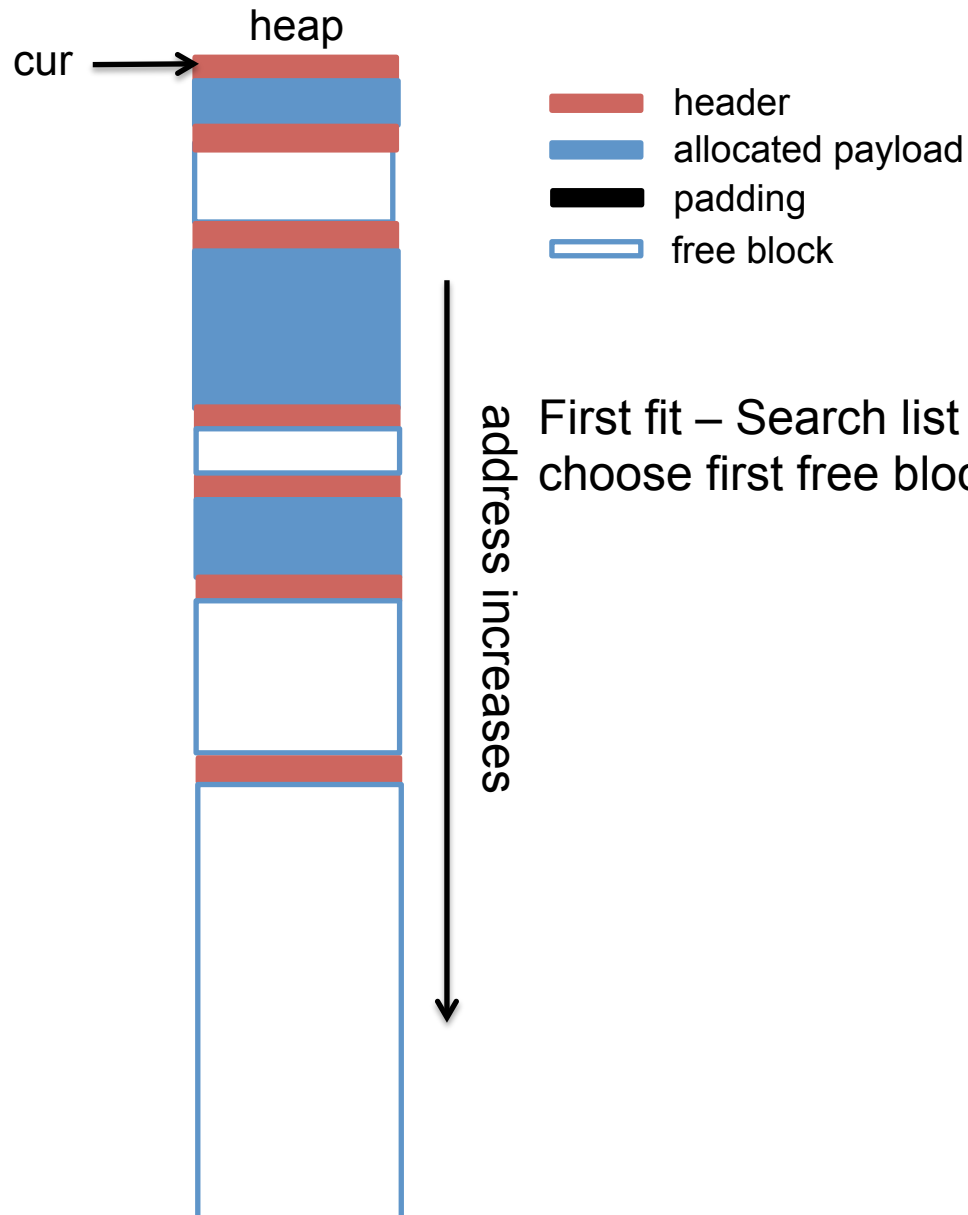
Where to place an allocation?

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
```



First fit

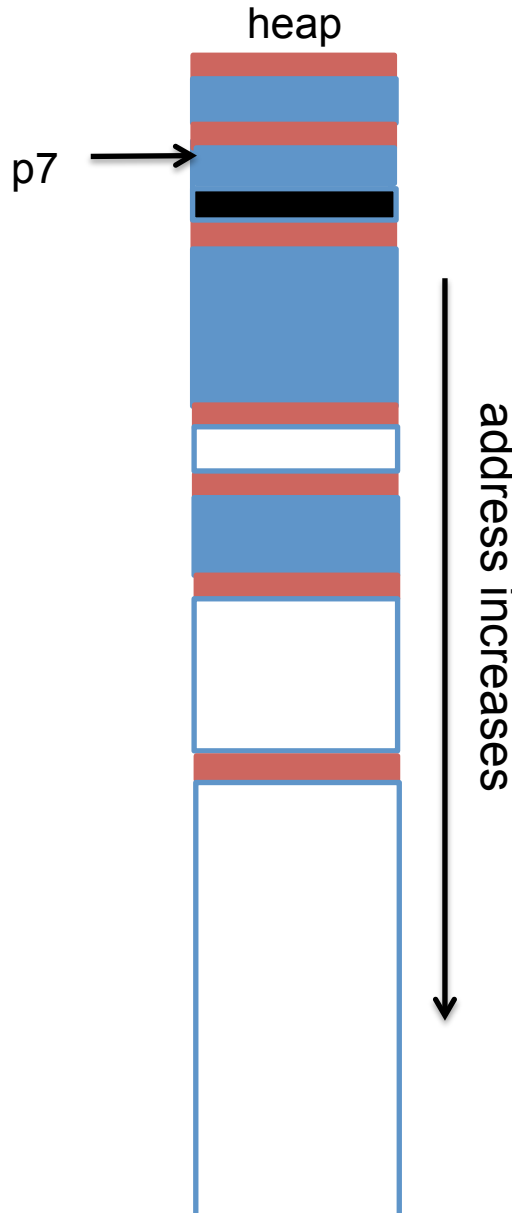
```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
```



First fit – Search list from beginning, choose first free block that fits

First fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
```

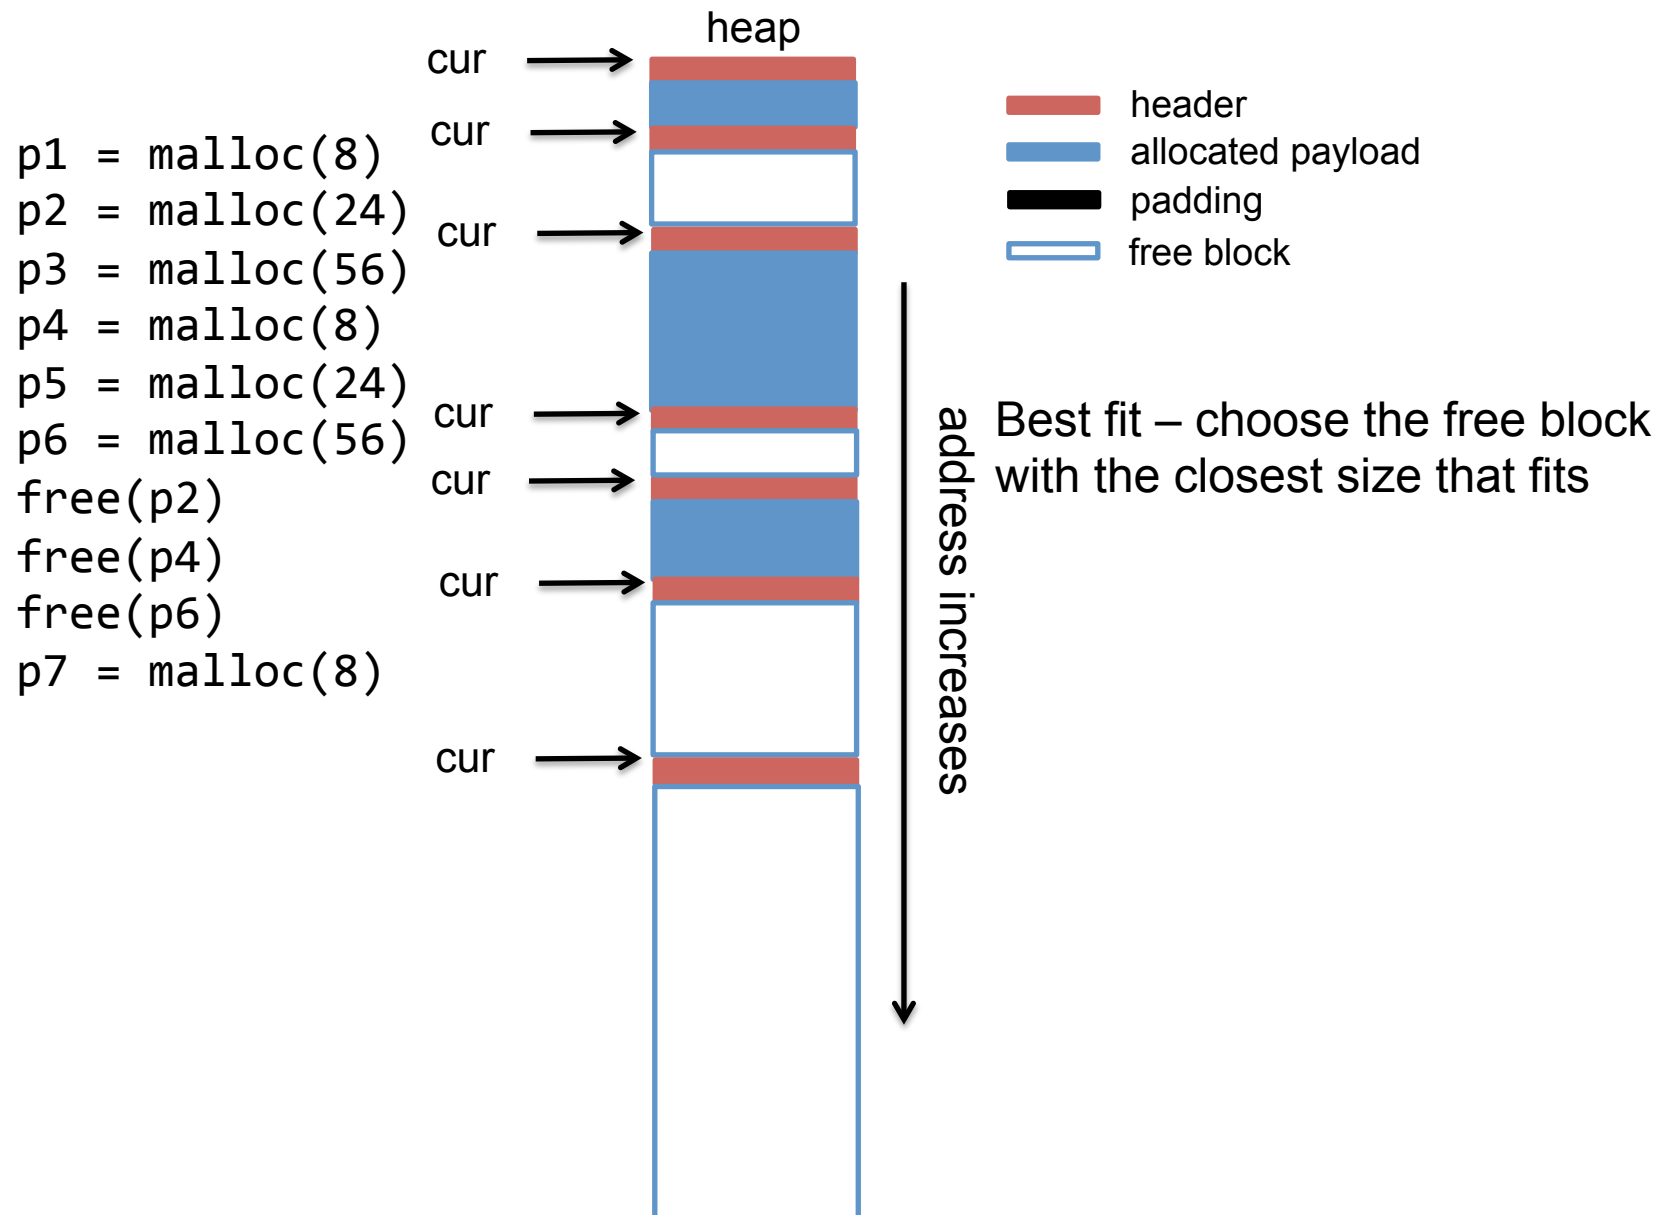


- header
- allocated payload
- padding
- free block

First fit – Search list from beginning, choose first free block that fits

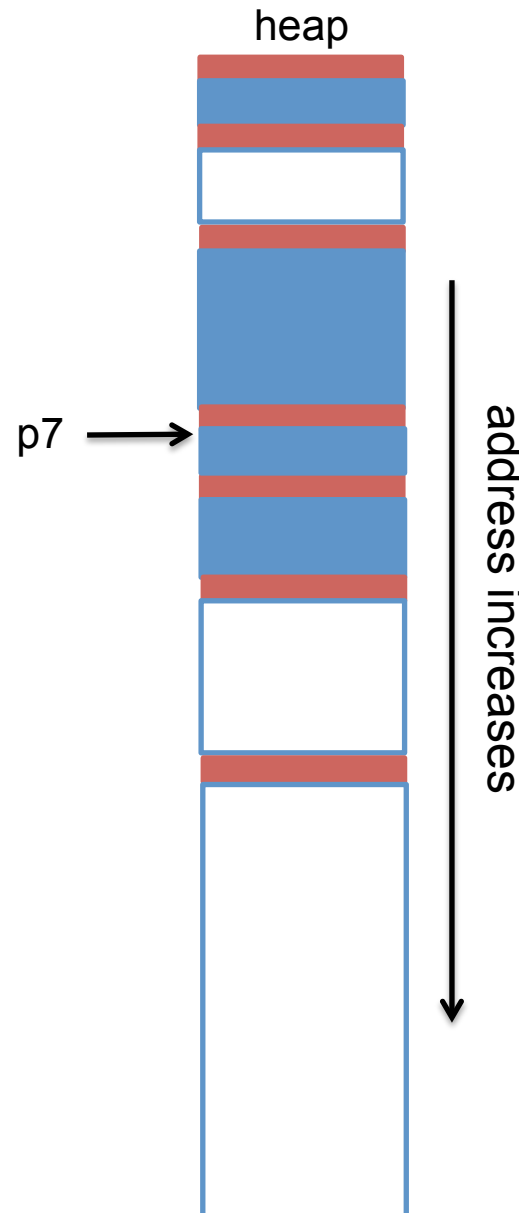
Downside: cause fragmentation at beginning of the heap

Best fit



Best fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
```



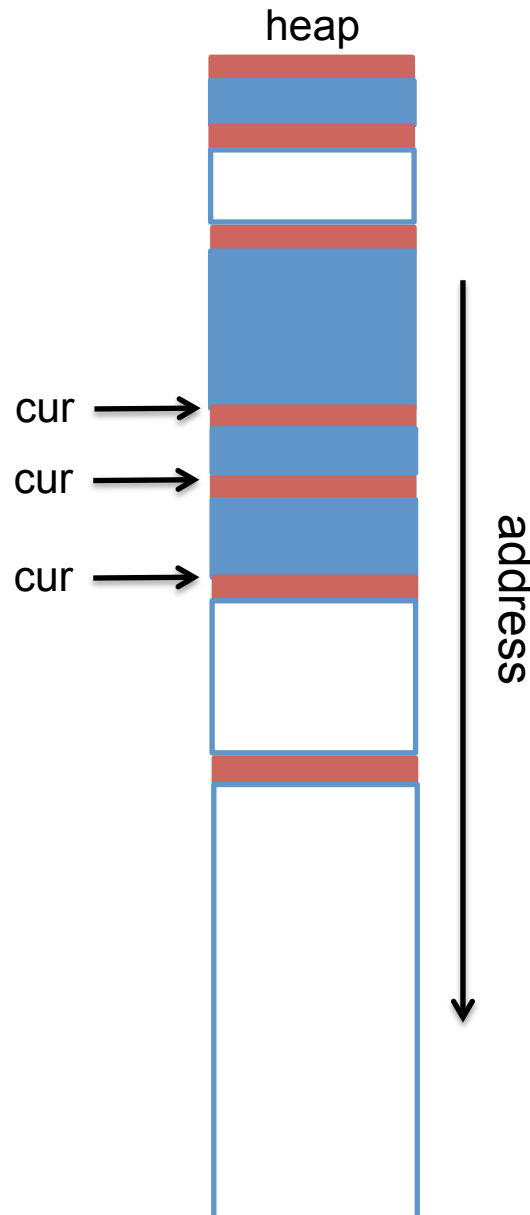
- header
- allocated payload
- padding
- free block

Best fit – choose the free block with the closest size that fits

Downside: run slower than first fit.

Next fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
p8 = malloc(56)
```

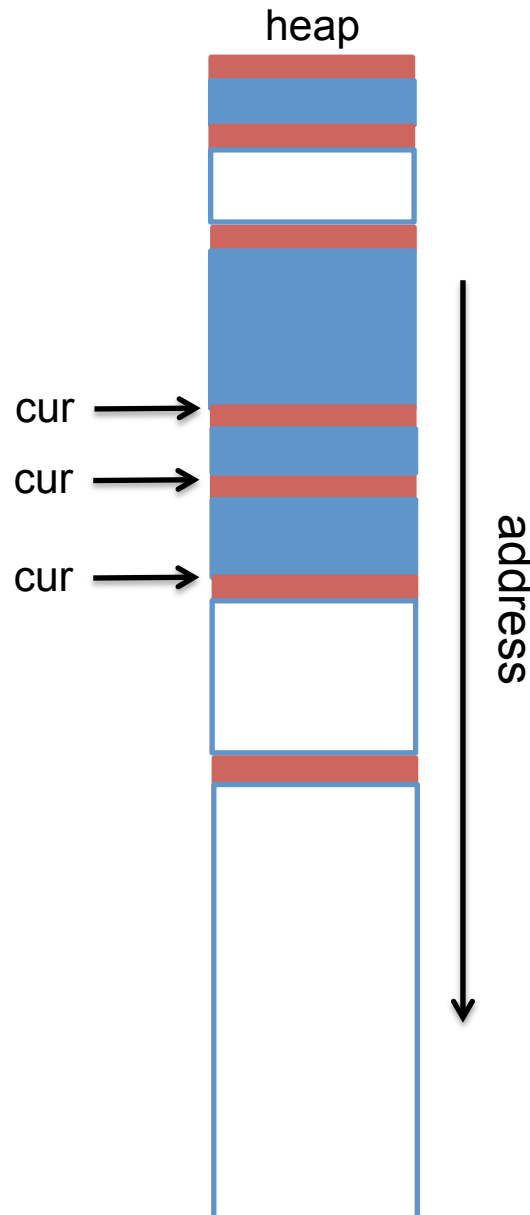


header
allocated payload
padding
free block

Next fit – like first-fit, but search starts from where the previous search left off.

Next fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
p8 = malloc(56)
```




header
allocated payload
padding
free block

Next fit – like first-fit, but search starts from where the previous search left off.

Next fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
p8 = malloc(56)
```

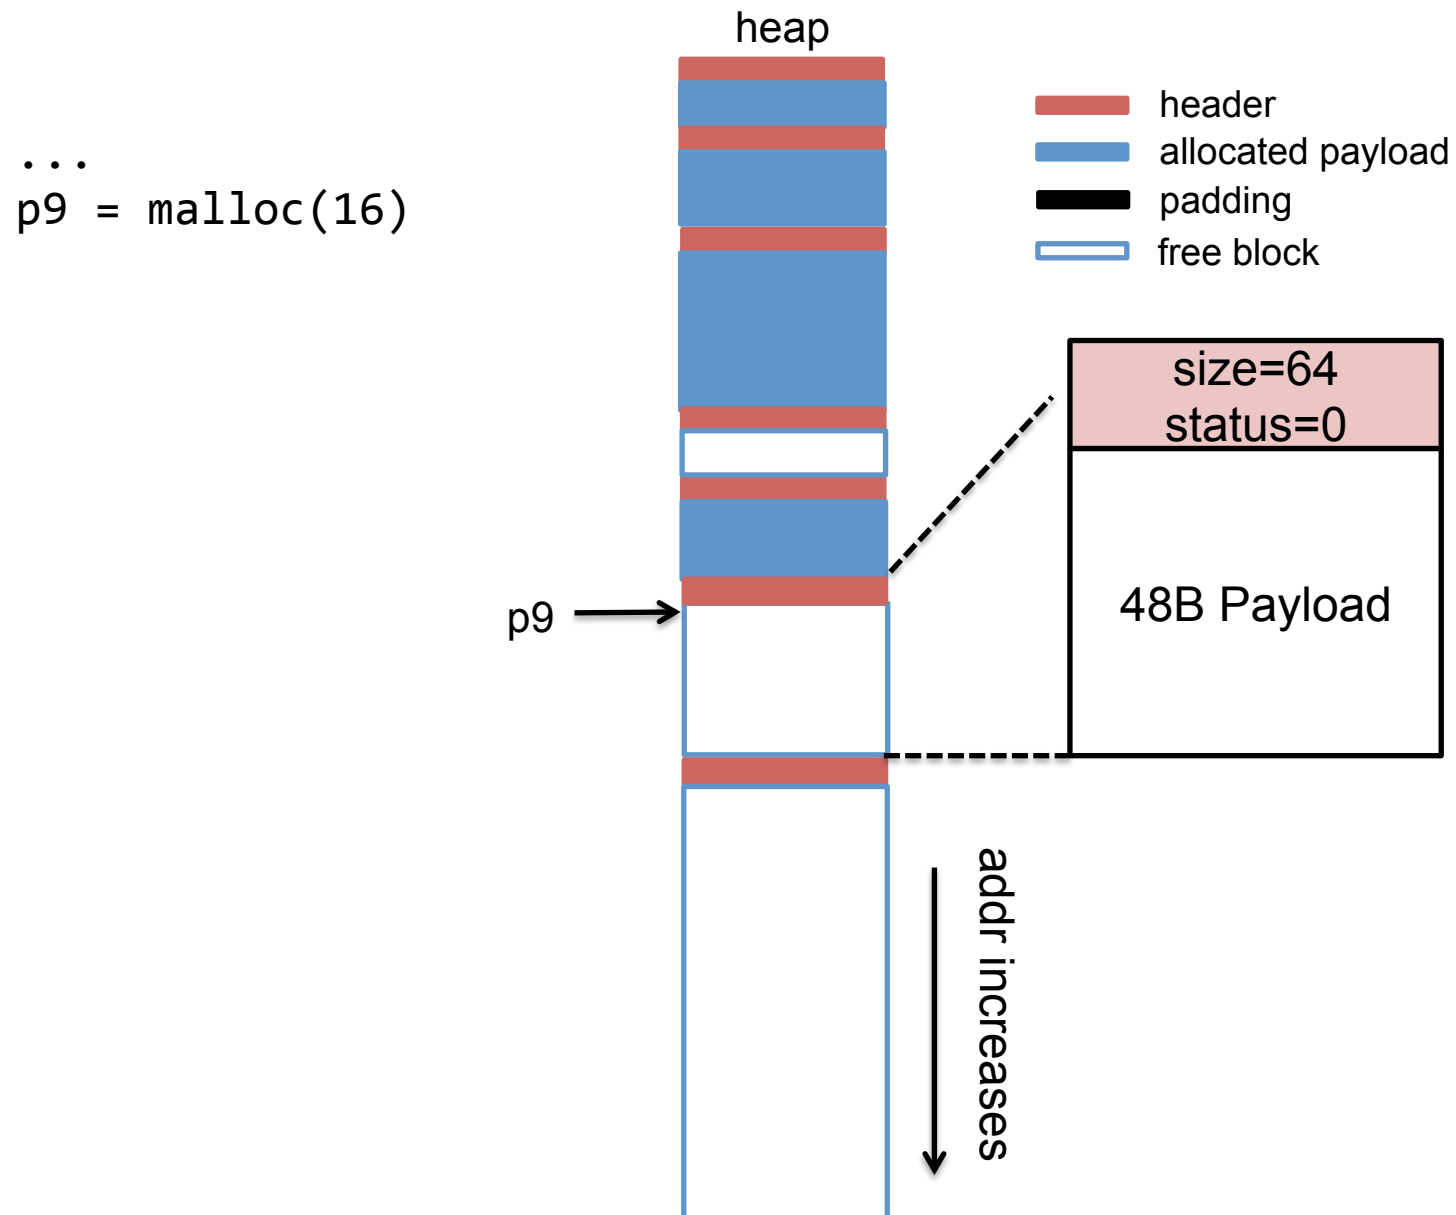


- header
- allocated payload
- padding
- free block

Next fit – like first-fit, but search starts from where the previous search left off.

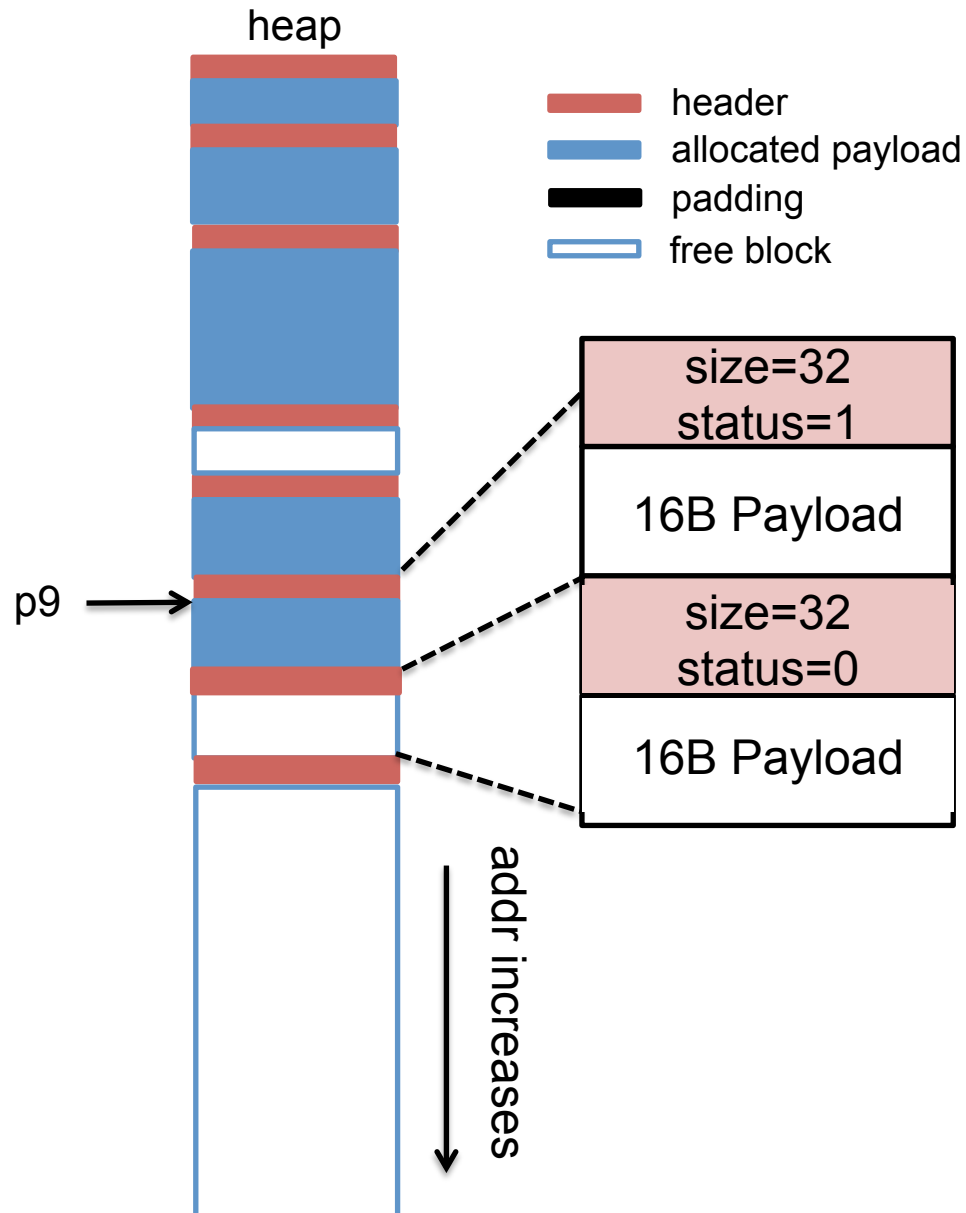
Next fit runs faster than first fit,
but fragmentation is worse.

Splitting a free block

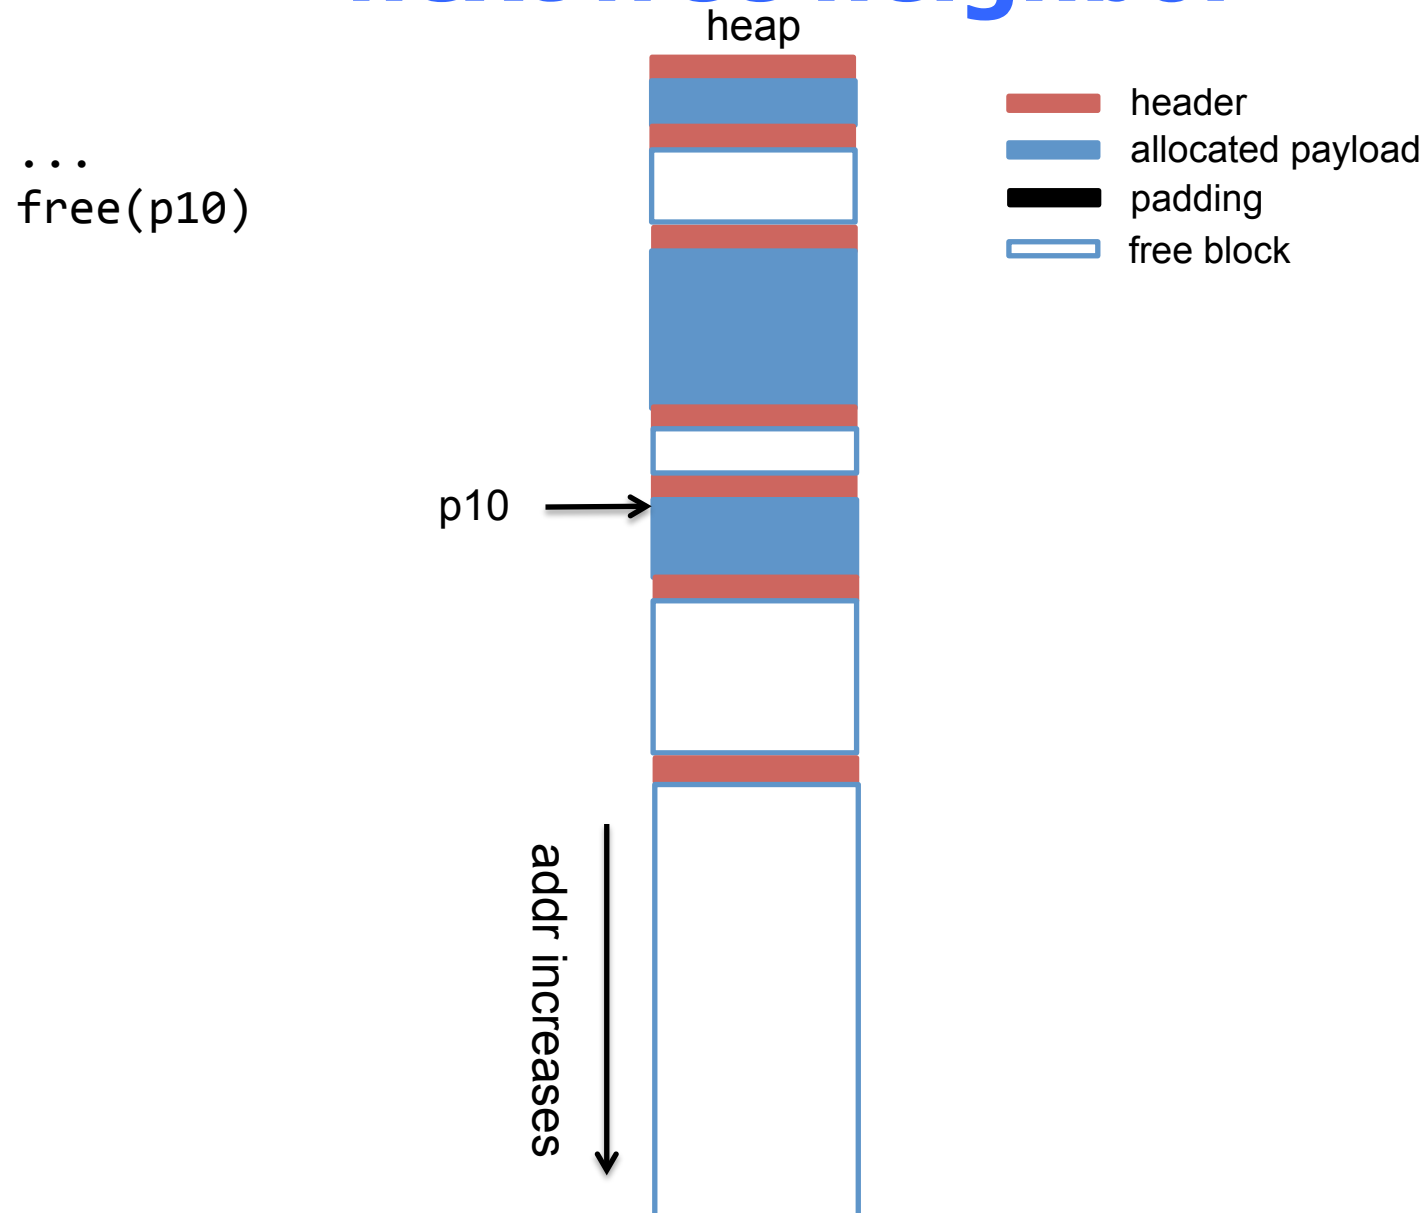


Splitting a free block

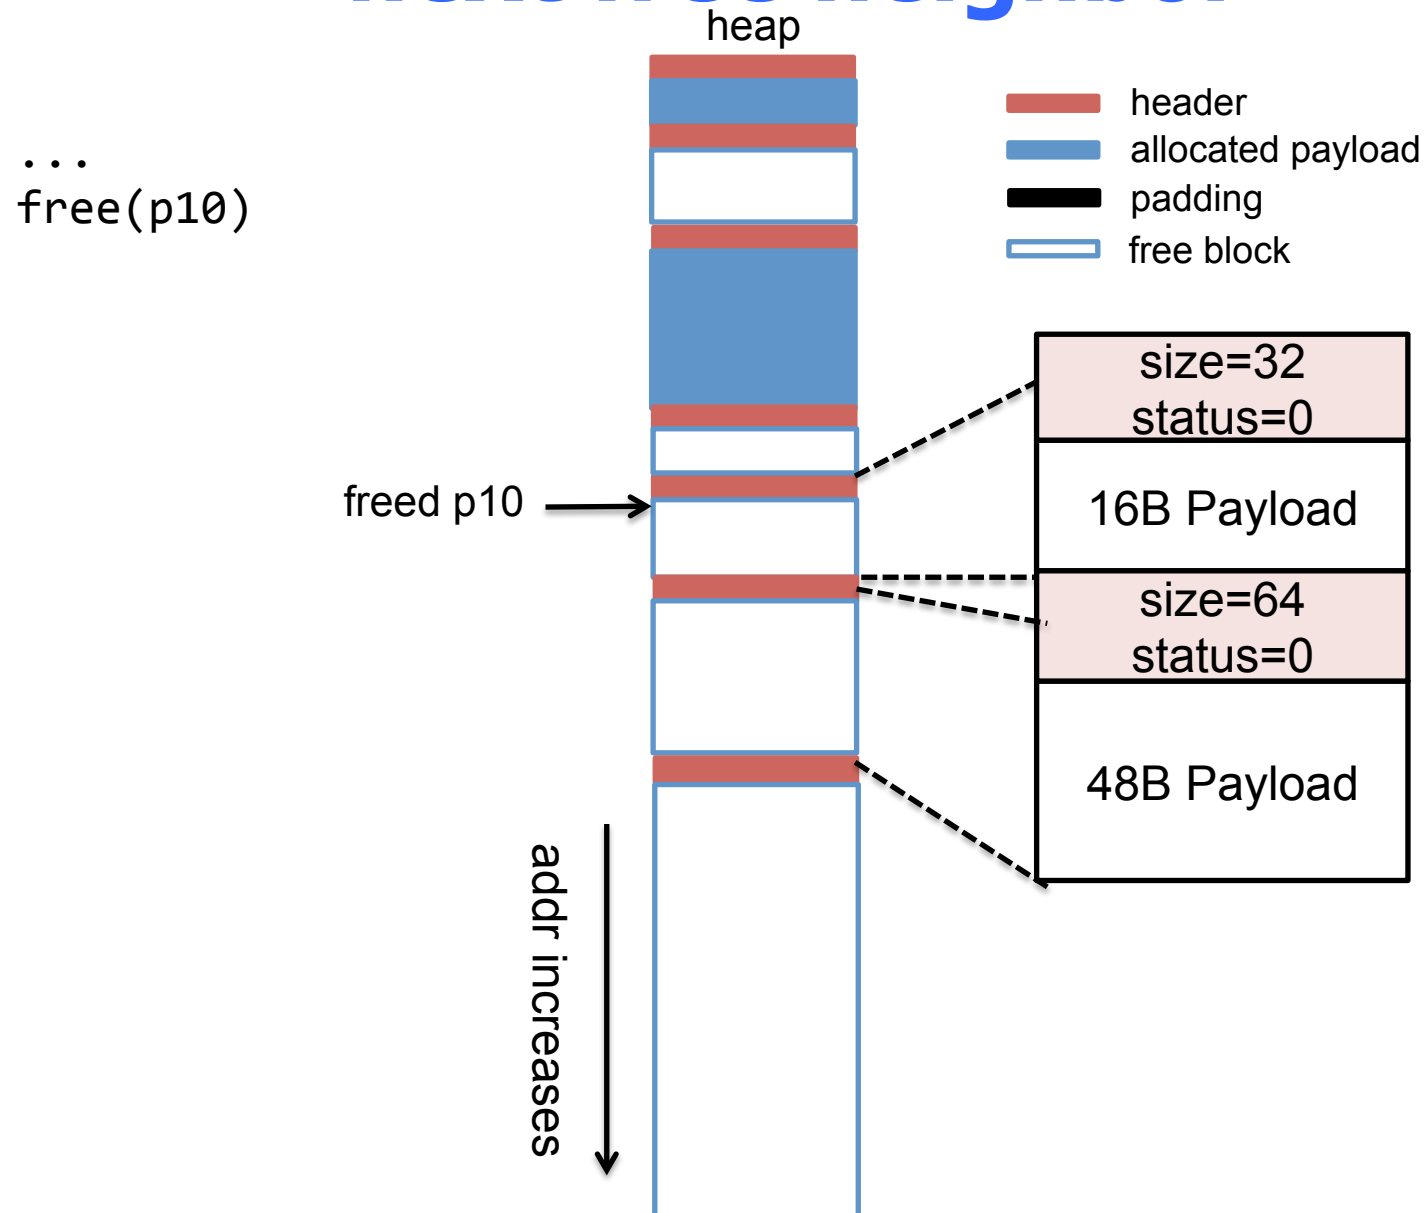
...
p9 = malloc(16)



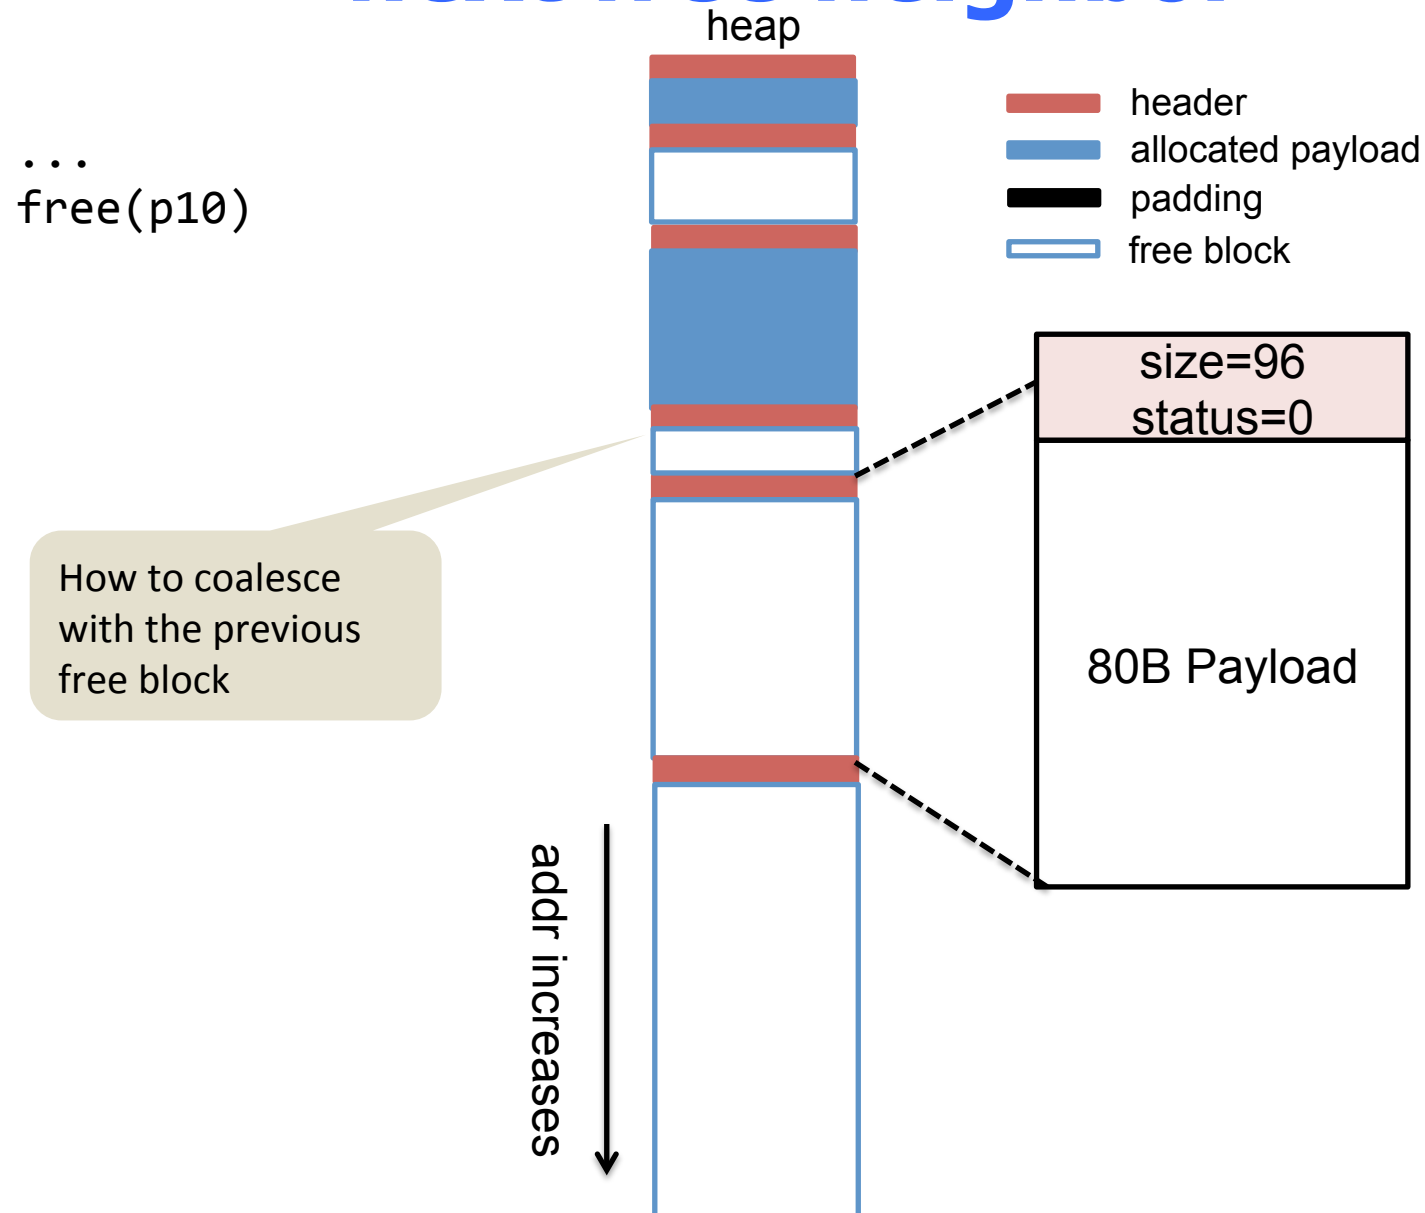
Coalescing a free block with its next free neighbor



Coalescing a free block with its next free neighbor



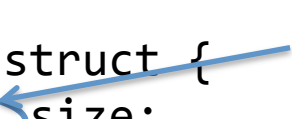
Coalescing a free block with its next free neighbor



implicit list impl (lab 4)

```
typedef struct {  
    size_t size;  
    size_t allocated;  
} header_t;
```

unsigned long



```
int mm_init();  
void *mm_malloc(size_t size);  
void mm_free(void *p);
```

implicit list implementation

```
size_t hdr_size = sizeof(header_t);
```

```
int init() {  
    return 0; //start with empty heap  
}
```

```
void *mm_malloc(size_t size) {  
  
    size = align(size);  
    size_t csz = size + hdr_size;  
    header_t *h = first_fit(csz); //find a free chunk  
    if (!h) {  
        h = ask_os_for_chunk(csz);  
    } else {  
        split(h, csz); //split if necessary  
    }  
    ... // set chunk status to be allocated  
}
```

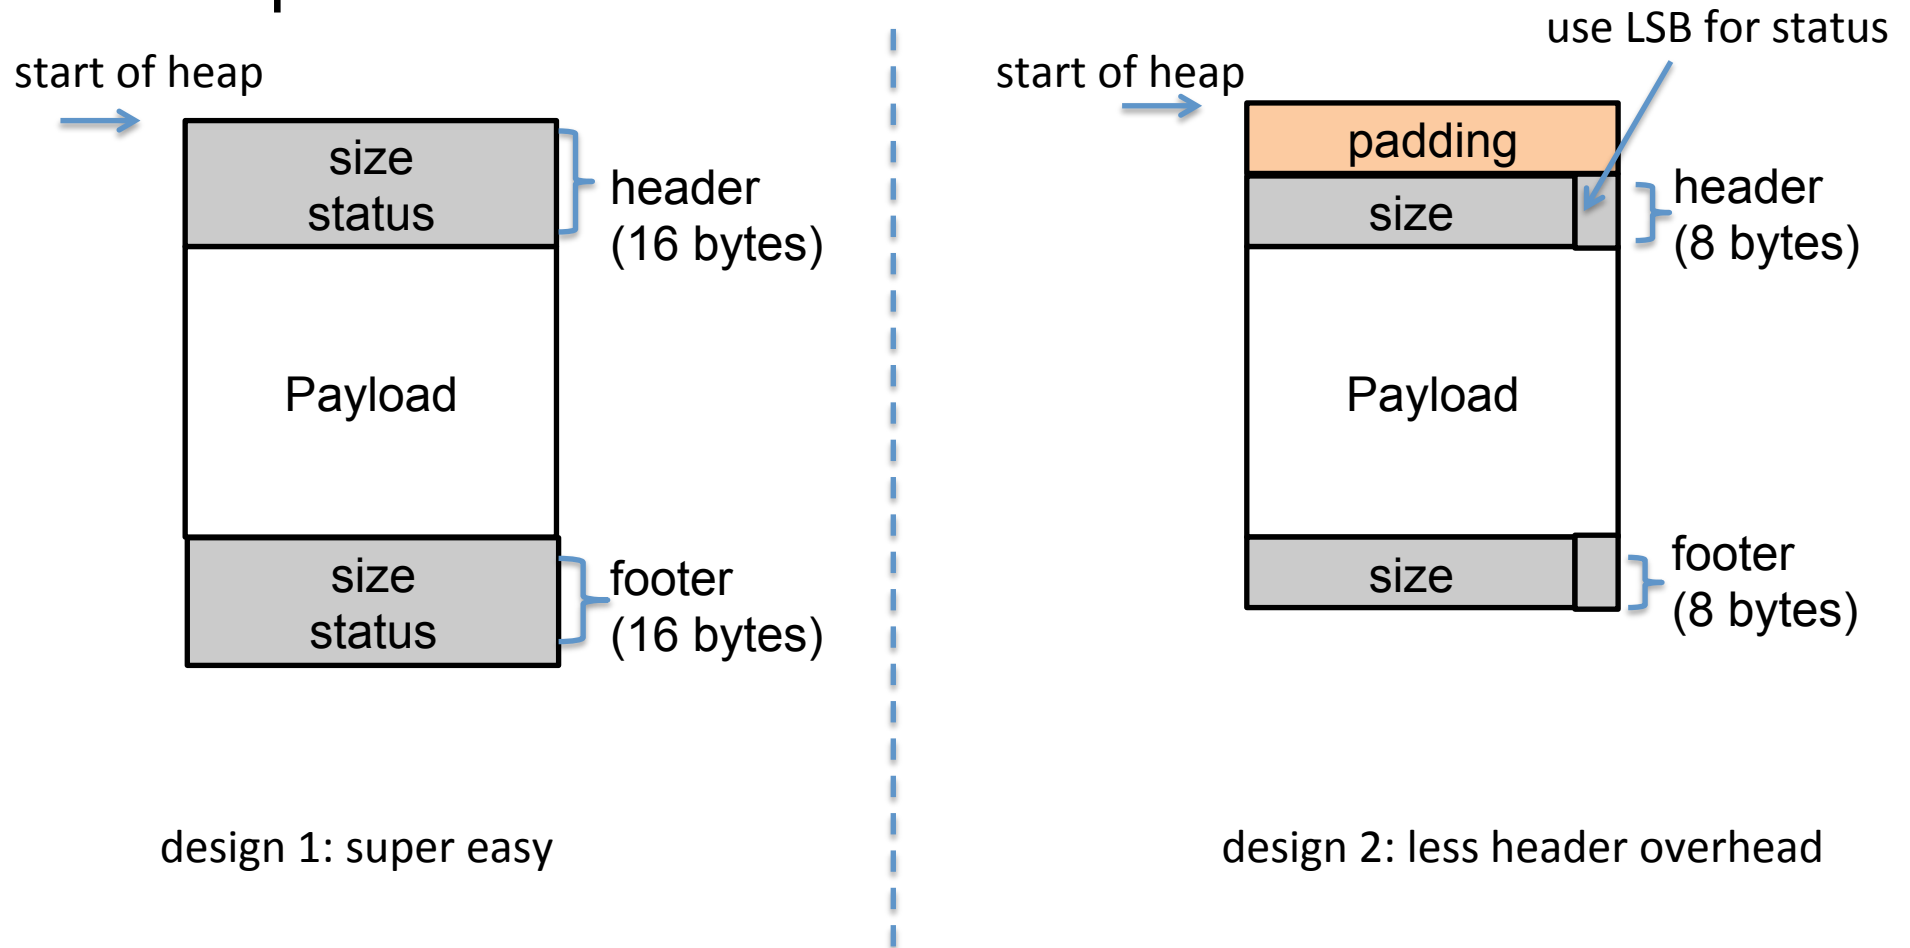
implicit list implementation

```
header_t *payload2header(void *p) {  
  
}
```

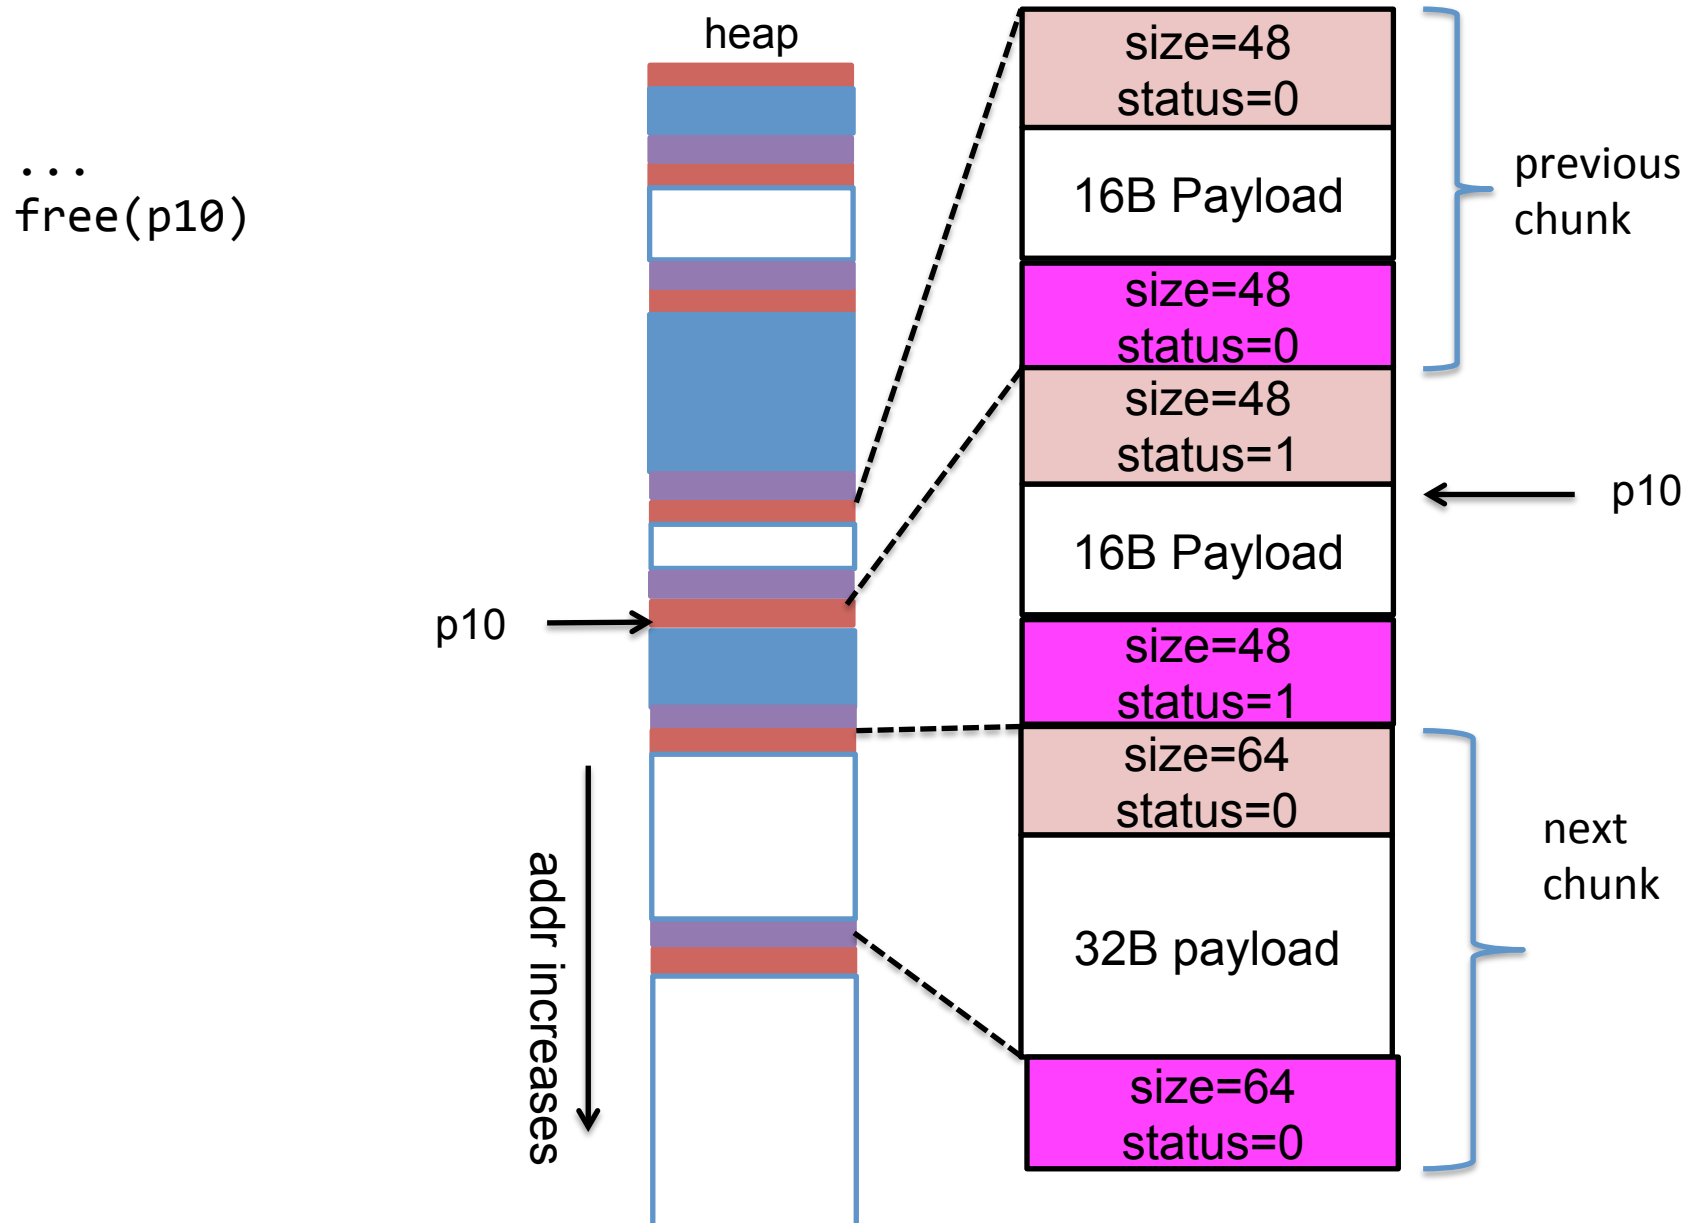
```
void *mm_free(void *p) {  
    header_t *h = payload2header(p);  
    ... //set chunk status to be free  
    coalesce(h);  
}
```

Use footer to coalesce with previous block

- Duplicate header information into the footer




Coalescing prev and next blocks




Coalescing prev and next blocks


...
free(p10)



Recap: malloc using implicit list



Design without footer



Design with footer

Explicit free lists

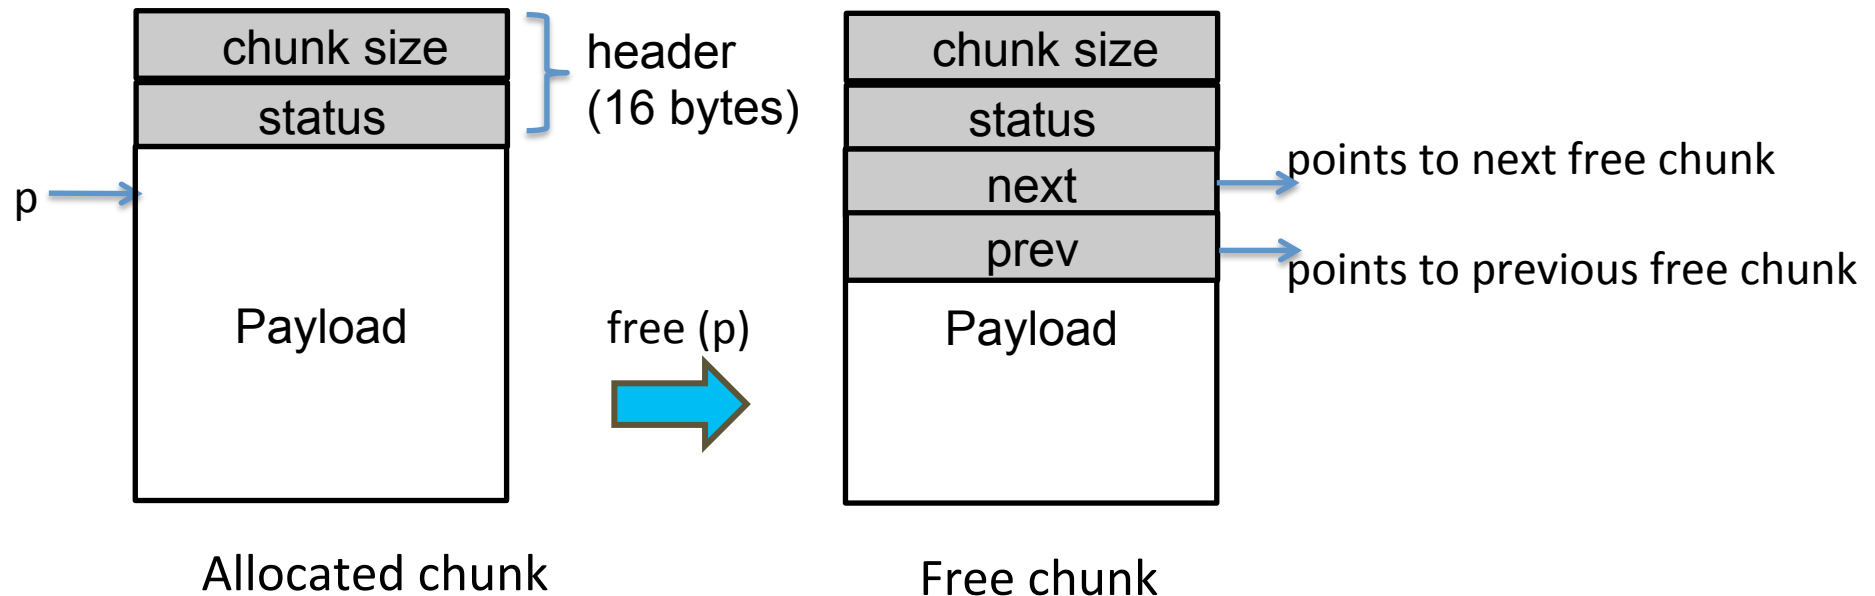
Problems of implicit list:

- Allocation time is linear in # of total (free and allocated) chunks

Explicit free list:

- Maintain a linked list of free chunks only.

Explicit free list



Add footer to coalesce with previous chunk during free

- Question: do we need next/prev fields for allocated blocks?

Answer: No. We do not need to chain together allocated blocks. We can still traverse all blocks (free and allocated) as in the case of implicit list.

- Question: what's the minimal size of a chunk?

Answer: 8 (header) + 8 (footer) + 8 (next pointer) + 8 (previous pointer) = 32 bytes

Explicit list: types, globals, initialization

```
typedef struct {  
    size_t size;  
    size_t status;  
}hdr_t;  
  
typedef struct free_hdr_t {  
    hdr_t header;  
    struct free_hdr *next;  
    struct free_hdr *prev;  
} free_hdr_t;
```

```
free_hdr_t *freelist;
```

```
void init() {  
    freelist = NULL;  
}
```

Explicit list: globals, initialization

```
void *malloc(size_t size) {  
    size = align(size);  
    size_t csz = sizeof(hdr_t) + size;           //size of chunk to allocate  
  
    free_hdr_t *h = first_fit(csz);  
    if (!h) {  
        h = ask_os_for_chunk(csz);  
    } else {  
        split(h);                               //split if necessary  
    }  
    return (char *)h + sizeof(hdr_t);           //return pointer to payload  
}
```


Explicit list: find a chunk

```
void list_insert(free_hdr_t **head, free_hdr_t *node) {  
    ...  
}
```

```
void list_delete(free_hdr_t **head, free_hdr_t *node) {  
    ...  
}
```

```
//find a chunk in freelist whose size is >= csz  
//detach the found chunk from freelist
```

```
free_hdr_t *  
first_fit(size_t csz) {  
    free_hdr_t *curr = freelist;  
    while (curr) {  
        if (curr->header.size) >= sz) {  
            list_delete(&freelist, curr);  
            return curr;  
        }  
        curr = curr->next;  
    }  
    return NULL;  
}
```


Explicit list: free

```
void free(void *p) {  
    header *h = payload2header(p);  
    init_free_chunk((free_hdr_t *)h, h->size);  
  
    coalesce(h);  
    list_insert(&freelist, h);  
}
```

```
void init_free_chunk(free_hdr_t *h, size_t sz) {  
    h->next = h->prev = NULL;  
    h->status = 0;  
    h->size = size;  
}
```

```
void coalesce(free_hdr_t *h) {  
  
    header_t *next = next_chunk((header_t)h);  
    if (!next->status) {  
        list_delete(&freelist, (free_hdr_t *)next);  
        h->size += next->size;  
    }  
}
```

Explicit list with footer



```
typedef size_t hdr_t;

typedef struct free_hdr_t {
    hdr_t header;
    struct free_hdr *next;
    struct free_hdr *prev;
} free_hdr_t;
```

```
bool get_status(hdr_t *h) {
    return (*h) & 0x1L;
}
```

```
size_t get_size(hdr_t *h) {
    return (*h) & ~0x1L;
}
```

```
void set_size_status(hdr_t *h,
    size_t sz, bool status) {
    *h = sz | status;
}
```

```
void set_status(header *h, bool status){
    ...
}
```

```
void set_size(header *h, size_t sz) {
    ...
}
```

Segregated list


- Idea: keep multiple freelists
 - each freelist contains chunks of similar sizes

Segregated list: initialize

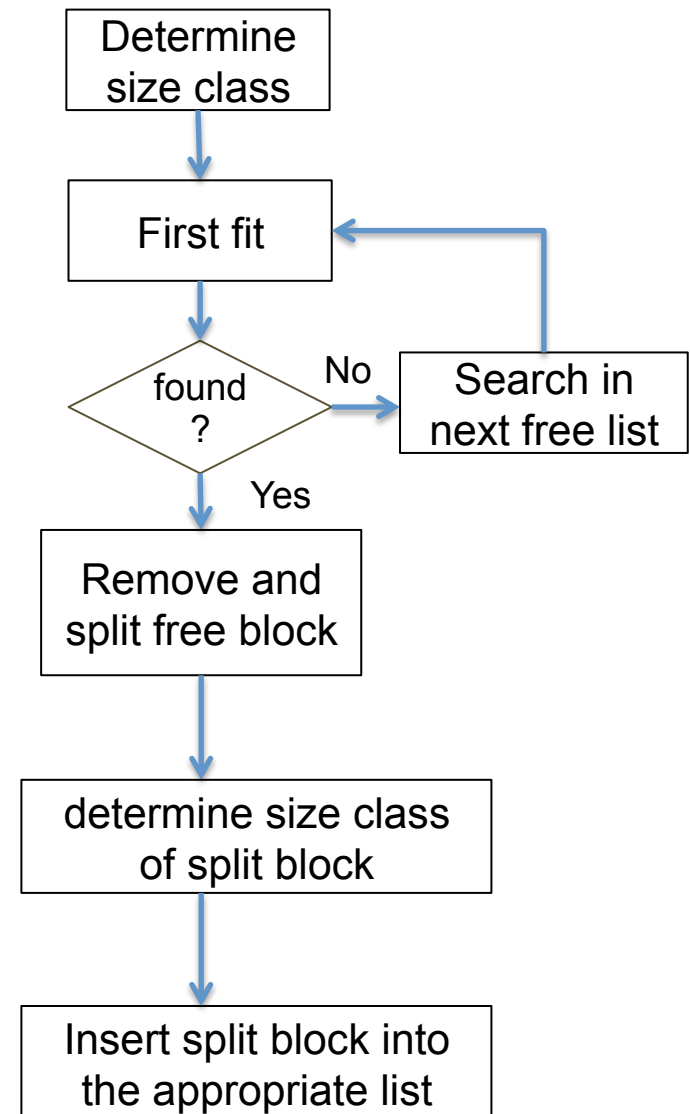

```
#define NLISTS 3
free_hdr* freelists[NLISTS];
size_t size_classes[NLISTS] = {32, 128, -1};
```

```
int which_freelist(size_t s) {
    int ind = 0;
    while (s > size_classes[ind])
        ind++;
    return ind;
}
```

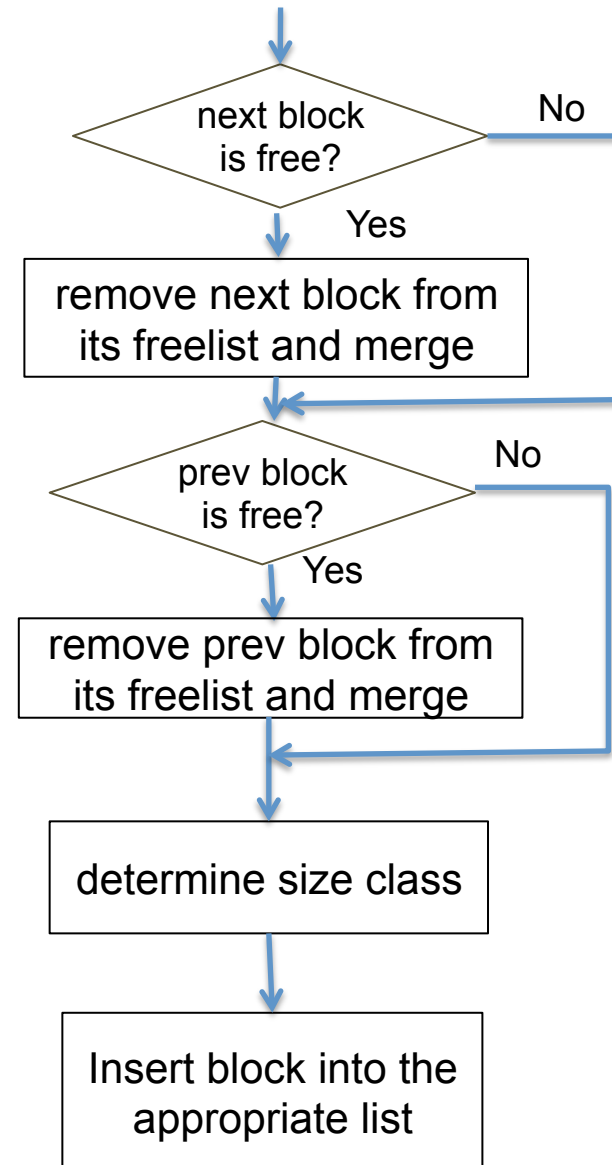

```
void init() {
    for (int i = 0; i < NLISTS; i++) {
        freelists[i] = NULL;
    }
}
```



Segregated list: allocation



Segregated list: free



Buddy System


- A special case of segregated list
 - each freelist has *identically-sized* blocks
 - block sizes are powers of 2
- Advantage over a normal segregated list?
 - Less search time (no need to search within a freelist)
 - Less coalescing time
- Adopted by Linux kernel and jemalloc

Simple binary buddy system

Initialize:

- assume heap starts at 00...00 (real heap address is just a constant offset from 00...00)

(0000 0000 0000 0000)₂



Binary buddy system: allocate

```
p = malloc(16000);
```


Recursive split in half until having the right size

- insert free buddy into appropriate freelist

Addresses of buddies at size 2^m differ in exactly 1-bit at position m (from right)

$(0000\ 0000\ 0000\ 0000)_2$

$(1000\ 0000\ 0000\ 0000)_2$

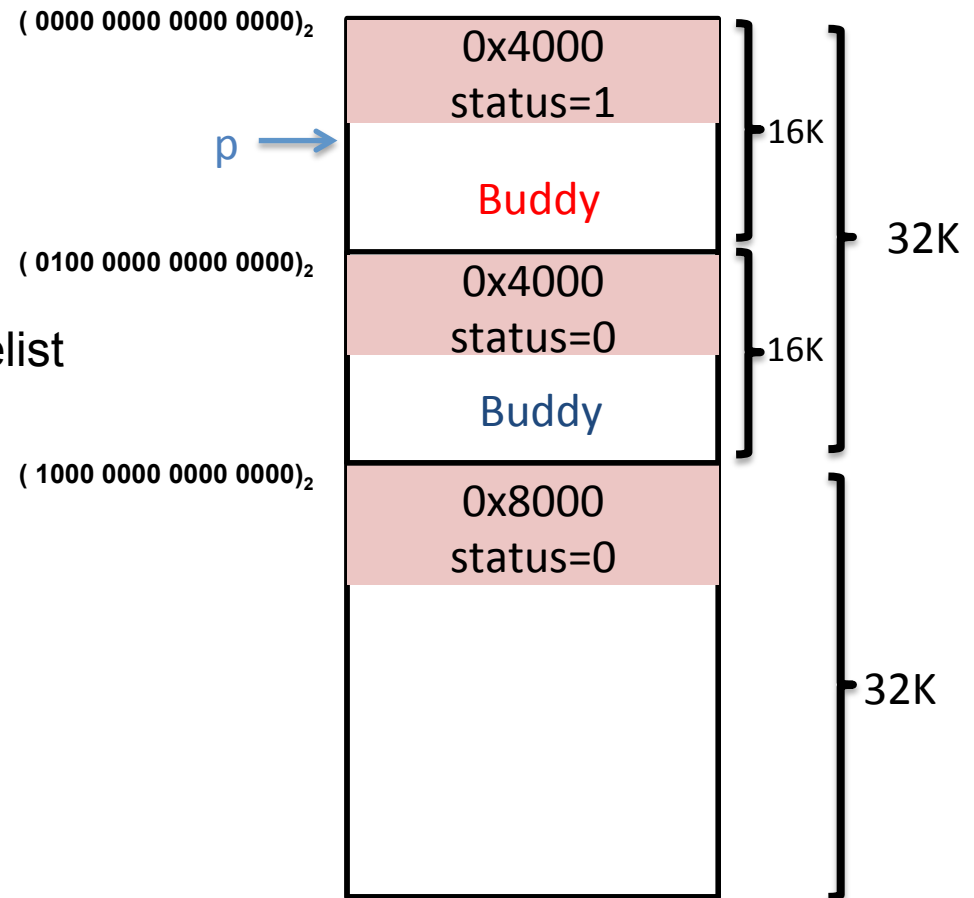


Binary buddy system: allocate

```
p = malloc(16000);
```

Recursive split in half until having the right size

- insert free buddy into appropriate freelist



Binary buddy system: free

`free(p);`

Recursively merge block with buddy

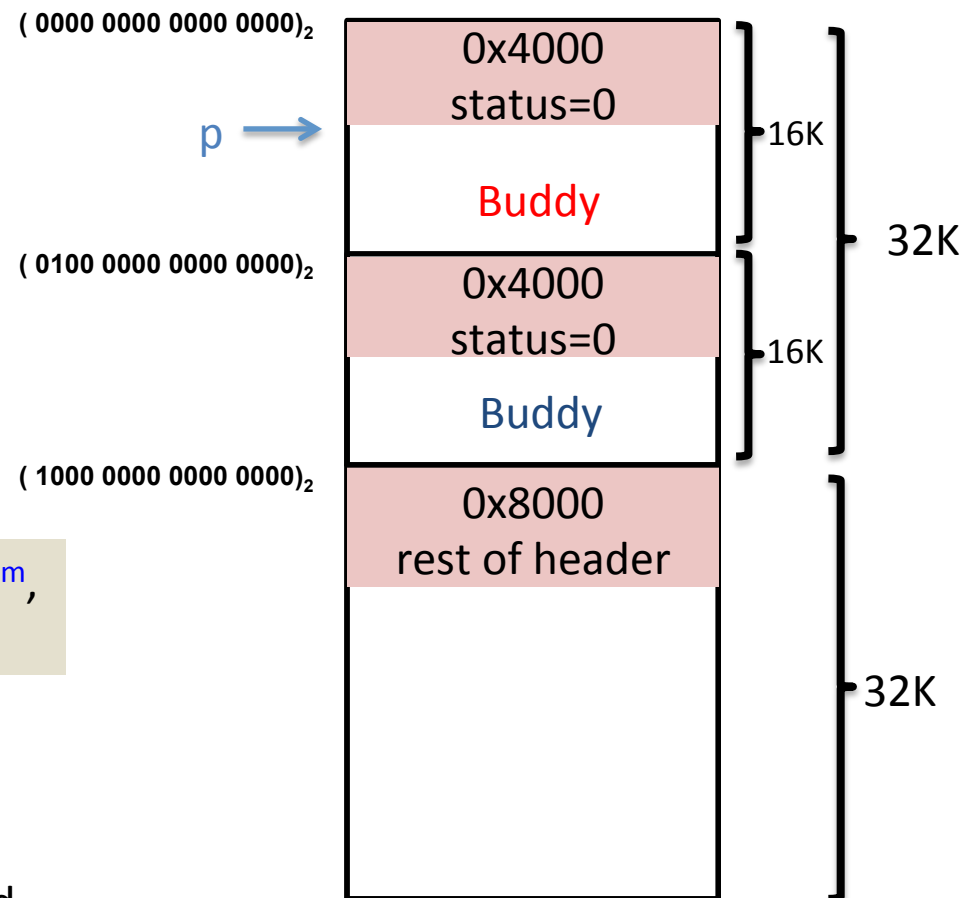
1. Calculate addr of buddy block,
determine buddy status

Question: given addr a of block with size 2^m ,
how to calculate its buddy's address?

$$a \wedge (1 \ll m)$$



any bit XOR 0 = unchanged
any bit XOR 1 = flipped




Binary buddy system: free

`free(p);`

If buddy is free:

2. Detach free buddy from its list
3. Combine with current block



Binary buddy system: free

`free(p);`

Repeat to merge with larger buddy
Insert final block into appropriate
freelist

