

# Dynamic Memory Allocation

Jinyang Li

based on Tiger Wang's slides

# Why dynamic memory allocation?

```
typedef struct node {
    int val;
    struct node *next;
} node;


void
list_insert(node *head, int v) {
    node *np = malloc(sizeof(node));
    np->next = head;
    np->val = v;
    *head = np;
}

int
main(void) {
    char buf[100];
    node *head = NULL;
    while (fgets(buf, 100, stdin)) {
        list_insert(&head, atoi(buf));
    }
}
```

How many nodes to allocate is only known at runtime (when the program executes)

# Dynamic allocation on heap

Question: can one dynamically allocate memory on stack?

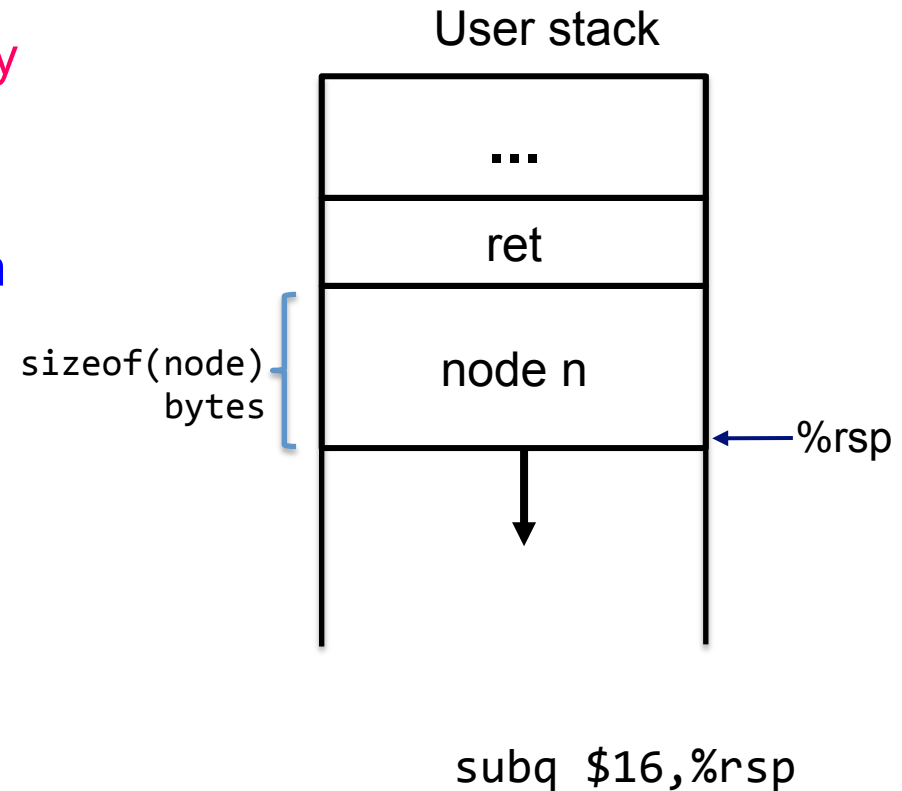


# Dynamic allocation on heap

Question: Is it possible to dynamically allocate memory on stack?

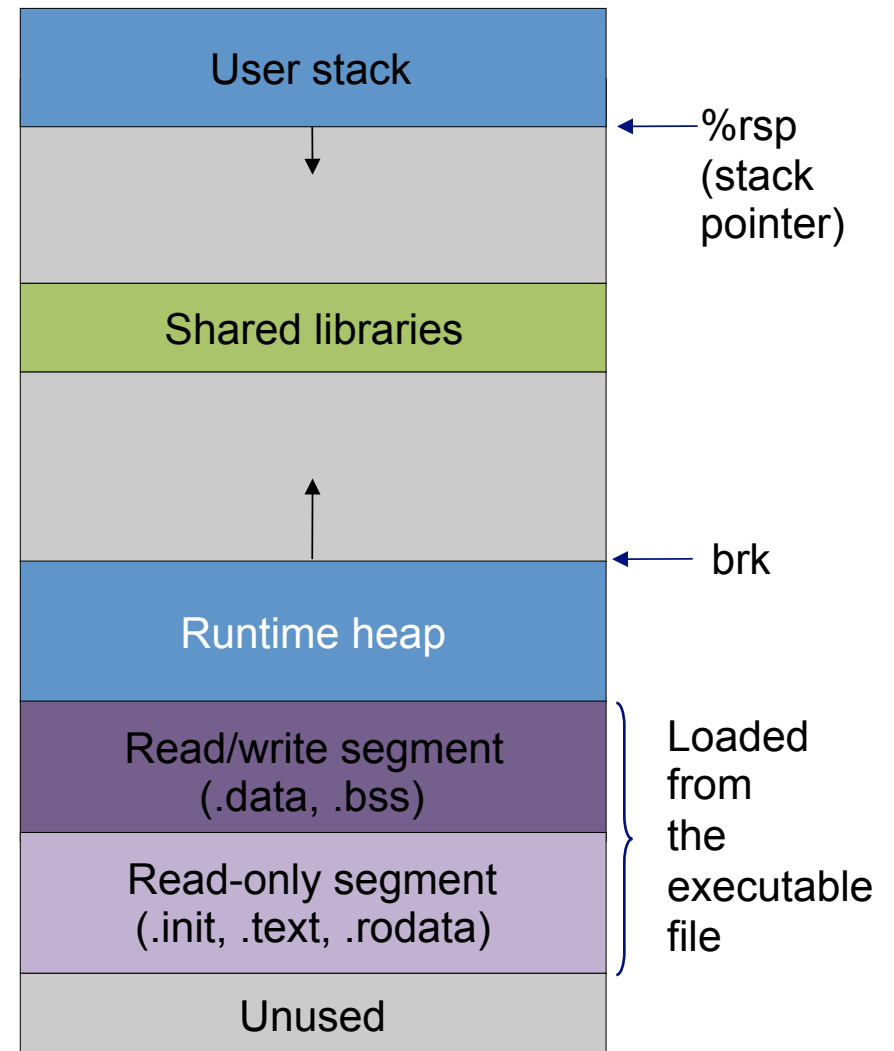
Answer: Yes, but space is freed upon function return

```
void  
list_insert(node *head, int v) {  
    node n;  
    node *np = &n;  
    np->next = head;  
    np->val = v;  
    *head = np;  
}
```



# Dynamic allocation on heap

Question: How to allocate memory on heap?



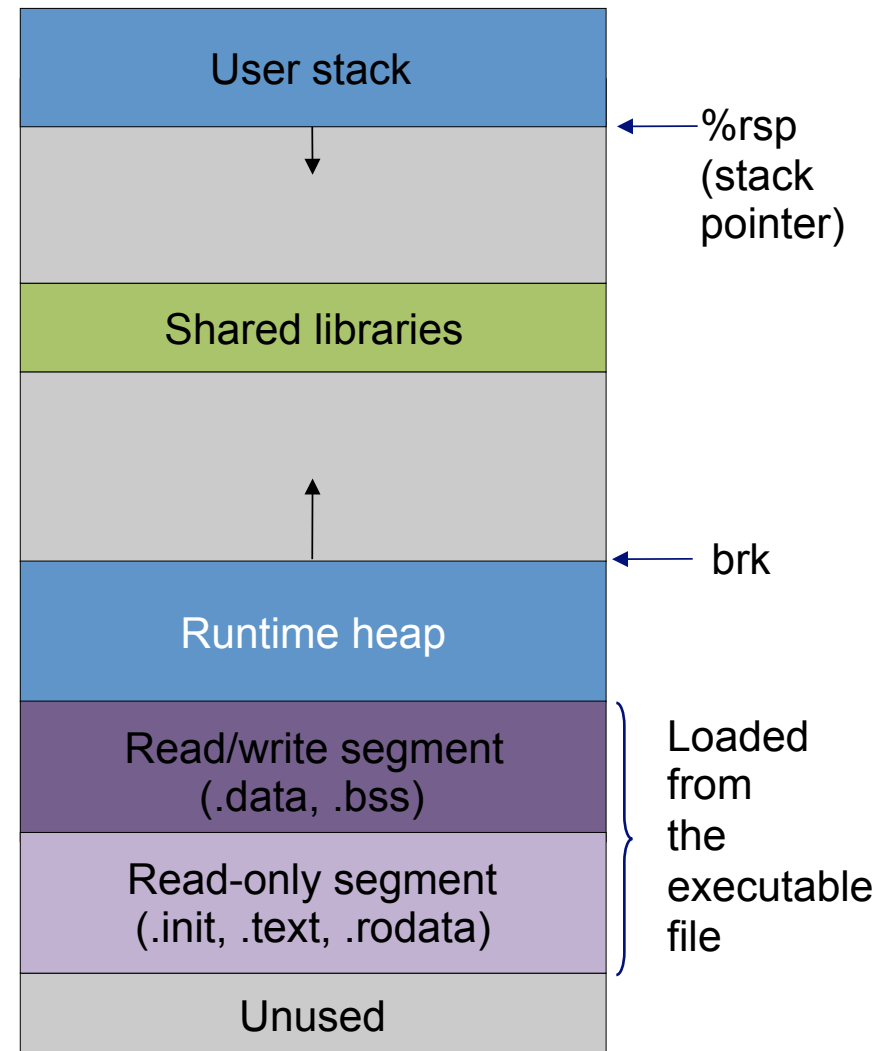
# Dynamic allocation on heap

Question: How to allocate memory on heap?

Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

It increases the top of heap by “size” and returns a pointer to the base of new storage. The “size” can be a negative number.



# Dynamic allocation on heap

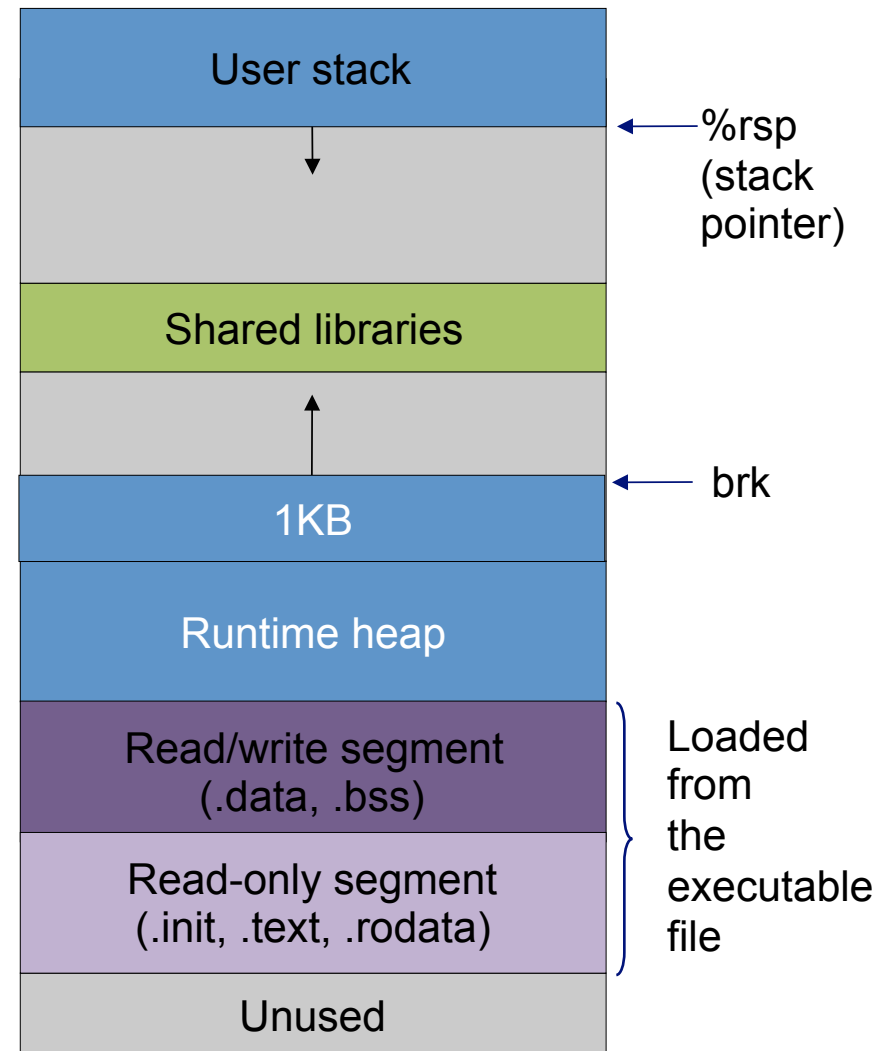
Question: How to allocate memory on heap?

Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

It increases the top of heap by “size” and returns a pointer to the base of new storage. The “size” can be a negative number.

```
p = sbrk(1024) //allocate 1KB
```



# Dynamic allocation on heap

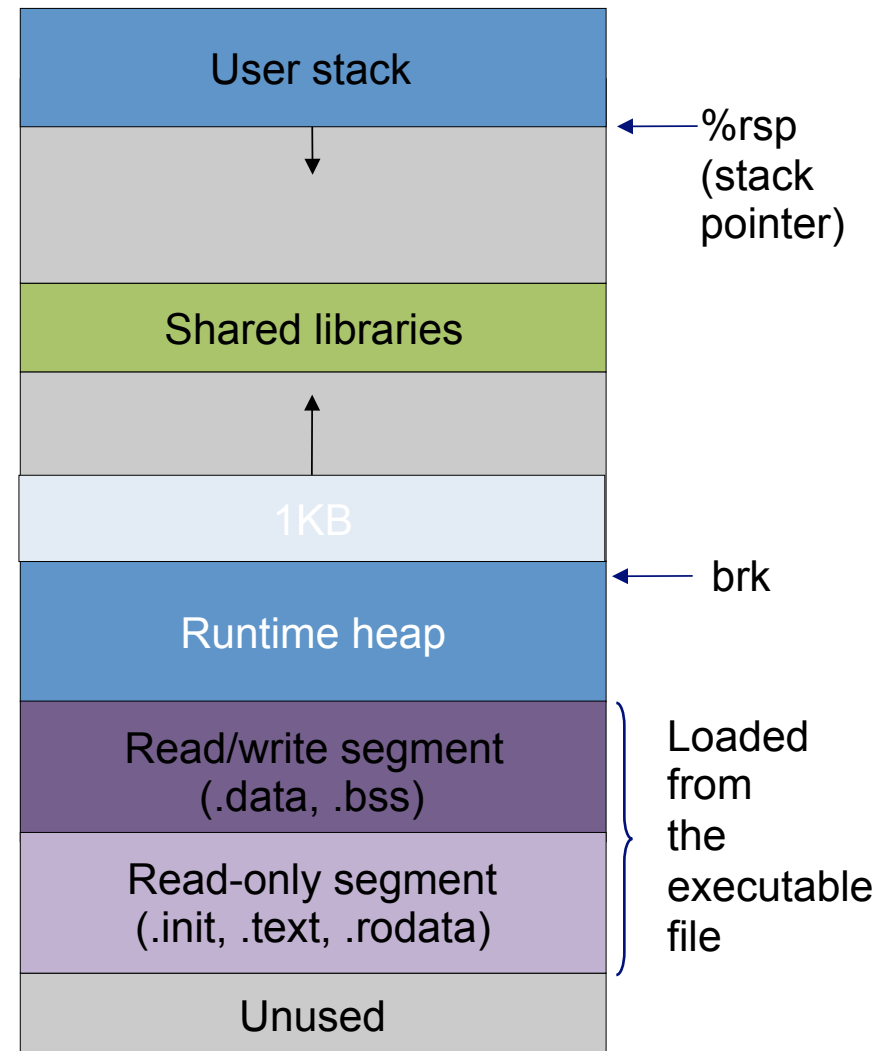
Question: How to allocate memory on heap?

Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

It increases the top of heap by “size” and returns a pointer to the base of new storage. The “size” can be a negative number.

```
p = sbrk(1024) //allocate 1KB  
sbrk(-1024) //free p
```





# Dynamic allocation on heap

Question: How to allocate memory on heap?

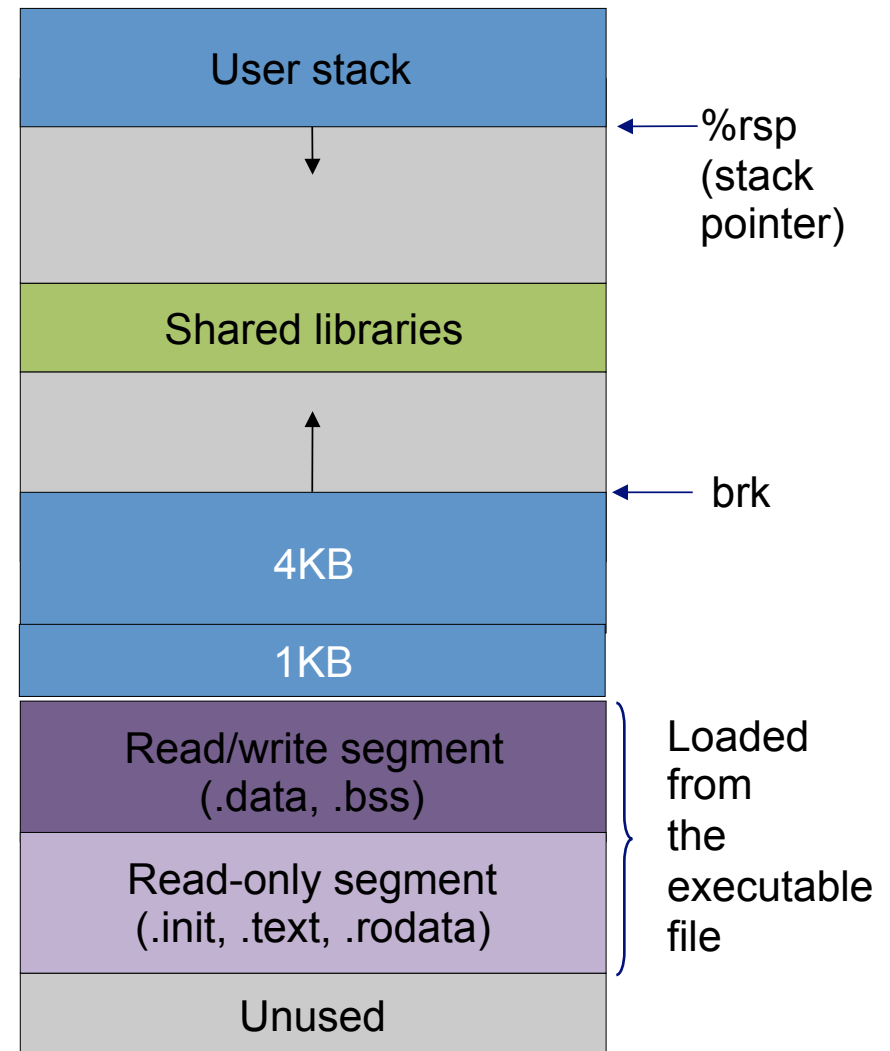
Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

Issue 1 – can only free the memory on the top of heap

```
p1 = sbrk(1024) //allocate 1KB  
p2 = sbrk(4096) //allocate 4KB
```

How to free p1?



# Dynamic allocation on heap

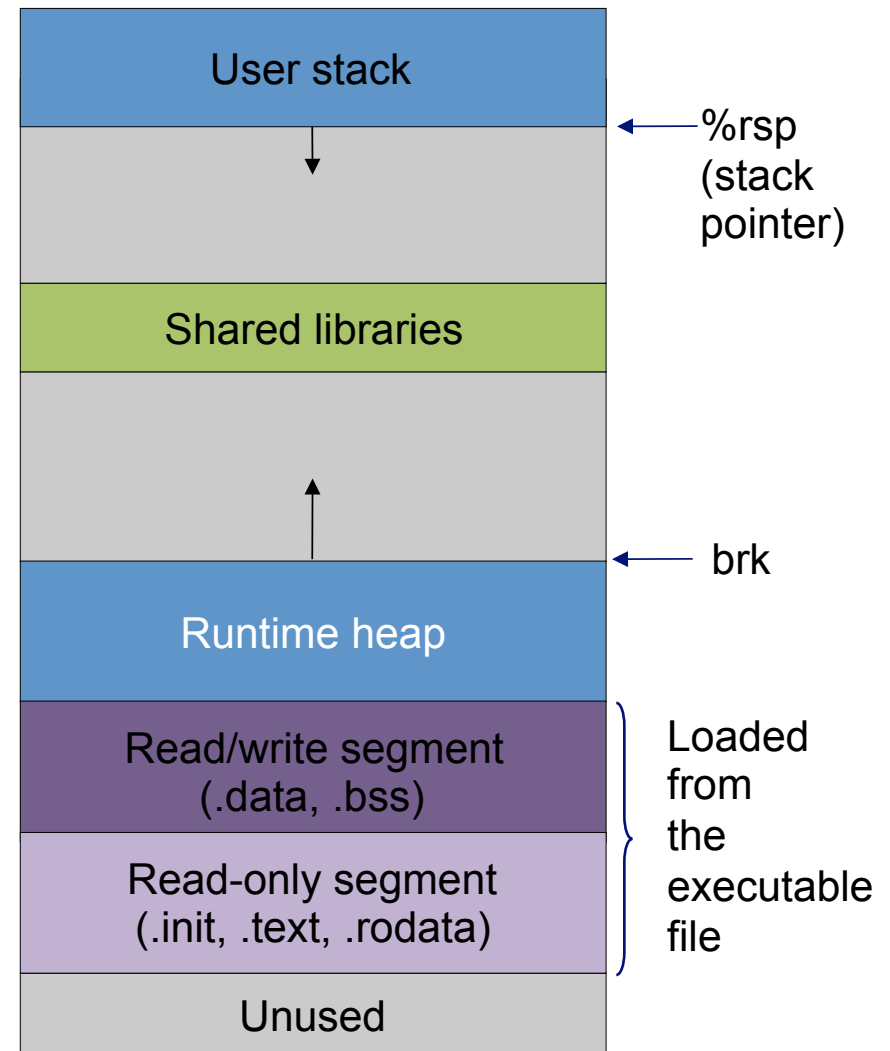
Question: How to allocate memory on heap?

Ask OS for allocation on the heap via **system calls**

```
void *sbrk(intptr_t size);
```

Issue I – can only free the memory on the top of heap

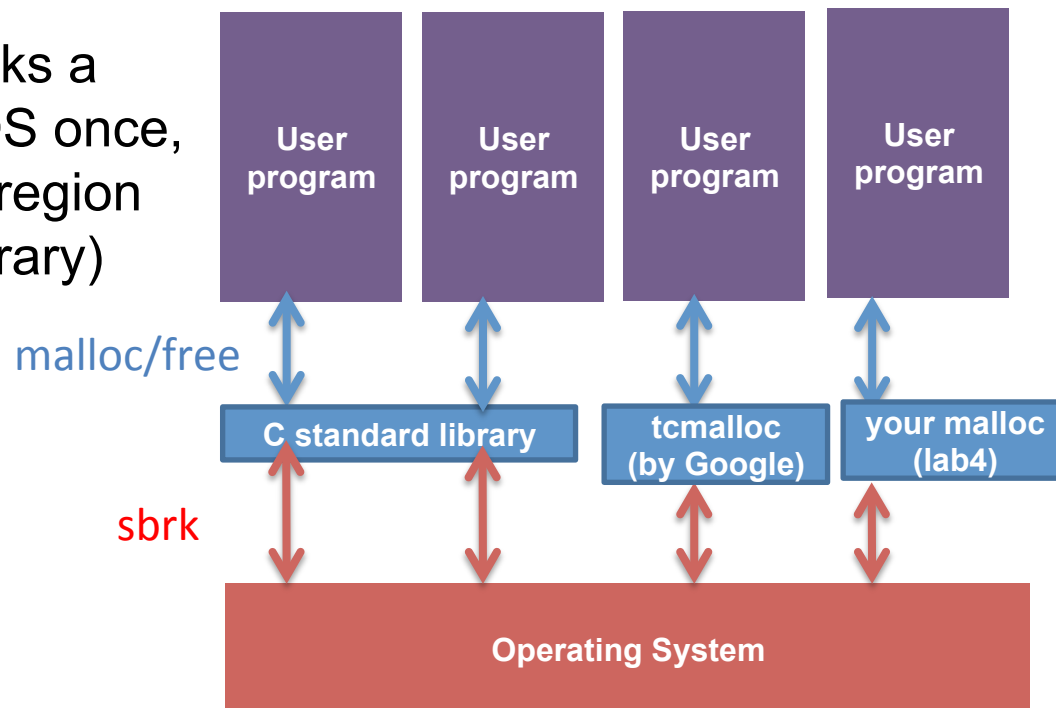
Issue II – system call has high performance cost > 10X



# Dynamic allocation on heap

Question: How to efficiently allocate memory on heap?

Basic idea: user program asks a large memory region from OS once, then manages this memory region by itself (using a “malloc” library)



# How to implement a memory allocator?

API:

- `void* malloc(size_t size);`
- `void free(void *ptr);`

Goal:

- Efficiently utilize acquired memory with high throughput
  - high throughput – how many mallocs / frees can be done per second
  - high utilization – fraction of allocated size / total heap size

# How to implement a memory allocator?

Assumed behavior of applications:

- Issue an arbitrary sequence of malloc/free
- Argument of free must be the return value of a previous malloc
- No double free

Restrictions on the allocator:

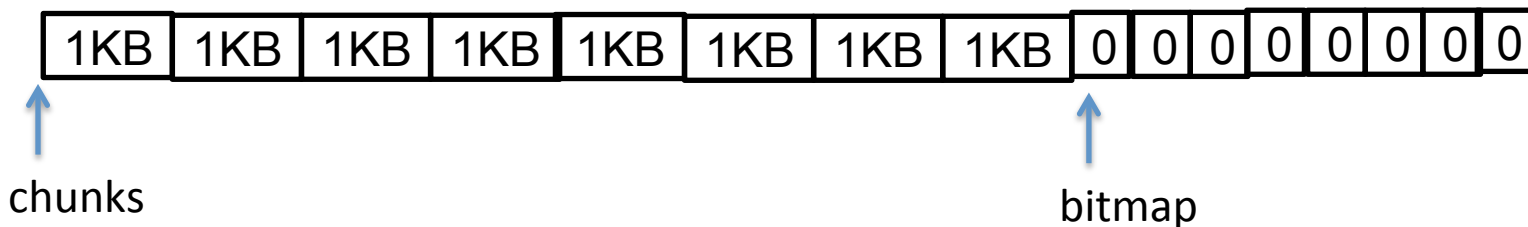
- Once allocated, space cannot be moved around

# Questions

1. (Basic book-keeping) How to keep track which bytes are free and which are not?
2. (Allocation decision) Which free chunk to allocate?
3. (API restriction) free is only given a pointer, how to find out the allocated chunk size?

# How to bookkeep? Strawman #1

- Structure heap as n 1KB chunks + n metadata

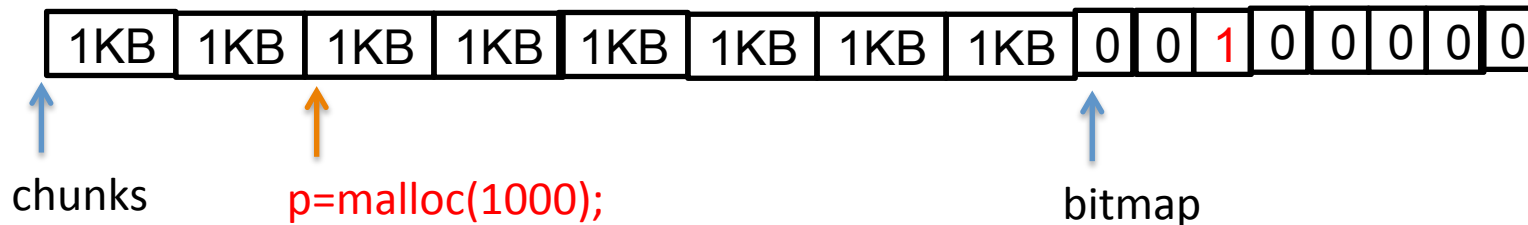


```
#define CHUNKSIZE 1<<10;
typedef char[CHUNKSIZE] chunk;
char *bitmap;
chunk *chunks;
size_t n_chunks;
```

```
void init() {
    n_chunks = 128;
    sbrk(n_chunks*sizeof(chunk)+ n_chunks/8);
    chunks = (chunk *)heap_lo();
    bitmap = heap_lo() + n_chunks *CHUNKSIZE;
}
```

Assume allocator asks for  
enough memory from OS  
in the beginning

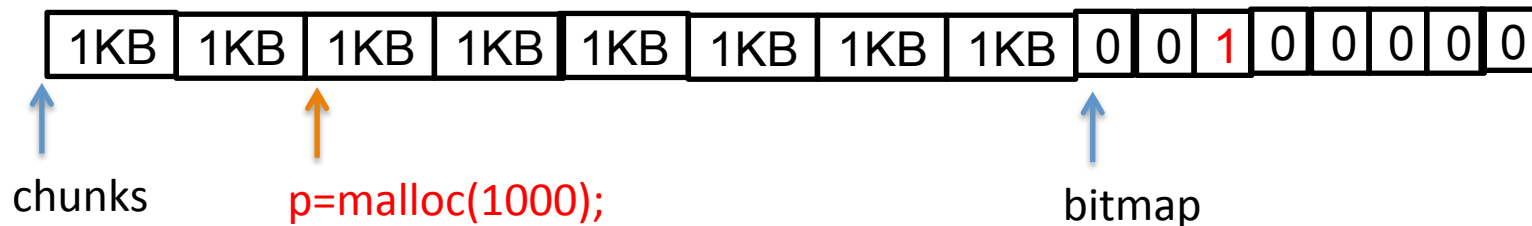
# How to bookkeep? Strawman #1



```
void* malloc(size_t sz) {  
    // find out # of chunks needed to fit sz bytes  
    CSZ = ...  
  
    //find csz consecutive free chunks according to bitmap  
    int i = find_consecutive_chunks(bitmap);  
  
    // return NULL if did not find csz free consecutive chunks  
    if (i < 0)  
        return NULL;  
  
    // set bitmap at positions i, i+1, ... i+csz-1  
    bitmap_set_pos(bitmap, i, csz);  
    return (void *)&chunks[i];  
}
```



# How to bookkeep? Strawman #1



```
void free(void *p) {  
    i = ((char *)p - (char *)chunks)/sizeof(chunk);  
    bitmap_clear_pos(bitmap, i); //how many bits to clear??  
}
```

- Problem with strawman?
  - free does not know how many chunks allocated
  - wasted space within a chunk (internal fragmentation)
  - wasted space for non-consecutive chunks (external fragmentation)

# How to bookkeep? Other Strawmans

- How to support a variable number of variable-sized chunks?
  - Idea #1: use a hash table to map address → [chunk size, status]
  - Idea #2: use a linked list in which each node stores [address, chunk size, status] information.

## Problems of strawmans?

Implementing a hash table and linked list requires use of a dynamic memory allocator!


# **How to implement a “linked list” without use of malloc**

# Implicit list

Embed chunk metadata in the chunks

- Chunk has a header storing size and status
- 16-byte aligned
  - Chunk size (metadata+payload) is multiple of 16
  - Header must be also aligned to 16 bytes

status: allocated or free




allocated: `size_and_status & 0x1L`  
size: `size_and_status & ~(0x1L)`

# Implicit list

Embed chunk metadata in the chunks

- Chunk has a header storing size and status
- Payload is 16-byte aligned

```
p = malloc(1024)
```




# Implicit list

Embed chunk metadata in the chunks

- Chunk has a header storing size and status
- Payload is 16-byte aligned

```
p = malloc(1)
```




# Implicit list

Embed chunk metadata in the chunks

- Chunk has a header storing size and status
- Payload is 16-byte aligned

```
p = malloc(1)
```



# How to traverse an implicit list

```
typedef struct {
    unsigned long size_and_status;
    unsigned long padding;
} header;

void traverse_implicit_list() {
    header *curr = (header *)heap_lo();
    while ((char *)curr < heap_high()) {
        bool allocated = get_status(curr);
        size_t csz = get_chunksz(curr);
        curr = (header *)((char *)curr + csz);
    }
}


bool get_status(header *h) {
    return h->size_and_status & 0x1L;
}

size_t get_size(header *h) {
    return h->size_and_status & ~(0x1L);
}
```




# Placing allocated blocks

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
```




# Where to place an allocation?

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
```



# First fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
```




header  
allocated payload  
padding  
free block

First fit – Search list from beginning,  
choose first free block that fits

# First fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
```




- header
- allocated payload
- padding
- free block

First fit – Search list from beginning, choose first free block that fits


Downside: cause fragmentation at beginning of the heap

# Best fit



# Best fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
```




- header
- allocated payload
- padding
- free block

Best fit – choose the free block with the closest size that fits

Downside: run slower than first fit.

# Next fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
p8 = malloc(56)
```




header  
allocated payload  
padding  
free block

Next fit – like first-fit, but search starts from where the previous search left off.

# Next fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
p8 = malloc(56)
```




- header
- allocated payload
- padding
- free block

Next fit – like first-fit, but search starts from where the previous search left off.



## Next fit

```
p1 = malloc(8)
p2 = malloc(24)
p3 = malloc(56)
p4 = malloc(8)
p5 = malloc(24)
p6 = malloc(56)
free(p2)
free(p4)
free(p6)
p7 = malloc(8)
p8 = malloc(56)
```




- header
- allocated payload
- padding
- free block

Next fit – like first-fit, but search starts from where the previous search left off.


Next fit runs faster than first fit,  
but fragmentation is worse.

# Splitting a free block

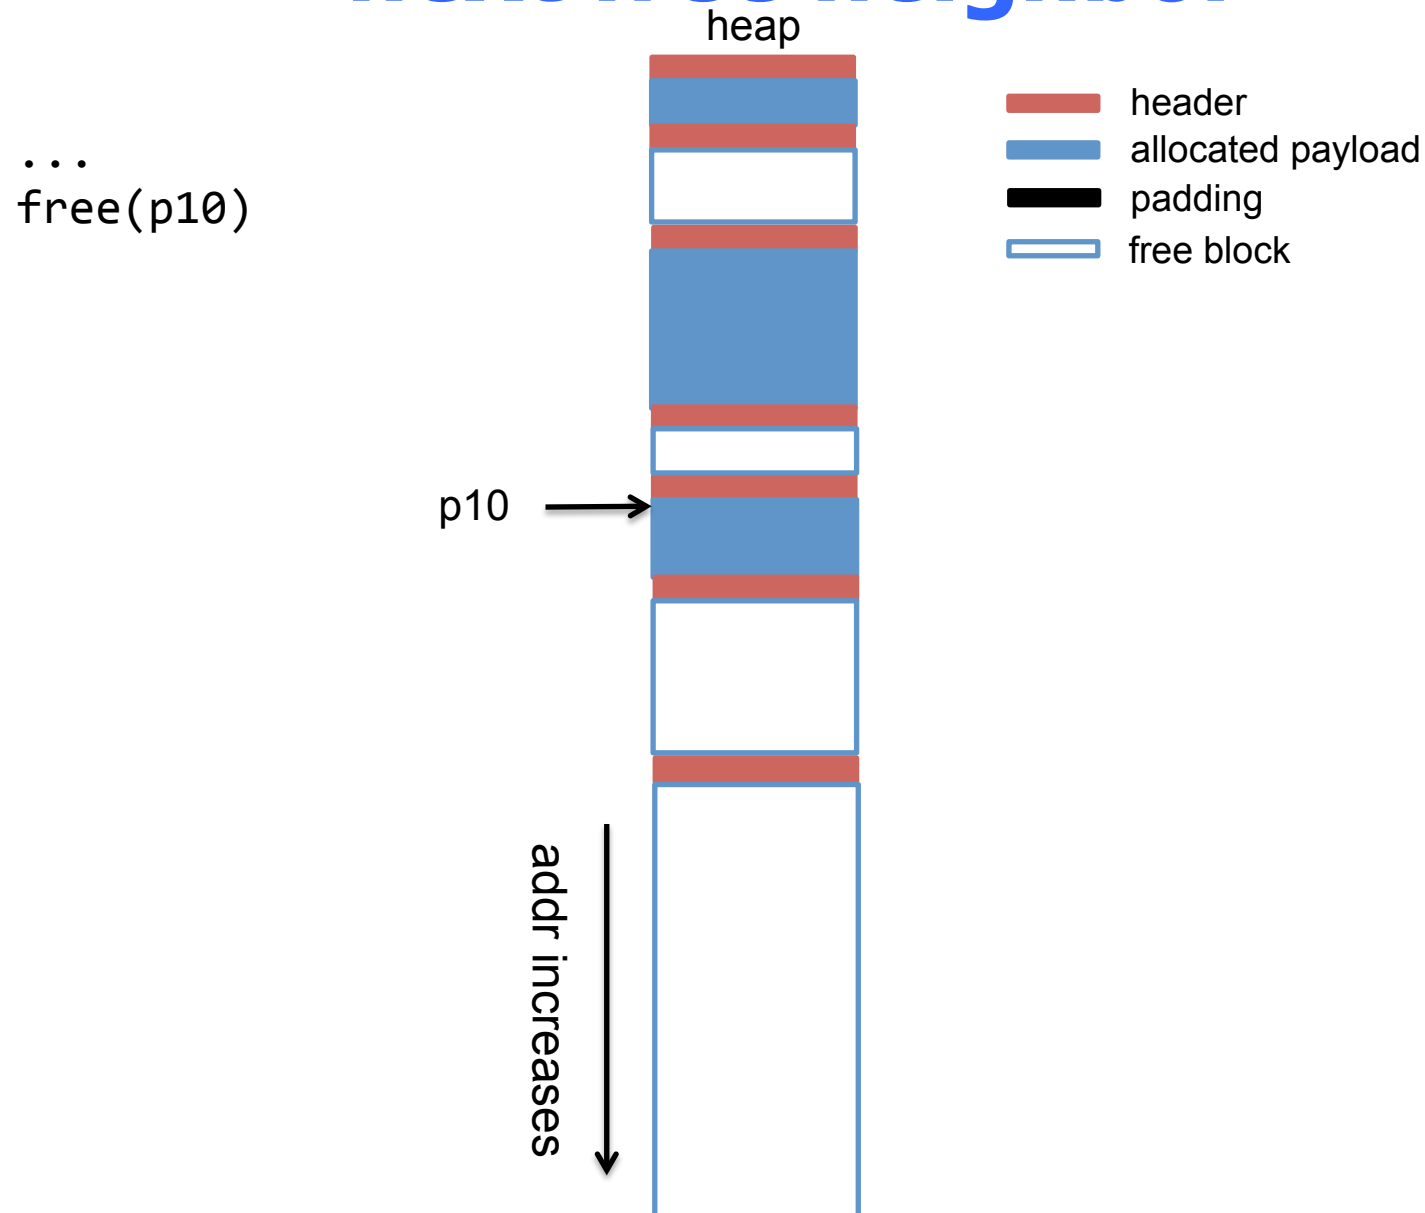


# Splitting a free block

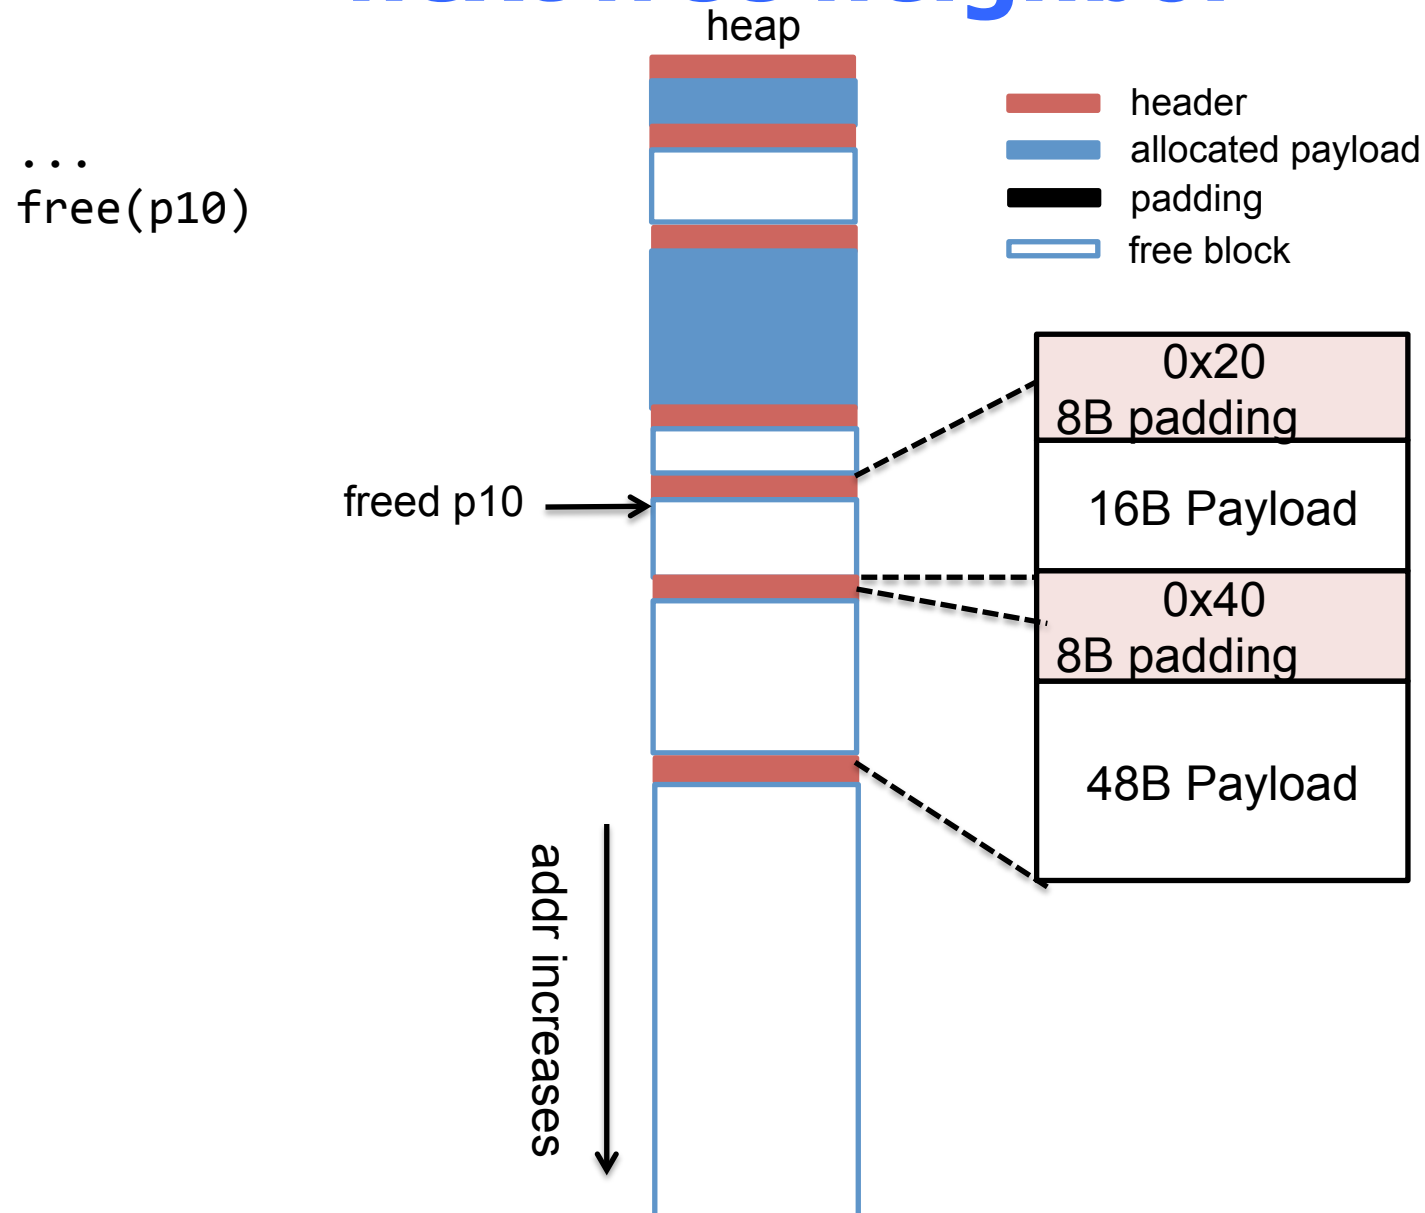
...  
p9 = malloc(16)



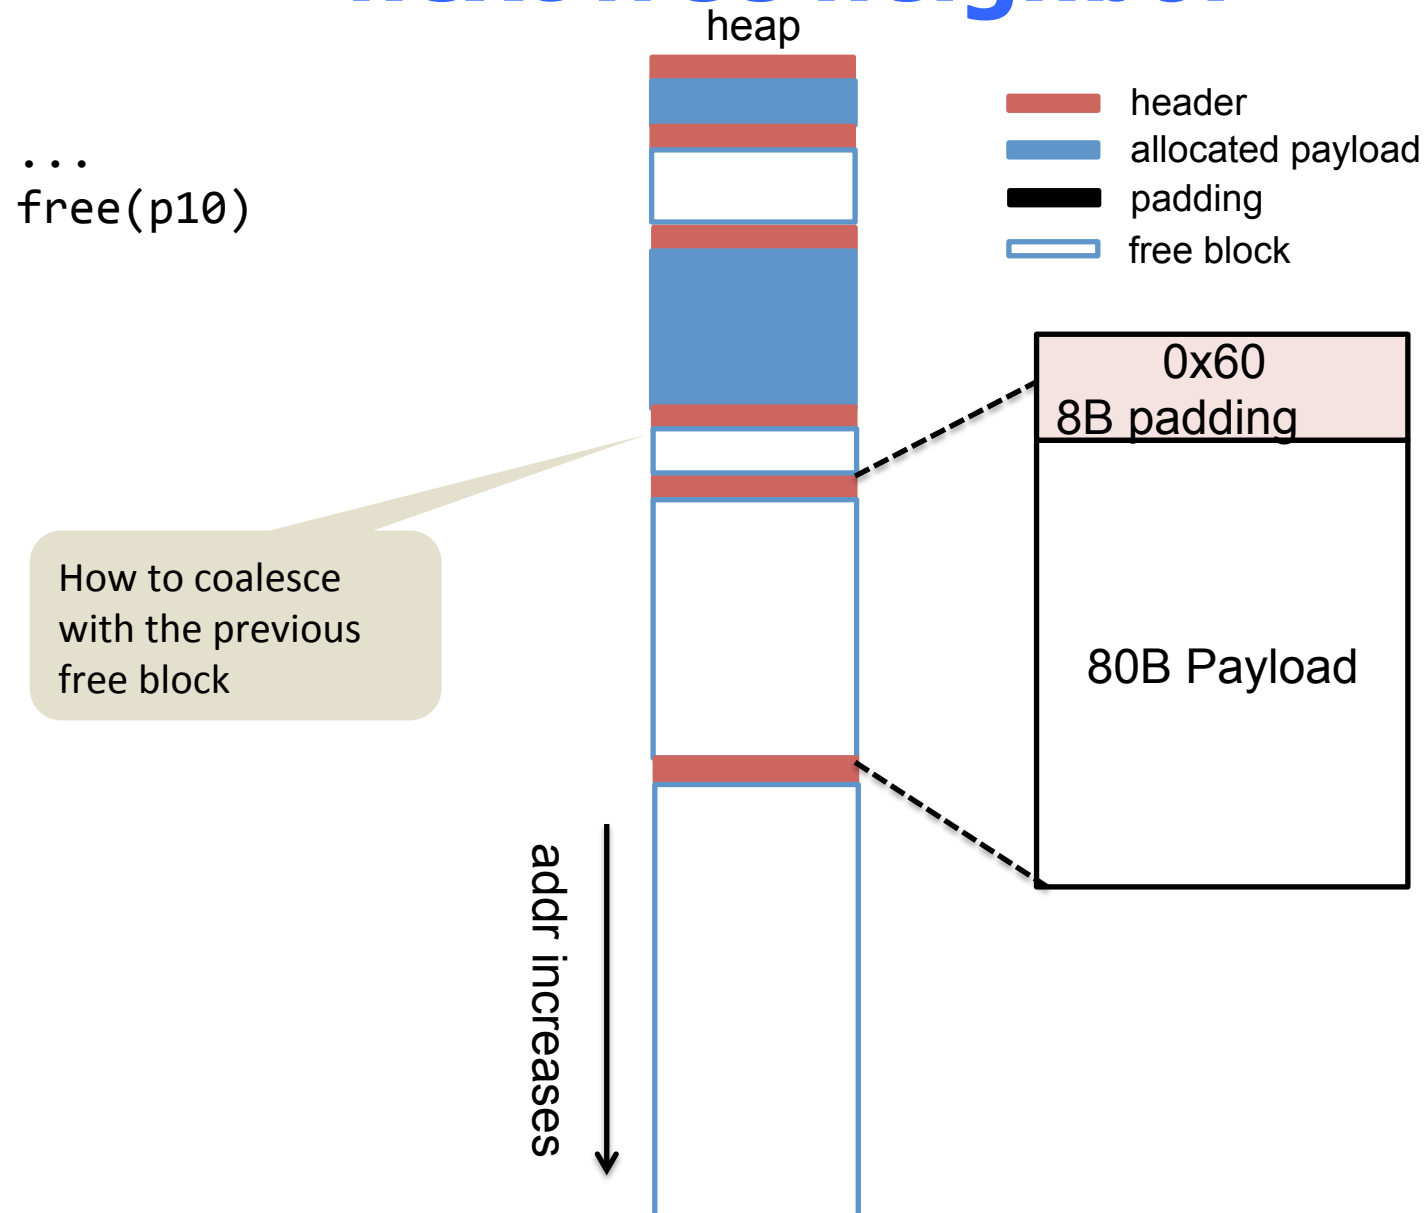
# Coalescing a free block with its next free neighbor



# Coalescing a free block with its next free neighbor

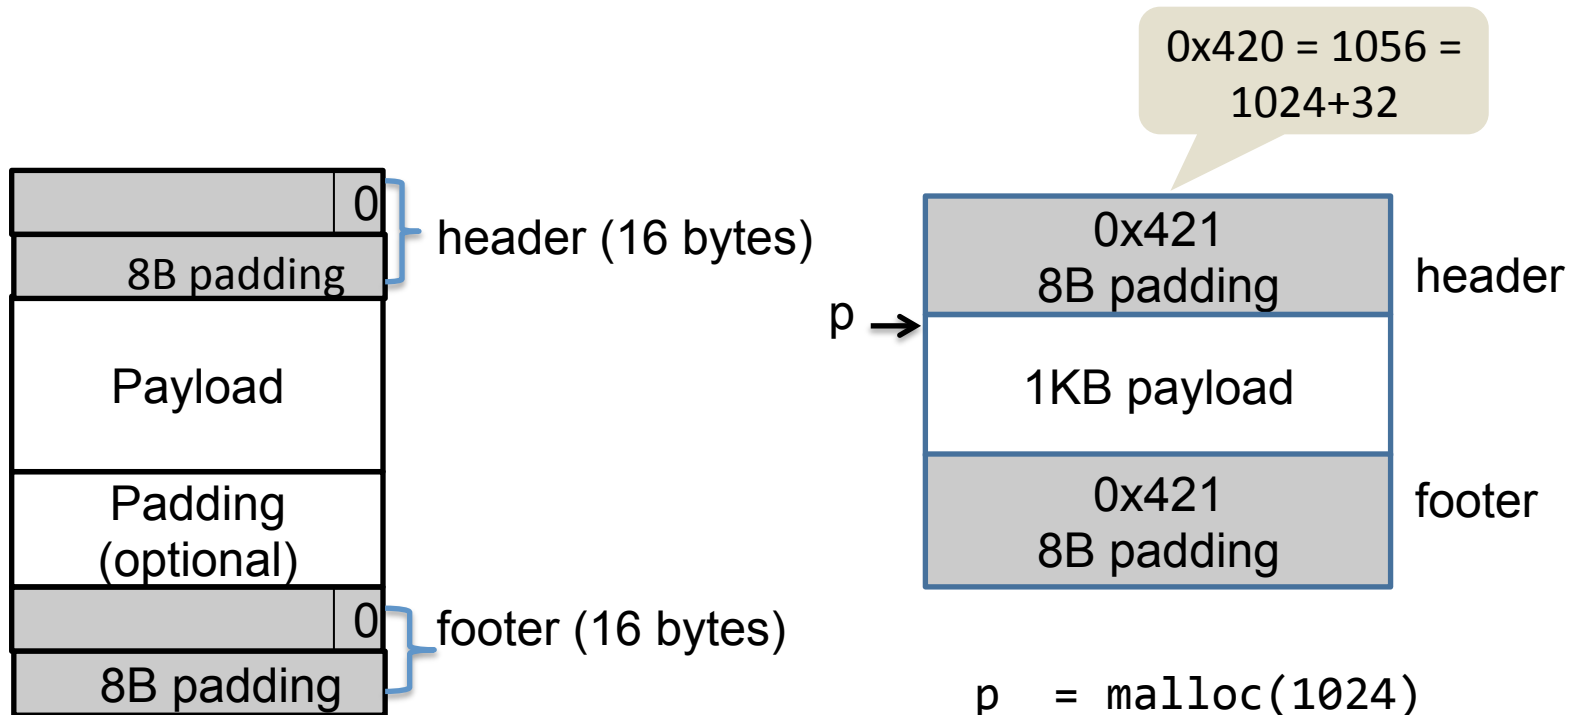


# Coalescing a free block with its next free neighbor

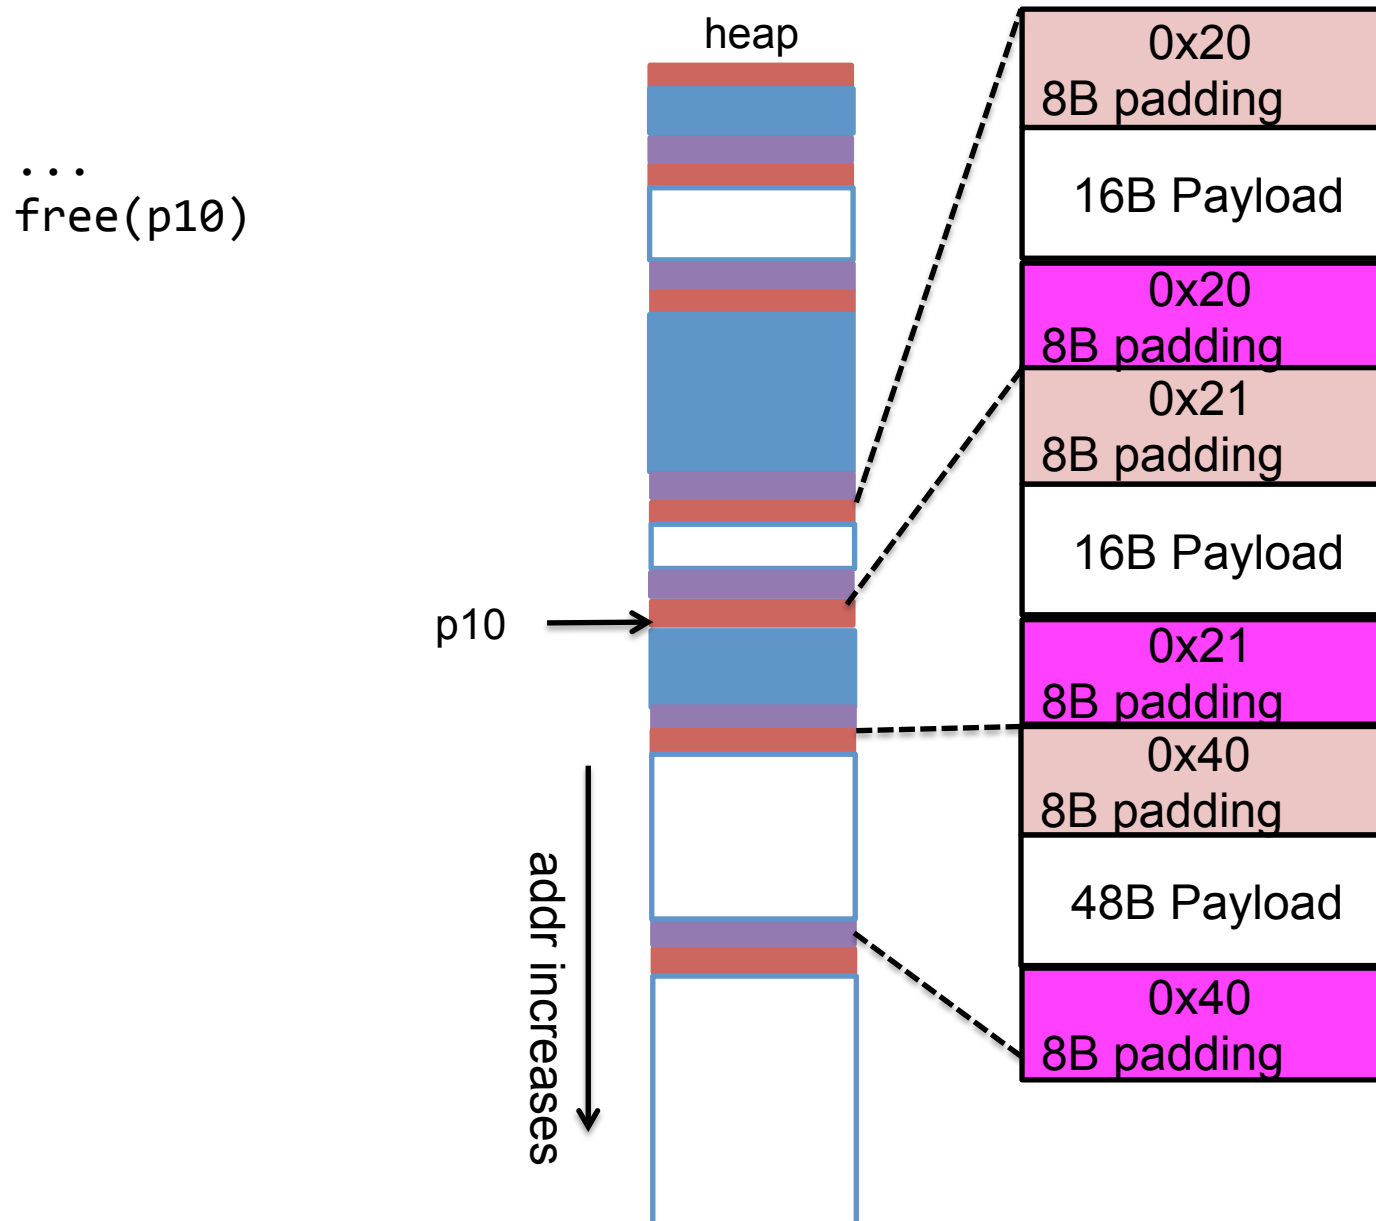


# Use footer to coalesce with previous block

- Duplicate header information into the footer




# Coalescing prev and next blocks






# Coalescing prev and next blocks

...  
free(p10)



# Recap: malloc using implicit list




- We can traverse the entire list of chunks on heap by incrementing pointer with chunk sizes,
- To allocate, find a block that fits, split if necessary

- Question: what's the minimal size of a chunk?


Answer:  $> 16 \text{ (header)} + 16 \text{ (footer)} + 16 \text{ (min payload)} = 48 \text{ bytes}$

# Coalescing prev and next blocks



# Coalescing prev and next blocks

...  
free(p10)



# Explicit free lists

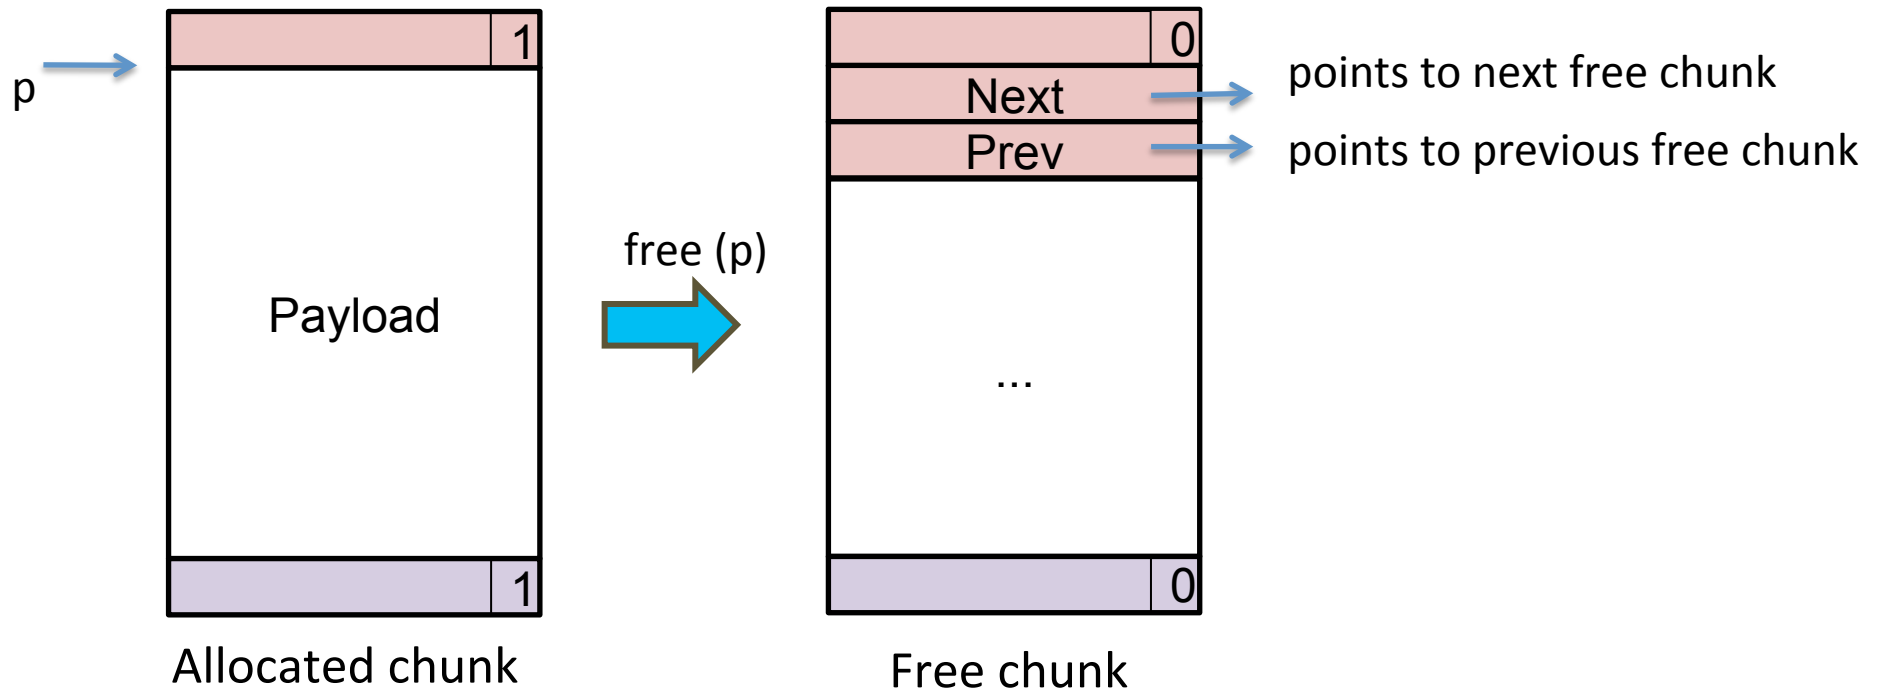
Problems of implicit list:

- Allocation time is linear in # of total (free and allocated) chunks

Explicit free list:

- Maintain a linked list of free chunks only.

# Explicit free list



- Question: do we need next/prev fields for allocated blocks?

Answer: No. We do not need to chain together allocated blocks. We can still traverse all blocks (free and allocated) as in the case of implicit list.

- Question: what's the minimal size of a chunk?

Answer: 16 (header) + 16 (footer) + 8 (next pointer) + 8 (previous pointer) = 48 bytes

# Explicit list: types, basic helpers

```
typedef struct {
    unsigned long size_and_status;
    unsigned long padding;
} header;

typedef struct free_hdr {
    header common_header;
    struct free_hdr *next;
    struct free_hdr *prev;
} free_hdr;

bool
get_status(header *h) {
    return h->size_and_status & 0x1L;
}

size_t
get_size(header *h) {
    return h->size_and_status & ~(0x1L);
}
```

```
void
set_size_status(header *h,
    size_t sz, bool status) {

    h->size_and_status = sz | status;
}

void
set_status(header *h, bool status){
    size_t sz = get_size(h);
    set_size_status(h, sz, status);
}

void
set_size(header *h, size_t sz) {
    status = get_status(h);
    set_size_status(h, sz, status);
}
```

# Explicit list: globals, initialization

```
free_hdr *freelist;
```

```
header*
```

```
get_footer_from_header(header *h) {  
    return (header *)((char *)h + get_size(h) - sizeof(header));  
}
```

```
init_free_chunk(free_hdr *h, size_t sz) {  
    set_size_status(&h->common_header, sz, false);  
    h->prev = h->next = NULL;  
    set_size_status(get_footer_from_header(h->common_header), sz, false);  
}
```

```
free_hdr *
```

```
get_block_from_OS(size_t sz) {  
    free_hdr *h = sbrk(sz);  
    init_free_chunk(h, sz); //init header and footer  
    return h;  
}
```

```
#define MIN_OS_ALLOC_SZ 1024
```

```
void init() {  
    freelist = get_block_from_OS(MIN_OS_ALLOC_SZ);  
}
```



# Explicit list: helpers to insert and detach from freelist

```
void insert(free_hdr **head, free_hdr *node) {
    if (*head)
        (*head)->prev = node;
    node->next = *head;
    *head = node;  //node becomes the new head
}

void delete(free_hdr **head, free_hdr *node) {
    if (node->prev) { //node is not the first node in the list
        node->prev->next = node->next;
        if (node->next)
            node->next->prev = node->prev;
    } else { //delete the first node in the list
        *head = node->next;
        if (node->next)
            node->next->prev = NULL;
    }
}
```

# Explicit list: allocate

```
void * malloc(size_t s) {  
    size_t csz = s + 2*sizeof(header); //min chunk size required  
    free_hdr *n = first_fit(csz);  
    if (!n)  
        n = get_block_from_OS(csz>MIN_OS_ALLOC_SIZE?csz:MIN_OS_ALLOC_SIZE);  
  
    free_hdr *newchunk = split(n, csz);  
    insert(&freelist, newchunk);  
    set_status(n, true);  
    return (char *)n+sizeof(header);  
}  
  
free_hdr * first_fit(size_t sz) {  
    free_hdr *n = freelist;  
    while (n) {  
        if (get_size(n->common_header)>= sz) {  
            delete(&freelist, n);  
            break;  
        }  
        n = n->next;  
    }  
    return n;  
}
```

# Explicit list: free

```
void free(void *p) {
    header *h = get_header_from_payload(p);
    init_free_chunk((free_hdr *)h, get_size(h));

    header *next = get_next_header(h);
    if (!get_status(next))
        h = coalesce((free_hdr *)h, (free_hdr *)next);
    header *prev = get_prev_header(h);
    if (!get_status(prev))
        h = coalesce((free_hdr *)h, (free_hdr *)prev);

    insert(&freelist, (free_hdr *)h);
}

free_hdr *
coalesce(free_hdr *me, free_hdr *other) {
    delete(&freelist, other);
    int sum = get_size(me->common_header)+get_size(other->common_header));
    free_hdr *h = me<other? me:other;
    set_size_status(h->common_header, sum, false);
    set_size_status(get_footer_from_header((header *)h, sum, false);
    h->next = h->prev = NULL;
    return h;
}
```

# Segregated list


- Idea: keep multiple freelists
  - each freelist contains chunks of similar sizes

# Segregated list: initialize

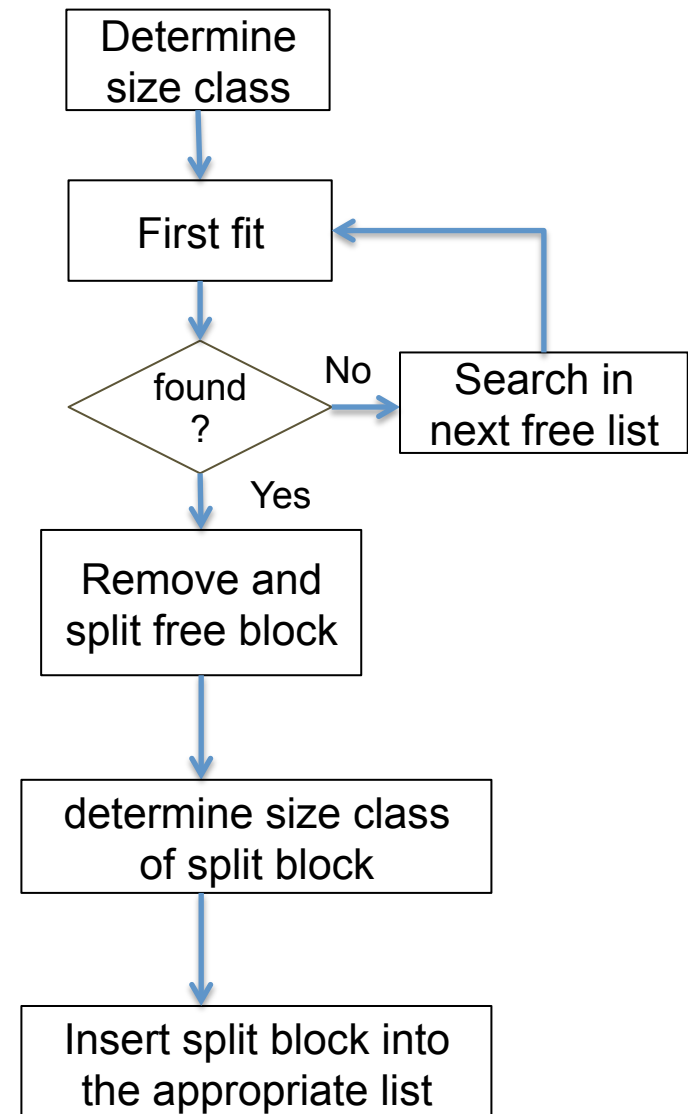

```
#define NLISTS 3
free_hdr* freelists[NLISTS];
size_t size_classes[NLISTS] = {32, 128, -1};
```

```
int which_freelist(size_t s) {
    int ind = 0;
    while (s > size_classes[ind])
        ind++;
    return ind;
}
```

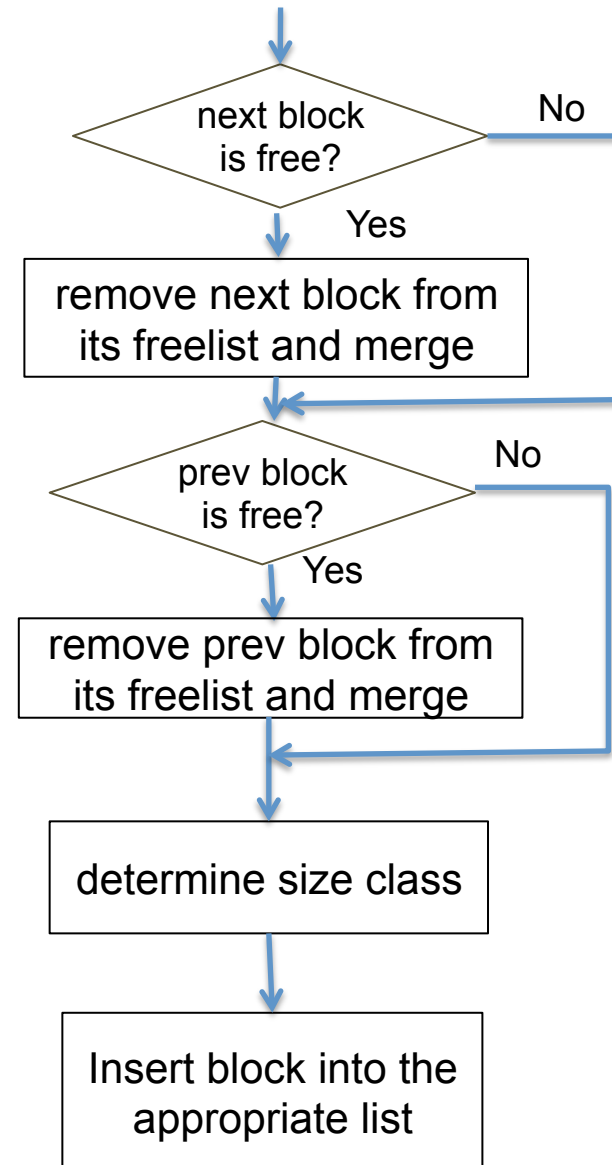

```
void init() {
    free_hdr *h = get_block_from_OS(1024);
    freelist[which_freelist(1024)] = h;
}
```



# Segregated list: allocation



# Segregated list: free



# Buddy System

- A special case of segregated list
  - each freelist has *identically-sized* blocks
  - block sizes are powers of 2
- Advantage over a normal segregated list?
  - Less search time (no need to search within a freelist)
  - Less coalescing time
- Adopted by Linux kernel and jemalloc




# Simple binary buddy system

Initialize:

- for simplicity, assume the initial  $2^m$  block is aligned at  $2^m$  (i.e. the least significant  $m$ -bits of its addr are zero)

(0000 0000 0000 0000)<sub>2</sub>



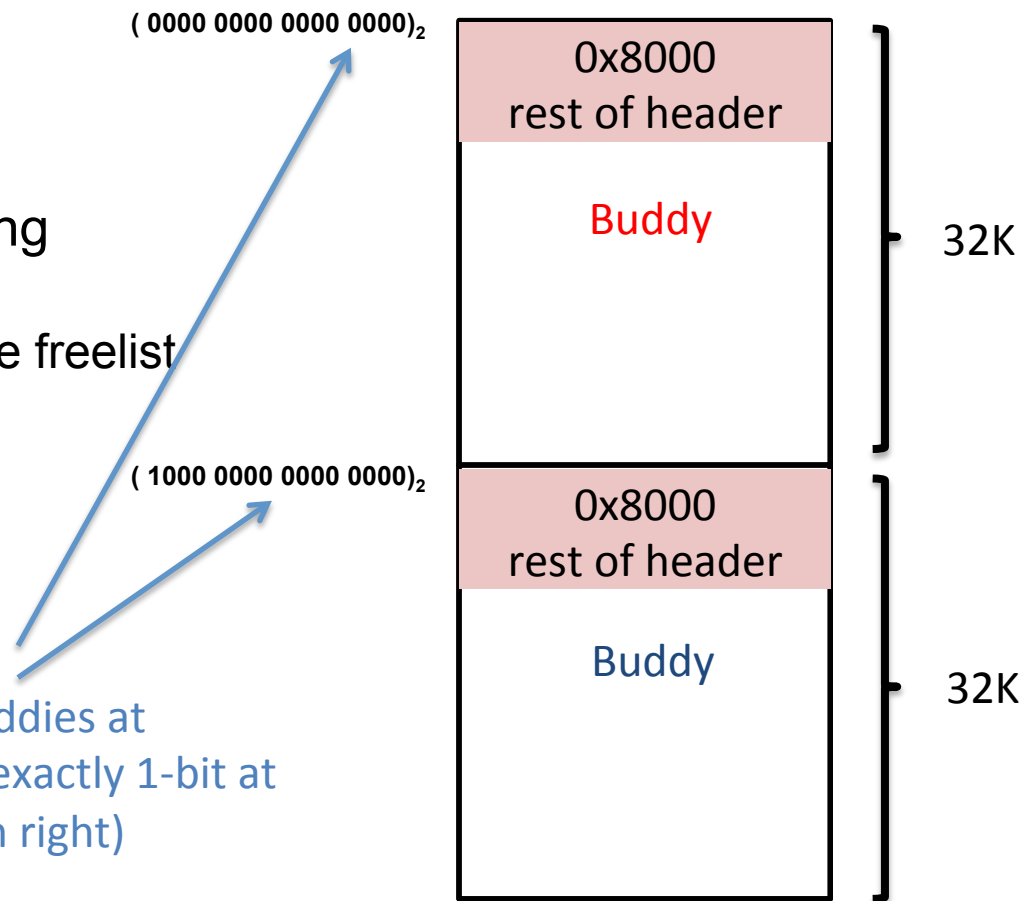
# Binary buddy system: allocate

```
p = malloc(16000);
```

Recursive split in half until having the right size

- insert free buddy into appropriate freelist

Addresses of buddies at size  $2^m$  differ in exactly 1-bit at position  $m$  (from right)

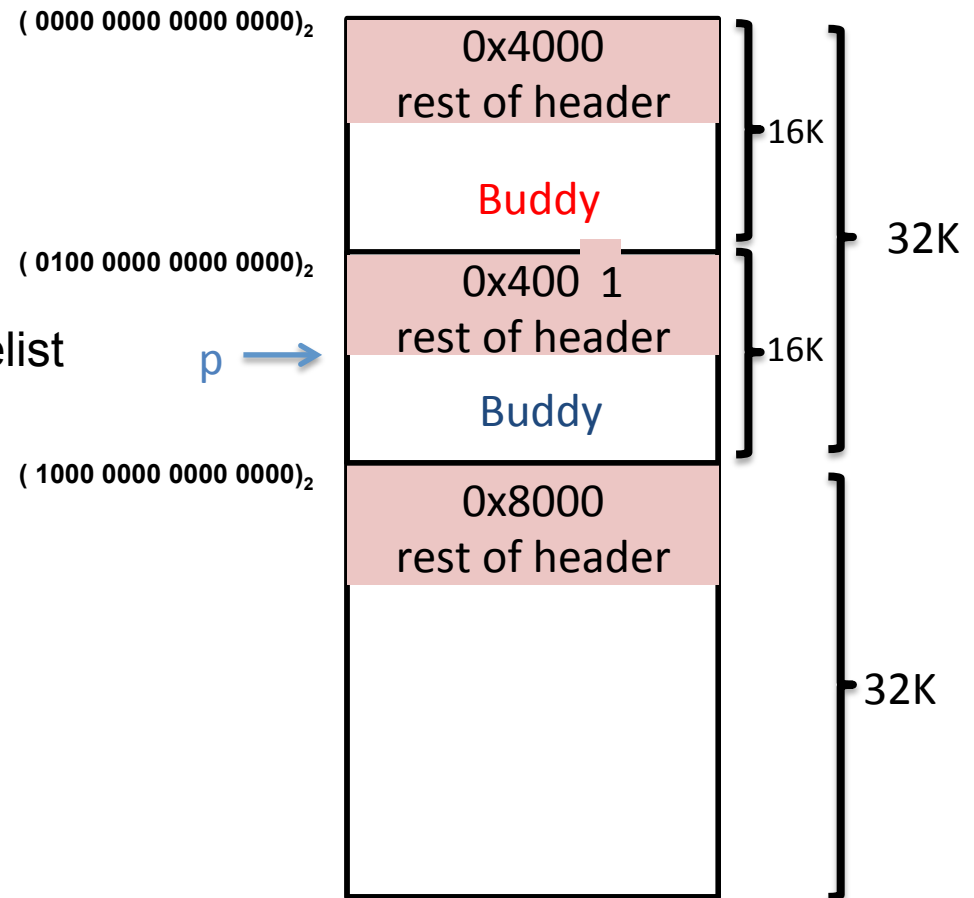


# Binary buddy system: allocate

```
p = malloc(16000);
```

Recursive split in half until having the right size

- insert free buddy into appropriate freelist



# Binary buddy system: free

`free(p);`

Recursively merge block with buddy

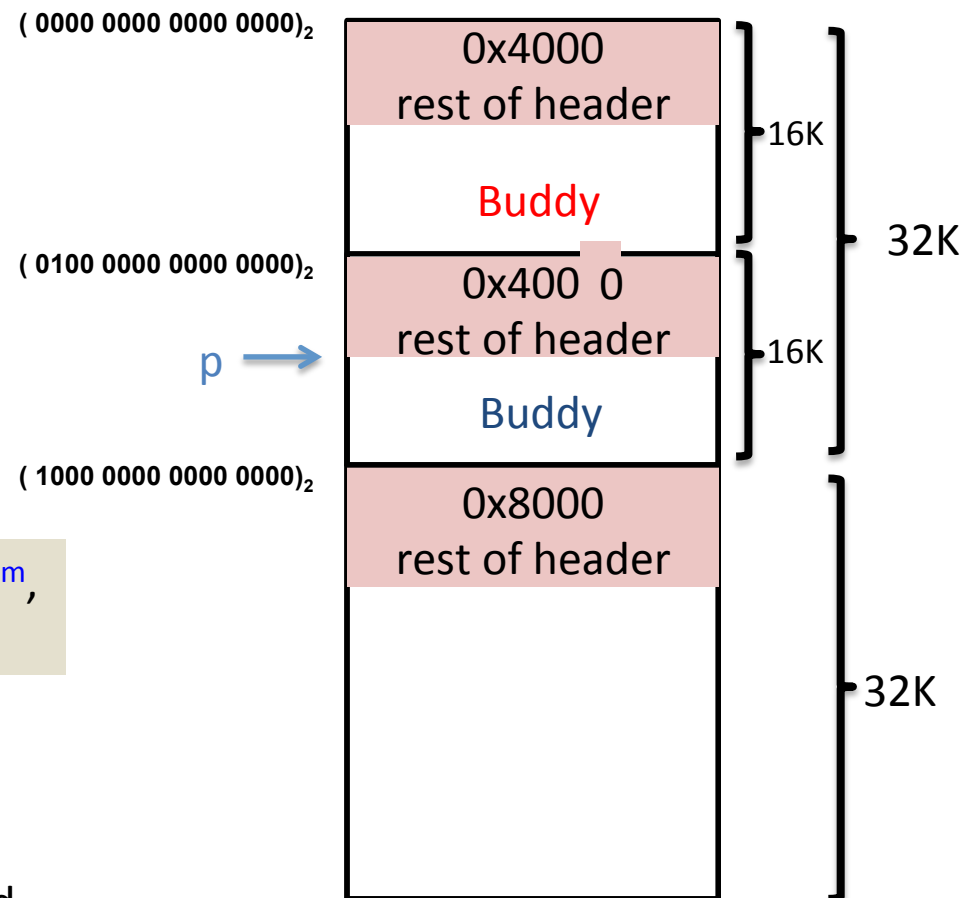
1. Calculate addr of buddy block,  
determine buddy status

Question: given addr  $a$  of block with size  $2^m$ ,  
how to calculate its buddy's address?

$$a \wedge (1 \ll m)$$



any bit XOR 0 = unchanged  
any bit XOR 1 = flipped

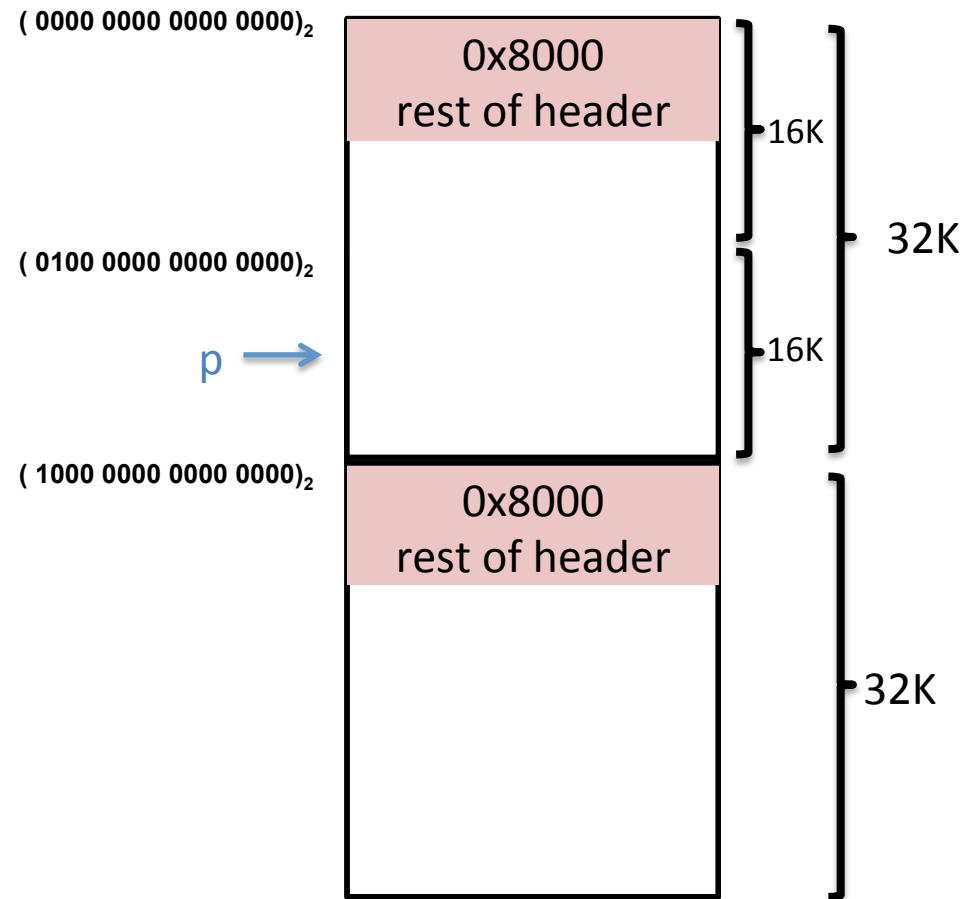


# Binary buddy system: free

`free(p);`

If buddy is free:

2. Detach free buddy from its list
3. Combine with current block



# Binary buddy system: free

`free(p);`

Repeat to merge with larger buddy  
Insert final block into appropriate  
freelist

