

# **User program and OS interaction**

## **Multiprocessing**

Jinyang Li

# What we've learnt so far

- Machine instructions
  - compiler translates C to x86 instructions
  - x86 instructions are executed by CPU hardware only
- Dynamic memory allocator
  - realized as a library implementation
- Virtual memory
  - each process has its own virtual address space
  - VM is realized by a combination of hardware mechanism and OS implementation
    - MMU performs address translation
    - OS populates page table

# Today's lesson plan

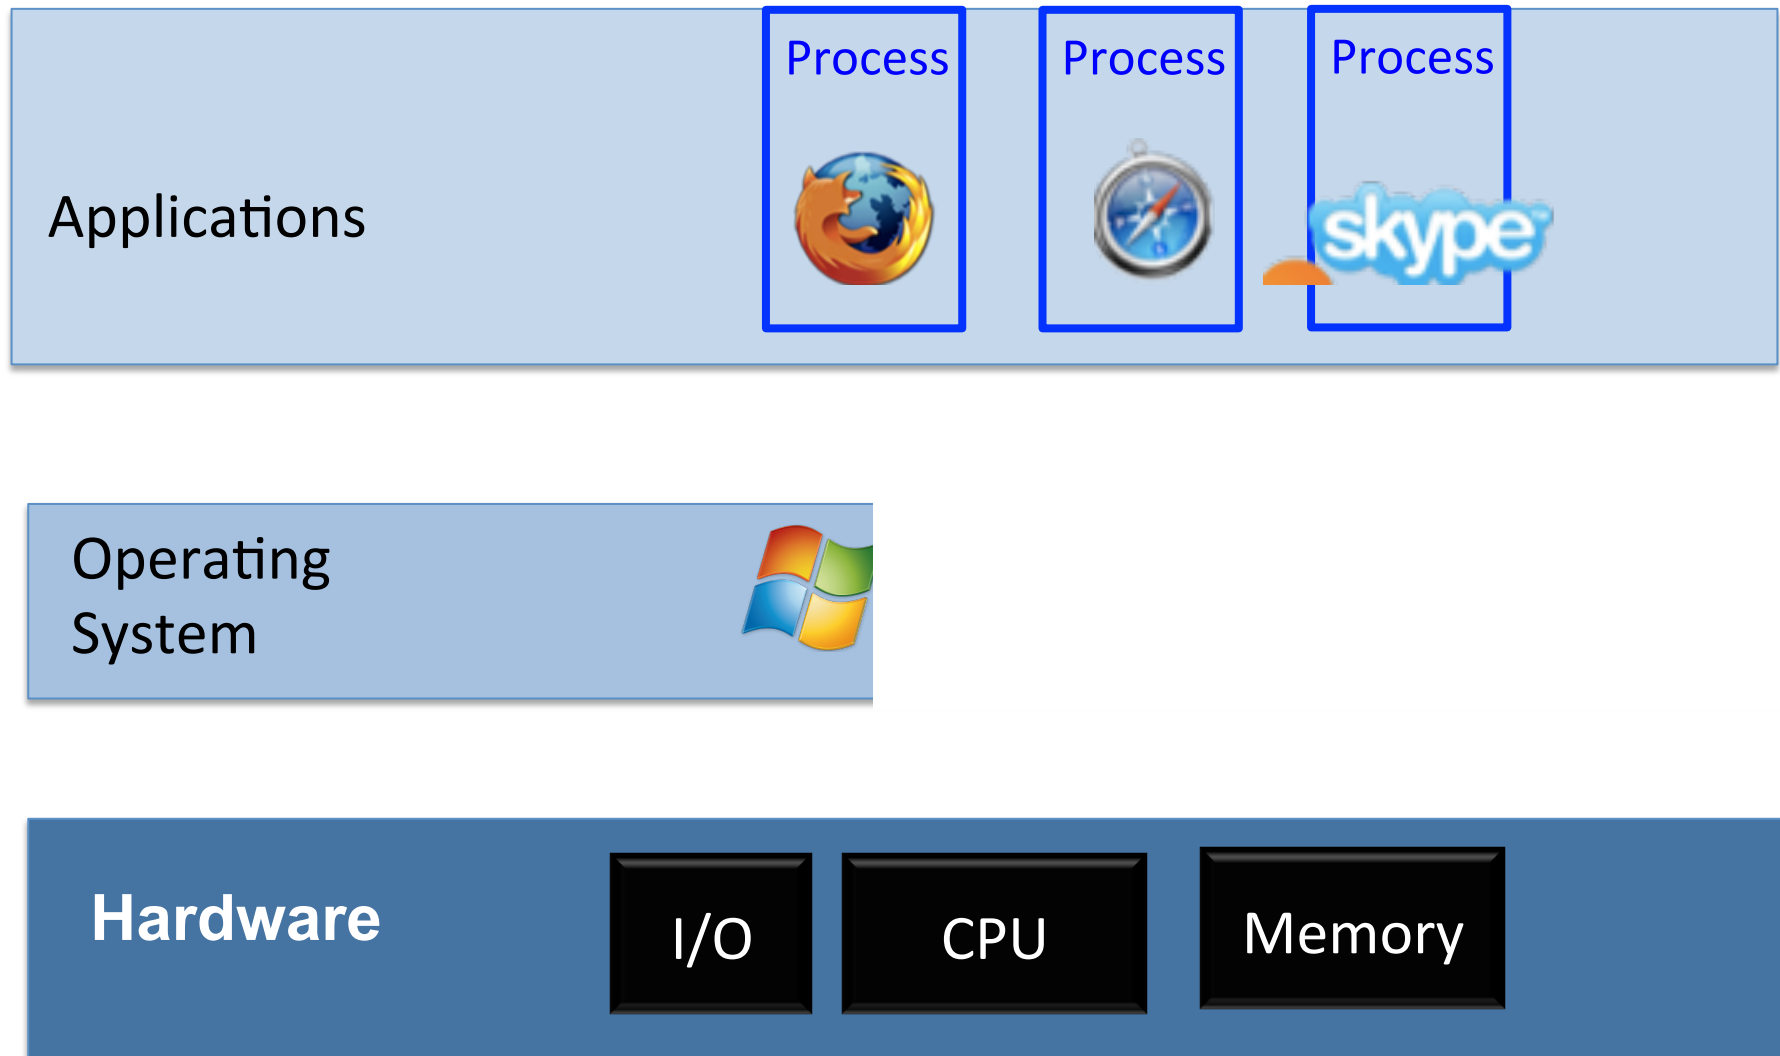
1. Interaction between user programs and OS
2. Multiprocessing

# Interaction between user programs and OS

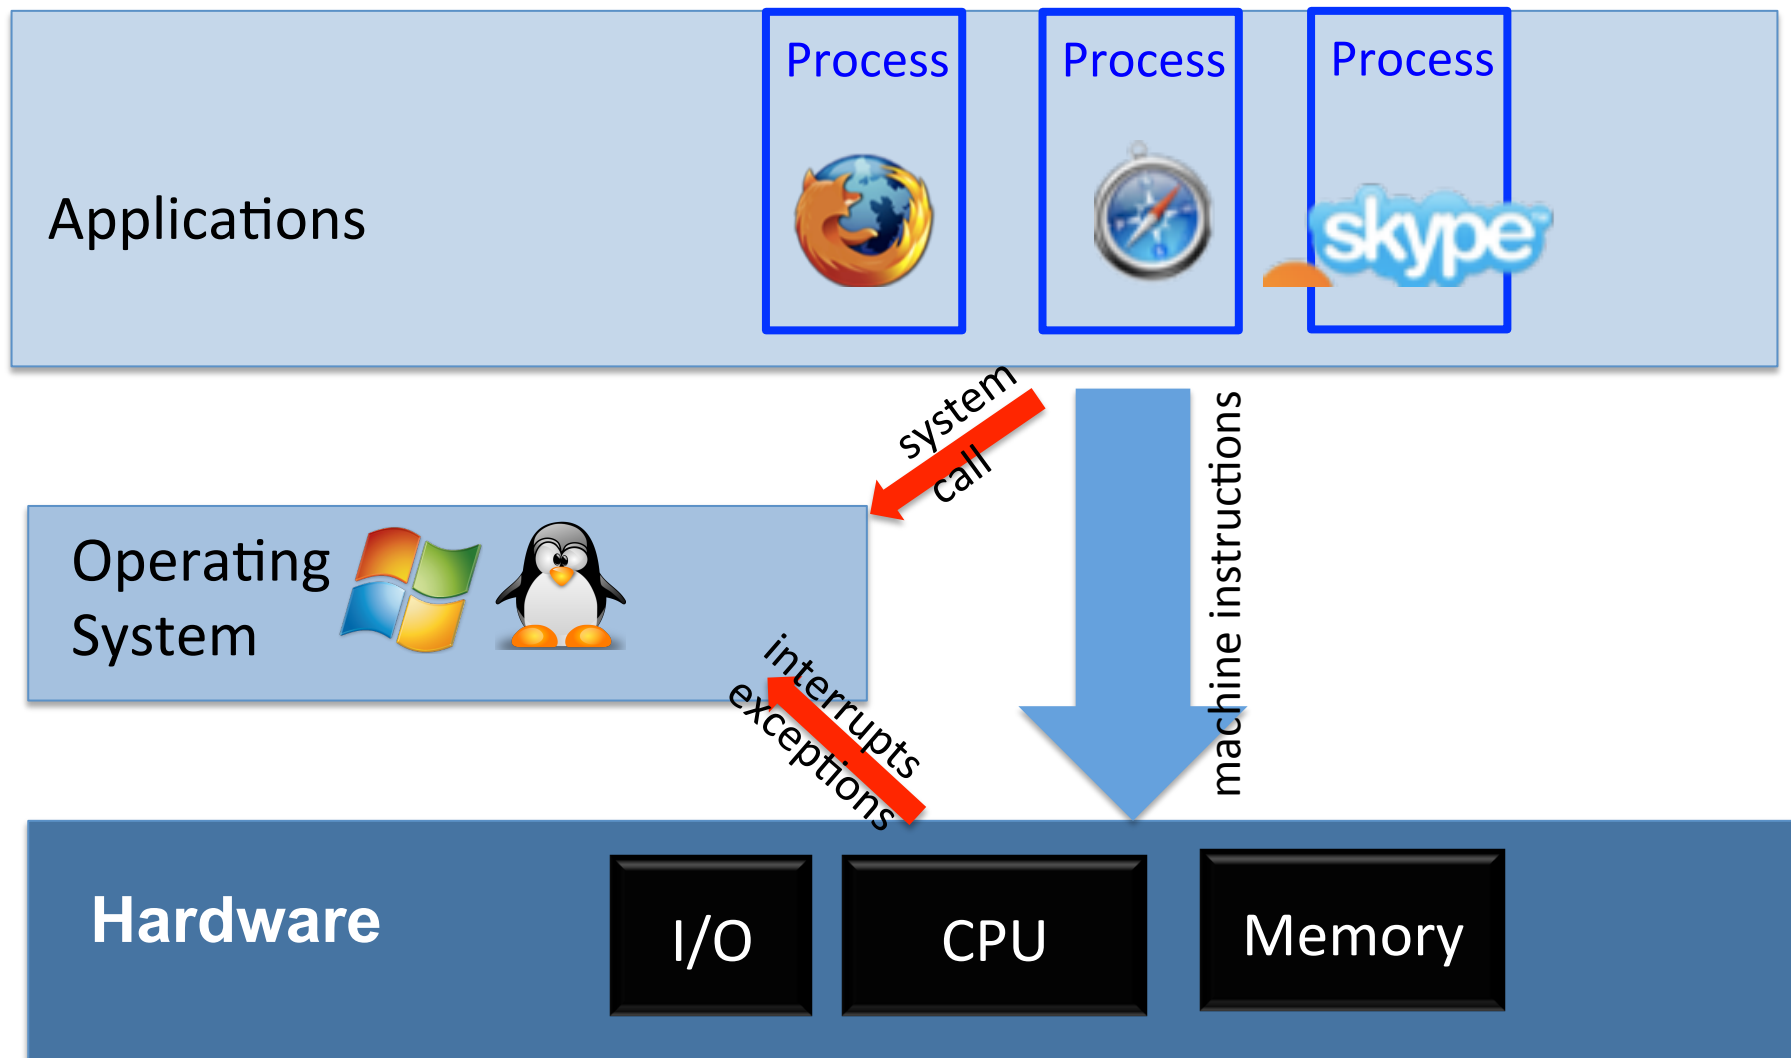


I mean OS kernel

# Applications, OS, Hardware



# Applications, OS, Hardware



# The role of OS



## What does the OS do?

### 1. Resource management


- **scheduling**: give each process illusion of exclusive CPU use
- **VM management**: give each process illusion of exclusive memory use

### 2. Hide messy hardware details

- file system
- networking

# Process

- Process is an instance of a running program
  - when you type `./a.out`, OS launches a process
  - when you type Ctrl-C, OS kills the process
- OS maintains some state for each process
  - process identifier (process id)
    - user id
    - status (e.g. runnable or blocked)
    - saved rip and other registers
    - VM structure (including its page table)



Only OS can  
modify these  
data



# How to protect the OS from user processes?

- Hardware provides privileged vs. non-privileged mode of execution
  - also called supervisor or kernel mode
  - also called user mode
- OS runs in privileged mode
  - can change content of CR3 (points to root page table)
  - can access VA marked as supervisor only
  - ...
- User programs run in non-privileged mode
  - cannot access kernel data structures because they are stored in VA marked as supervisor only

# How to get into privileged mode?

Hardware provides 3 controlled mechanisms to switch from non-privileged to privileged execution:

## 1. Traps

- syscalls (user programs explicitly ask for OS help)

## 2. Exception (caused by the current running program)

- e.g. divide by zero, page fault

## 3. Interrupt (caused by external events)

- timer, keyboard press, packet arrival

# How to get out of privileged mode?

- OS uses the special hardware instruction `iret`
- OS may return to the same program or context switch to execute a different program

# #1 Traps:

## Syscall: User → OS

- User programs ask for OS services using syscalls
  - it's like invoking a function in OS
- Each syscall has a known number

0	read
1	write
2	open
3	close
...	
57	fork
59	execve
60	exit
62	kill

C library wraps these syscalls to provide file I/O

linux syscall number

# Syscall: user → OS

user code

```
movq %rax, %rbx
inc %rbx
...
syscall 2
movq %rax, %r8
add %r8, %r9
...
```

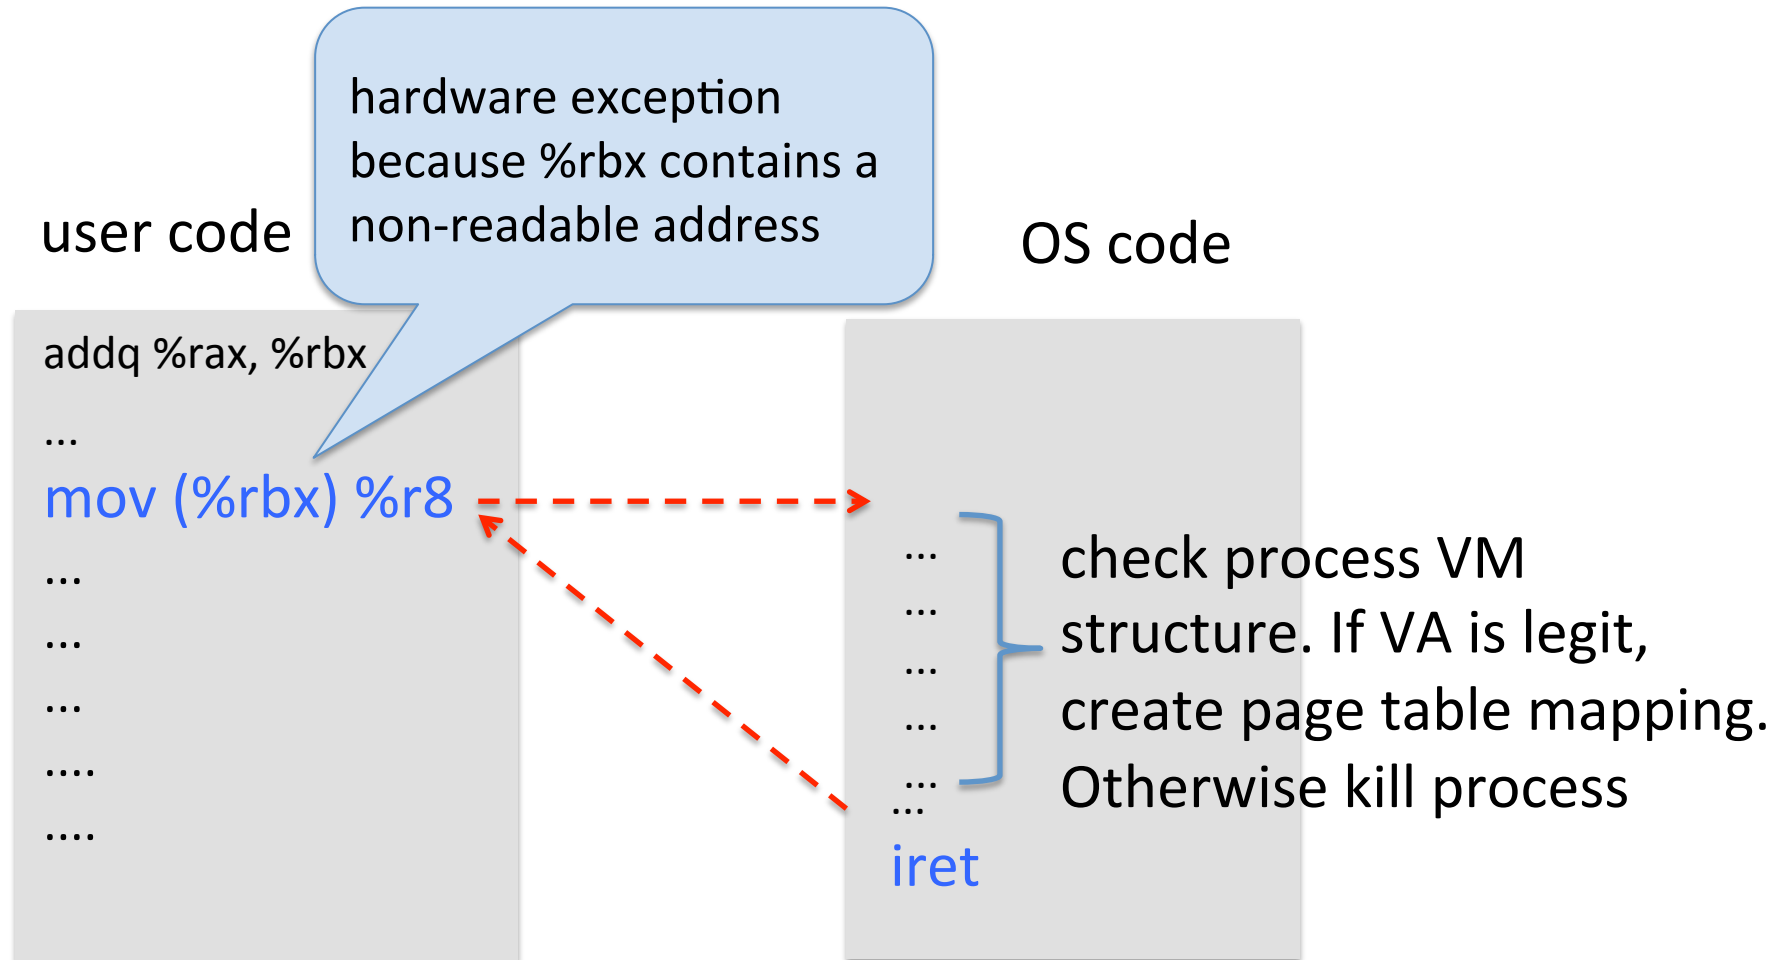
OS code

```
...
...
...
...
...
...
...
...
...
iret
```

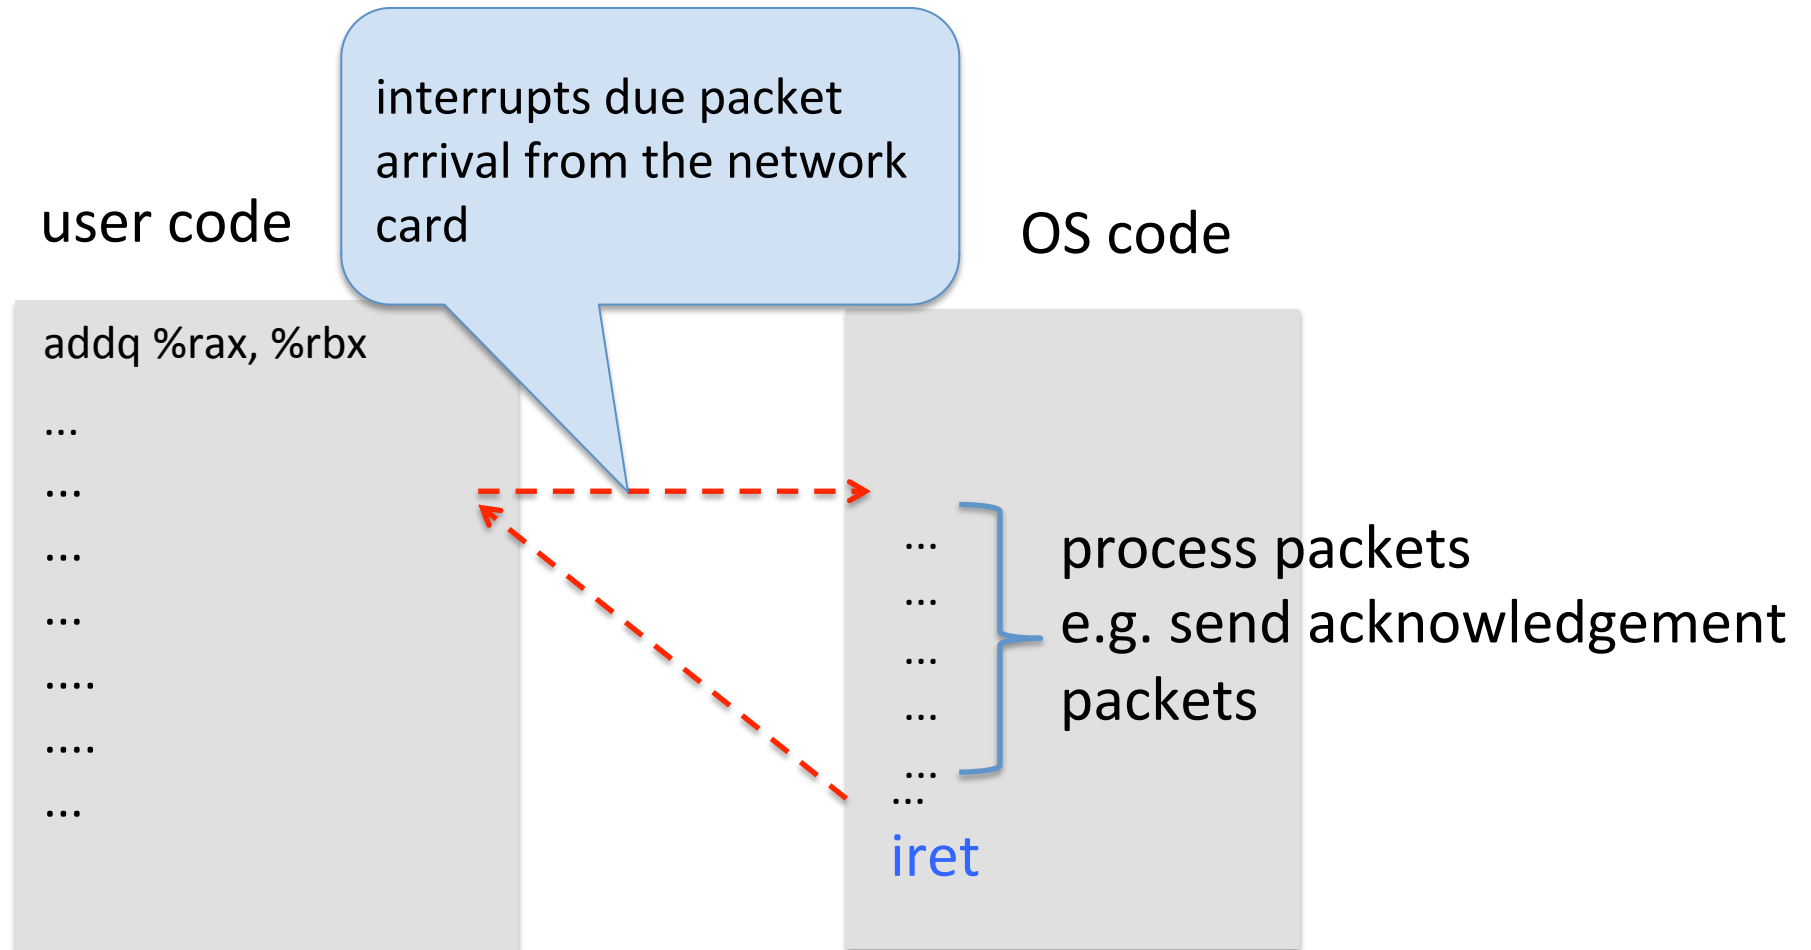
code to open  
the requested file

Assuming OS wants to execute  
the same process next; it does  
not have to

## #2 exceptions: OS takes control upon exceptions



# #3 interrupts: OS takes control upon interrupts



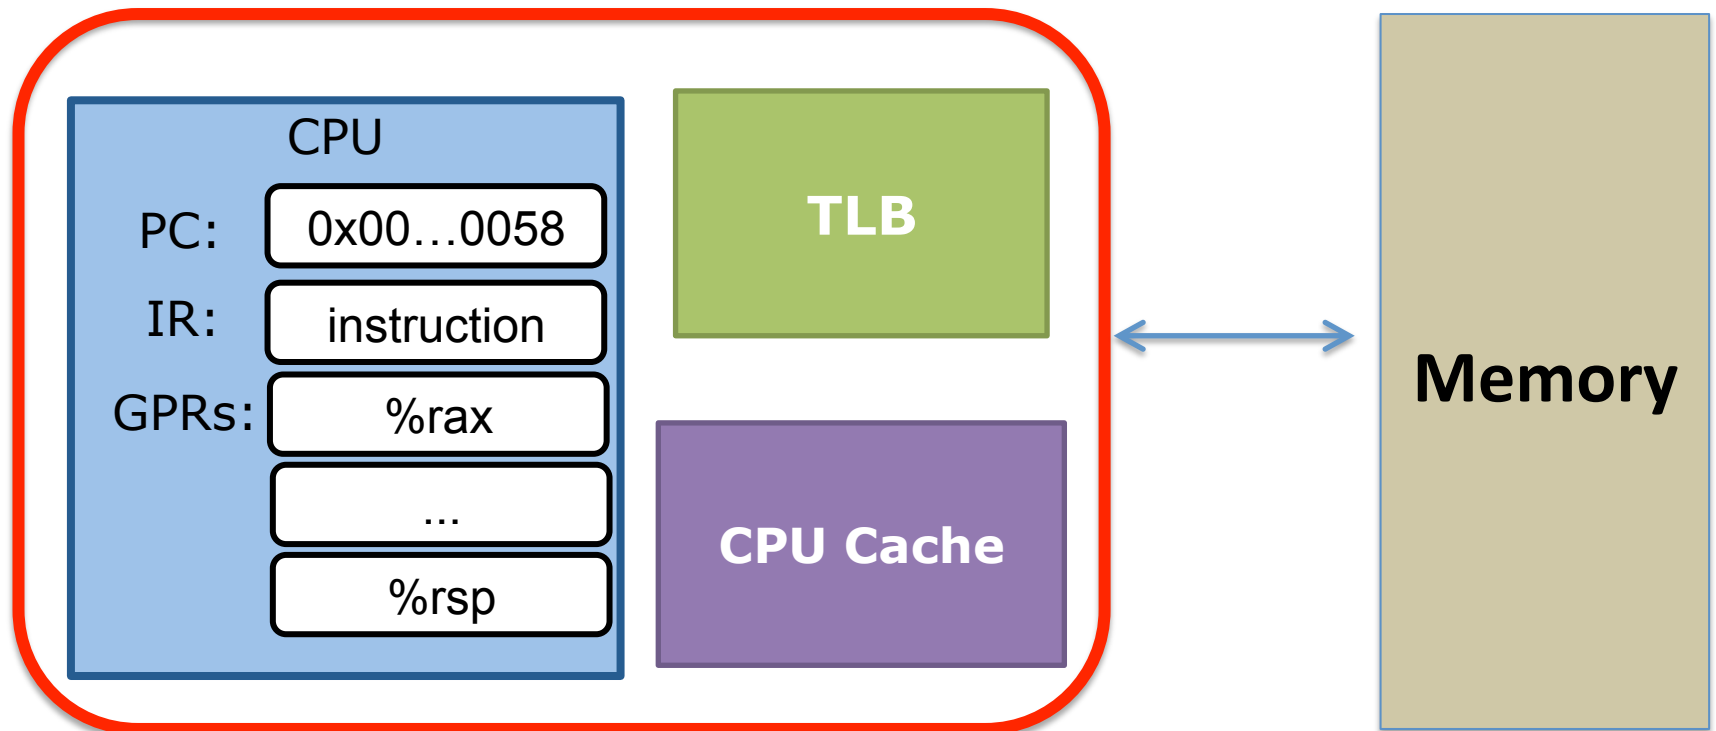
# Multi-processing



# Goal of multi-processing

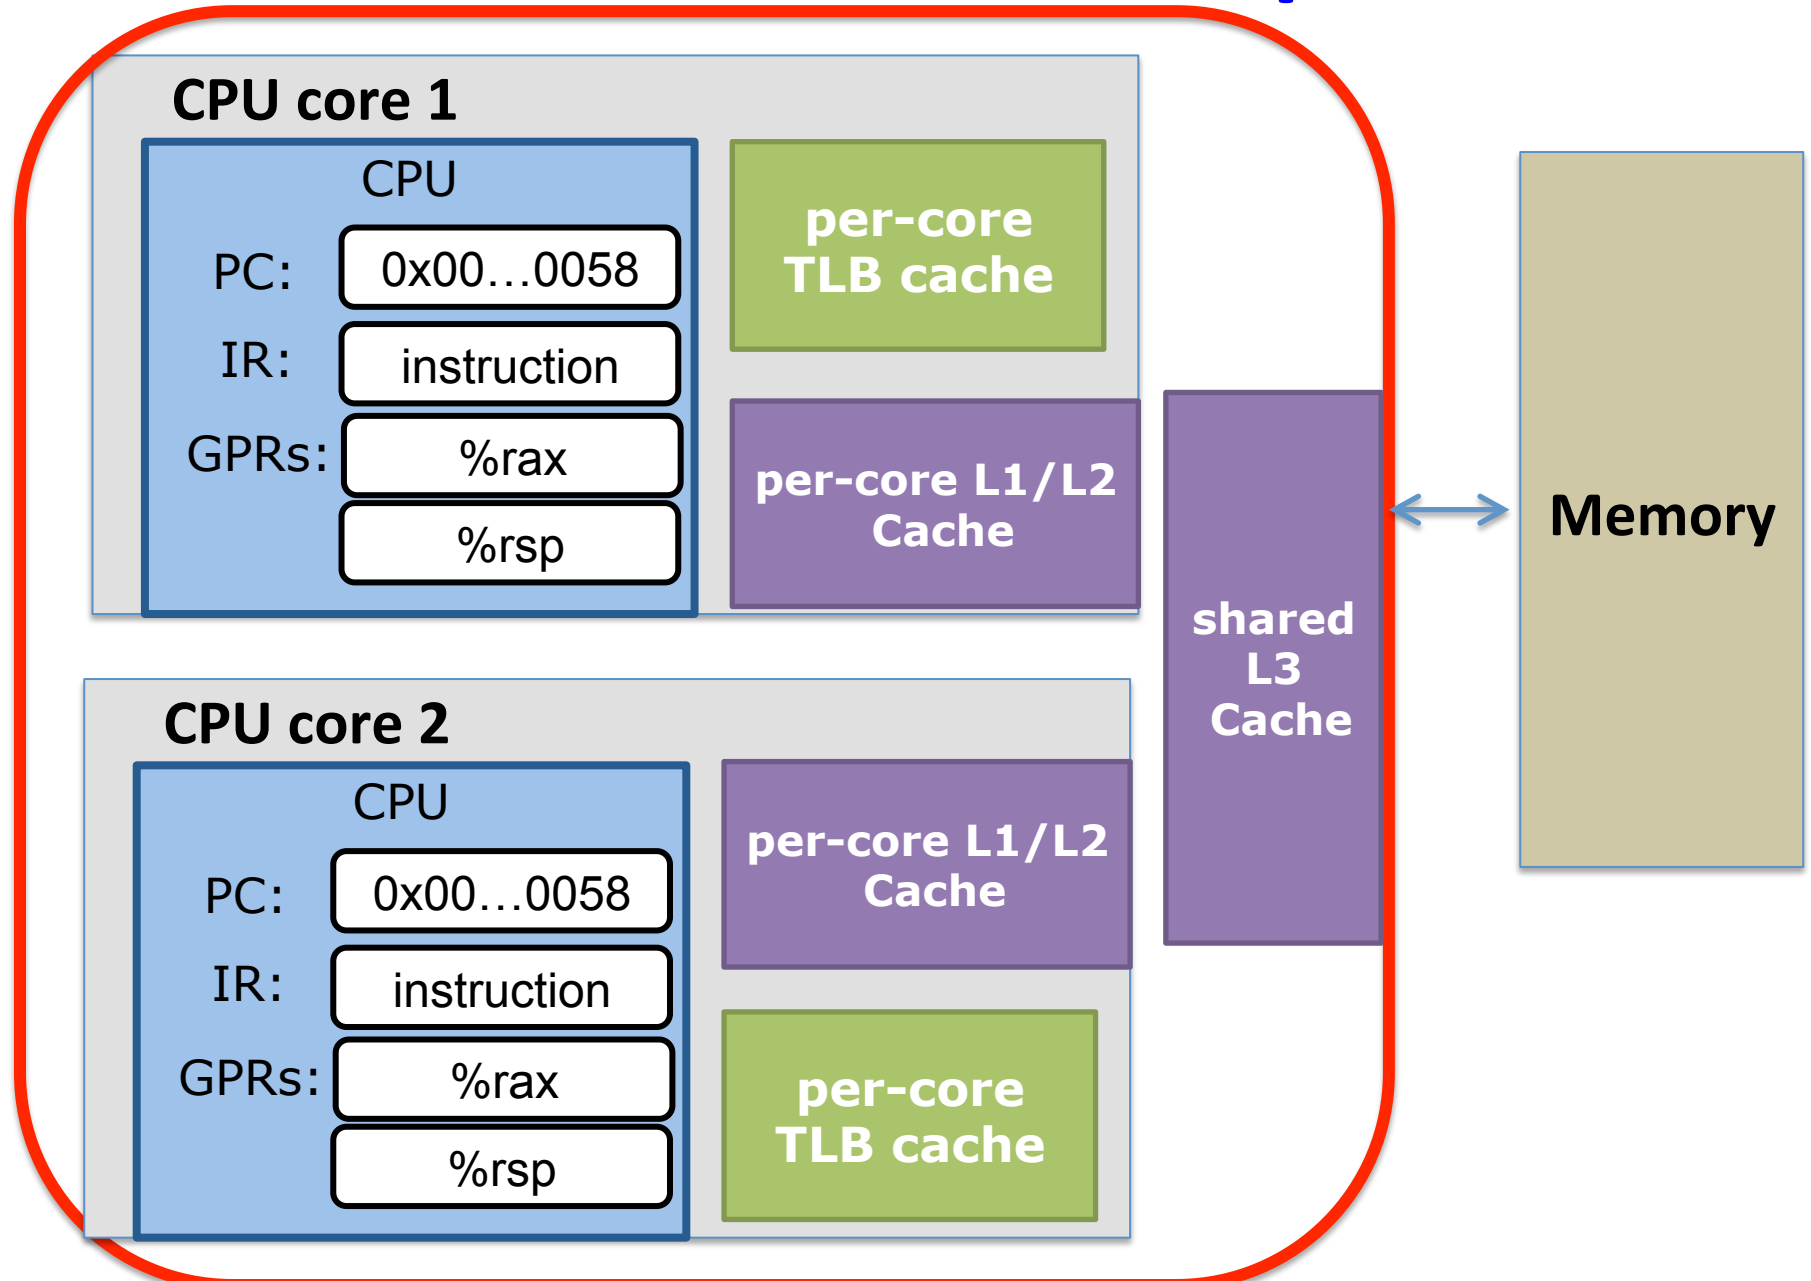
- Run multiple processes “simultaneously”
- Why?
  - listening to music while writing your lab
  - Running a web server, a database server, a PHP program together

# Modern CPUs have multiple cores




Your mental model of the CPU as a single core machine

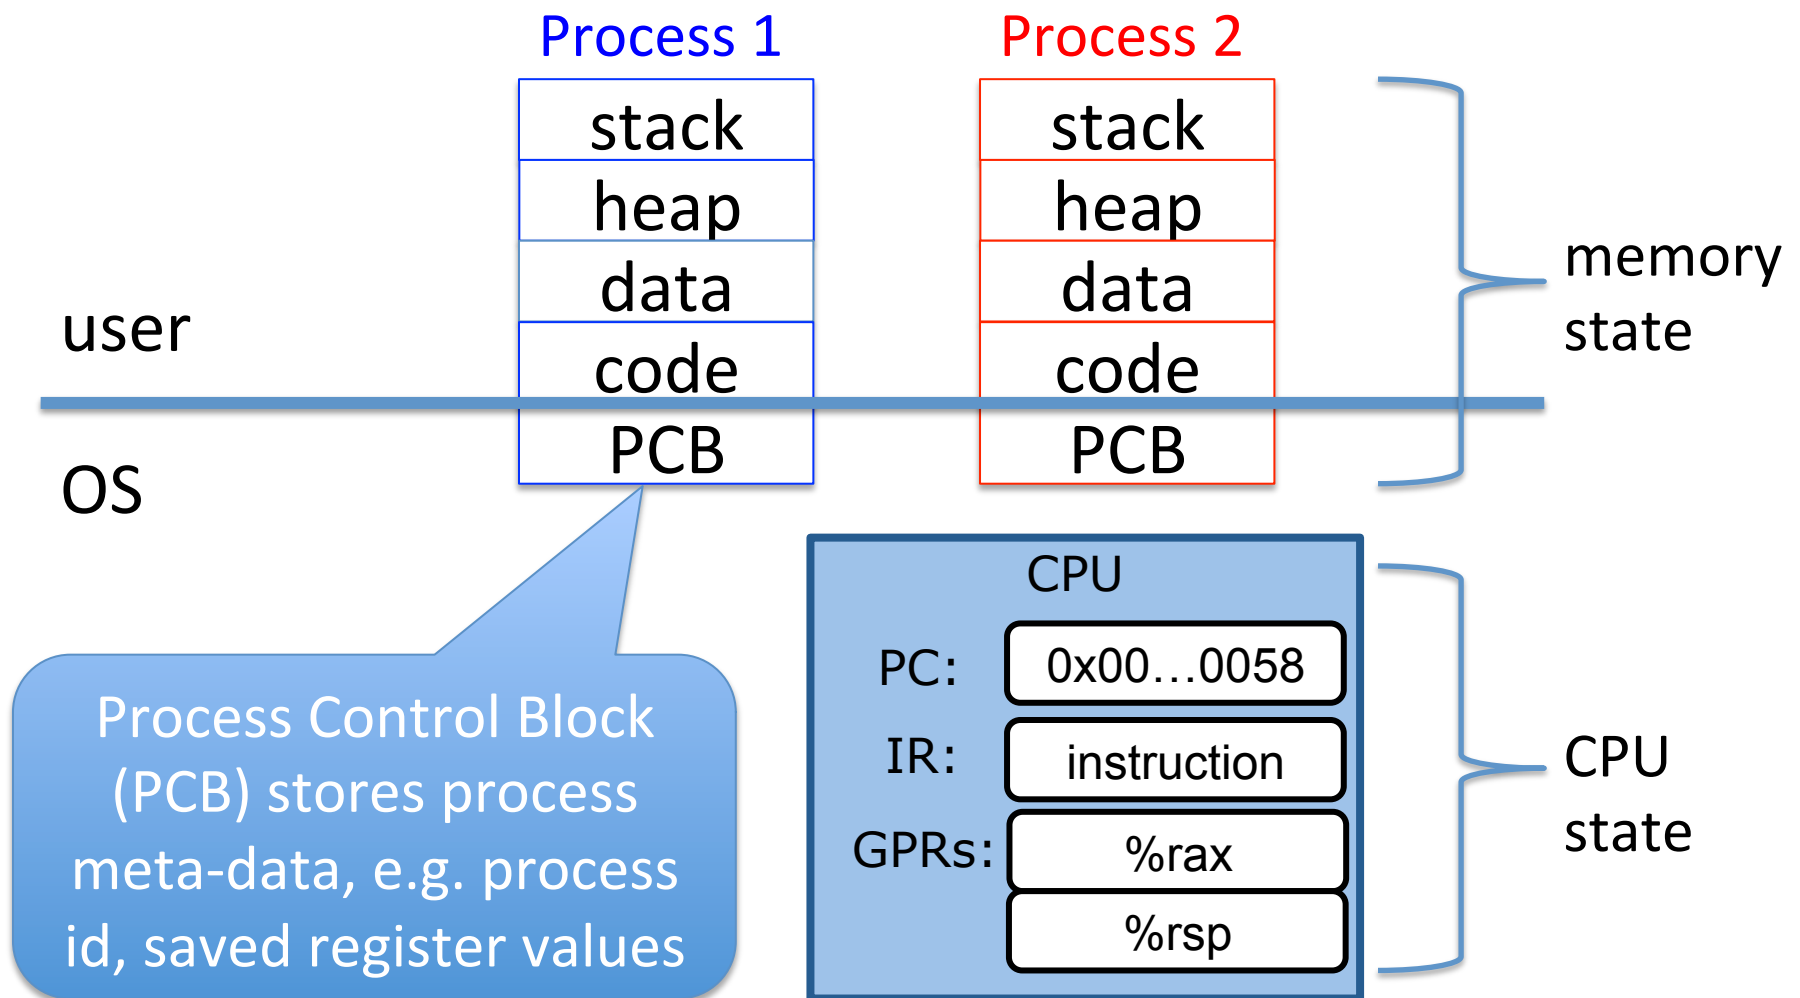
# Modern CPUs have multiple cores



# How to multi-process?

- Execute one process exclusive on each core?
  - 2 cores → 2 processes only 
- How to “simultaneously” execute more processes than there are cores?

# Multiprocessing (e.g. on a single core machine)



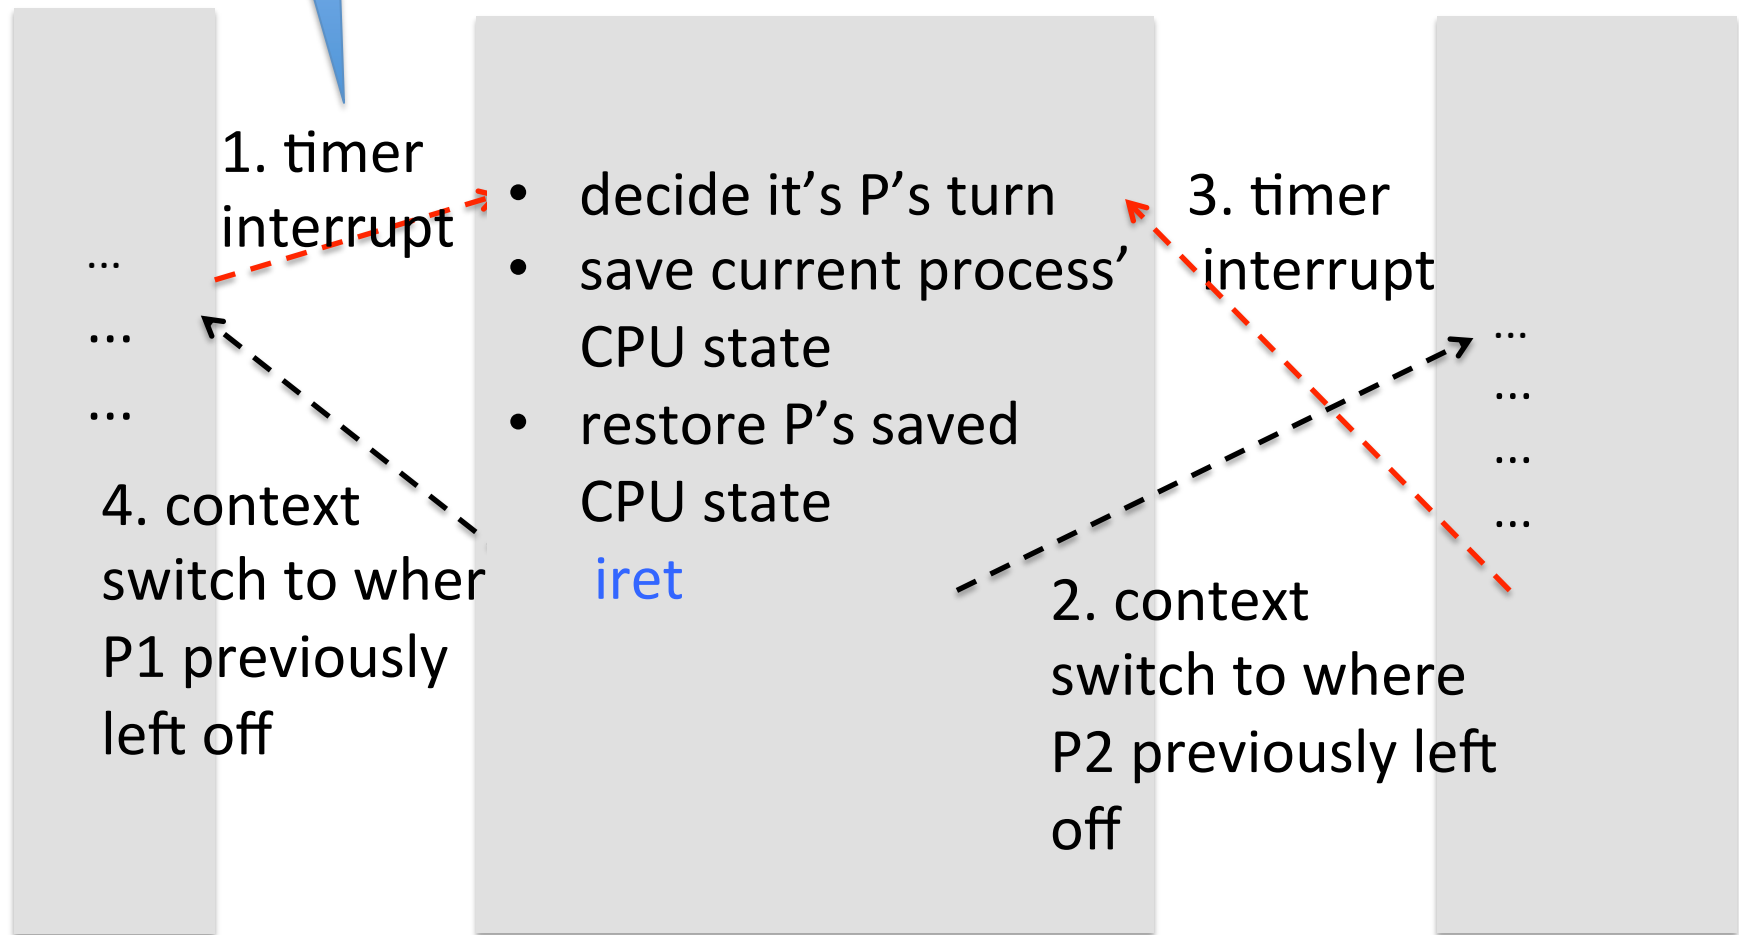
every  
10ms

# Context switch

Process P1

OS code

Process P2



# Creating and killing processes

- One process creates another process via syscall `fork()`
  - All processes are created by some processes (a tree).
  - The first process is a special one (`init`) and is created by OS.
  - When launching a program via command-line, the shell program creates the process

# The fork syscall

- OS creates a new child process (almost completely) identical to the parent process
- Same code, data, heap, stack, register state except different return values of the fork syscall
- Returns child process's id in parent process
- Returns zero in the child process



“called once, returned twice”




# Example fork call

```
void main()
{
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent, child pid=%d\n", pid);
    }
}
```

# Example fork call

process 1




```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```

# Example fork call


## process 1

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



## process 2


```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



# Example fork call


## process 1

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



## process 2


```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



# Example fork call


## process 1

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



## process 2


```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```




# Example fork call

## process 1

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



## process 2

```
void
main() {
    pid_t pid = fork();
     assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```


output:

In parent...

# Example fork call


## process 1

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



## process 2

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```




output:

In parent...

# Example fork call


## process 1

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



## process 2

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



output:


In parent...



# Example fork call


## process 1

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



## process 2

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("In child\n");
    } else {
        printf("In parent...\n");
    }
}
```



output:

In parent...

In child

# Notes on fork

- Execution of parent and child are concurrent
  - interleaving is non-deterministic.
  - In the example, both outputs are possible

In parent...

In child

In child

In parent...

- Parent and child have separate address space (but their contents immediately after fork are identical)

# Execution of parent and child are concurrent

```
void main()  
{  
1:    printf("hello\n");  
2:    fork();  
3:    printf("world\n");  
4:    fork();  
5:    printf("bye\n");  
}
```

How many processes are created in total?

# Execution of parent and child are concurrent

```
void main()  
{  
    L1: printf("hello\n");  
    L2: fork();  
    L3: printf("world\n");  
    L4: fork();  
    L5: printf("bye\n");  
}
```

What are the possible printouts?




hello  
world  
world  
bye  
bye  
bye  
bye



hello  
world  
bye  
bye  
world  
bye  
bye



hello  
world  
bye  
bye  
bye  
world  
bye



# Yet another example


```
void main()
{
    L1: printf("hello\n");
    L2: if (fork() == 0) {
        L3:   printf("big\n");
        L4:   if (fork() == 0) {
        L5:     printf("world\n");
                }
            }
    L6:   printf("bye\n");
}
```

What are the possible printouts?

✓  
hello  
big  
world  
bye  
bye  
bye

✓  
hello  
bye  
big  
bye  
world  
bye

✗  
hello  
bye  
big  
bye  
bye  
world



# Parent and child have separate address space with (initially) identical content

```
void main()
{
    int total = 0;
    → pid_t pid = fork();
    assert(pid >= 0);
    total++;
    if (pid == 0)
        printf("child %d\n", total);
    else
        printf("parent %d\n", total);
}
```

What are the possible printouts?



child 1  
parent 1



child 1  
parent 2



parent 1  
child 2

parent



# Parent and child have separate address space with (initially) identical content


```
void main()
{
    int total = 0;
    pid_t pid = fork();
    → assert(pid >= 0);
    total++;
    if (pid == 0)
        printf("child %d\n");
    else
        printf("parent %d\n");
}
```

What are the possible printouts?

✓ child 1  
parent 1

✗ child 1  
parent 2

✗ parent 1  
child 2



# Parent and child have separate address space with (initially) identical content

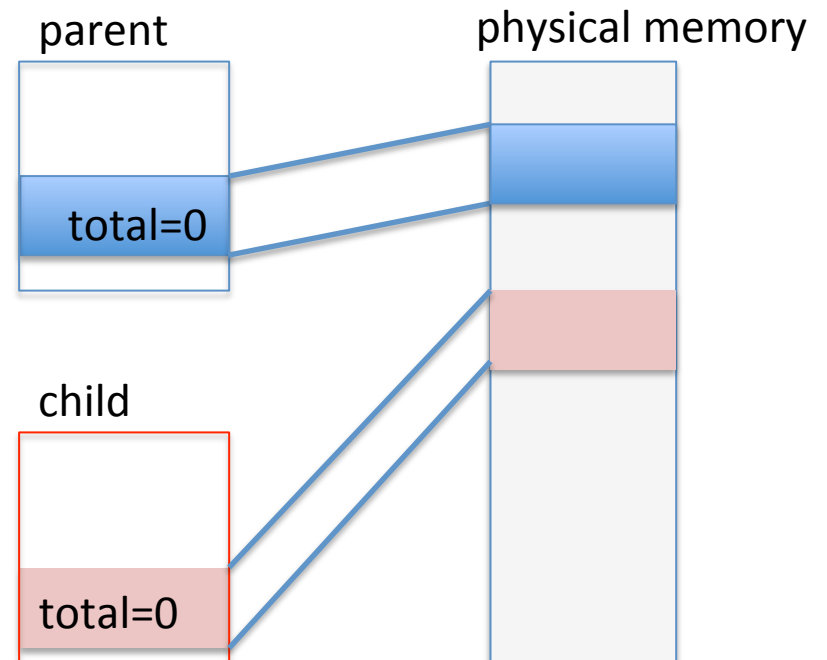
```
void main()
{
    int total = 0;
    pid_t pid = fork();
    → assert(pid >= 0);
    total++;
    if (pid == 0)
        printf("child %d\n");
    else
        printf("parent %d\n");
}
```

What are the possible printouts?

✓ child 1  
parent 1

✗ child 1  
parent 2

✗ parent 1  
child 2





# wait: synchronize with child

- Parent process could wait for the exit of its child process(es).
  - `int waitpid(pid_t pid, int * child_status, ...)`
- Good practice for parent to wait
  - Otherwise, some OS process state about the child cannot be freed even after child exits
  - leaks memory

# Exercise

What are the possible printouts?

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("child\n");
    } else {
        printf("parent\n");
    }
}
```

✓ child  
parent

✓ parent  
child

# Exercise

```
void
main() {
    pid_t pid = fork();
    assert(pid >= 0);
    if (pid == 0) {
        printf("child\n");
    } else {
        waitpid(pid, NULL, 0);
        printf("parent\n");
    }
}
```

What are the possible printouts?

✓	child	✗	parent
	parent		child

# **execv: load program in current process**

- `int execv(char *filename, char *argv[])`
  - overwrites code, data, heap, stack of existing process (retains process pid)
- called once, never returns

# Example

```
void main() {  
    pid_t pid;  
    pid = fork();  
    if (pid == 0) {  
        execv("/bin/echo", "hello");  
        printf("world\n");  
    }  
    waitpid(pid, NULL, 0);  
    printf("bye\n");  
}
```

Never executed  
because execv has  
replaced process's  
memory with that  
of the echo program

How many processes are created in total? output?

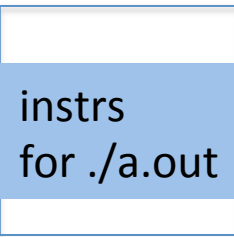
2

hello bye


# Example

```
void main() {  
    pid_t pid;  
    → pid = fork();  
    if (pid == 0) {  
        execv("/bin/echo", "hello");  
        printf("world\n");  
    }  
    waitpid(pid, NULL, 0);  
    printf("bye\n");  
}
```

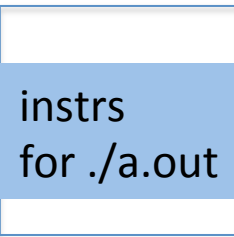
parent



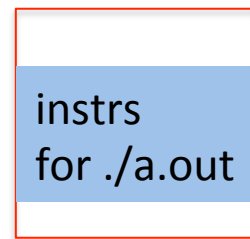
# Example

```
void main() {  
    pid_t pid;  
    pid = fork();  
     if (pid == 0) {  
        execv("/bin/echo", "hello");  
        printf("world\n");  
    }  
    waitpid(pid, NULL, 0);  
    printf("bye\n");  
}
```

parent



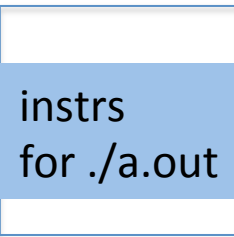
child



# Example

```
void main() {  
    pid_t pid;  
    pid = fork();  
    → if (pid == 0) {  
        → execv("/bin/echo", "hello");  
        printf("world\n");  
    }  
    waitpid(pid, NULL, 0);  
    printf("bye\n");  
}
```

parent



child

