

# Concurrency – Multithreading

Jinyang Li

based on slides by Tiger Wang

# Example


```
long bigloop(int *arr, int sz) {  
    long r = 0;  
    for(int i = 0; i < sz; i++)  
        r += arr[i];  
    return r;  
}
```

How to improve the performance  
with multicore?


```
int main() {  
    ...  
    long r = bigloop(arr, 1000000);  
    ...  
}
```

# Parallelization

bigloop: 0→7



# Parallelization



Performance can be improved by 4X

# Concurrency

## What's concurrency?

- things happening "simultaneously"
  1. multiple CPU cores concurrently executing instructions
  2. CPU and I/O devices concurrently doing processing

## Why write concurrent programs?

- speed up programs using multiple CPUs
- speed up programs by concurrently doing CPU processing and I/O.

# How to write concurrent programs?

## Use multiple processes

- Each process uses a different CPU
- Different processes runs different tasks
  - They have separate address spaces
  - Elaborate to communicate with each other

## Use multiple threads

# In this lecture


## Use multiple processes

- Each process uses a different CPU
- Different processes runs different tasks
  - They have separated address space
  - Elaborate to communicate with each other

Use multiple threads

# Multiple threads (Multithreading)

Process



```
long bigloop(int *arr, int sz) {  
    long r = 0;  
    for(int i = 0; i < sz; i++)  
        r += arr[i];  
    return r;  
}
```

```
int main() {  
    ...  
    long r = bigloop(arr, 8);  
    ...  
}
```

CPU0

CPU1

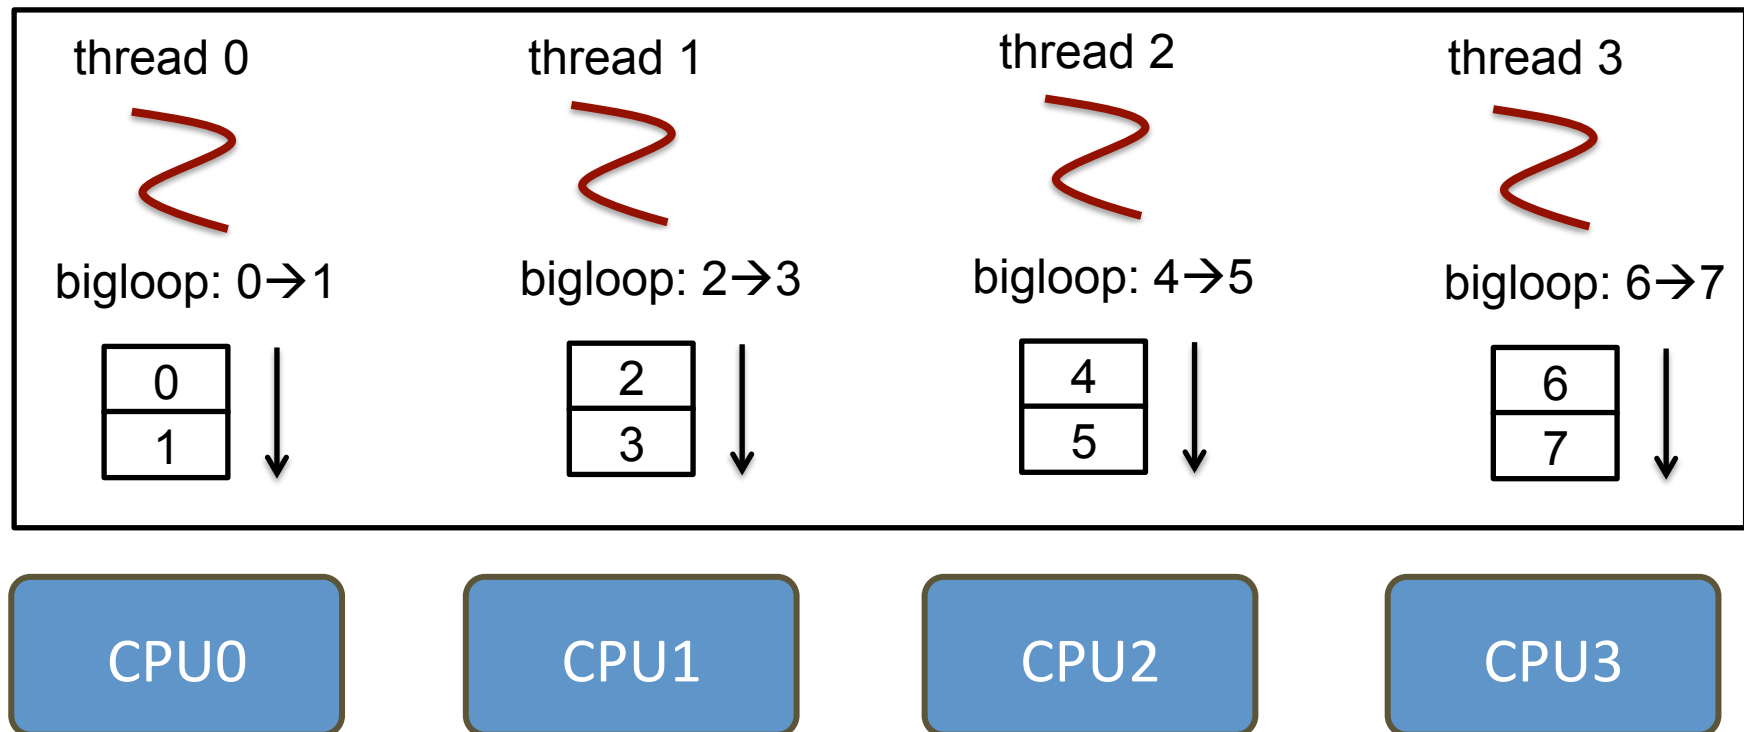
CPU2

CPU3



# Multiple threads (Multithreading)

Process




# Multiple threads (Multithreading)


Single process, multiple threads

- Share the same memory space
- Has its own stack
- Has its own control flow

Process




# Share the memory space

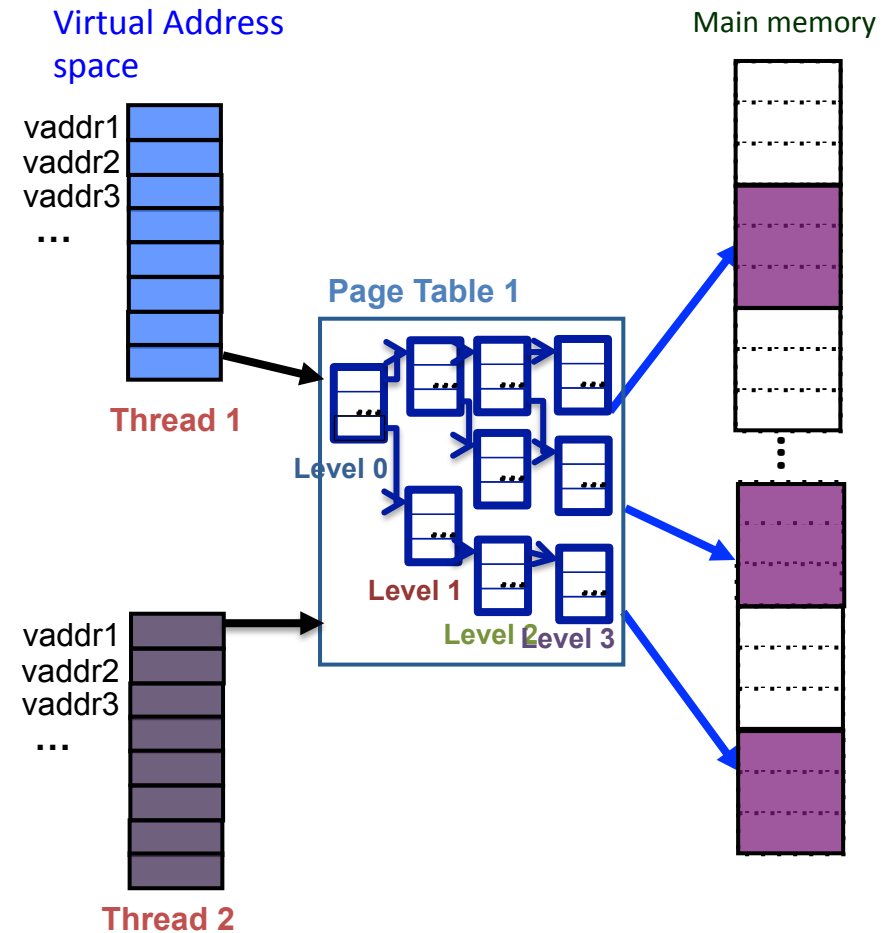


Different processes have different page tables

# Share the memory space

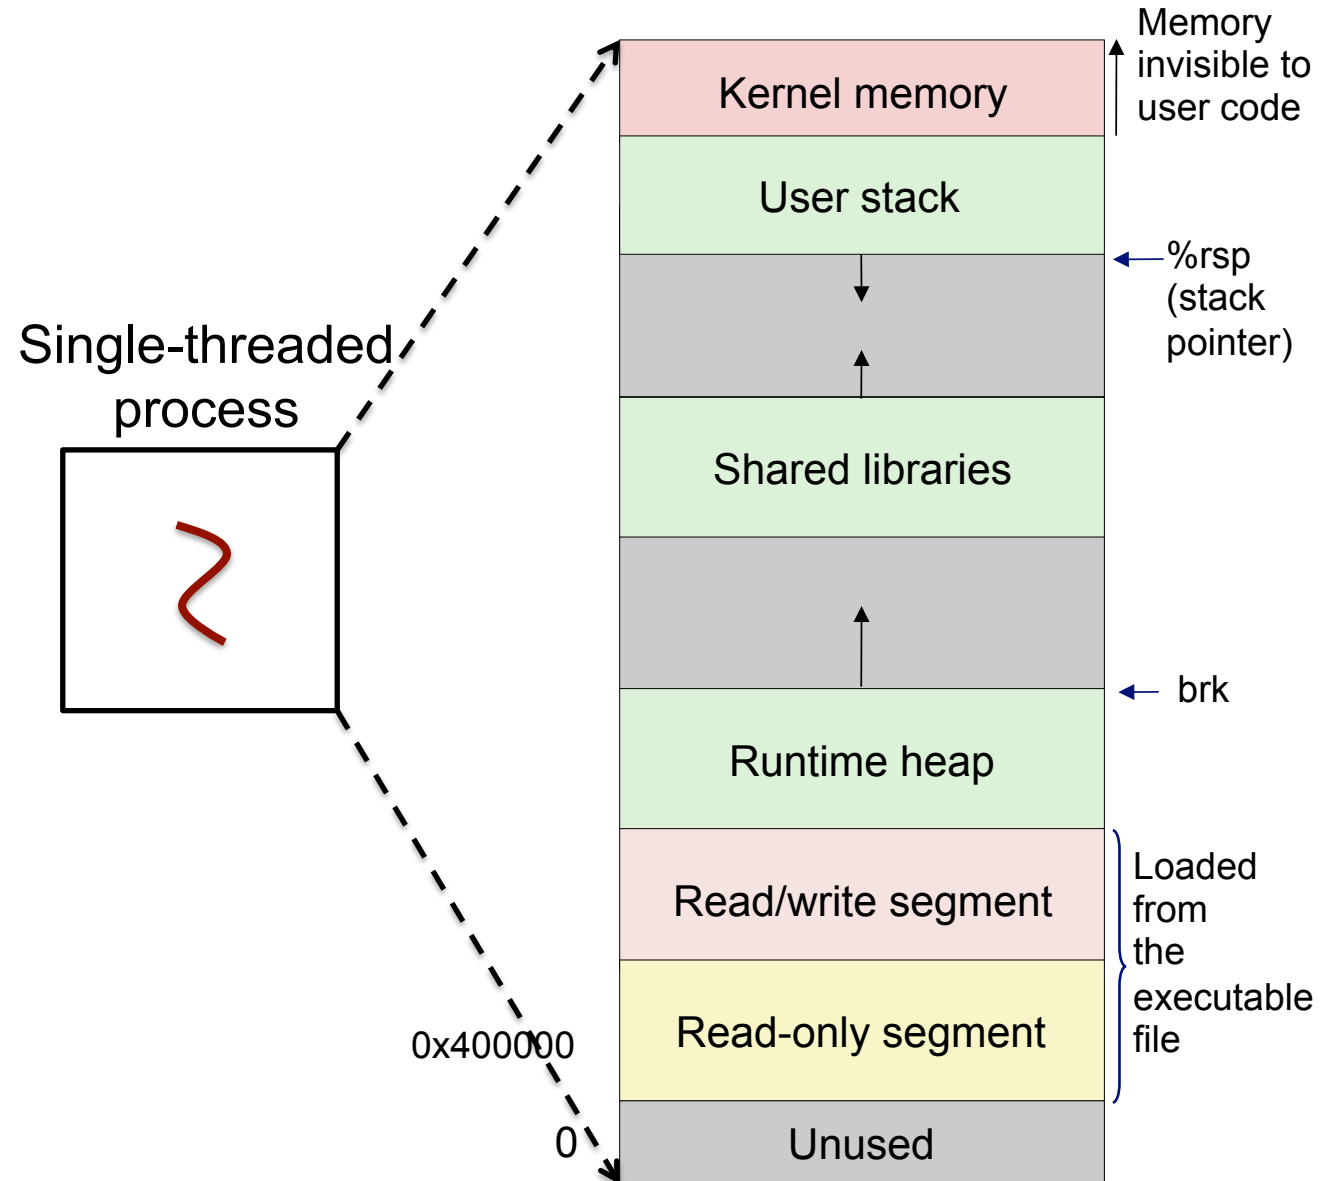


Different processes have different page tables



Different threads of the same process share the same page table

# Single threaded process

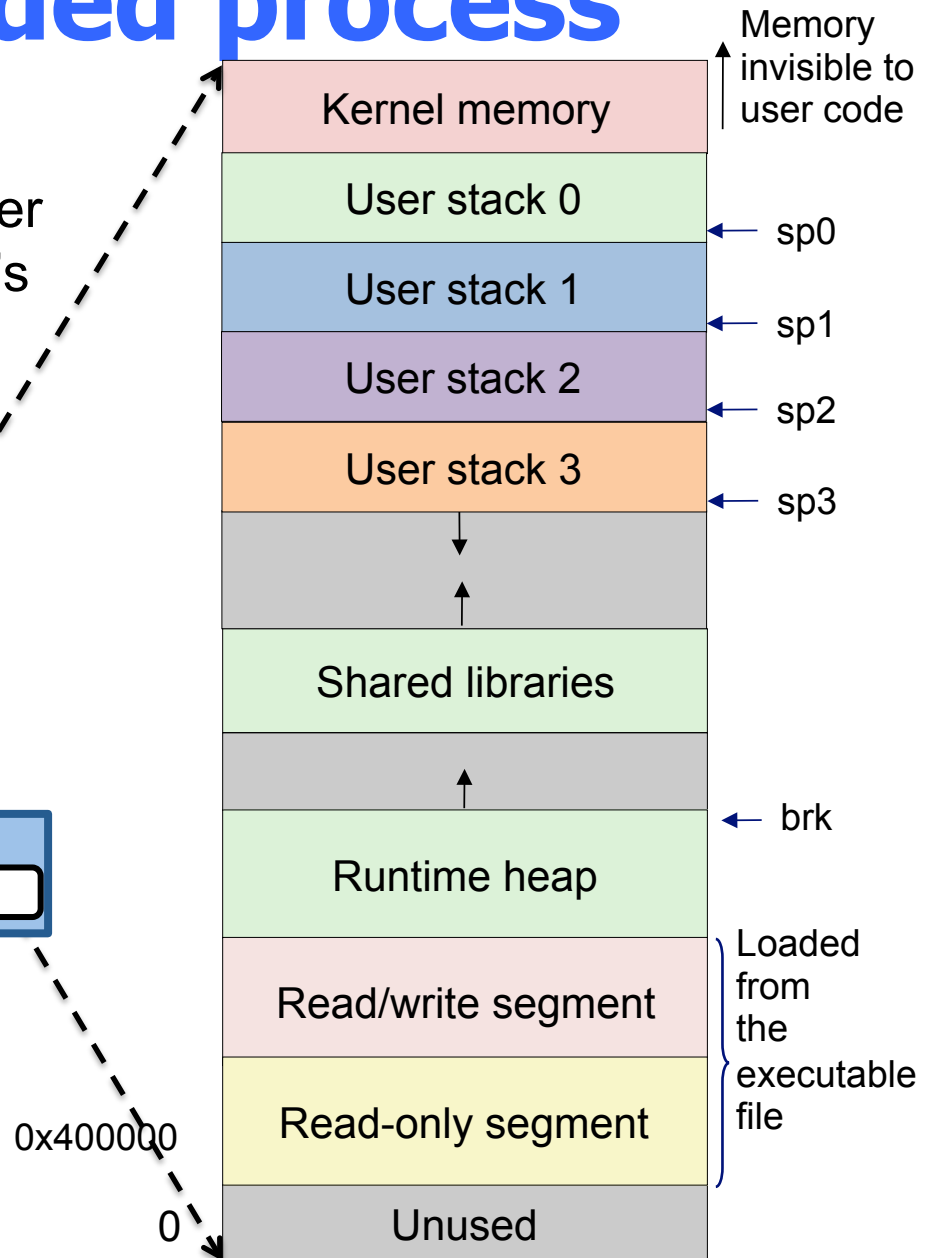



# Multi-threaded process

Each thread has its own stack

- Each thread has its own stack pointer
- Store stack pointer into a CPU core's %rsp before running

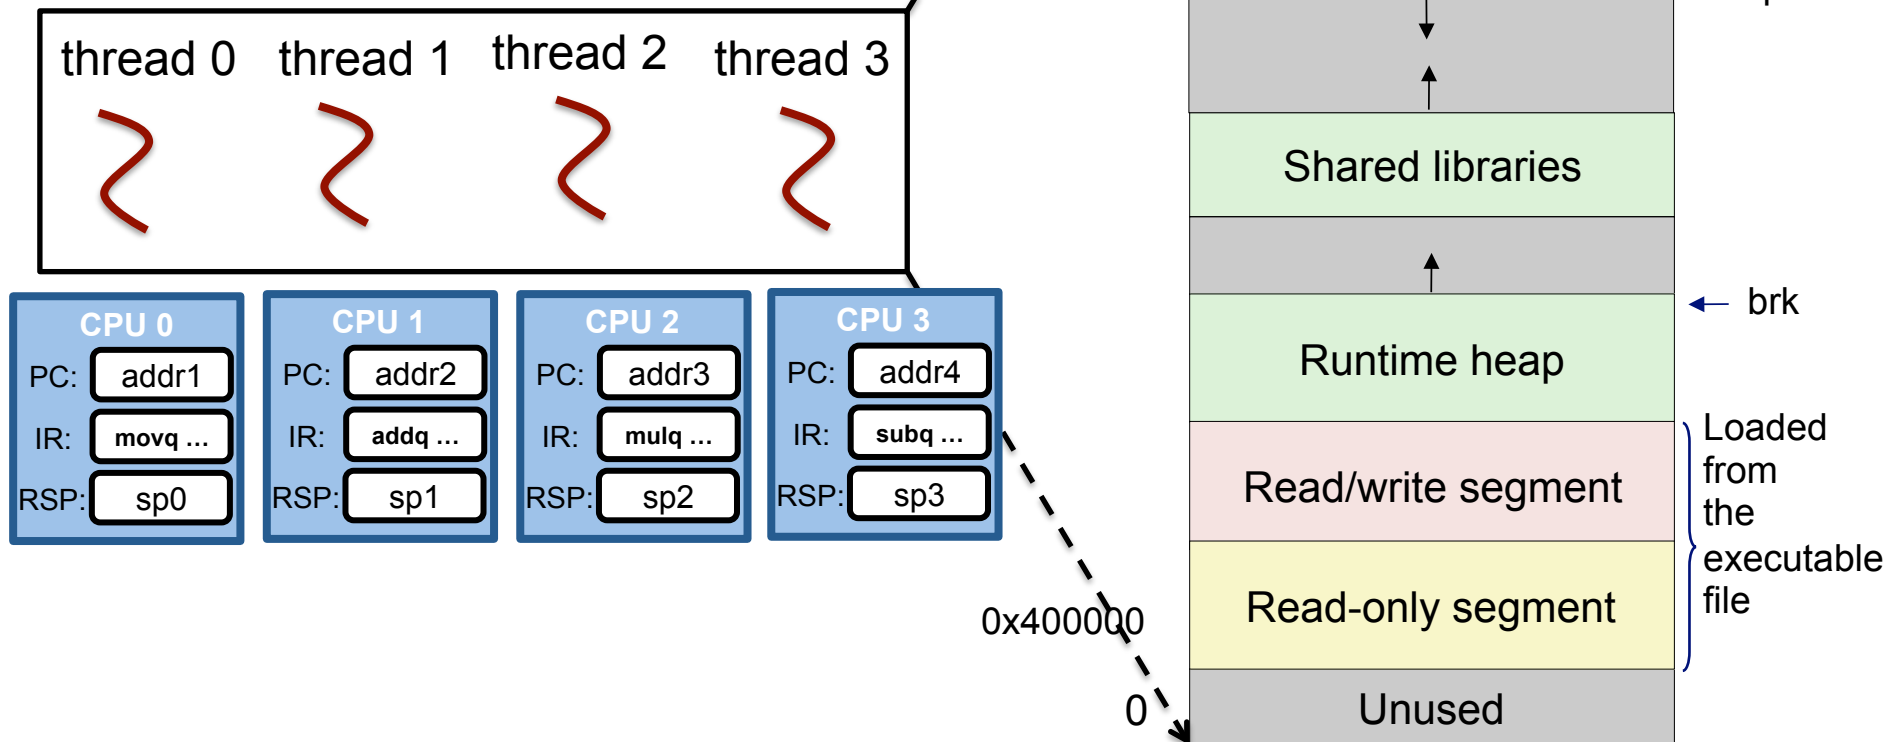
Process 1



# Own control flow

Each thread has its own CPU state (registers, RFLAGS). It loads its CPU state to a CPU core's registers before running.

Process 1



# POSIX thread interface

POSIX: Portable Operating System Interface

- POSIX defines the API for variants of Unix

Thread interface defined by POSIX

- pthread\_create: create a new thread
- pthread\_join: wait for the target thread terminated



# pthread\_create

```
#include <pthread.h>
int pthread_create(pthread_t *thread_id,
                  const pthread_attr_t *attr,
                  void *(*start_routine)(void*),
                  void *arg);
```

## Create a new thread

- It executes `start_routine` with `arg` as its sole argument.
- Its attribute is specified by `attr`
- Upon successful completion, it will store the ID of the created thread in the location referenced by `thread_id`.

## Return value

- zero: success
- non-zero (error number): fail

# Example 1 – Create

```
void* func(void* arg) {  
    printf("This is the created thread\n");  
    return NULL;  
}  
  
int main(int argc, char* argv[]) {  
  
    pthread_t tid;  
    int r = pthread_create(&tid, NULL, &func, NULL);  
    if(r != 0) {  
        printf("create thread failed");  
        return 1;  
    }  
  
    return 0;  
}
```

```
gcc create.c -lpthread
```

# Example 1 – Create

```
void* func(void* arg) {  
    printf("This is the created thread\n");  
    return NULL;  
}  
  
int main(int argc, char* argv[]) {  
    pthread_t tid;  
    pthread_create(&tid, NULL, &func, NULL);  
    return 0;  
}
```


```
gcc create.c -lpthread
```

Process finishes when its main thread exits.

- All created threads are terminated

# pthread\_join

```
#include <pthread.h>
int pthread_join(pthread_t thread_id, void **ret_ptr);
```



Wait for the target thread to finish

- Upon success, the return value of the target thread is stored at the location pointed to by ret\_ptr.

Return value

- zero: success
- non-zero (error number): fail

## Example 2 – Join

```
void* func(void* arg) {  
    printf("This is the created thread\n");  
    return NULL;  
}  
  
int main(int argc, char* argv[]) {  
  
    pthread_t tid;  
    pthread_create(&tid, NULL, &func, NULL);  
    pthread_join(tid, NULL);  
    return 0;  
}
```

## Example 3 – Parameter

```
void* func(void* arg) {  
    int p = *(int *)arg;  
    p = p + 1;  
    return &p;  
}
```

Question – what is expected result ?

```
int main(int argc, char* argv[]) {  
  
    int param = 100;  
  
    pthread_t tid;  
    pthread_create(&tid, NULL, &func, (void *)&param);  
    ...  
  
    int *res = NULL;  
    pthread_join(tid, &res);  
    ...  
  
    printf("result: addr %lx val %d\n", res, *res);  
    return 0;  
}
```

## Example 3 – Parameter

```
void* func(void* arg) {  
    int p = *(int *)arg;  
    p = p + 1;  
    return &p;  
}
```

p is on the stack of the created thread  
-- it is destroyed when the thread terminates

```
int main(int argc, char* argv[]) {  
  
    int param = 100;  
  
    pthread_t tid;  
    pthread_create(&tid, NULL, &func, (void *)&param);  
    ...  
  
    int *res = NULL;  
    pthread_join(tid, &res);  
    ...  
  
    printf("result: addr %lx val %d\n", res, *res);  
    return 0;  
}
```

# Example 3 – Parameter

```
void* func(void* arg) {
    int p = *(int *)arg;
    p = p + 1;
    int *r = malloc(sizeof(int));
    *r = p;
    return (void *)r;
}

int main(int argc, char* argv[]) {

    int param = 100;

    pthread_t tid;
    pthread_create(&tid, NULL, &func, (void *)&param);
    ...

    int *res = NULL;
    pthread_join(tid, &res);
    ...

    printf("result: addr %lx val %d\n", res, *res);
    return 0;
}
```



# Example 3 – Parameter

```
void* func(void* arg) {
    int p = *(int *)arg;
    p = p + 1;
    int *r = malloc(sizeof(int));
    *r = p;
    return (void *)r;
}

int main(int argc, char* argv[]) {

    int param = 100;

    pthread_t tid;
    pthread_create(&tid, NULL, &func, (void *)&param);
    ...

    int *res = NULL;
    pthread_join(tid, &res);
    ...

    printf("result: addr %lx val %d\n", res, *res);
    free(res);
    return 0;
}
```

## Example 4 – Interleave

```
void* func(void* arg) {  
    printf("1");  
}  
  
int main(int argc, char* argv[]) {  
  
    printf("0");  
  
    pthread_t tid;  
    int r = pthread_create(&tid, NULL, &func, NULL);  
    ...  
    printf("2");  
  
    ...  
    return 0;  
}
```

Question – what is the expected result ?

## Example 4 – Interleave

```
void* func(void* arg) {  
    printf("1");  
}
```

Question – what is the expected result ?

Answer: 012 or 021


```
int main(int argc, char* argv[]) {  
  
    printf("0");  
  
    pthread_t tid;  
    int r = pthread_create(&tid, NULL, &func, NULL);  
    ...  
    printf("2");  
  
    ...  
    return 0;  
}
```

# Example 4 – Interleave

```
void* func(void* arg) {  
    printf("1");  
}  
  
int main(int argc, char* argv[]) {  
    printf("0");  
  
    pthread_t tid;  
    int r = pthread_create(  
        &tid, NULL, &func, NULL);  
    ...  
    printf("2");  
    ...  
    return 0;  
}
```

Question – what is the expected result ?

Answer: 012 or 021

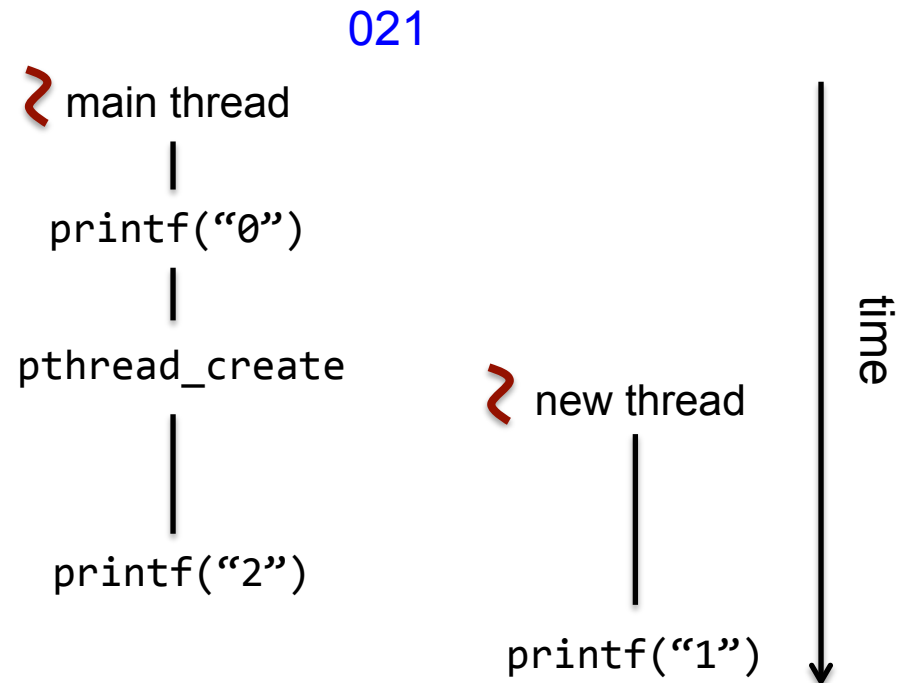


# Example 4 – Interleave

```
void* func(void* arg) {  
    printf("1");  
}  
  
int main(int argc, char* argv[]) {  
    printf("0");  
  
    pthread_t tid;  
    int r = pthread_create(  
        &tid, NULL, &func, NULL);  
    ...  
    printf("2");  
    ...  
    return 0;  
}
```

Question – what is the expected result ?

Answer: 012 or 021



# Example 5 – Stack, Heap, Global

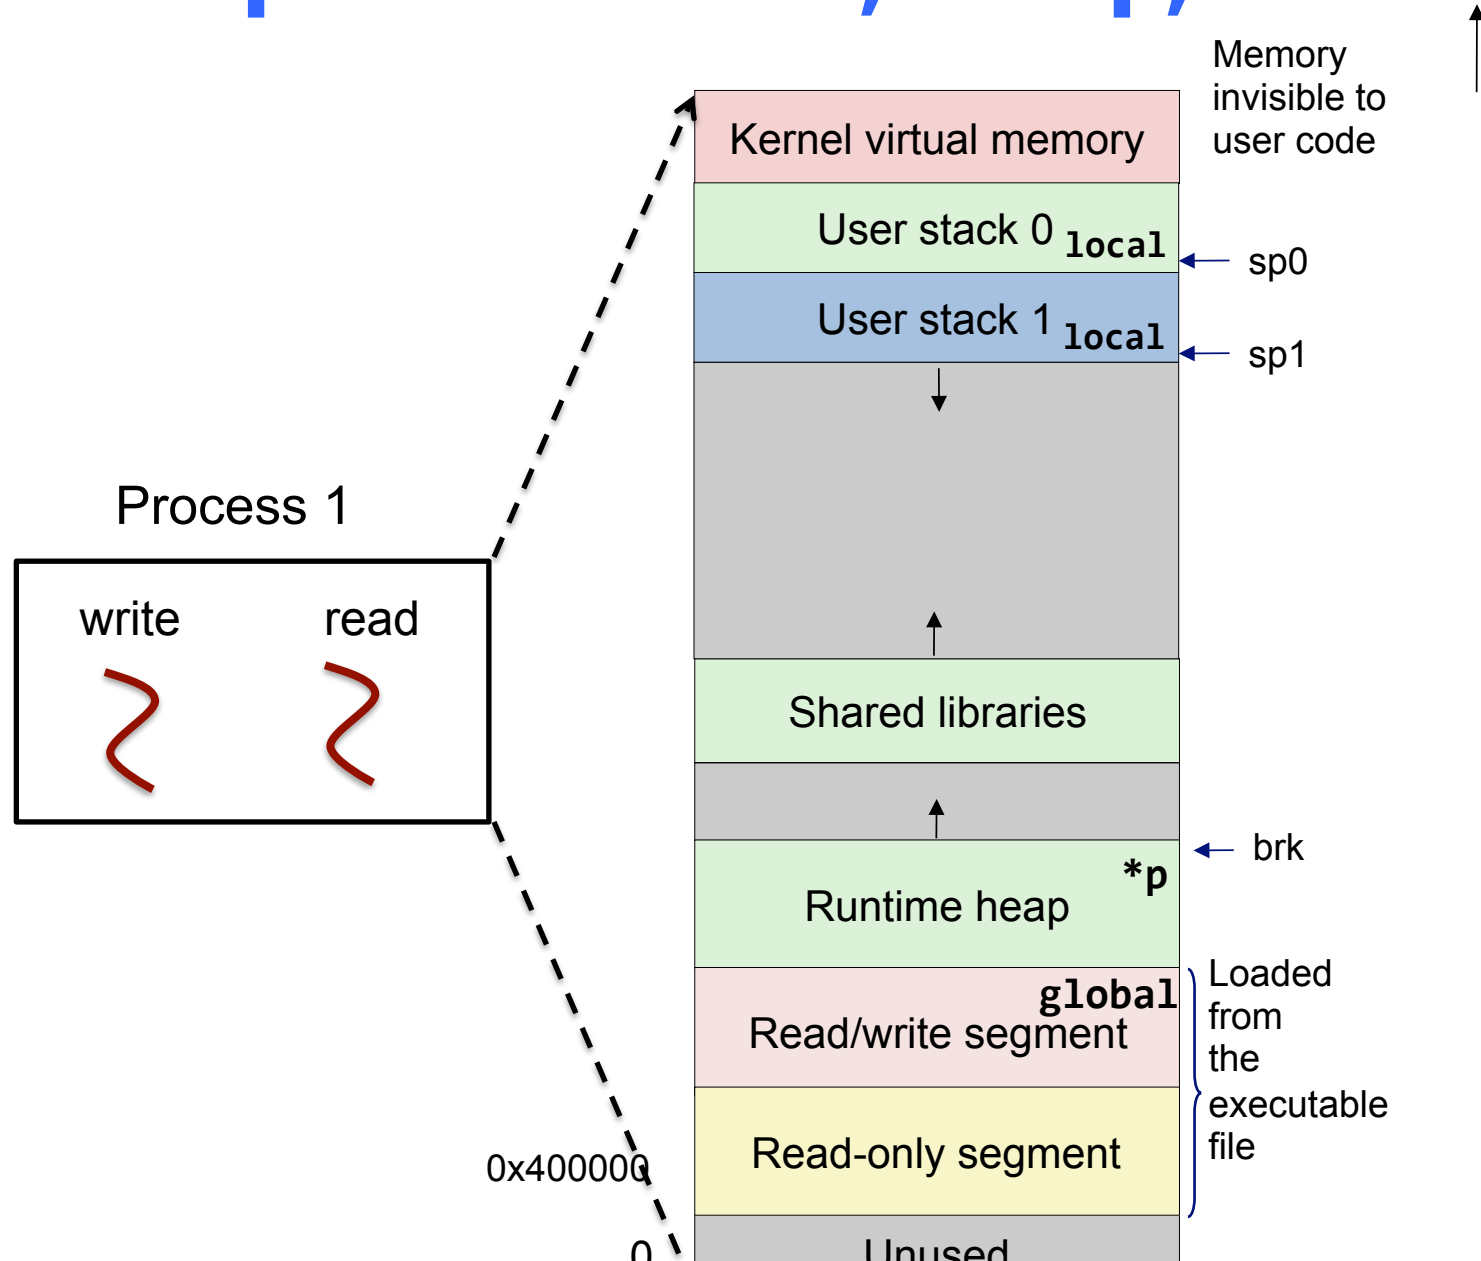
```
int global = 0;

void* write(void* arg) {
    int local = 100;
    global = 100;
    *(int *)arg = 100;
}

void* read(void* arg) {
    int local = 0;
    printf("local %d global %d heap %d\n",
        local, global, *(int *)arg);
    return NULL;
}

int main(int argc, char* argv[]) {
    int *p = (int *)malloc(sizeof(int));
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, &write, (void *)p);
    ...
    pthread_join(tid1, NULL);
    pthread_create(&tid2, NULL, &read, (void *)p);
    ...
    return 0;
}
```

# Example 5 – Stack, Heap, Global



# Example 5 – Stack, Heap, Global

```
int global = 0;

void* write(void* arg) {
    int local = 0;
    local = 100;
    global = 100;
    int *ptr = (int *)arg;
    (*ptr) = 100;
}

void* read(void* arg) {
    int local = 0;
    printf("local %d global %d heap %d\n",
           local, global, *(int *)arg);
    return NULL;
}

int main(int argc, char* argv[]) {
    int *p = (int *)malloc(sizeof(int));
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, &write, (void *)p);
    ...
    pthread_join(tid1, NULL);
    pthread_create(&tid2, NULL, &read, (void *)p);
    ...
    return 0;
}
```

What are the output?

local 0 global 100 heap 100



# Example 5 – Stack, Heap, Global

```
int global = 0;

void* write(void* arg) {
    int local = 0;
    local = 100;
    global = 100;
    int *ptr = (int *)arg;
    (*ptr) = 100;
}

void* read(void* arg) {
    int local = 0;
    printf("local %d global %d heap %d\n",
           local, global, *(int *)arg);
    return NULL;
}

int main(int argc, char* argv[]) {
    int *p = (int *)malloc(sizeof(int));
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, &write, (void *)p);
    ...
pthread_join(tid1, NULL);
    pthread_create(&tid2, NULL, &read, (void *)p);
    ...
    return 0;
}
```

What are the output?

local 0 global 0 heap 0

local 0 global 100 heap 0

local 0 global 100 heap 100

# Example 6 – bigloop

```
#define LEN 1000000000
```

Parallelize bigloop into two threads

```
long bigloop(int *arr) {  
    long r = 0;  
    for(int i = 0; i < LEN; i++)  
        r += arr[i];  
    return r;  
}  
  
int main() {  
    int *arr = malloc(LEN * sizeof(int));  
    ...  
    long r = bigloop(arr);  
    ...  
}
```

# Example 6 – bigloop

```
#define LEN 1000000000
```

```
void* loop_thr1(void *arg){
    long *r = malloc(sizeof(long));
    int *arr = (int *)arg;

    for(int i = 0; i < LEN/2; i++)
        (*r) += arr[i];
    return (void *)r;
}
```

```
int main() {
    int *arr = malloc(LEN * sizeof(int));
    ...
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, &loop_thr1, (void *)arr);
    pthread_create(&tid2, NULL, &loop_thr2, (void *)arr);
    long *res1, *res2;
    pthread_join(tid1, &res1);
    pthread_join(tid2, &res2);
    printf("result is %ld\n", (*res1) + (*res2));
}
```

```
void* loop_thr2(void *arg){
    long *r = malloc(sizeof(long));
    int *arr = (int *)arg;

    for(int i = LEN/2; i < LEN; i++)
        (*r) += arr[i];
    return (void *)r;
}
```

Can we merge loop\_thr1 with loop\_thr2?

# Example 6 – bigloop

```
#define LEN 1000000000

typedef struct {
    int *arr;
    int len;
} loop_info;

void* loop(void *arg){
    loop_info *info = (loop_info *)arg;
    long *r = malloc(sizeof(long));
    for(int i = 0; i < info->len; i++)
        (*r) += info->arr[i];
    return (void *)r;
}

int main() {
    int *arr = malloc(LEN * sizeof(int));
    ...
    pthread_t tids[2];
    for (int i = 0; i < 2; i++) {
        loop_info *info = (loop_info *)malloc(sizeof(loop_info));
        info->arr = arr + i * LEN/2;
        info->len = LEN/2;
        pthread_create(&tids[i], NULL, &loop, (void *)info);
    }
    for (int i = 0; i < 2; i++) {
        long *res;
        pthread_join(tids[i], &res);
        result += (*res);
    }
}
```