

# Concurrency – Locking

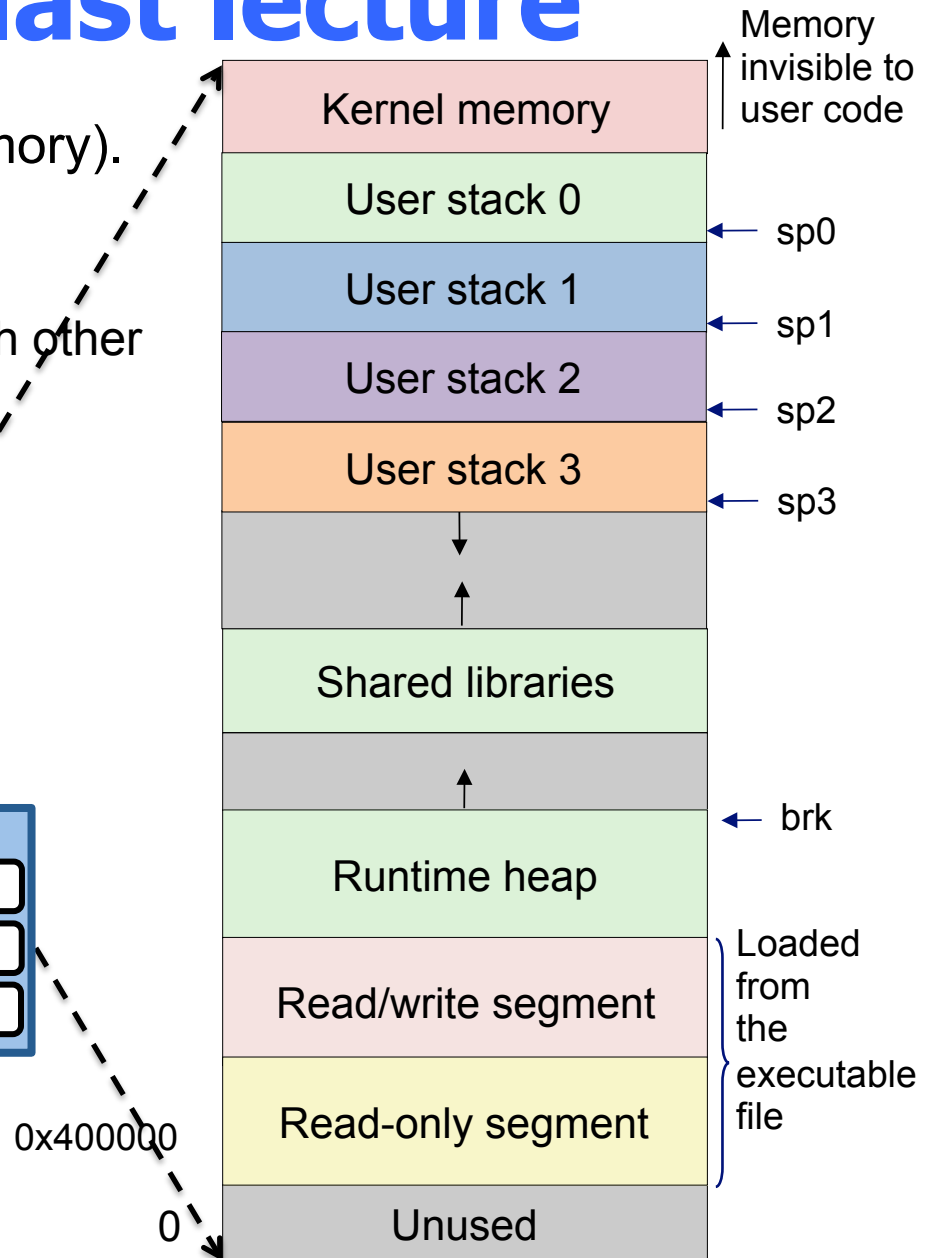

Jinyang Li

based on Tiger Wang's slides

# Review of last lecture

- Threads share address space (memory).
- Threads have separate CPU state.
- Threads have separate stacks
  - stacks are not protected from each other

Process



# Review of last lecture

- Create threads.
  - `pthread_create(&tid, NULL, &start, &args);`
- Wait for threads to finish.
  - `pthread_join(tid, &res);`
- The interleaving of threads are non-deterministics

# This lecture

- Races lead to buggy execution
- Preventing races using locks

# What are races? An example

global++



```
mov 0x20072d(%rip),%eax // load global into %eax
add $0x1,%eax           // update %eax by 1
mov %eax,0x200724(%rip) // restore global with %eax
```

Thread 1



global++

global: 0

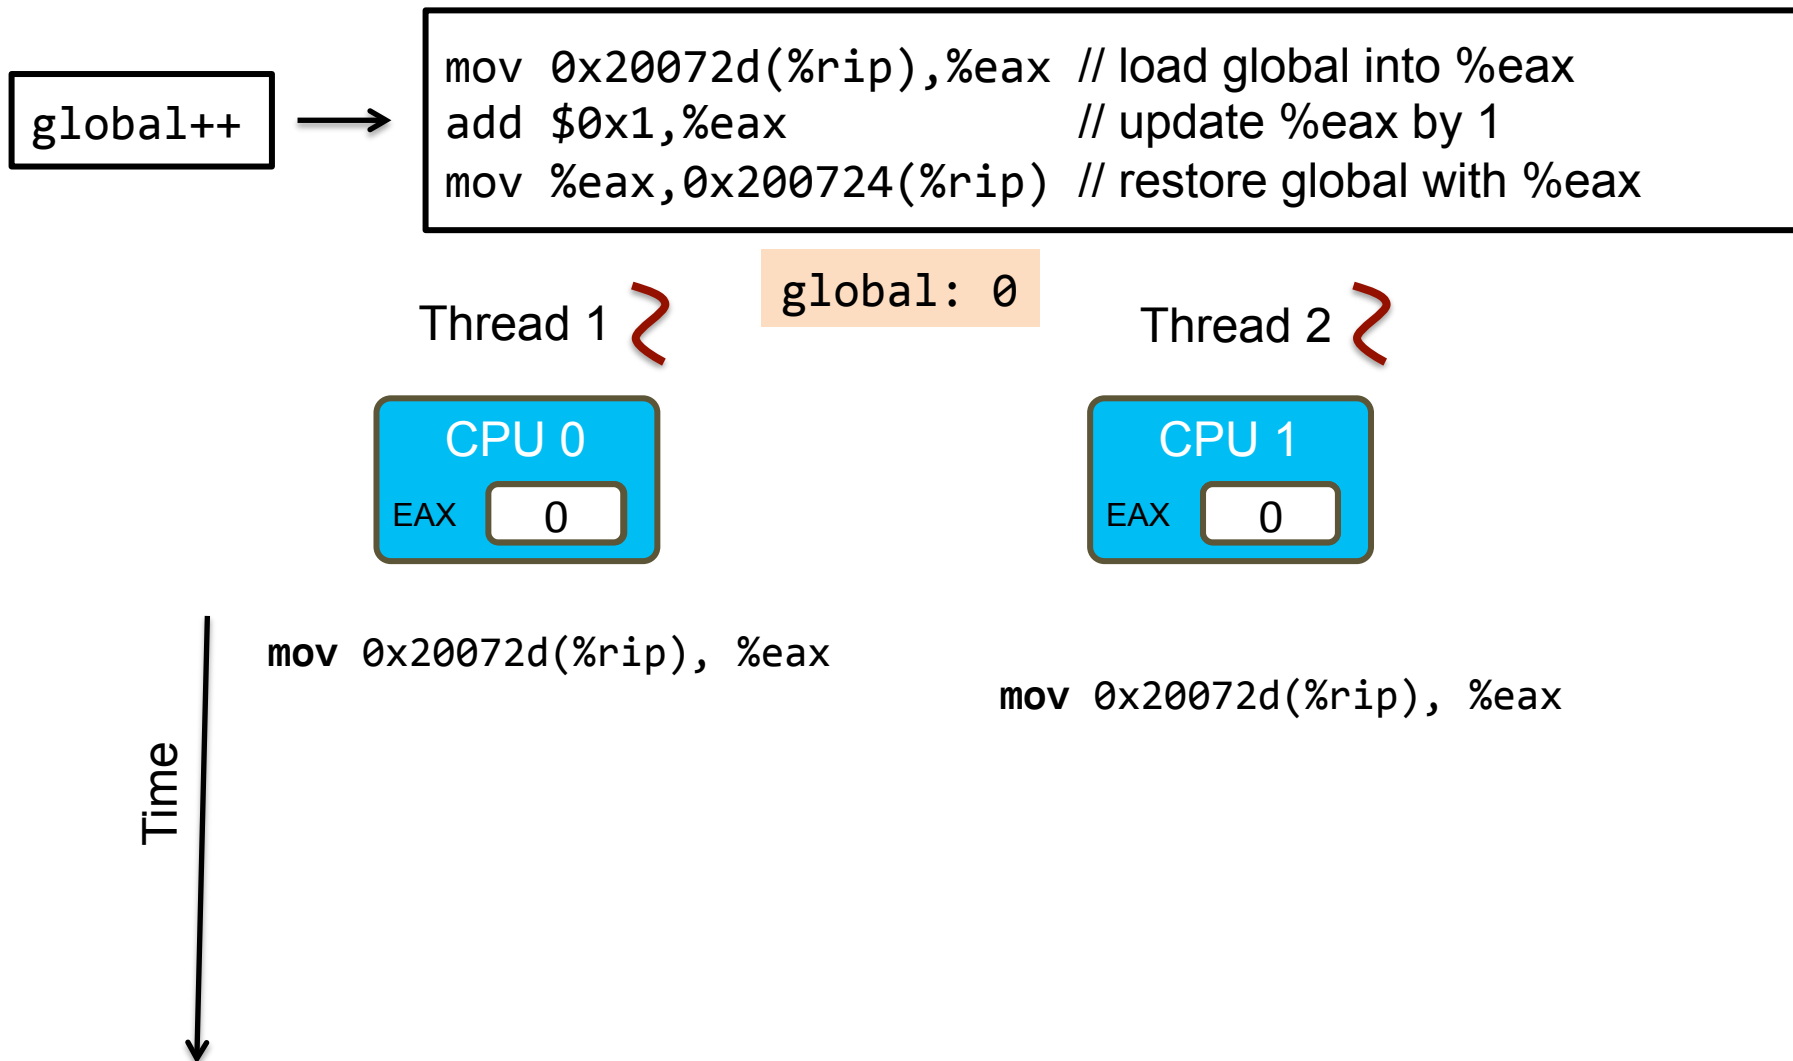
Thread 2



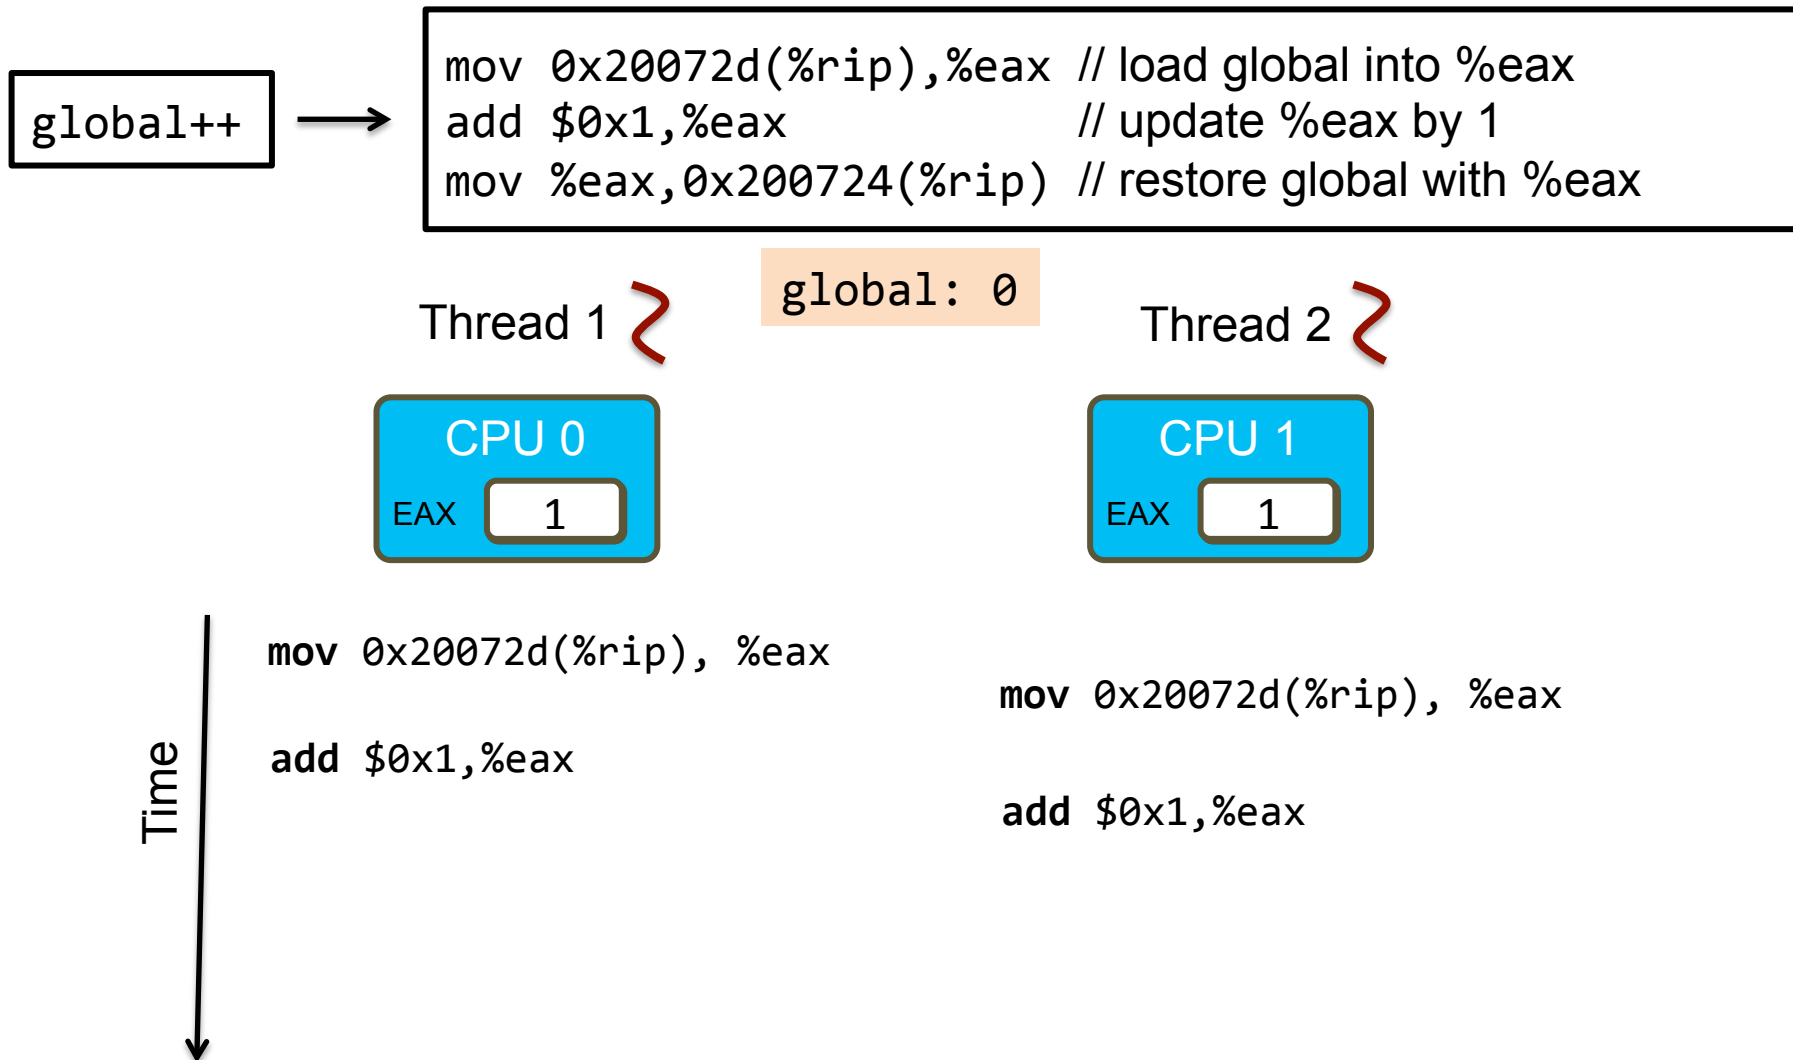
global++

What should be the value of `global`  
after both threads finish?

# What are races? An example



# What are races? An example



# What are races? An example

global++

```
mov 0x20072d(%rip),%eax // load global into %eax
add $0x1,%eax           // update %eax by 1
mov %eax,0x200724(%rip) // restore global with %eax
```

Thread 1

global: 1

Thread 2

CPU 0

EAX 1

CPU 1

EAX 1

Time

mov 0x20072d(%rip), %eax

add \$0x1,%eax

mov %eax, 0x20072d(%rip)

mov 0x20072d(%rip), %eax

add \$0x1,%eax

mov %eax, 0x20072d(%rip)

global = 1 ?!!



Worse. Sometimes global=2,  
sometimes global=1



# How to prevent race conditions?

- Mutual exclusion: only one thread enters critical section at any time.

```
void *thread_start(void *args) {  
    global++;  
} } critical section  
  
void main() {  
    for (int i = 0; i < 10; i++) {  
        pthread_create(.., thread_start, NULL);  
    }  
    pthread_join(...)  
}
```


# How to prevent race conditions?

- Use locks/mutexes to enforce mutual exclusion
- In pthread library, use `pthread_mutex_lock/unlock`

```
pthread_mutex_t mu;
```

You can also malloc a mutex

```
void *thread_start(void *args) {
```

```
 pthread_mutex_lock(&mu);  
    global++;
```

Blocks caller until mutx becomes unlocked (returns 0 on success)

```
 pthread_mutex_unlock(&mu);
```

Unlocks mutex (returns 0 on success)

```
}
```

```
void main() {
```

```
    pthread_mutex_init(&mu, NULL);
```

You must initialize a mutex before using it

```
    for (int i = 0; i < 10; i++) {  
        pthread_create(.., thread_start, NULL);
```

```
    }
```

```
    pthread_join(...)
```

```
}
```

# Lock enforces mutual exclusion

Thread 1 


```
pthread_mutex_lock(&mu);  
global++;  
pthread_mutex_unlock(&mu);
```

```
int global = 0;  
pthread_mutex_t mu;
```

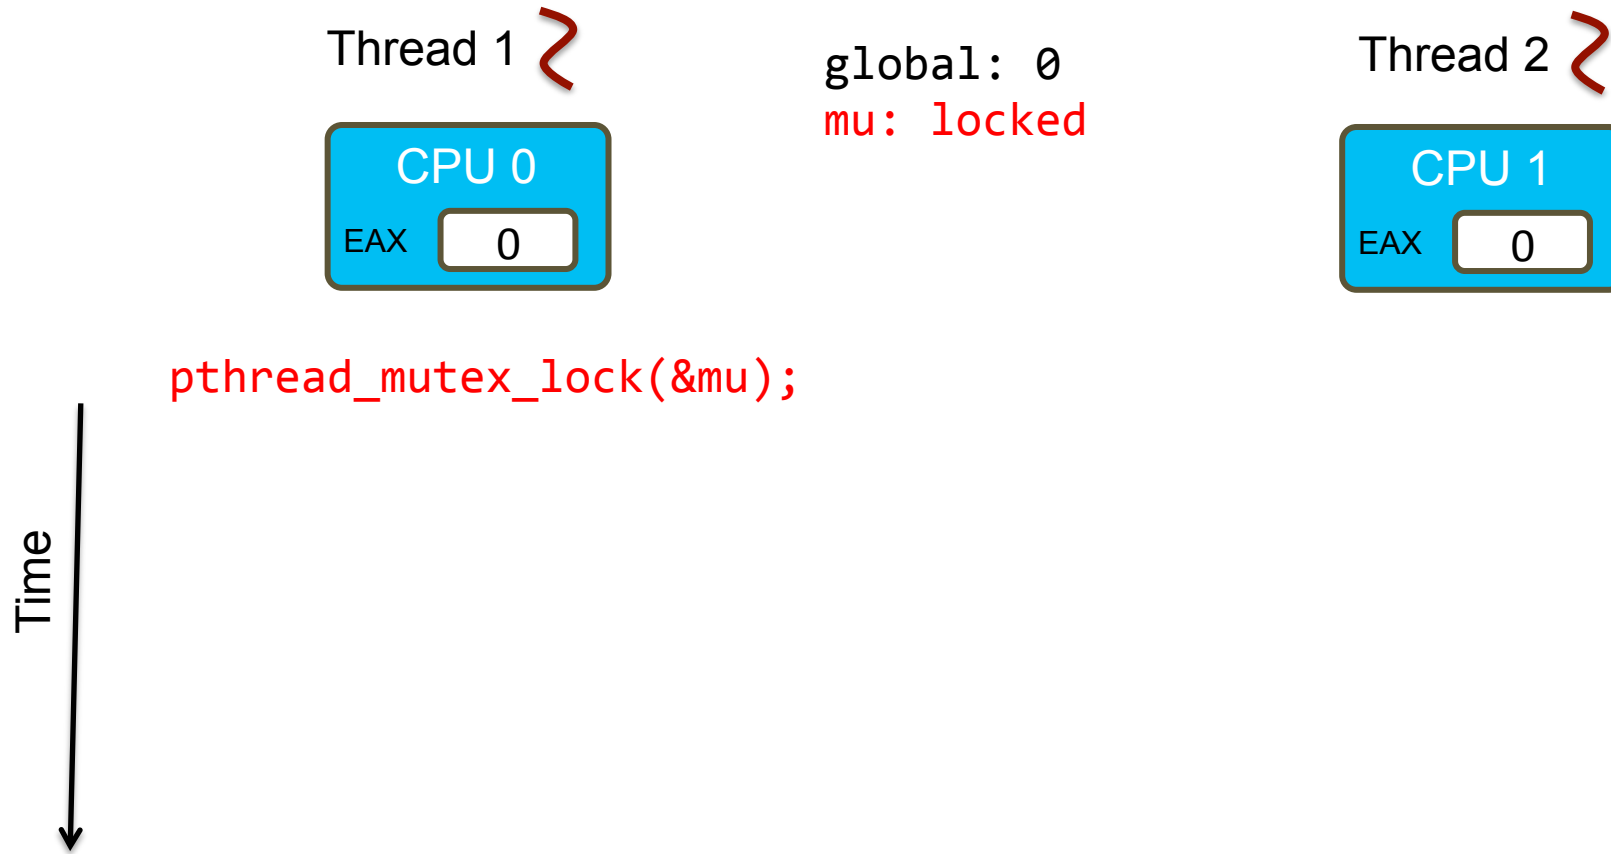
Thread 2 

```
pthread_mutex_lock(&mu);  
global++;  
pthread_mutex_unlock(&mu);
```

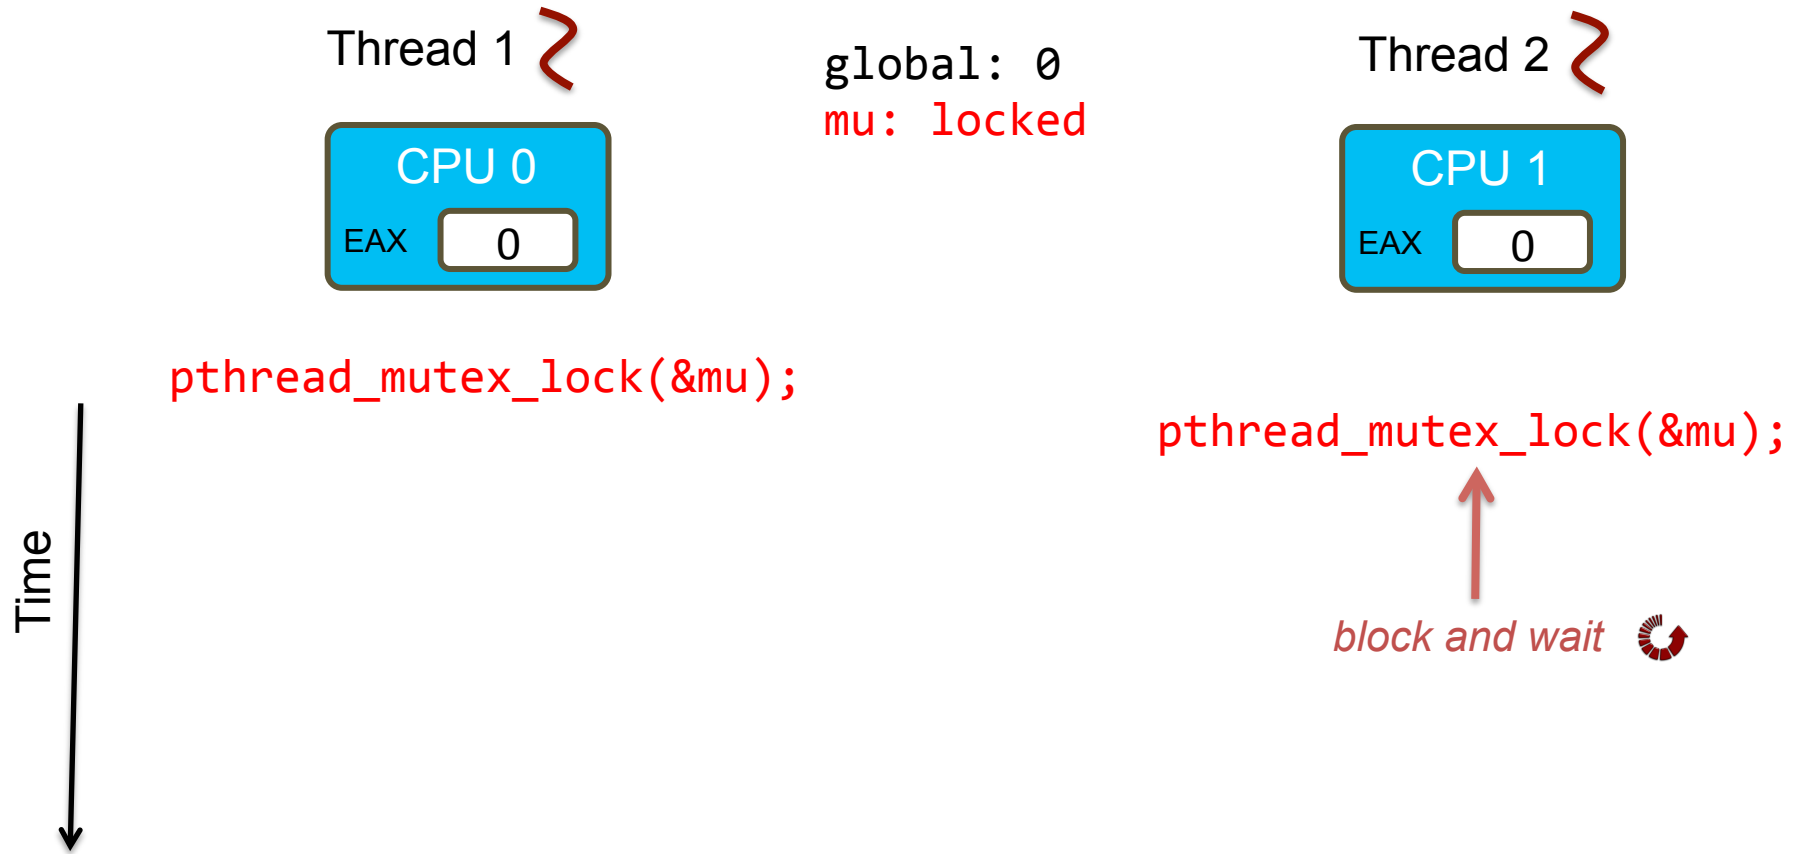
# Lock enforces mutual exclusion



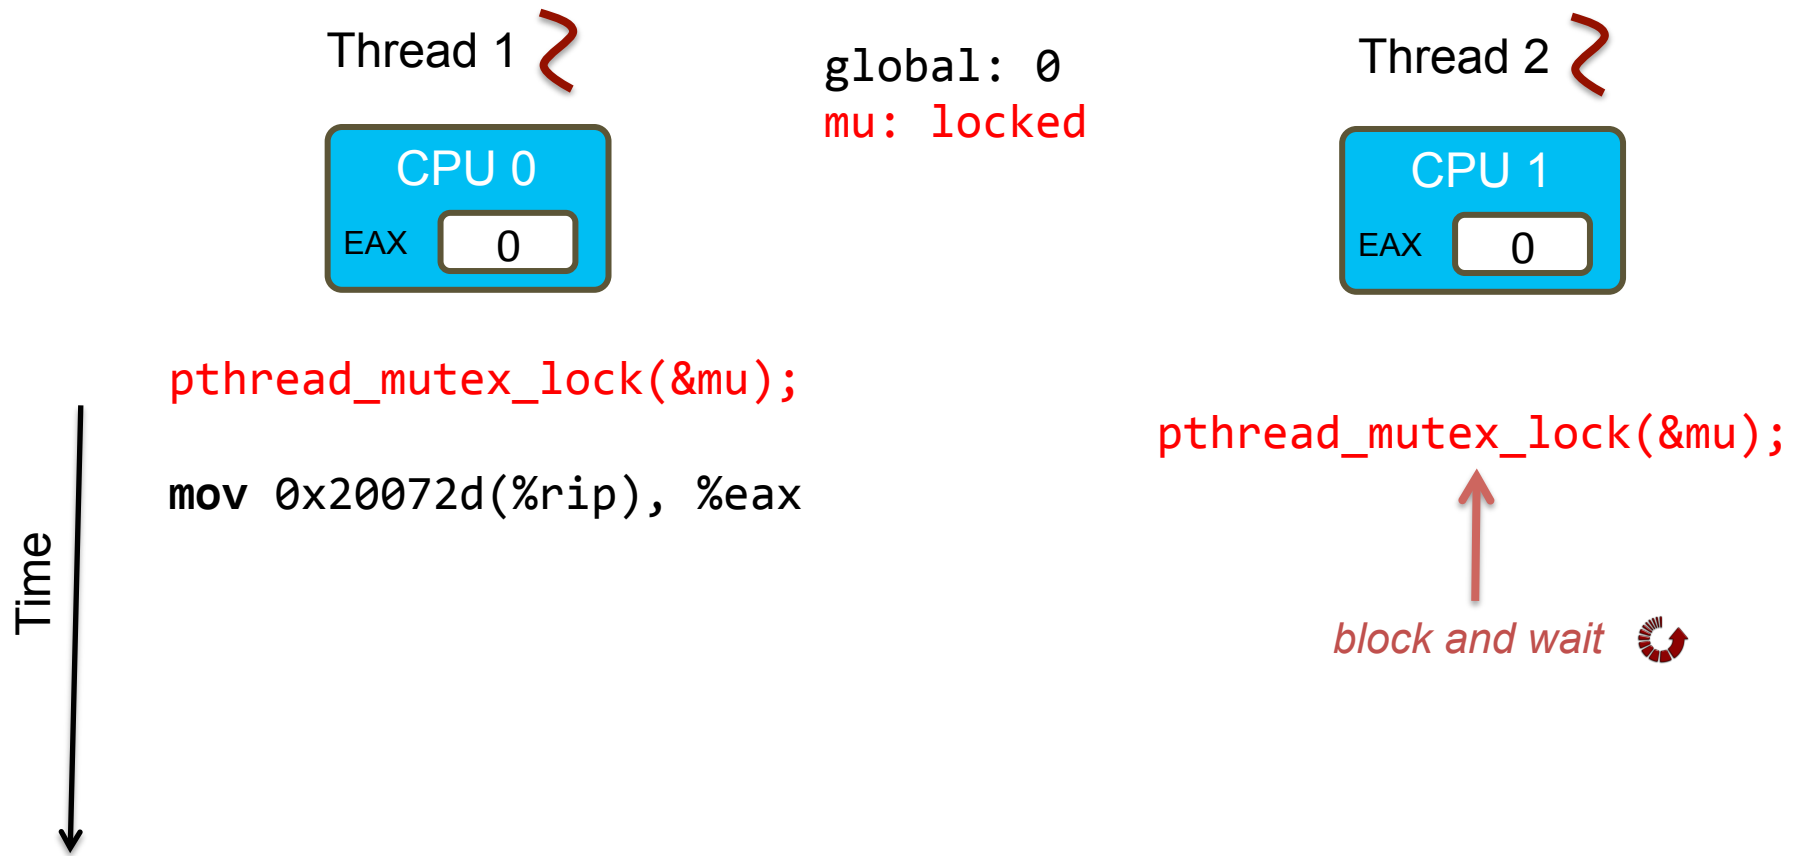
# Lock enforces mutual exclusion



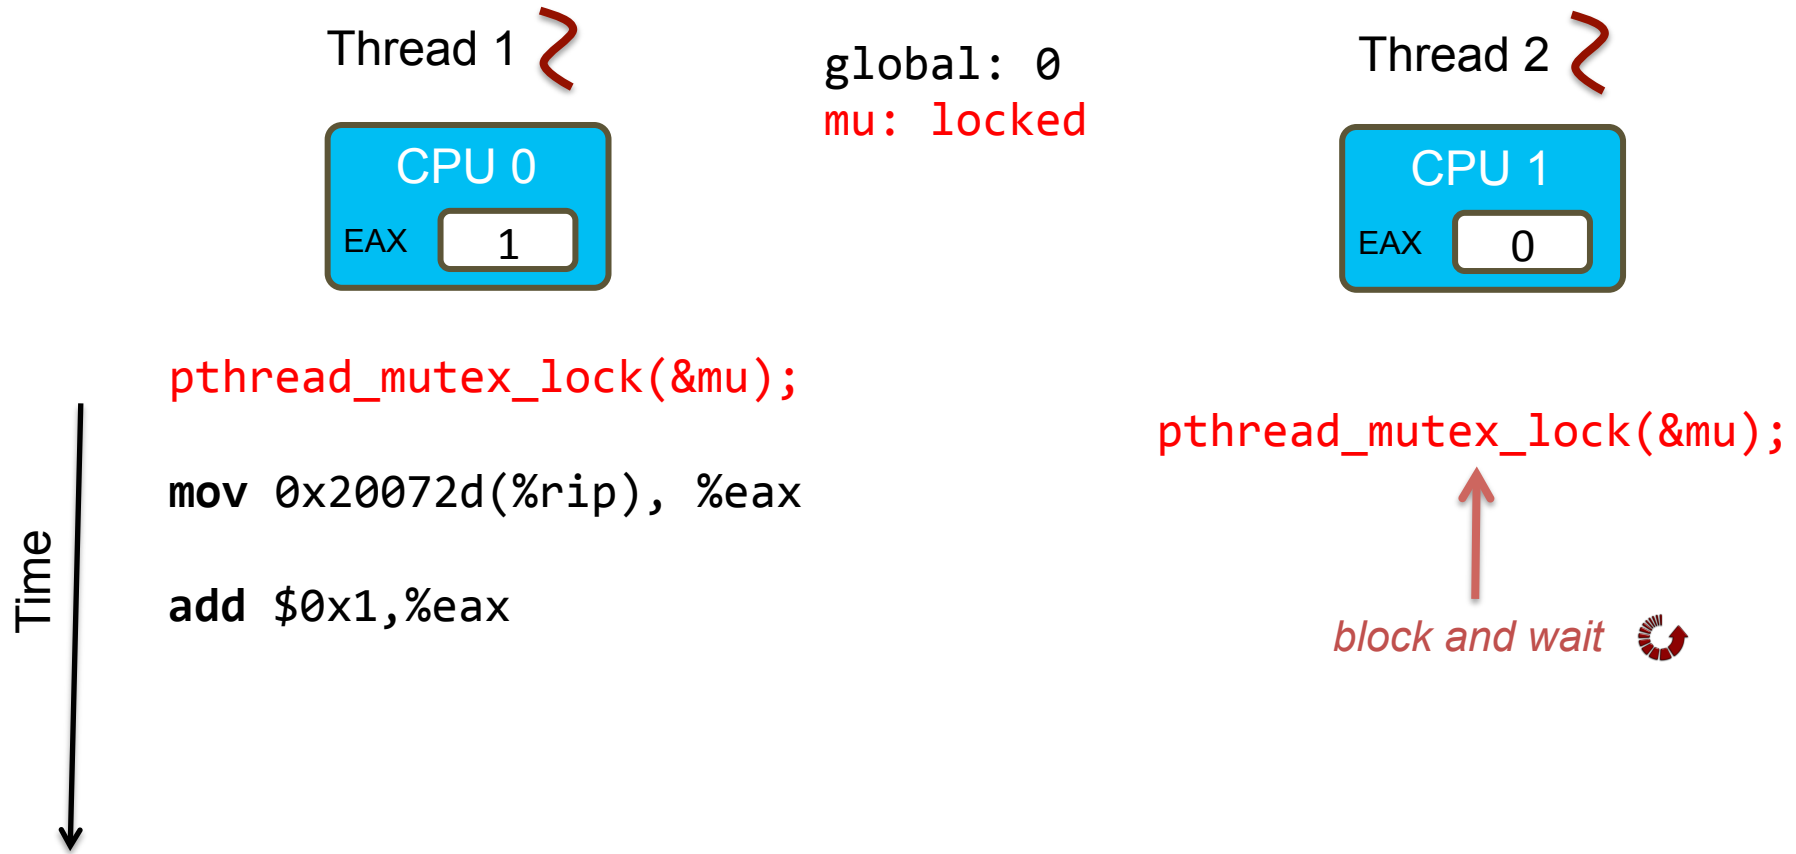
# Lock enforces mutual exclusion



# Lock enforces mutual exclusion

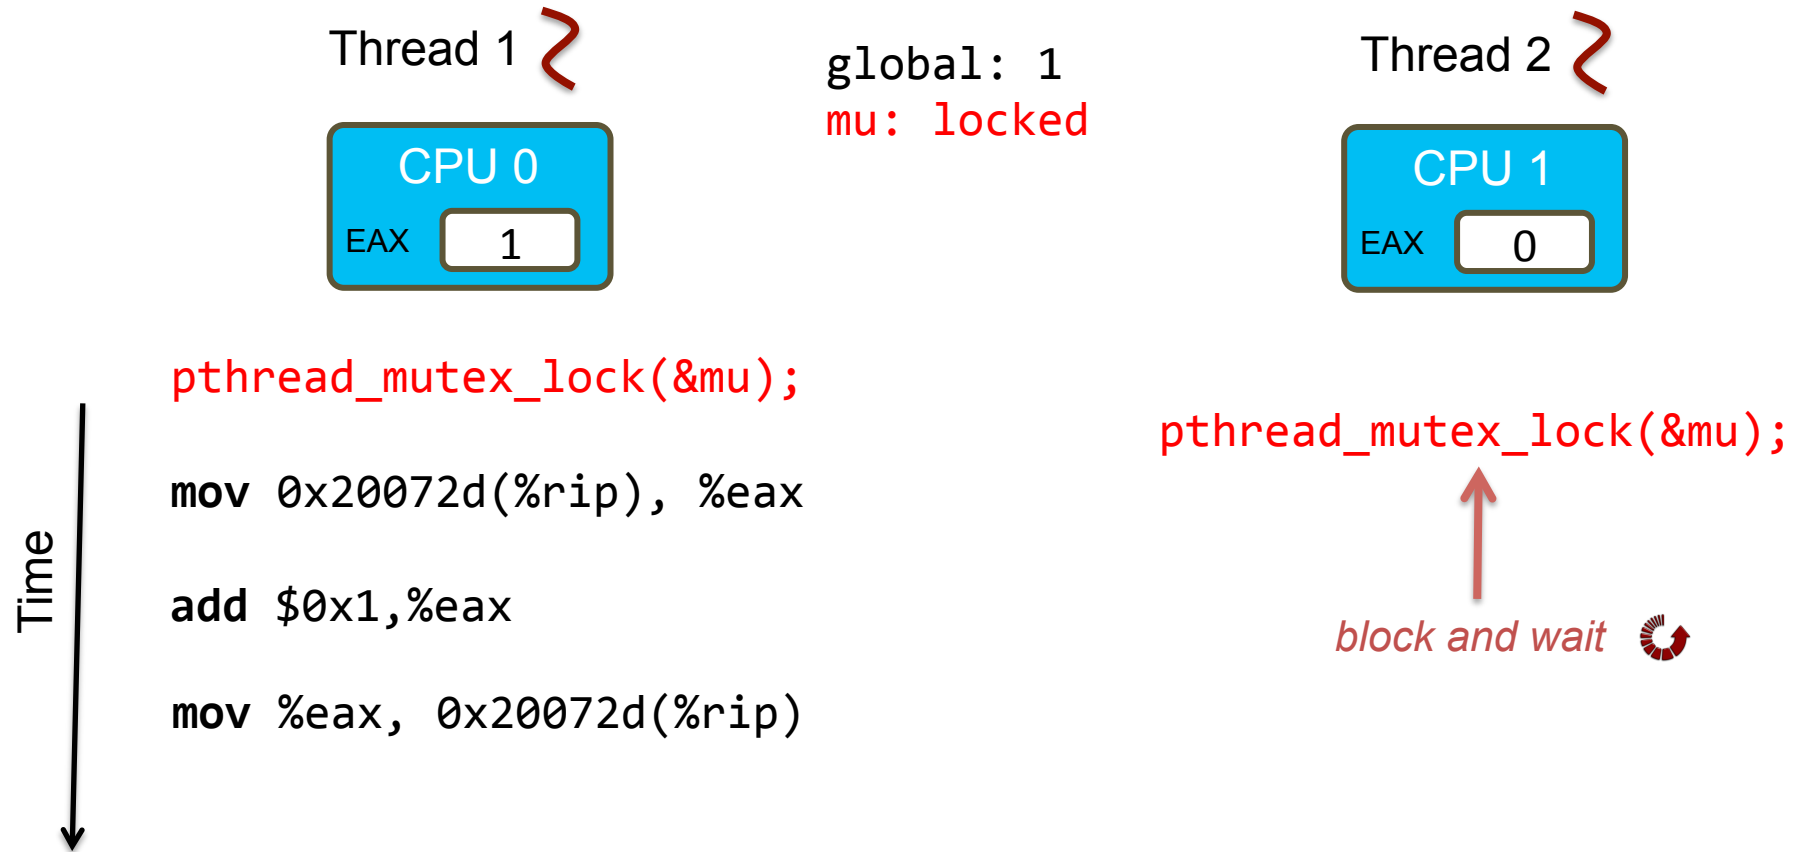


# Lock enforces mutual exclusion

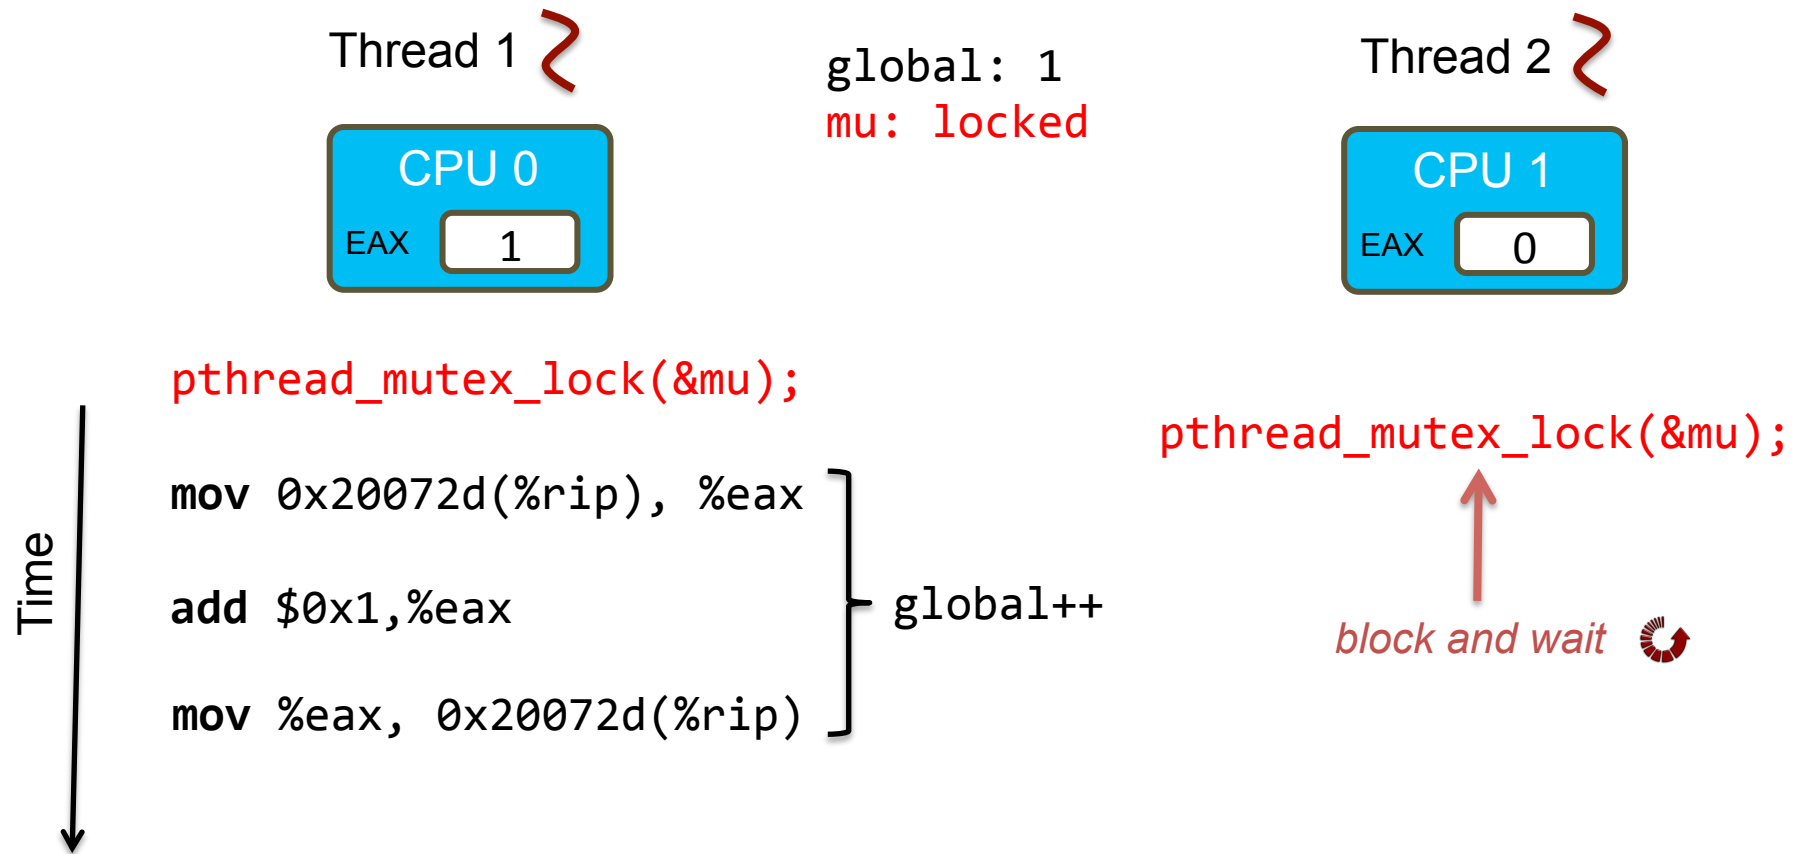




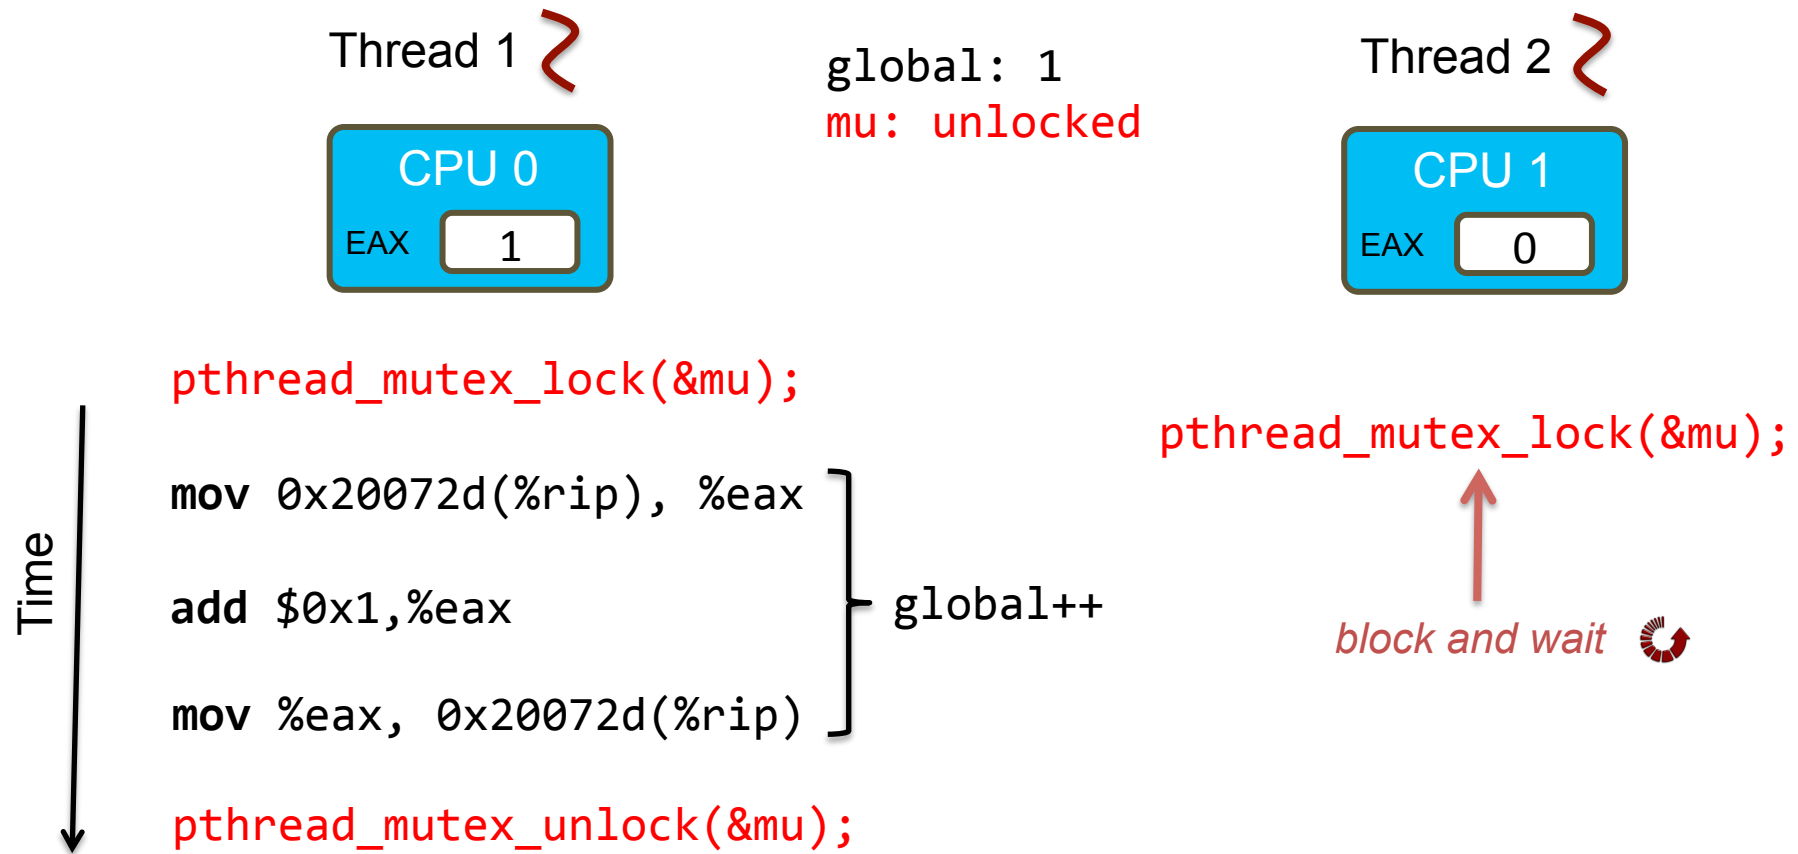
# Lock enforces mutual exclusion



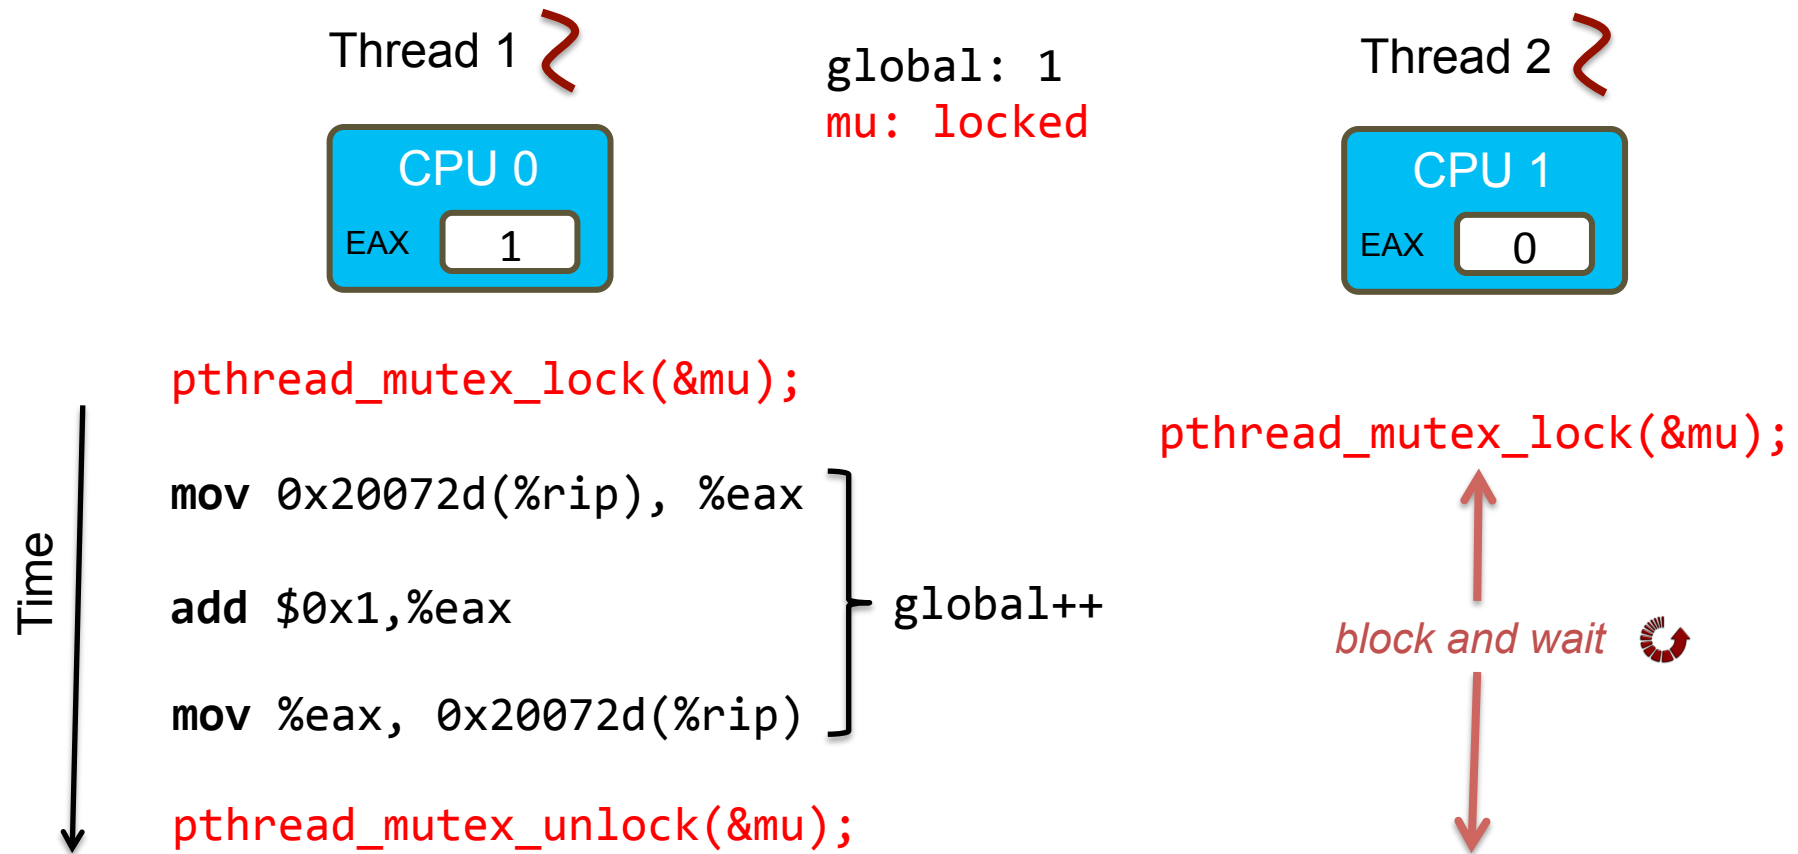
# Lock enforces mutual exclusion



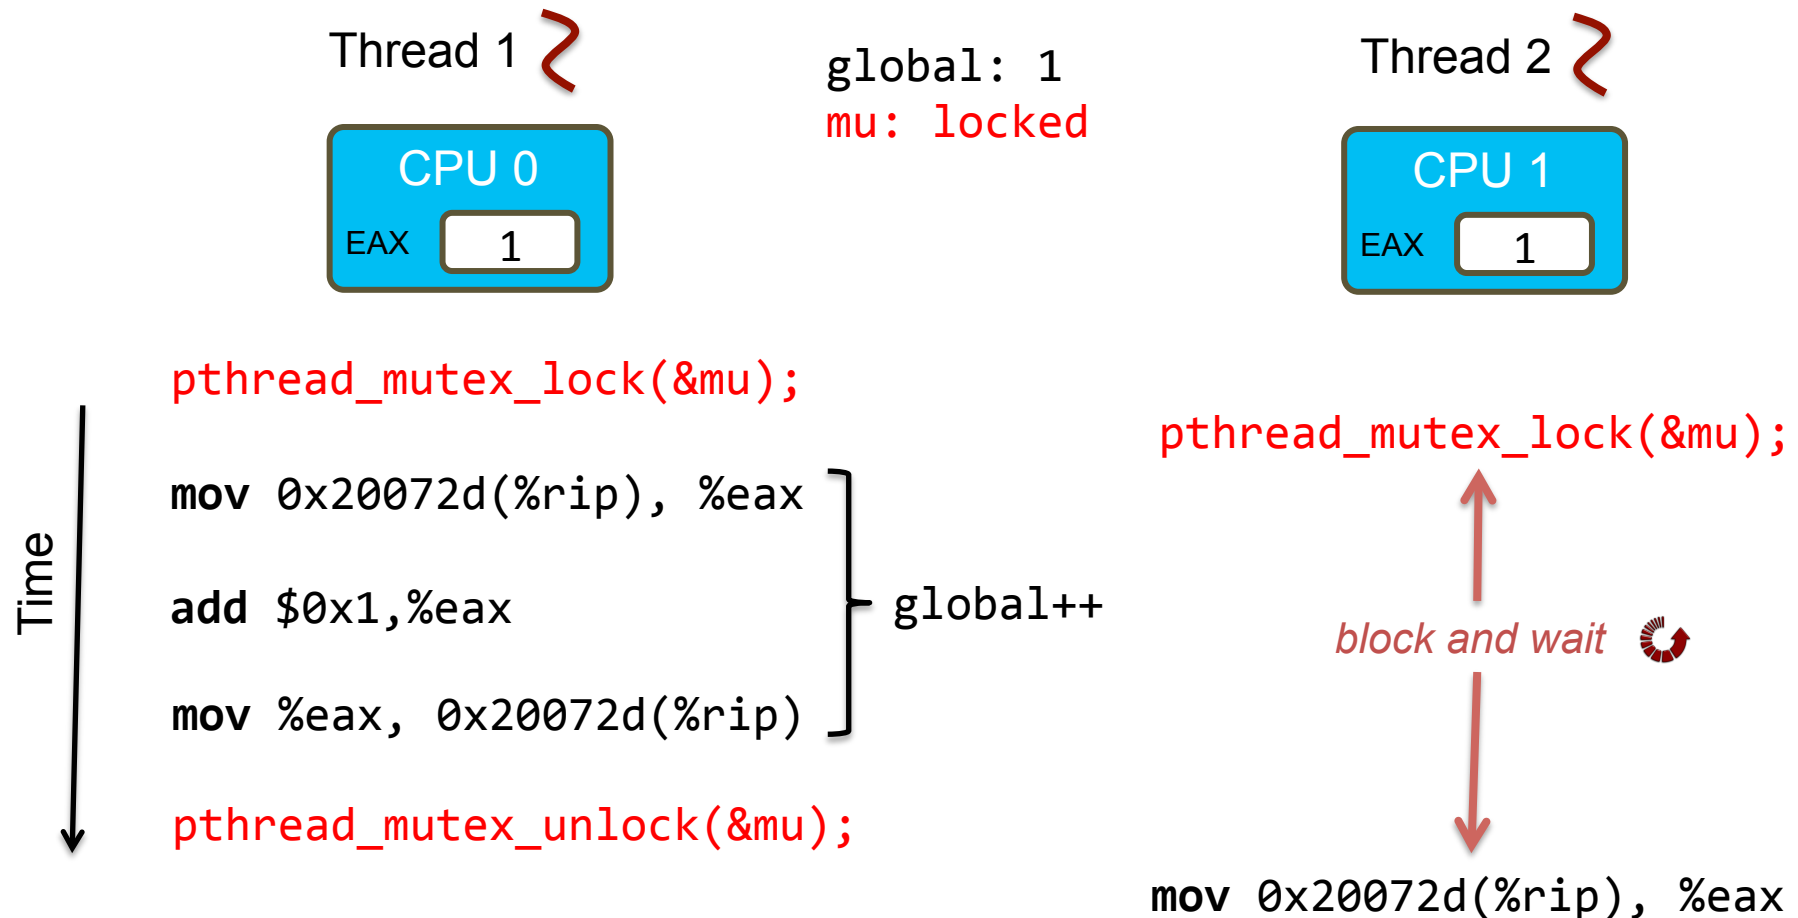
# Lock enforces mutual exclusion



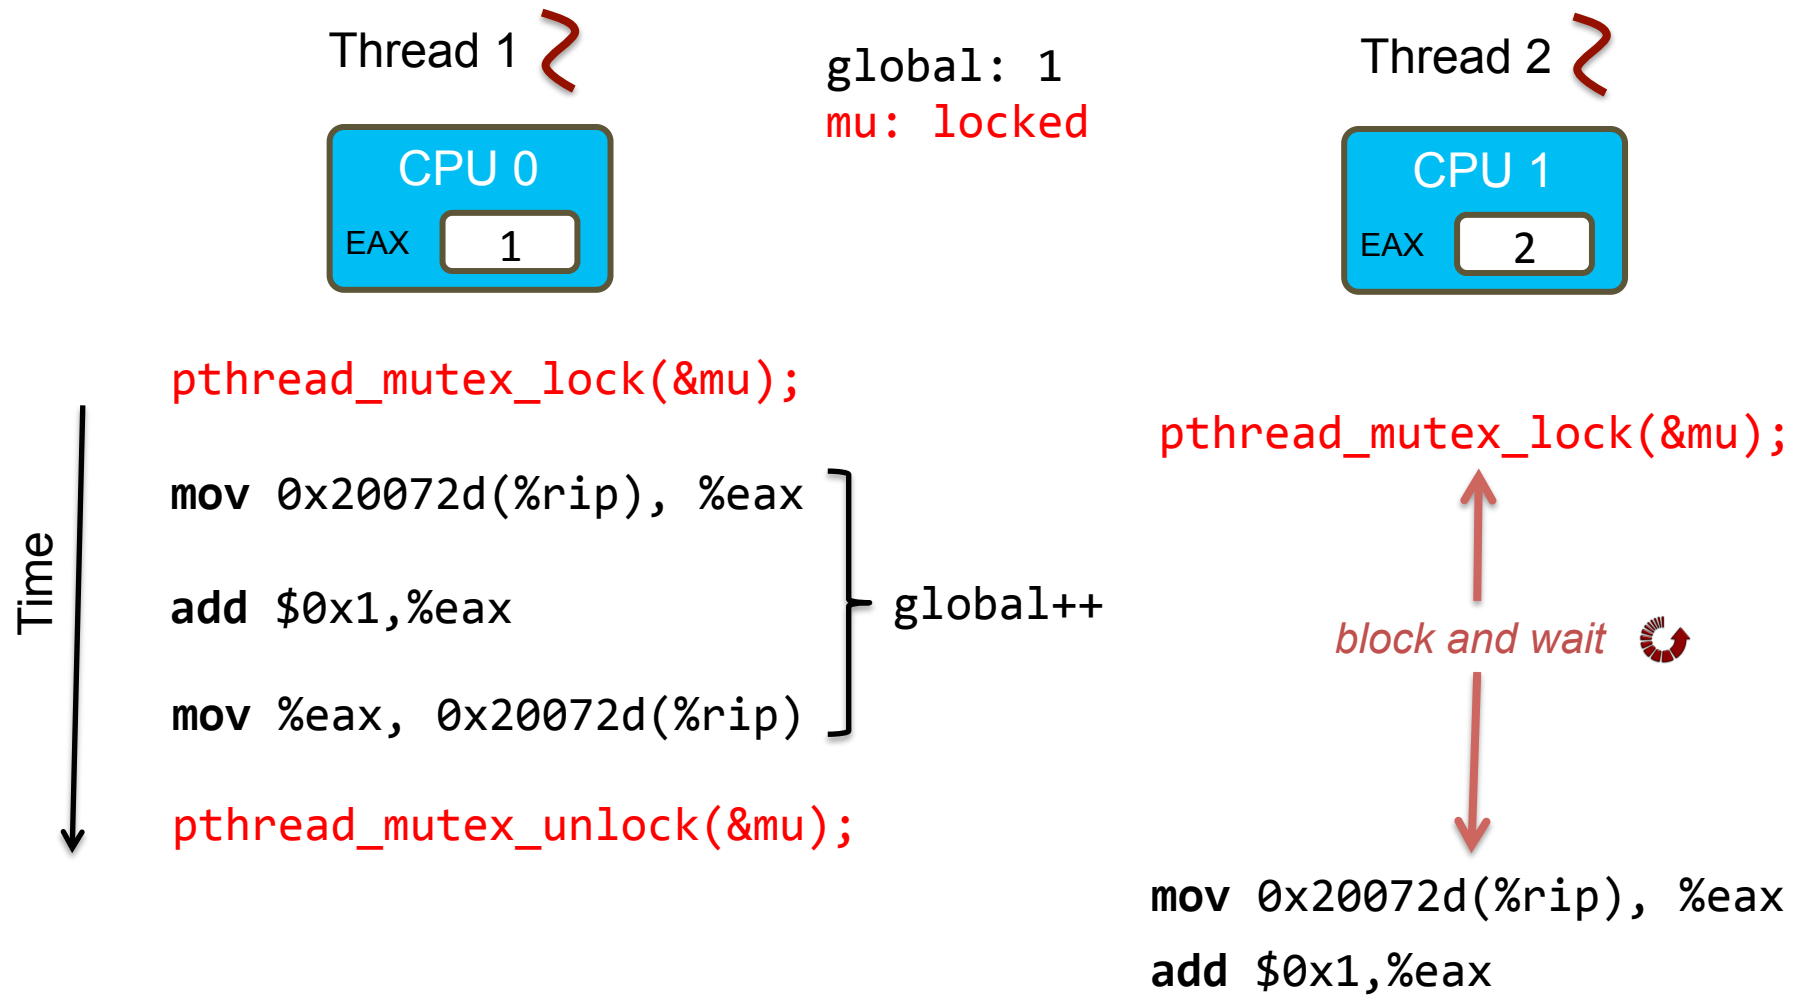
# Lock enforces mutual exclusion



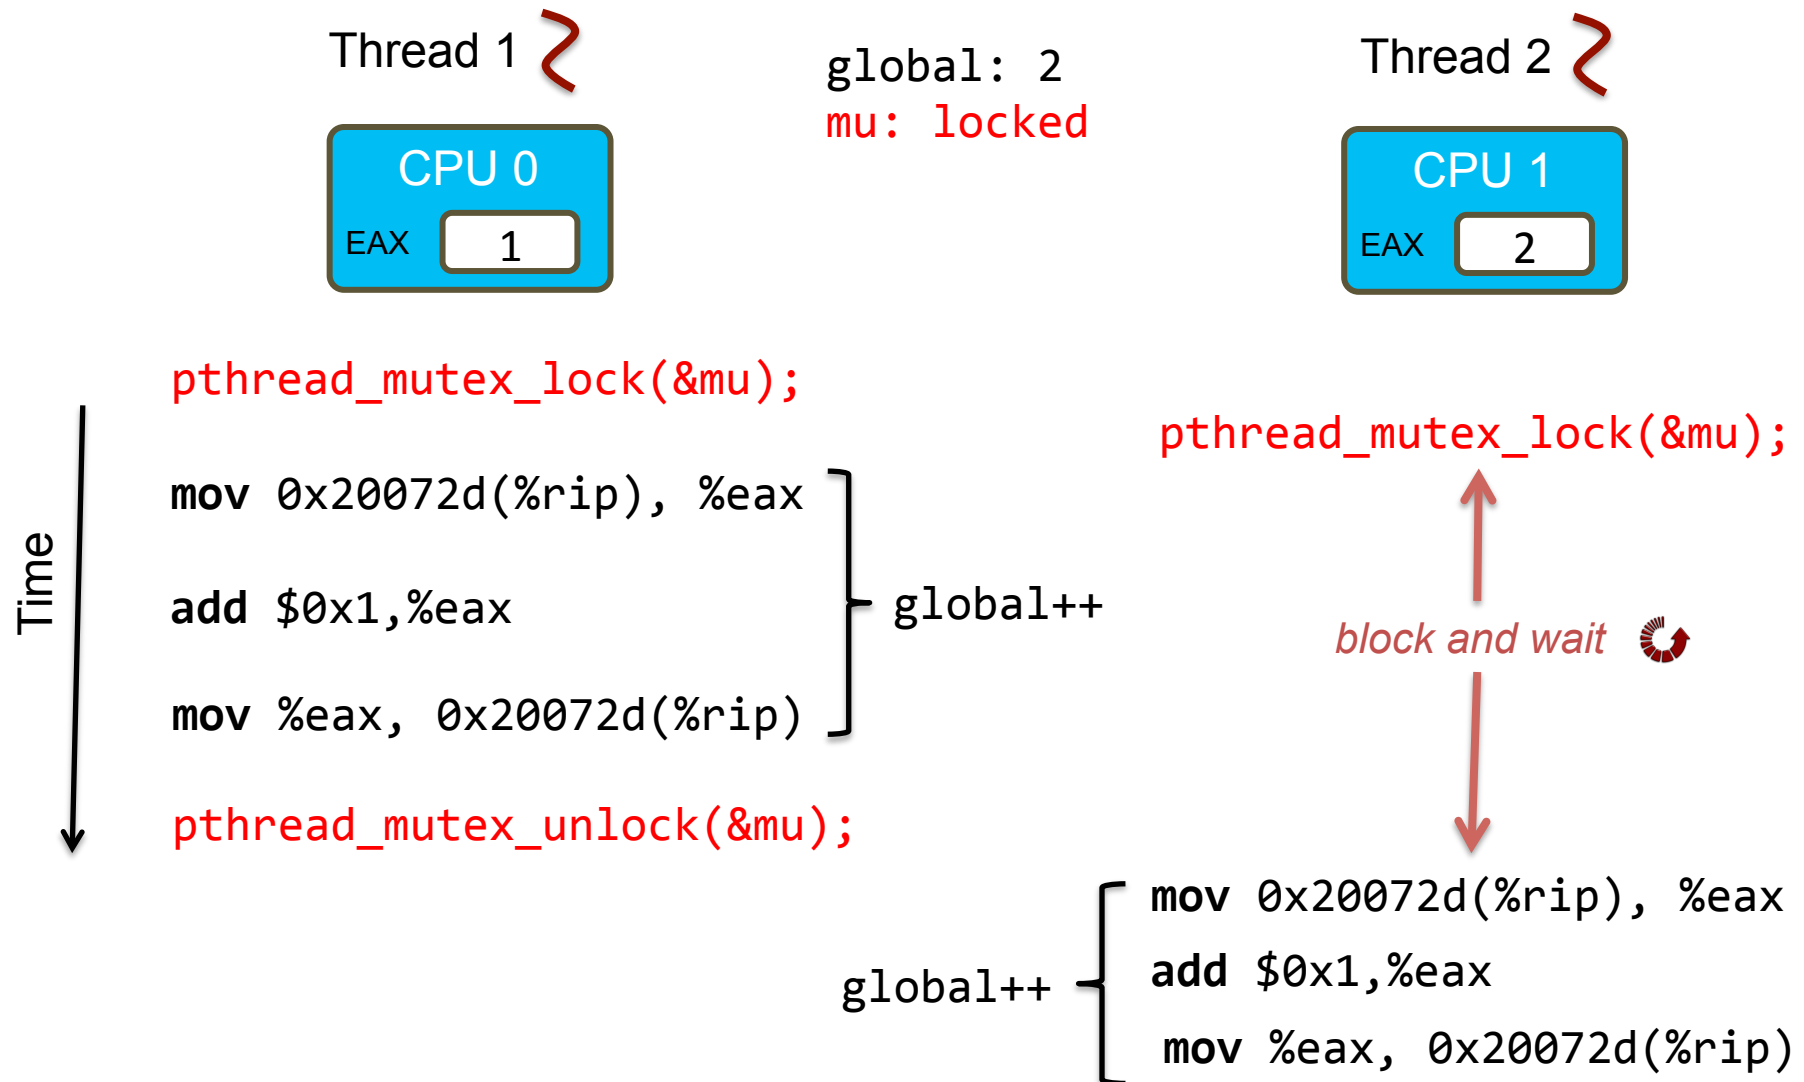
# Lock enforces mutual exclusion



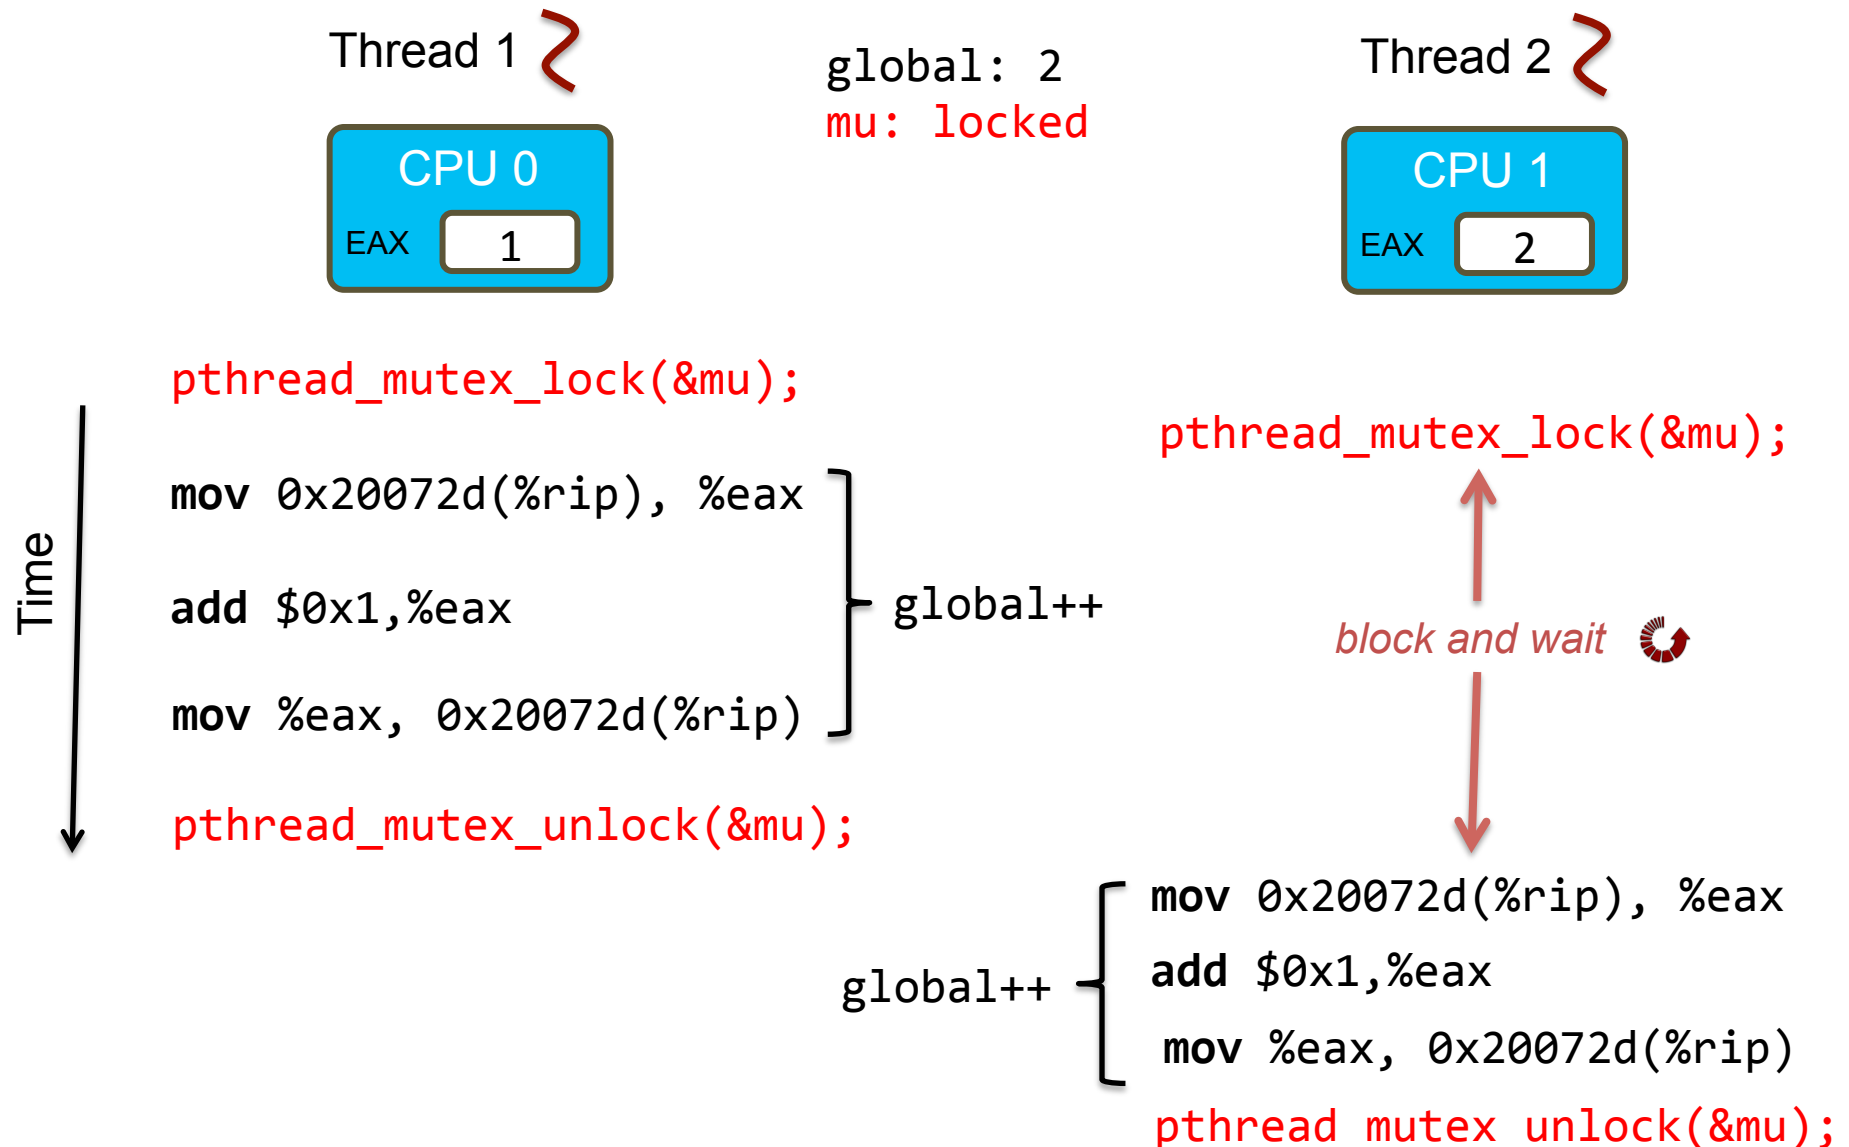
# Lock enforces mutual exclusion



# Lock enforces mutual exclusion



# Lock enforces mutual exclusion







# How to lock?

- Any vulnerable state must be locked before access
  - Vulnerable: state can be accessed by  $\geq 2$  threads, one of them is a writer.
- We mentally associate a separate lock to protect a different piece of vulnerable state

# Old example: What does mu protect?

```
pthread_mutex_t mu;  
int global = 0;  
void *thread_start(void *args) {  
 pthread_mutex_lock(&mu);  
    global++;  
 pthread_mutex_unlock(&mu);  
}
```



```
void main() {  
    for (int i = 0; i < 10; i++) {  
        pthread_create(.., thread_start, NULL);  
    }  
    pthread_join(...)  
}
```

# Lock granularity: an example

```
const int n=10;  
int accounts[total];
```

```
void transfer(int x, int y, int amount) {  
    accounts[x] -= amount;  
    accounts[y] += amount;  
}
```

```
int balance(User u) {  
    return accounts[u.checking]+  
           accounts[u.saving];  
}
```

Thread 1 	Thread 2 
%eax=accounts[1]=10	%eax=accounts[1]=10
%eax=%eax-10=0	%eax=%eax-10= 0
accounts[1]=%eax=0	accounts[1]=%eax=0
%eax=accounts[2]=10	
%eax=%eax+10=20	
accounts[2]=%eax=20	
	%eax=accounts[2]=20
	%eax=%eax+10=30
	accounts[2]=%eax=30

# A real hack



## Bitcoin Bank Flexcoin to Close After \$600k Bitcoin Theft

"The hacker discovered that if you place several withdrawals all in practically the same instant, they will get processed at more or less the same time," a user named busoni, who identified himself as the **owner** of the Poloniex exchange, [said on the BitcoinTalk forum](#).

# Lock granularity: an example

```
const int n=10;
int accounts[total];

void transfer(int x, int y, int amount) {
    accounts[x] -= amount;
    accounts[y] += amount;
}

int balance(User u) {
    return accounts[u.checking]+
           accounts[u.saving];
}
```

How to lock?

# One big lock is simple

```
const int n=10;  
int accounts[total];  
pthread_mutex_t mu;
```

What does mu protect?



```
void transfer(int x, int y, int amount) {  
    pthread_mutex_lock(&mu);  
    accounts[x] -= amount;  
    accounts[y] += amount;  
    pthread_mutex_unlock(&mu);  
}
```

```
int balance(User u) {  
    pthread_mutex_lock(&mu);  
    int bal = accounts[u.checking]+  
               accounts[u.saving];  
    pthread_mutex_unlock(&mu);  
    return bal;  
}
```

## What's the downside?

- Serializes execution.
- Only one thread can execute transfer() or balance() at a time

# Fine-grained locks have better performance, but...

```
const int n=10;
int accounts[total];
pthread_mutex_t mu[total];

void transfer(int x, int y, int amount) {
    pthread_mutex_lock(&mu[x]);
    accounts[x] -= amount;
    pthread_mutex_unlock(&mu[x]);
    pthread_mutex_lock(&mu[y]);
    accounts[y] += amount;
    pthread_mutex_unlock(&mu[y]);
}

int balance(User u) {
    pthread_mutex_lock(&mu[u.checking]);
    int bal = accounts[u.checking];
    pthread_mutex_unlock(&mu[u.checking]);
    pthread_mutex_lock(&mu[u.saving]);
    bal += accounts[u.saving];
    pthread_mutex_unlock(&mu[u.saving]);
    return bal;
}
```


Easy to lock incorrectly

# Fine-grained locks have better performance, but...

```
const int n=10;
int accounts[total];
pthread_mutex_t mu[total];

void transfer(int x, int y, int amount) {
    pthread_mutex_lock(&mu[x]);
    accounts[x] -= amount;
    pthread_mutex_unlock(&mu[x]);
    pthread_mutex_lock(&mu[y]);
    accounts[y] += amount;
    pthread_mutex_unlock(&mu[y]);
}

int balance(User u) {
    pthread_mutex_lock(&mu[u.checking]);
    int bal = accounts[u.checking];
    pthread_mutex_unlock(&mu[u.checking]);
    pthread_mutex_lock(&mu[u.saving]);
    bal += accounts[u.saving];
    pthread_mutex_unlock(&mu[u.saving]);
    return bal;
}
```

Thread 1 



accounts[1] -= 10

Thread 2 



bal = accounts[1] = 0



bal += accounts[2] = 10



accounts[2] += 10



bal=10  
missing \$10!!



# Fixing Fine-grained locks

```
int n=10;
int accounts[total];
pthread_mutex_t mu[total];

void transfer(int x, int y, int amount) {
    pthread_mutex_lock(&mu[x]);
    accounts[x] -= amount;
    pthread_mutex_lock(&mu[y]);
    accounts[y] += amount;
    pthread_mutex_unlock(&mu[x]);
    pthread_mutex_unlock(&mu[y]);
}

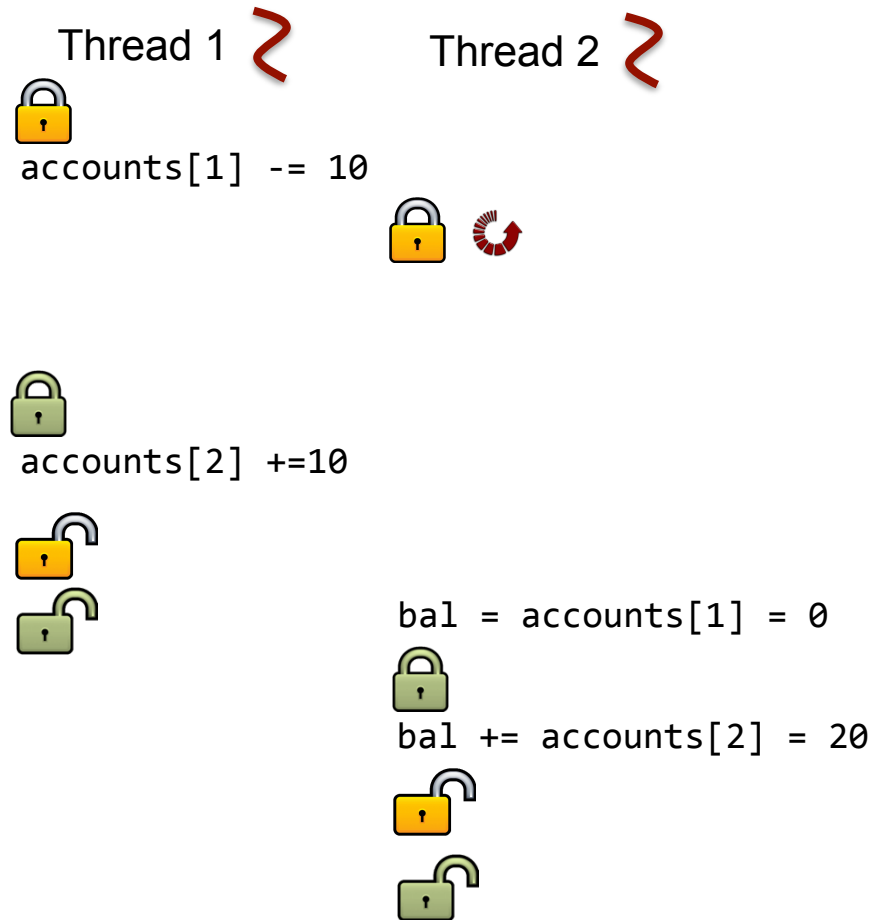
int balance(User u) {
    pthread_mutex_lock(&mu[u.checking]);
    int bal = accounts[u.checking];
    pthread_mutex_lock(&mu[u.saving]);
    bal += accounts[u.saving];
    pthread_mutex_unlock(&mu[u.checking]);
    pthread_mutex_unlock(&mu[u.saving]);
    return bal;
}
```

# Fixing Fine-grained locks

```
const int n=10;
int accounts[total];
pthread_mutex_t mu[total];

void transfer(int x, int y, int amount) {
    pthread_mutex_lock(&mu[x]);
    accounts[x] -= amount;
    pthread_mutex_lock(&mu[y]);
    accounts[y] += amount;
    pthread_mutex_unlock(&mu[x]);
    pthread_mutex_unlock(&mu[y]);
}

int balance(User u) {
    pthread_mutex_lock(&mu[u.checking]);
    int bal = accounts[u.checking];
    pthread_mutex_lock(&mu[u.saving]);
    bal += accounts[u.saving];
    pthread_mutex_unlock(&mu[u.checking]);
    pthread_mutex_unlock(&mu[u.saving]);
    return bal;
}
```

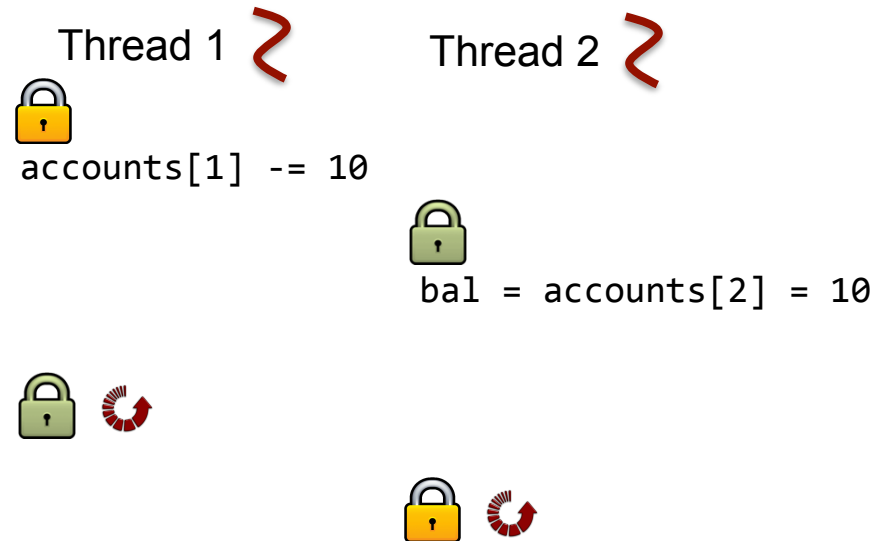


# Fine-grained locks may have deadlocks

```
const int n=10;
int accounts[total];
pthread_mutex_t mu[total];

void transfer(int x, int y, int amount) {
    pthread_mutex_lock(&mu[x]);
    accounts[x] -= amount;
    pthread_mutex_lock(&mu[y]);
    accounts[y] += amount;
    pthread_mutex_unlock(&mu[x]);
    pthread_mutex_unlock(&mu[y]);
}

int balance(User u) {
    pthread_mutex_lock(&mu[u.checking]);
    int bal = accounts[u.checking];
    pthread_mutex_lock(&mu[u.saving]);
    bal += accounts[u.saving];
    pthread_mutex_unlock(&mu[u.checking]);
    pthread_mutex_unlock(&mu[u.saving]);
    return bal;
}
```



# Deadlocks sound scary, but they are not

- Stack traces indicate why deadlocks happen
  - no data corruption

(gdb) thread apply all bt

Thread 3 (Thread 0x7fd57746e700 (LWP 9450)):

```
#0 __lll_lock_wait () at ../sysdeps/unix/sysv/linux/x86_64/lowlevellock.S:135
#1 0x00007fd57806b023 in __GI___pthread_mutex_lock (mutex=0x5627aad9f0a8 <mu+40>) at ../nptl/pthread_mutex_lock.c:78
#2 0x00005627aab9dace in balance (u=...) at thread.c:38
#3 0x00005627aab9db5f in th_balance (args=0x0) at thread.c:53
#4 0x00007fd5780686db in start_thread (arg=0x7fd57746e700) at pthread_create.c:463
#5 0x00007fd577d9188f in clone () at ../sysdeps/unix/sysv/linux/x86_64/clone.S:95
```

Thread 2 (Thread 0x7fd577c6f700 (LWP 9449)):

```
#0 __lll_lock_wait () at ../sysdeps/unix/sysv/linux/x86_64/lowlevellock.S:135
#1 0x00007fd57806b023 in __GI___pthread_mutex_lock (mutex=0x5627aad9f0d0 <mu+80>) at ../nptl/pthread_mutex_lock.c:78
#2 0x00005627aab9d9aa in transfer (x=1, y=2, amt=10) at thread.c:20
#3 0x00005627aab9da4d in th_transfer (args=0x0) at thread.c:29
#4 0x00007fd5780686db in start_thread (arg=0x7fd577c6f700) at pthread_create.c:463
#5 0x00007fd577d9188f in clone () at ../sysdeps/unix/sysv/linux/x86_64/clone.S:95
```

Thread 1 (Thread 0x7fd57847e740 (LWP 9448)):

```
#0 0x00007fd578069d2d in __GI___pthread_timedjoin_ex (threadid=140554814289664, thread_return=0x0, abstime=0x0, block=<optimized out>)
#1 0x00005627aab9dc7b in main (argc=1, argv=0x7ffdda8e2f38) at thread.c:67
```

(gdb) █

# Deadlocks sound scary, but they are not

- Stack traces indicate why deadlocks happen
  - no data corruption

```
(gdb) thread 3
[Switching to thread 3 (Thread 0x7fd57746e700 (LWP 9450))]
#0  __lll_lock_wait () at ../sysdeps/unix/sysv/linux/x86_64/lowlevellock.S:135
135      in ../sysdeps/unix/sysv/linux/x86_64/lowlevellock.S
(gdb) up
#1  0x00007fd57806b023 in __GI___pthread_mutex_lock (mutex=0x5627aad9f0a8 <mu+40>) at ../nptl/pthread_mutex_lock.c:78
78      ../nptl/pthread_mutex_lock.c: No such file or directory.
(gdb)
#2  0x00005627aab9dace in balance (u=...) at thread.c:38
38      pthread_mutex_lock(&mu[u.saving]);
(gdb) p u
$4 = {checking = 2, saving = 1}
(gdb) l
33      int
34      balance(User u)
35      {
36          pthread_mutex_lock(&mu[u.checking]);
37          int bal = accounts[u.checking];
38          pthread_mutex_lock(&mu[u.saving]);
39          bal += accounts[u.saving];
40          pthread_mutex_unlock(&mu[u.checking]);
41          pthread_mutex_unlock(&mu[u.saving]);
42          return bal;
(gdb)
```

# Deadlocks sound scary, but they are not


- Stack traces indicate why deadlocks happen
  - no data corruption

```
(gdb) thread 2
[Switching to thread 2 (Thread 0x7fd577c6f700 (LWP 9449))]
#0  __lll_lock_wait () at ../sysdeps/unix/sysv/linux/x86_64/lowlevellock.S:135
135  ../sysdeps/unix/sysv/linux/x86_64/lowlevellock.S: No such file or directory.
(gdb) up
#1  0x00007fd57806b023 in __GI___pthread_mutex_lock (mutex=0x5627aad9f0d0 <mu+80>) at ../nptl/pthread_mutex_lock.c:78
78  ../nptl/pthread_mutex_lock.c: No such file or directory.
(gdb)
#2  0x00005627aab9d9aa in transfer (x=1, y=2, amt=10) at thread.c:20
20  pthread_mutex_lock(&mu[y]);
(gdb) l
15  void
16  transfer(int x, int y, int amt)
17  {
18      pthread_mutex_lock(&mu[x]);
19      accounts[x] -= amt;
20      pthread_mutex_lock(&mu[y]);
21      accounts[y] += amt;
22      pthread_mutex_unlock(&mu[x]);
23      pthread_mutex_unlock(&mu[y]);
24  }
(gdb)
```

# Use lock ordering to prevent deadlock

## Observation:

- Deadlock occurs only if concurrent threads try to acquire locks in different order



## Technique:

- Each thread acquires lock in the same order

# Fine-grained locking with no deadlocks

```
const int n=10;  
int accounts[total];  
pthread_mutex_t mu[total];
```

```
void lock_in_order(int x, int y) {  
    if (x > y)  
        swap(&x, &y);  
    pthread_mutex_lock(&mu[x]);  
    pthread_mutex_lock(&mu[y]);  
}
```

```
void transfer(int x, int y, int amount) {  
    lock_in_order(x,y);  
    accounts[x] -= amount;  
    accounts[y] += amount;  
    pthread_mutex_unlock(&mu[x]);  
    pthread_mutex_unlock(&mu[y]);  
}
```

```
int balance(User u) {  
    lock_in_order(u.checking, u.saving);  
    int bal = accounts[u.checking] + accounts[u.saving];  
    pthread_mutex_unlock(&mu[u.checking]);  
    pthread_mutex_unlock(&mu[u.saving]);  
    return bal;  
}
```