

Pipelined CPU

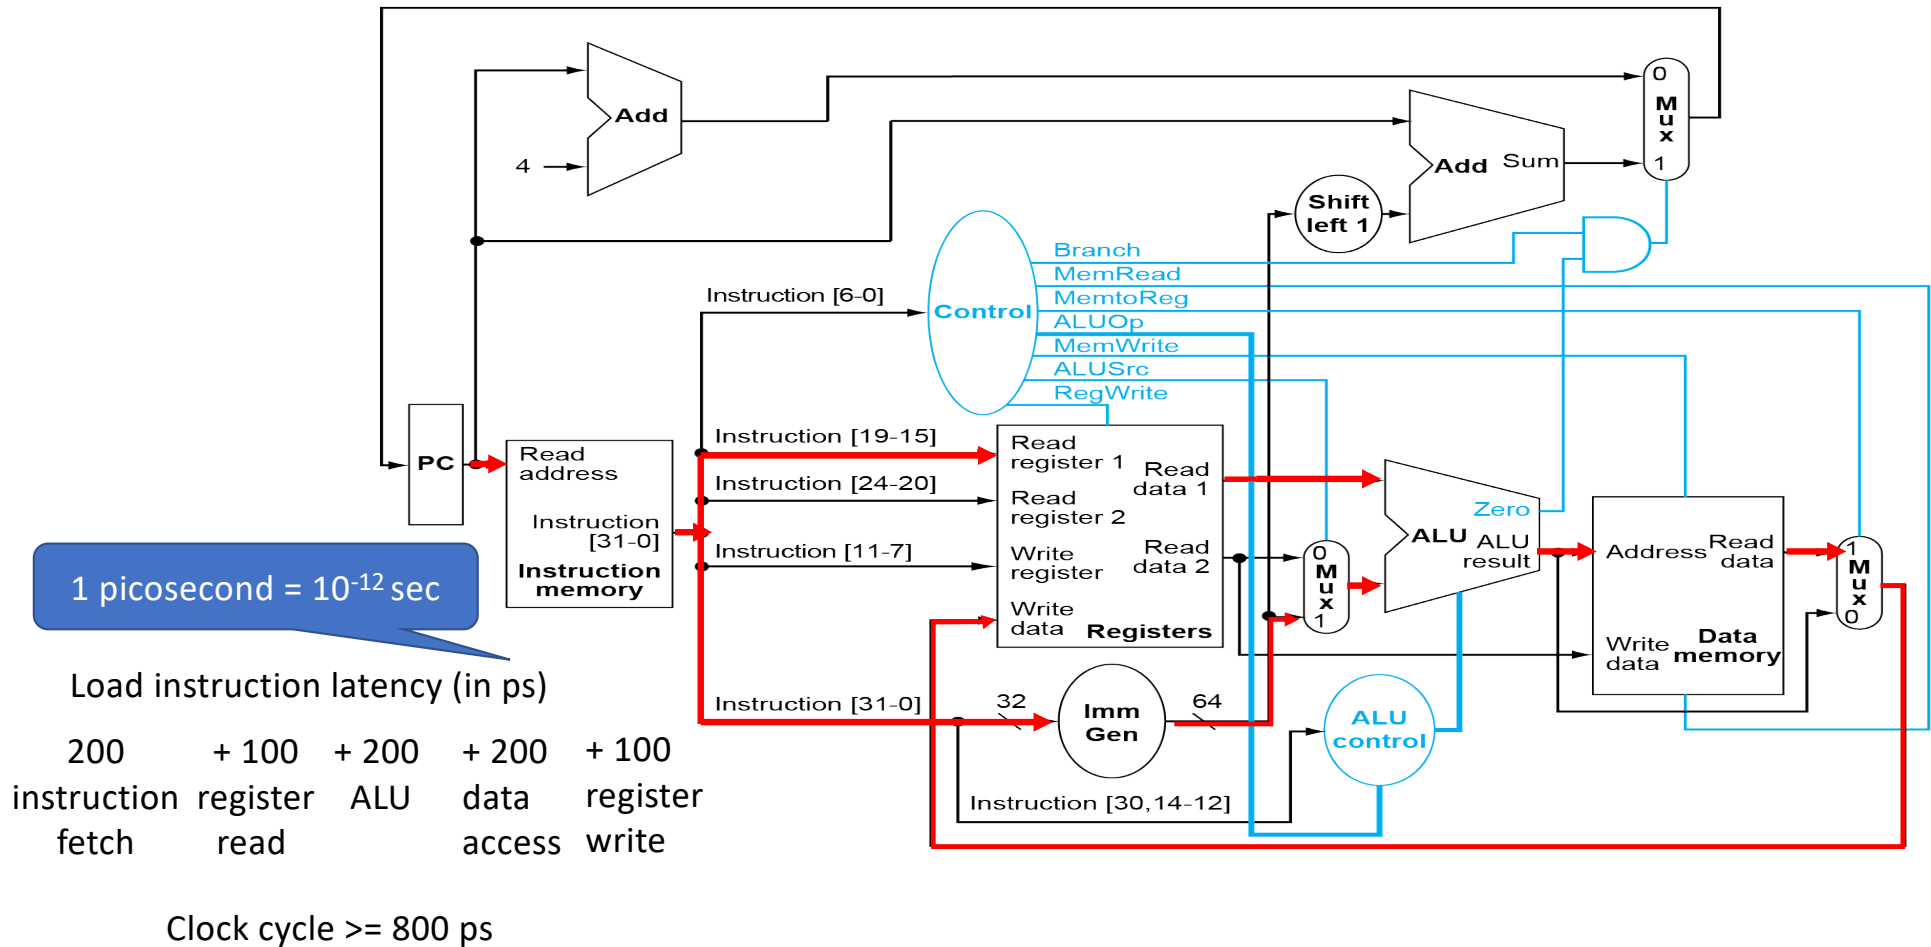
Jinyang Li

Based on the slides of Patterson and Hennessy



What we've learnt so far

- Combinatorial logic
 - ALU
- Sequential logic
 - Clocks and basic state elements (SR latch, D latch, flip-flop)
- The single-cycle CPU design

Single-cycle CPU uses a slow clock



Pipelining: a laundry analogy



Pipelined laundry: overlapping execution

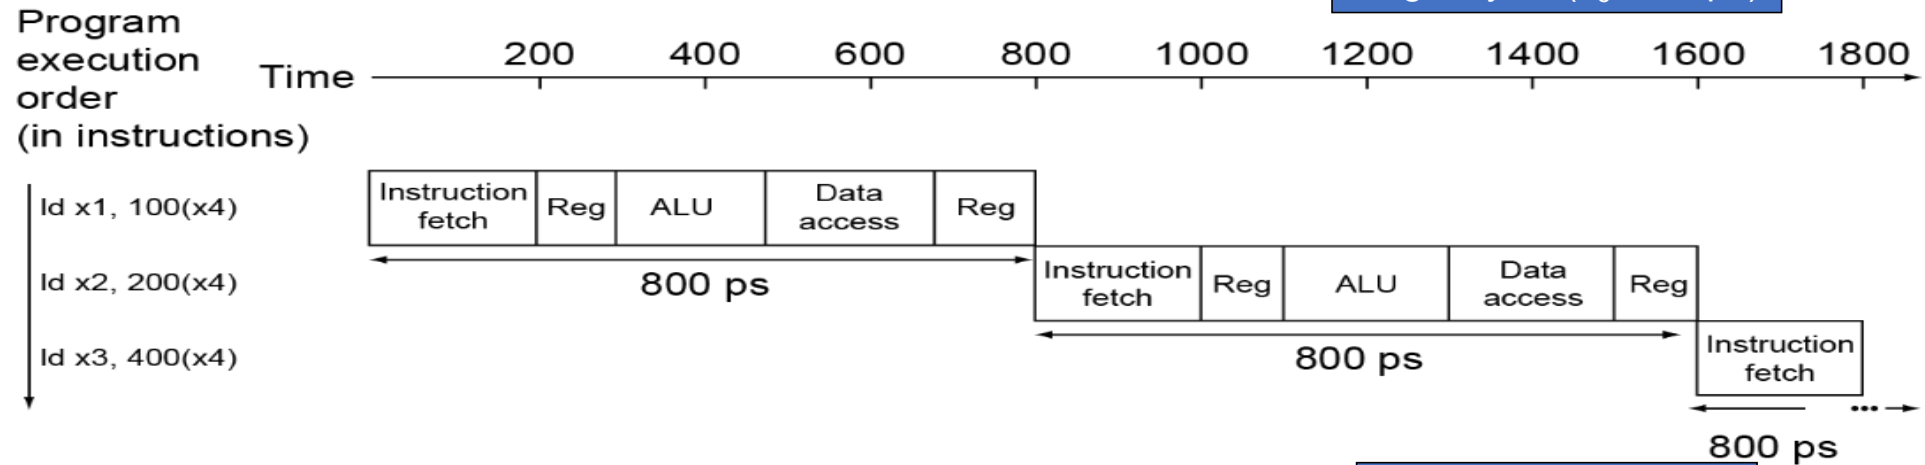
- Parallelism improves performance

RISC-V Pipeline

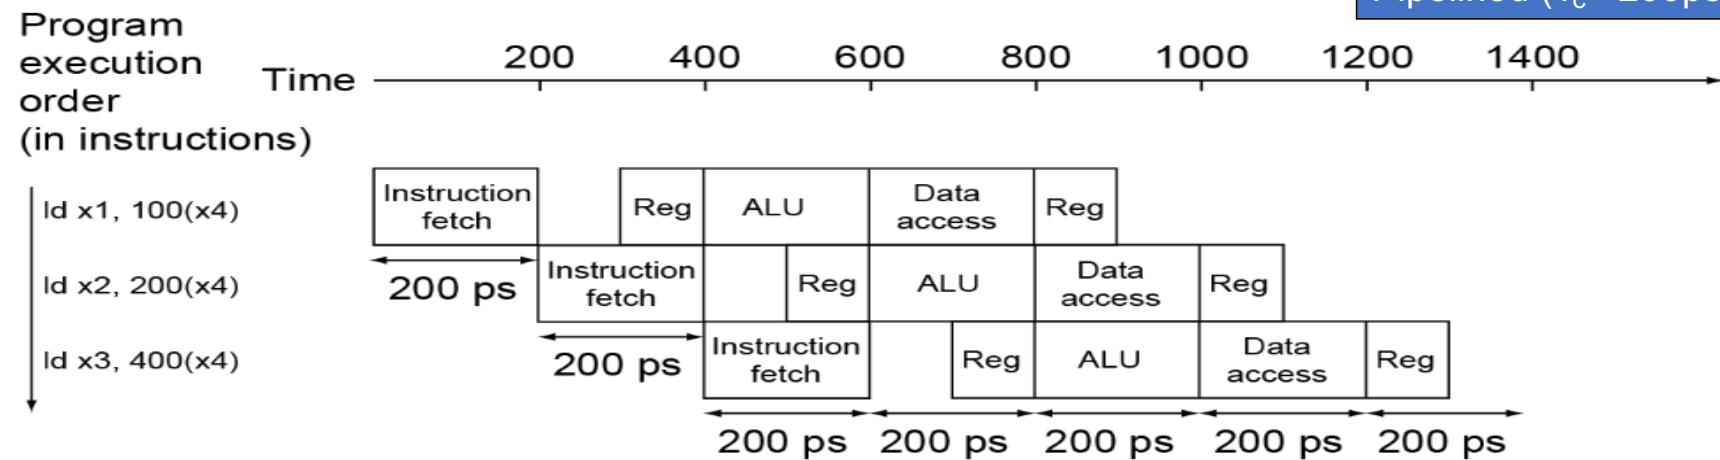
- Five stages, one step per stage
 1. IF: Instruction fetch from memory
 2. ID: Instruction decode & register read
 3. EX: Execute operation or calculate address
 4. MEM: Access memory operand
 5. WB: Write result back to register

Pipeline Performance

Single-cycle ($T_c = 800\text{ps}$)



Pipelined ($T_c = 200\text{ps}$)



Pipeline Speedup

- Pipelining increases throughput (instructions/sec)
 - Latency (time for each instruction) does not decrease
- If all stages are balanced (i.e., all take the same time)
 - $\text{throughput}_{\text{pipelined}} = \text{number-of-stages} * \text{throughput}_{\text{nonpipelined}}$
 - If not balanced, speedup is less

Throughput = $1/\text{time between instructions}$



Pipelining and ISA Design

- RISC-V ISA designed for pipelining
 - All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - c.f. x86: 1- to 17-byte instructions
 - Few and regular instruction formats
 - Can decode and read registers in one step
 - Load/store addressing
 - Can calculate address in 3rd stage, access memory in 4th stage

Pipeline challenges: hazards

- Situations that prevent starting the next instruction in the next cycle
- Structure hazards
 - A required resource is busy
- Data hazard
 - Need to wait for previous instruction to complete its data read/write
- Control hazard
 - Deciding on control action depends on previous instruction

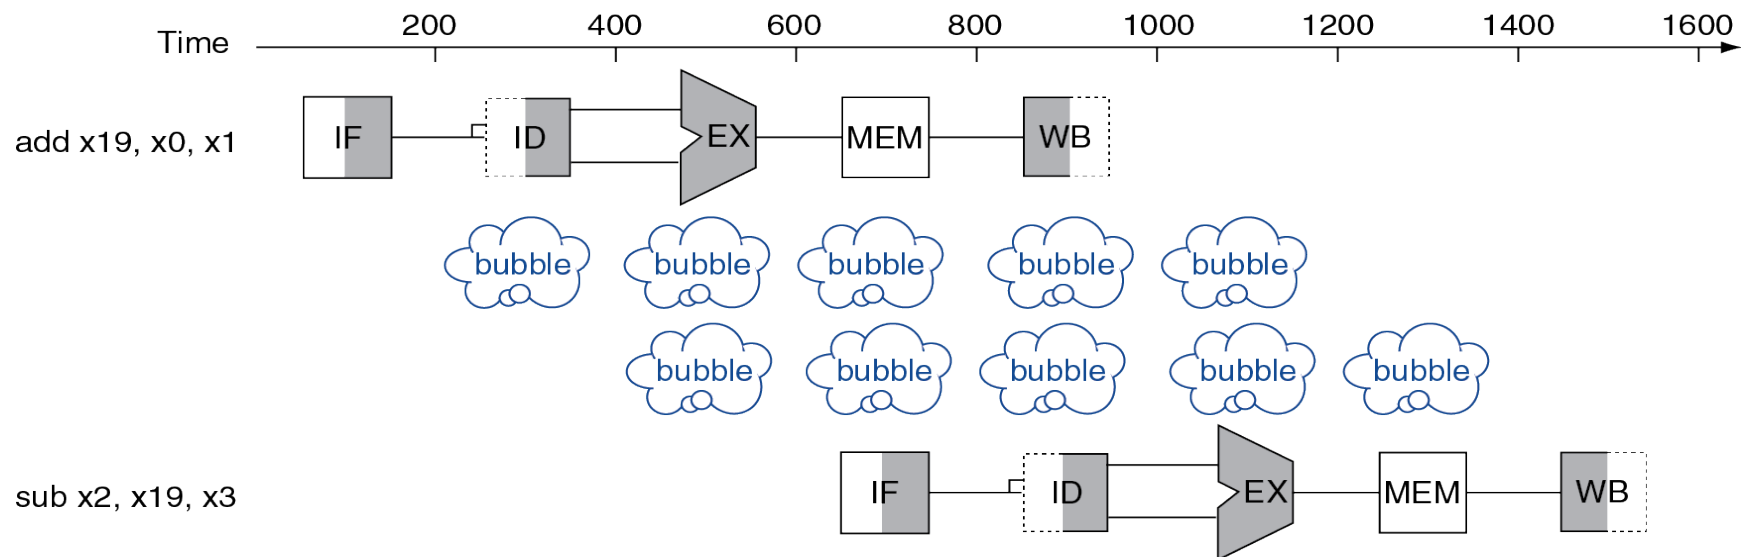
Structure Hazards

- Conflict use of a single resource
- Example: Suppose CPU uses a single memory
 - Load/store requires data access
 - Instruction fetch would have to *stall* for that cycle
 - Would cause a pipeline “bubble”
- Solution: Use separate instruction/data memories

Data Hazards

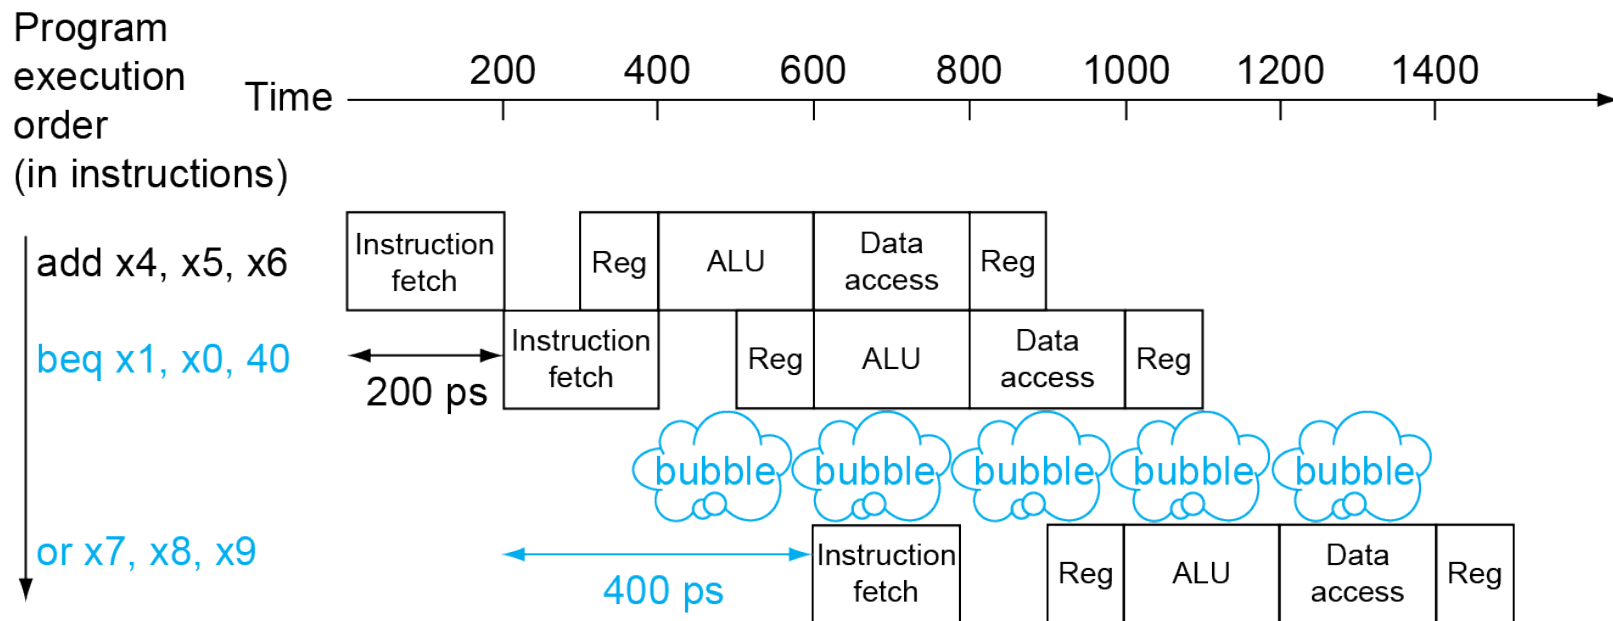
- An instruction depends on the previous instruction to complete its data read/write

add **x19**, x0, x1
sub x2, **x19**, x3



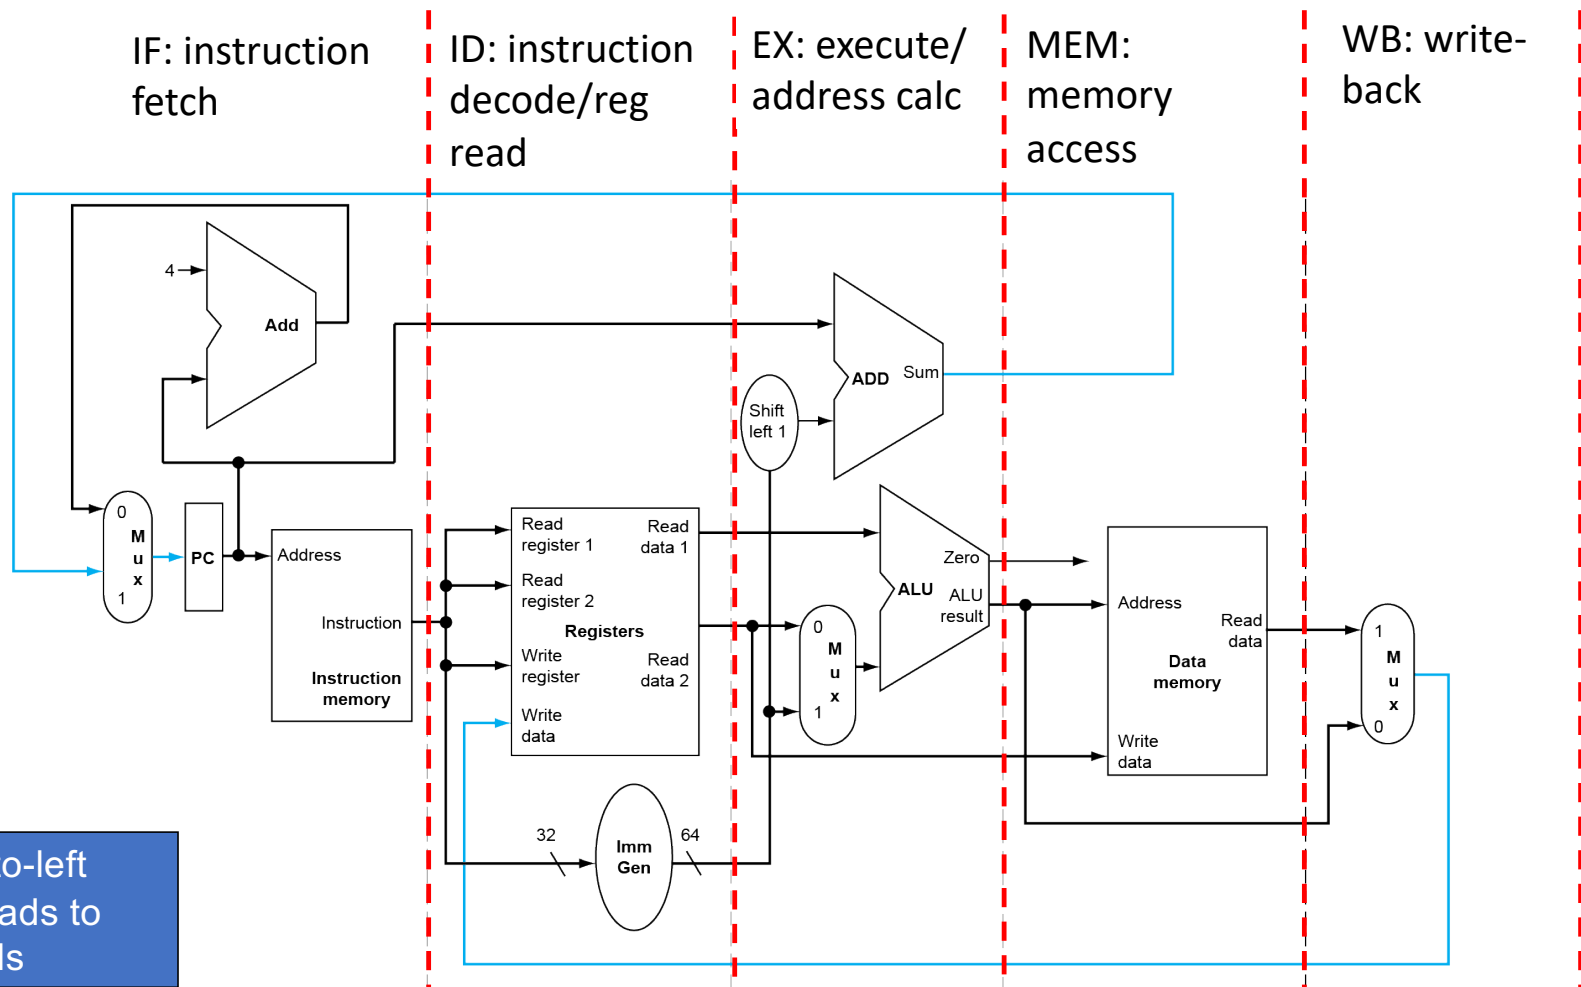
Control hazard

- Wait until branch outcome determined before fetching next instruction



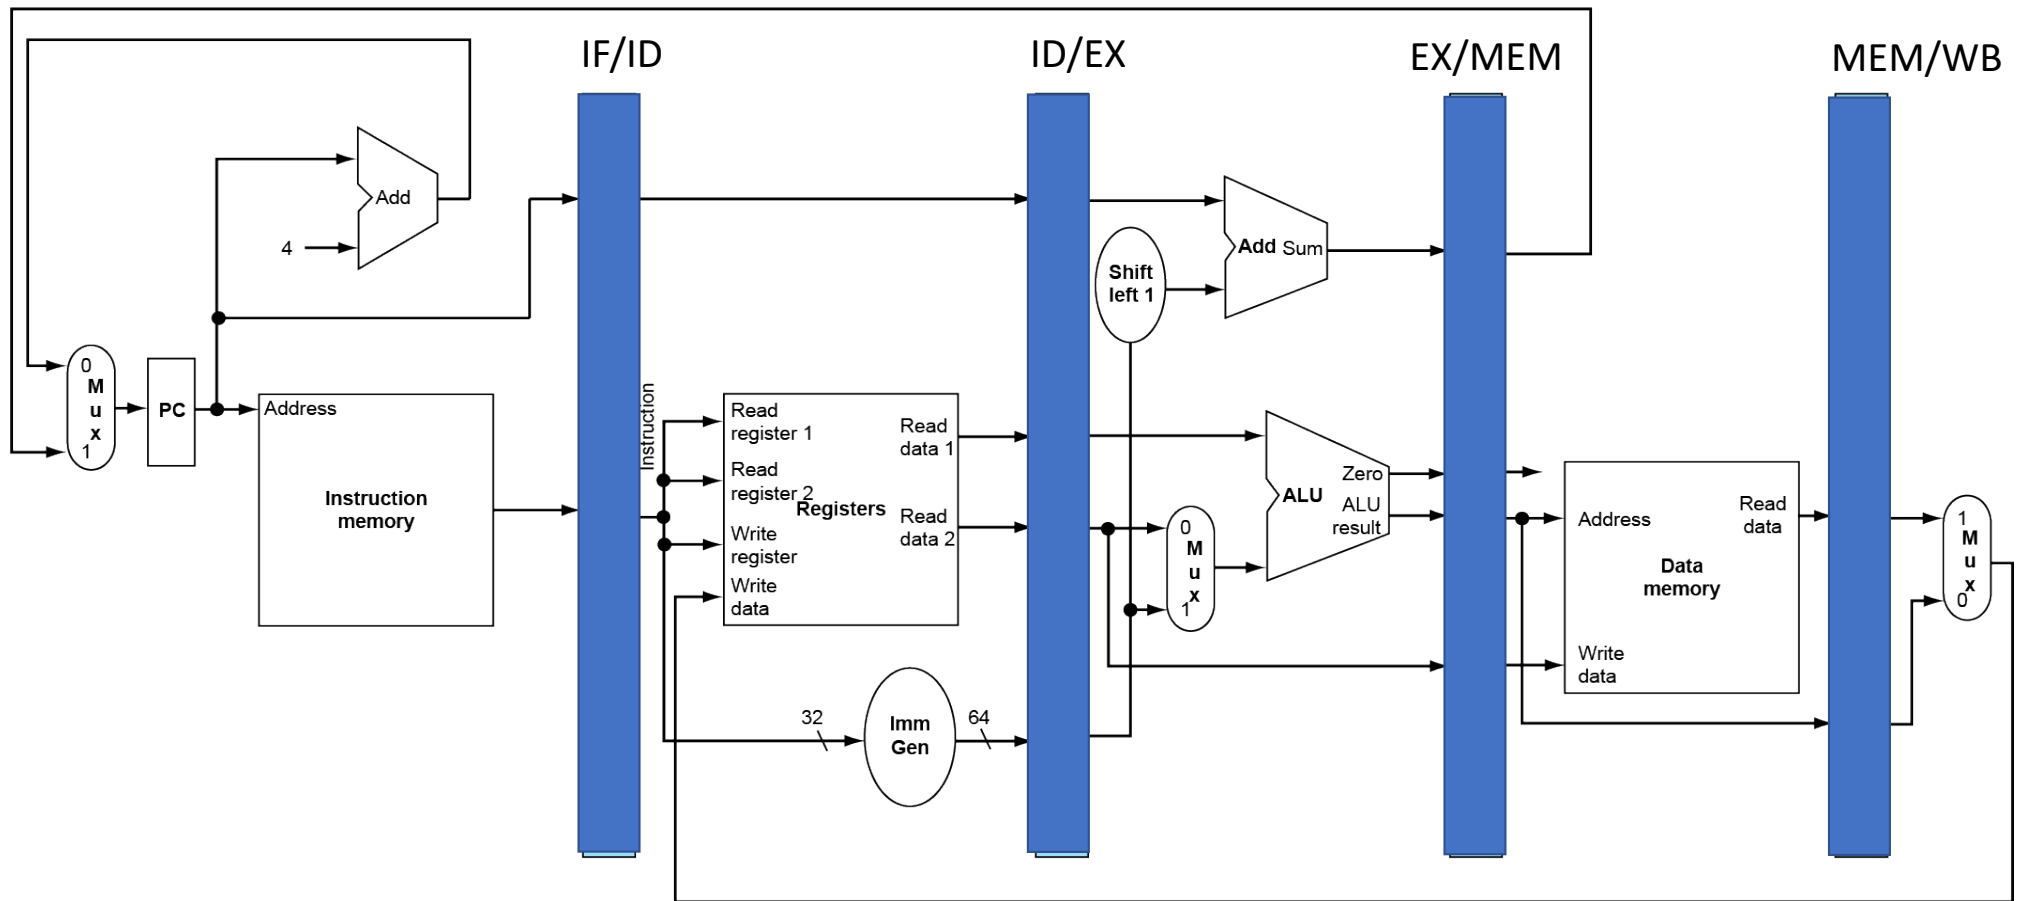
A basic pipelined RISC-V CPU

Pipelined Datapath




Pipeline registers

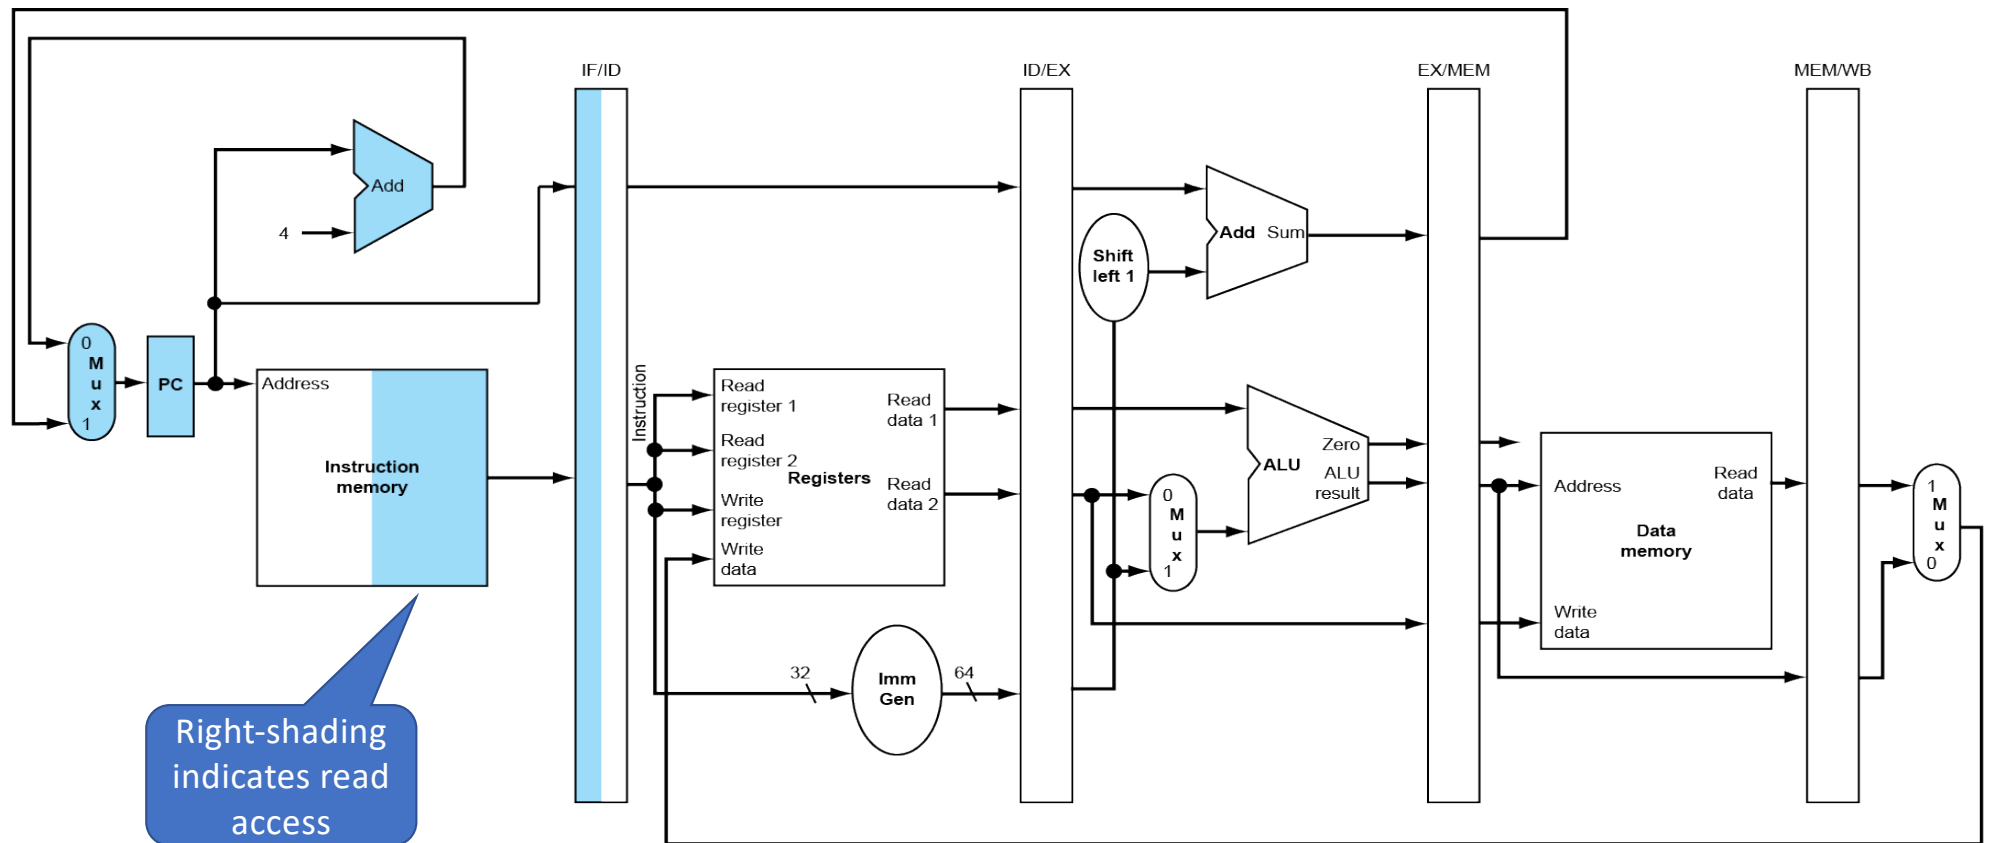
Needed to hold information produced in previous cycle



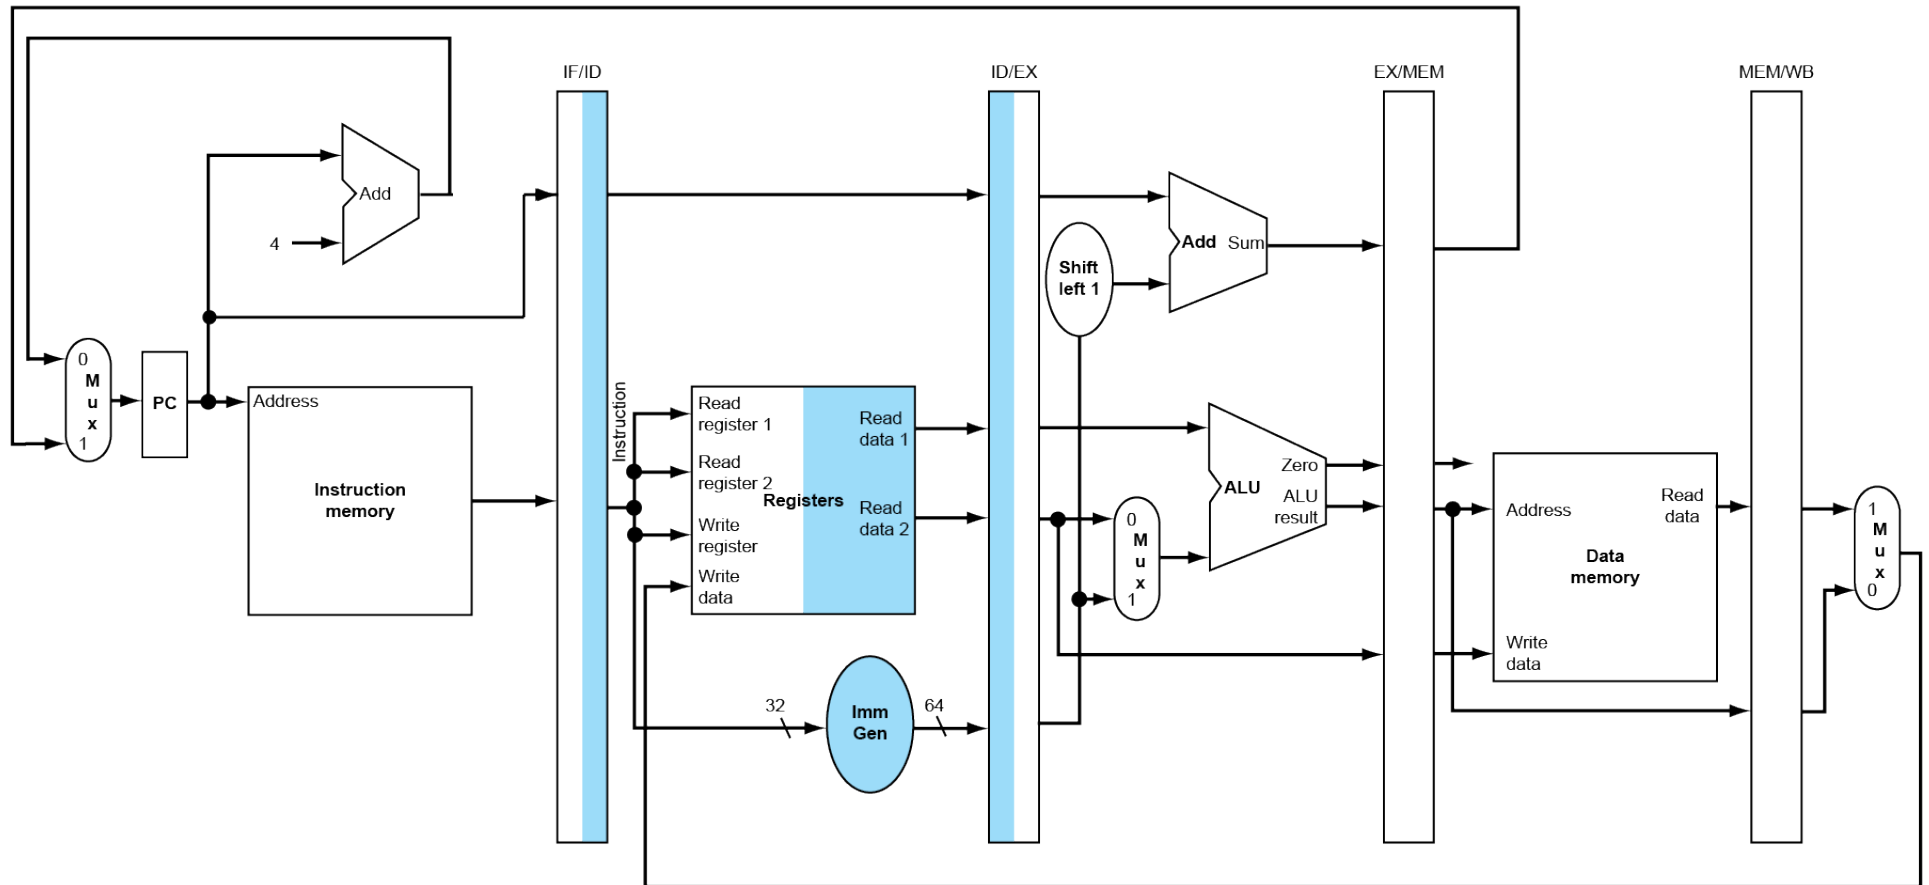

IF for Load, Store



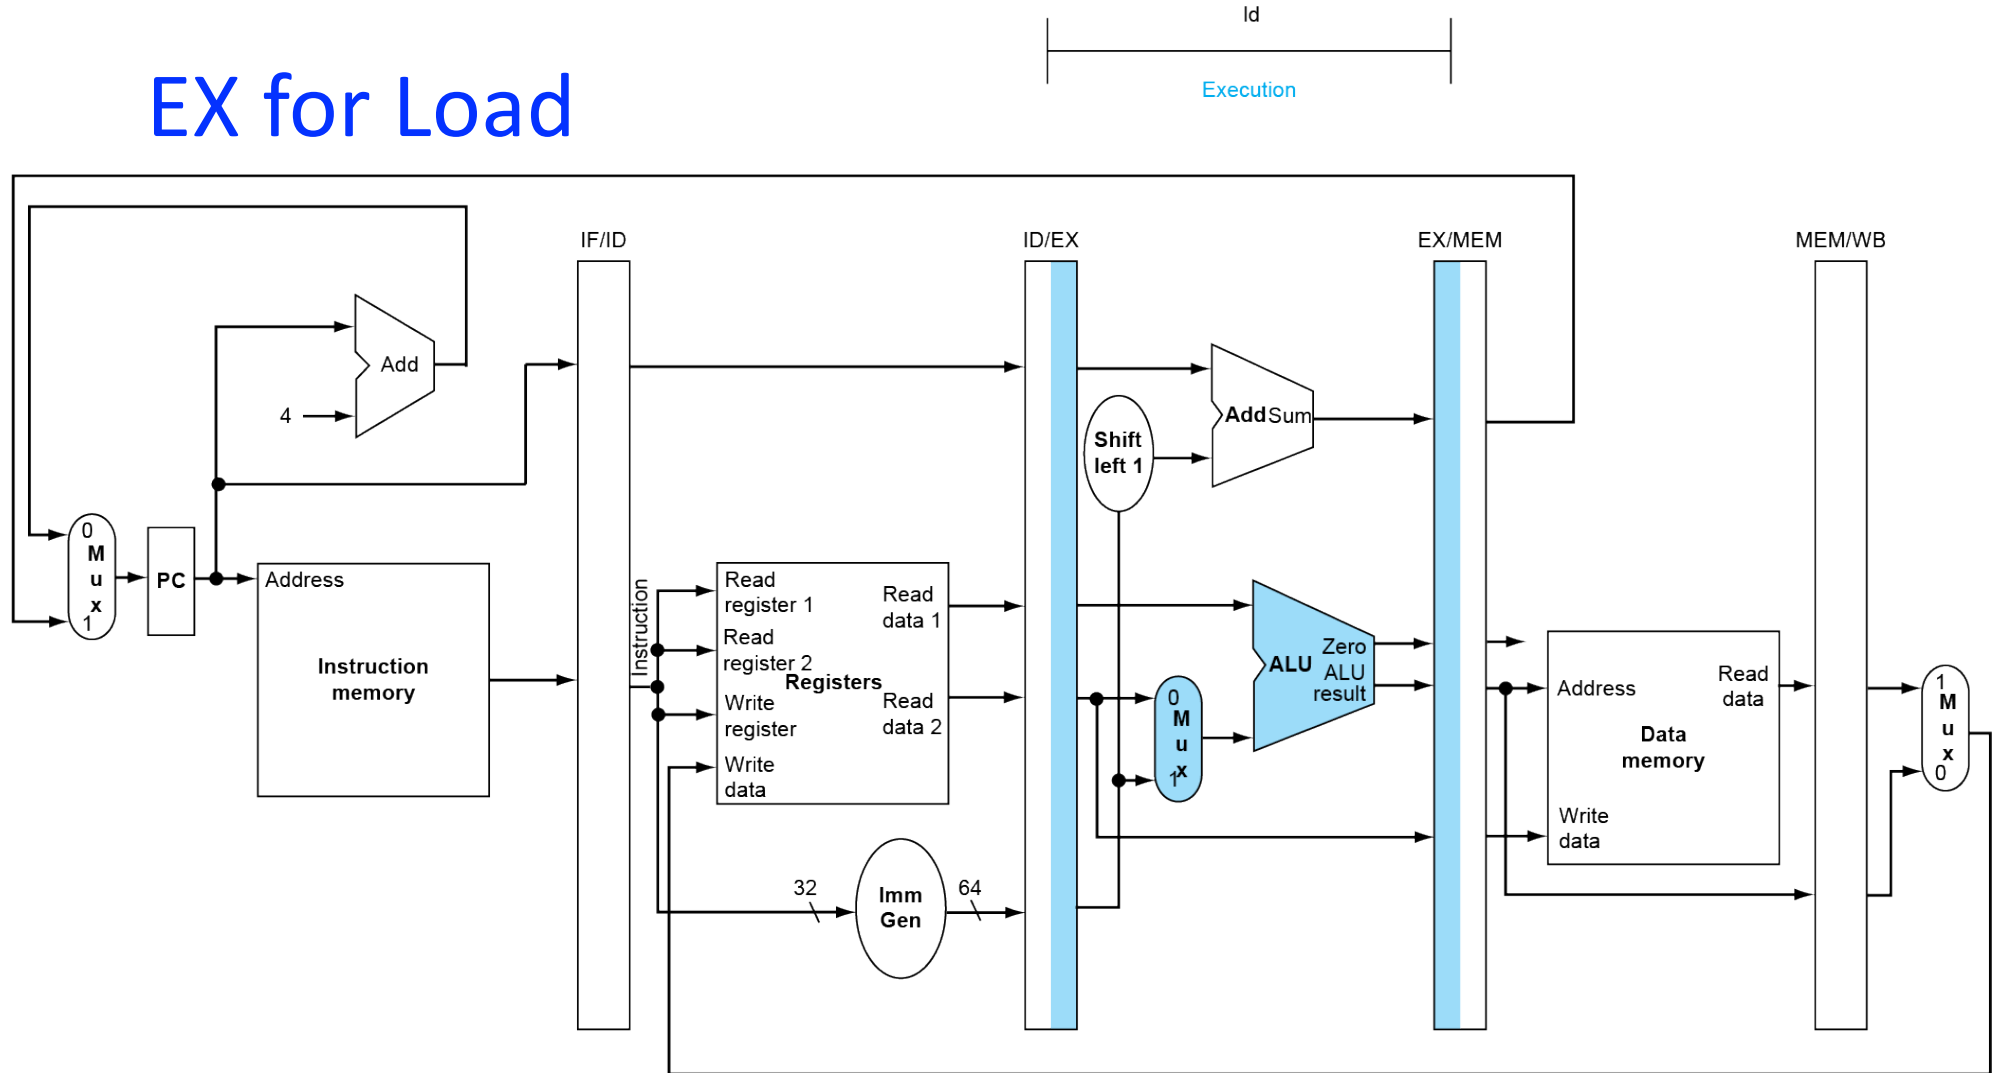
Single-clock-cycle diagram shows the state of an entire datapath during a single clock cycle



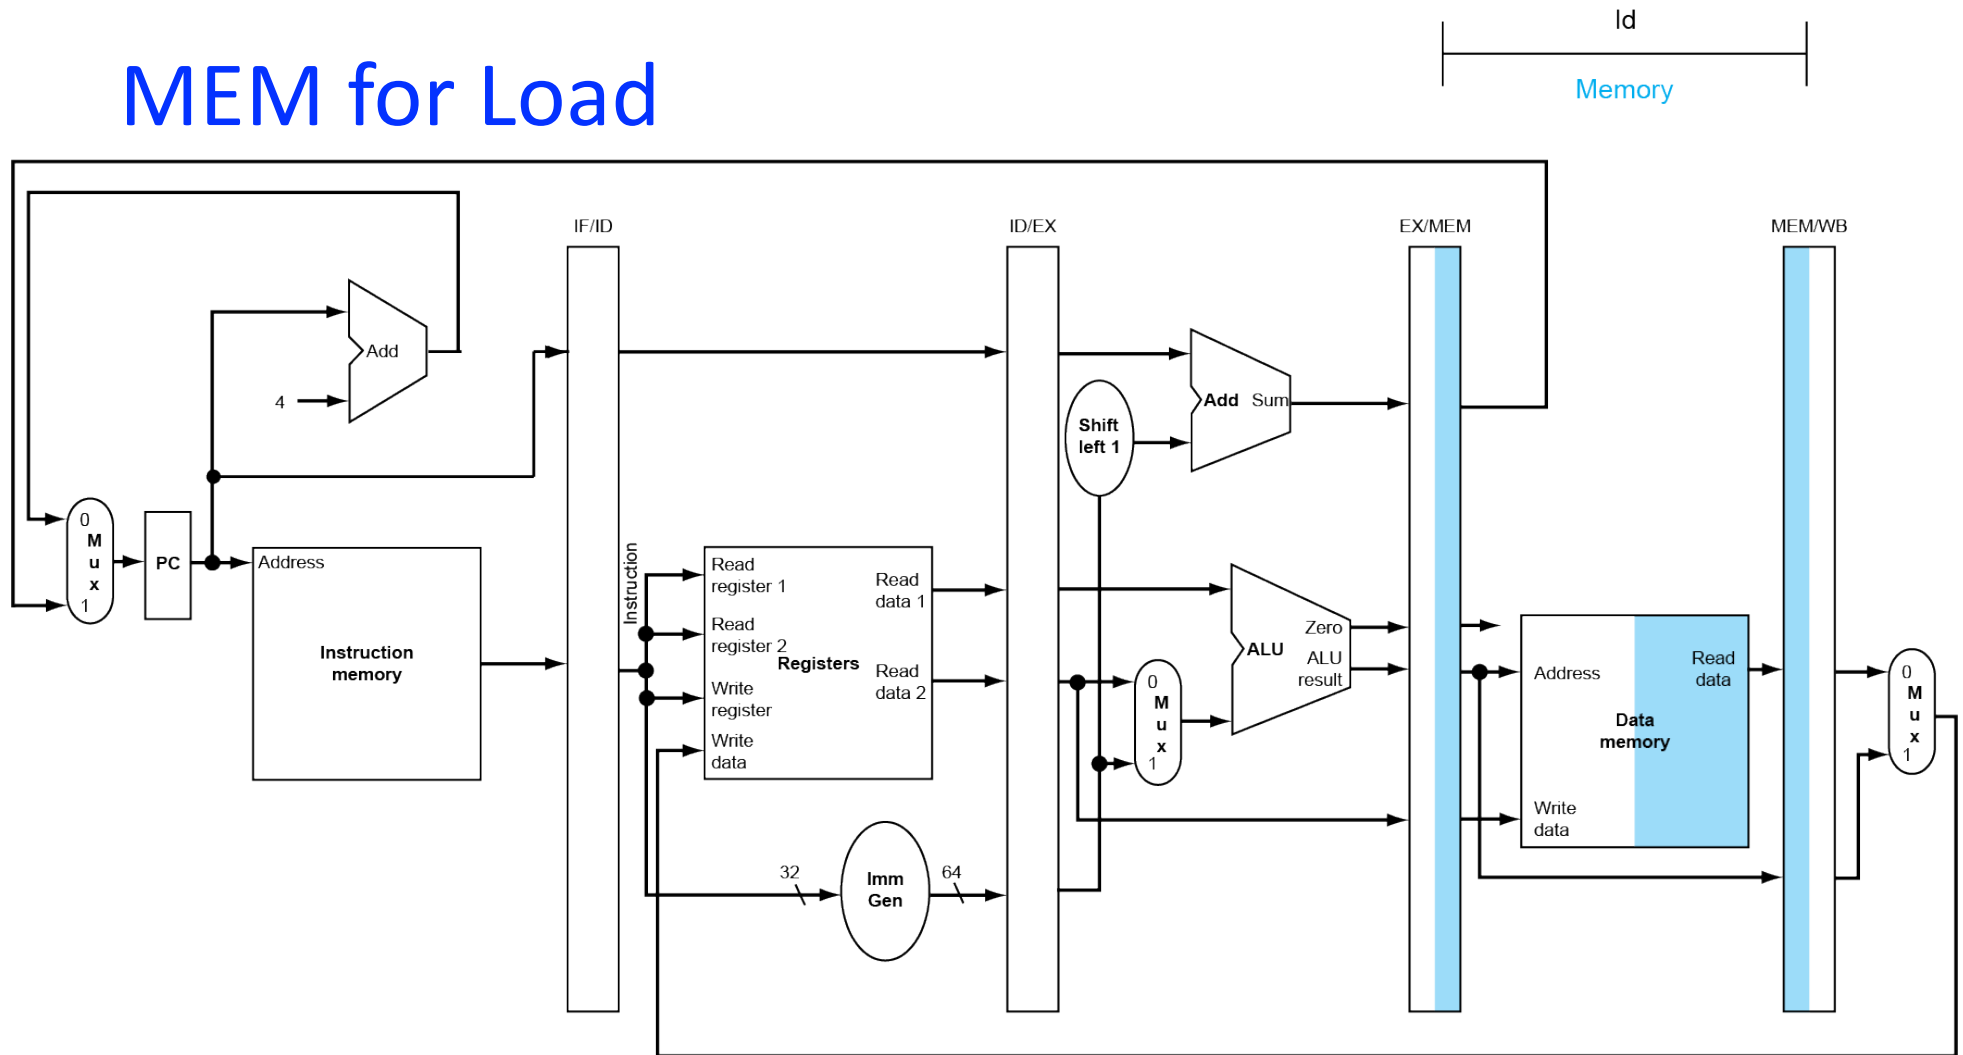
ID for Load, Store, ...



EX for Load

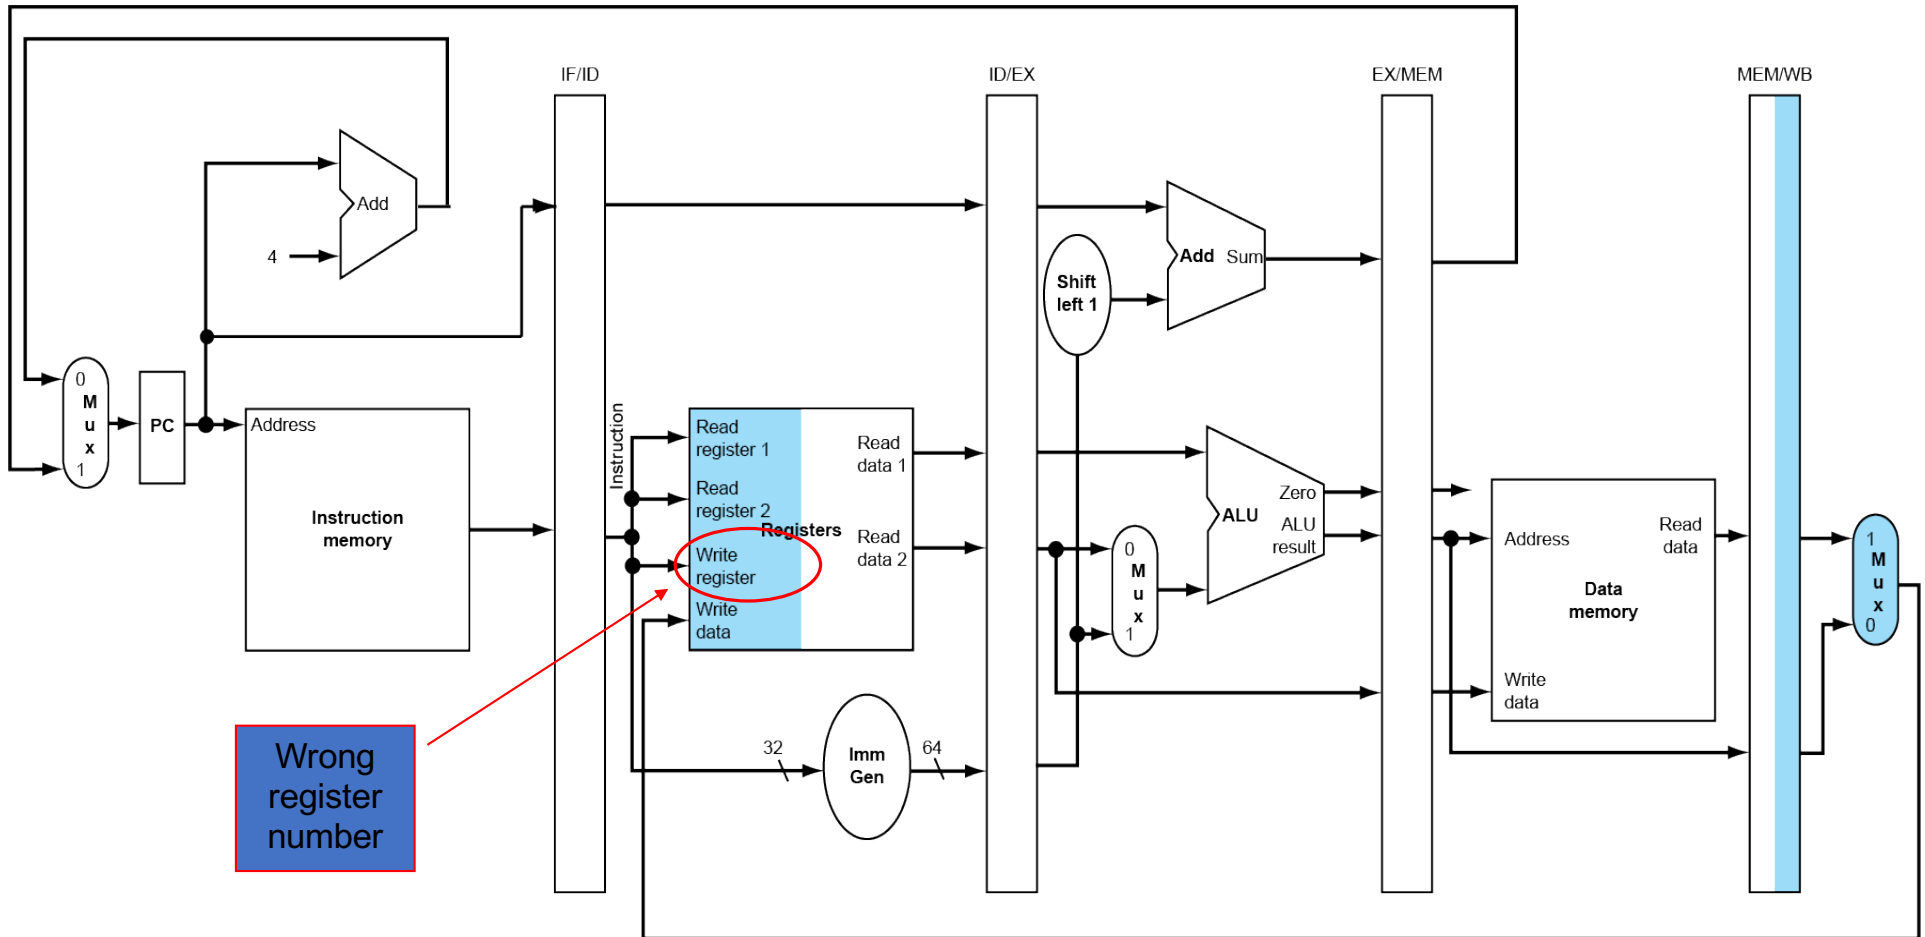


MEM for Load

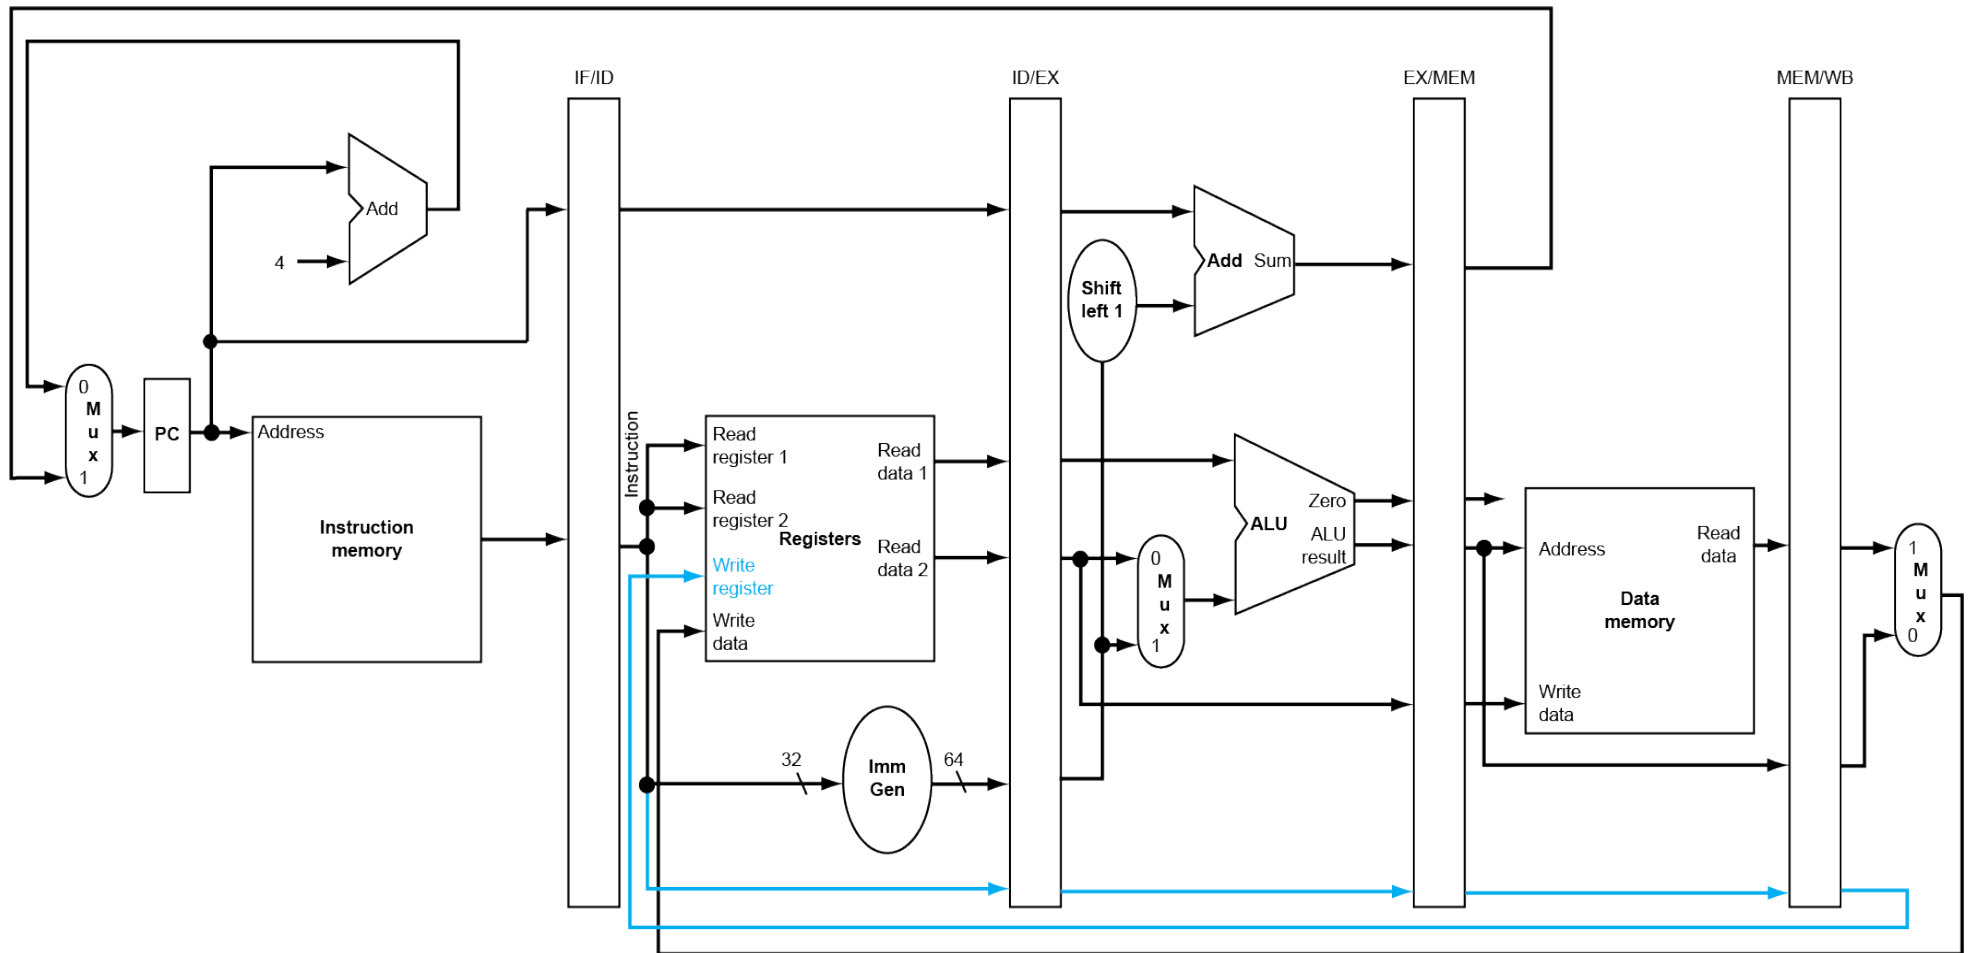


WB for Load

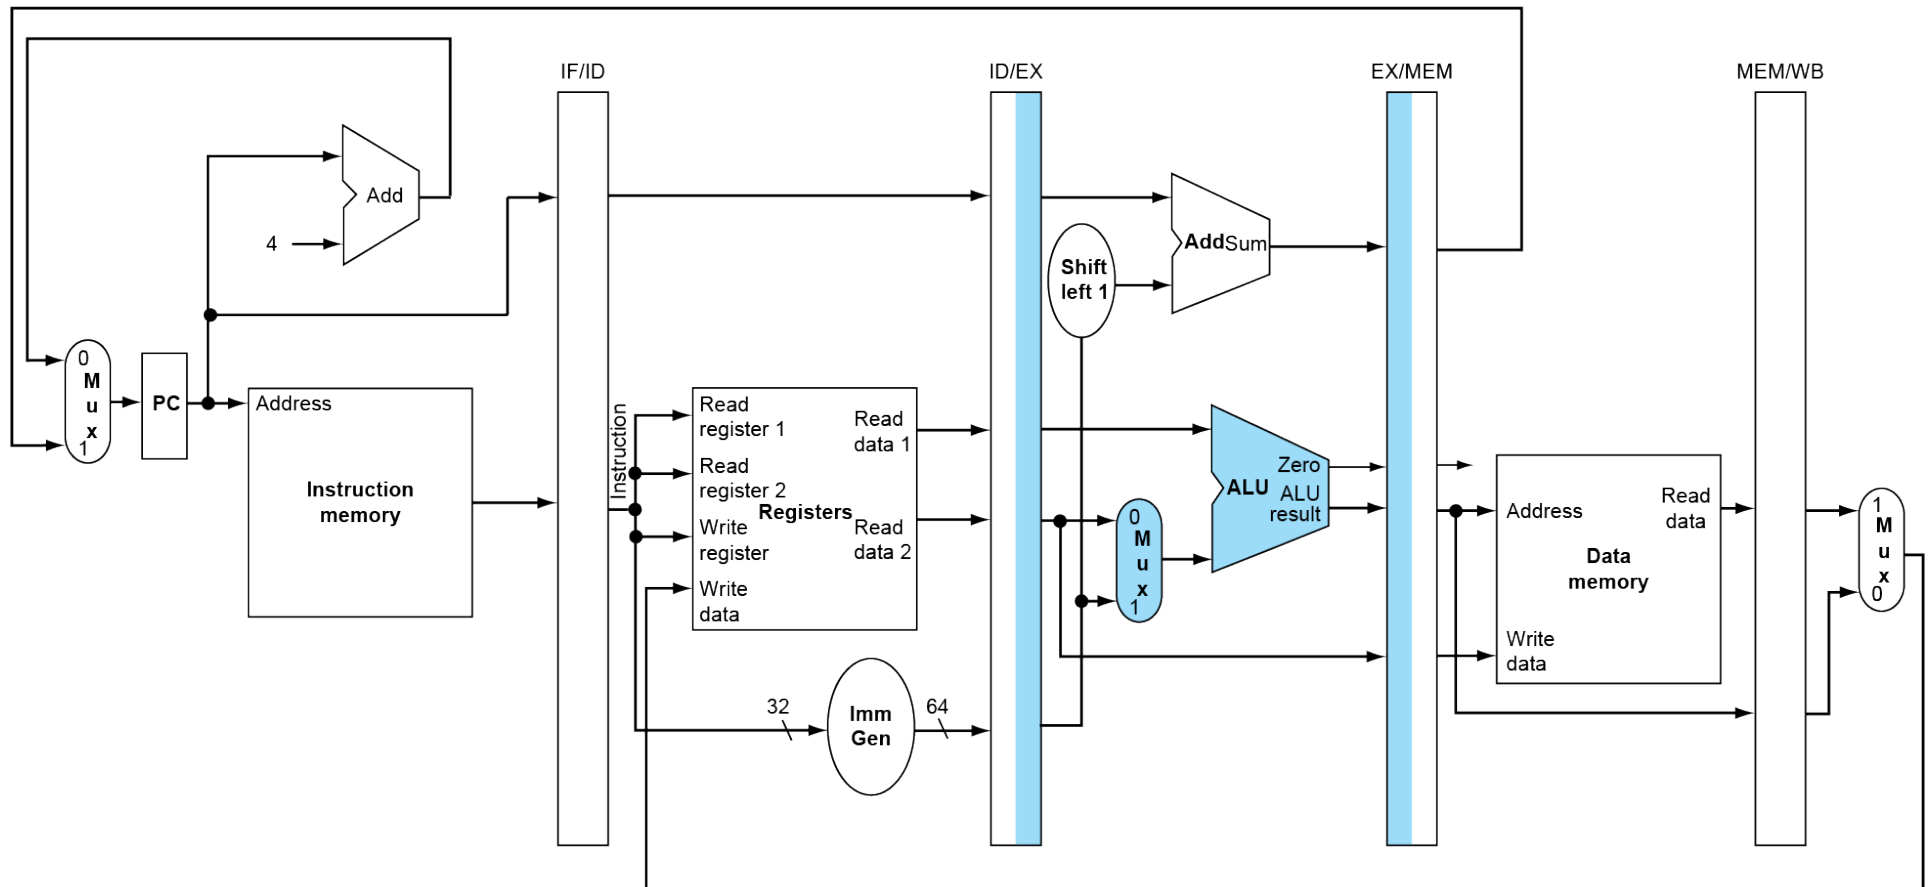

Id
Write-back



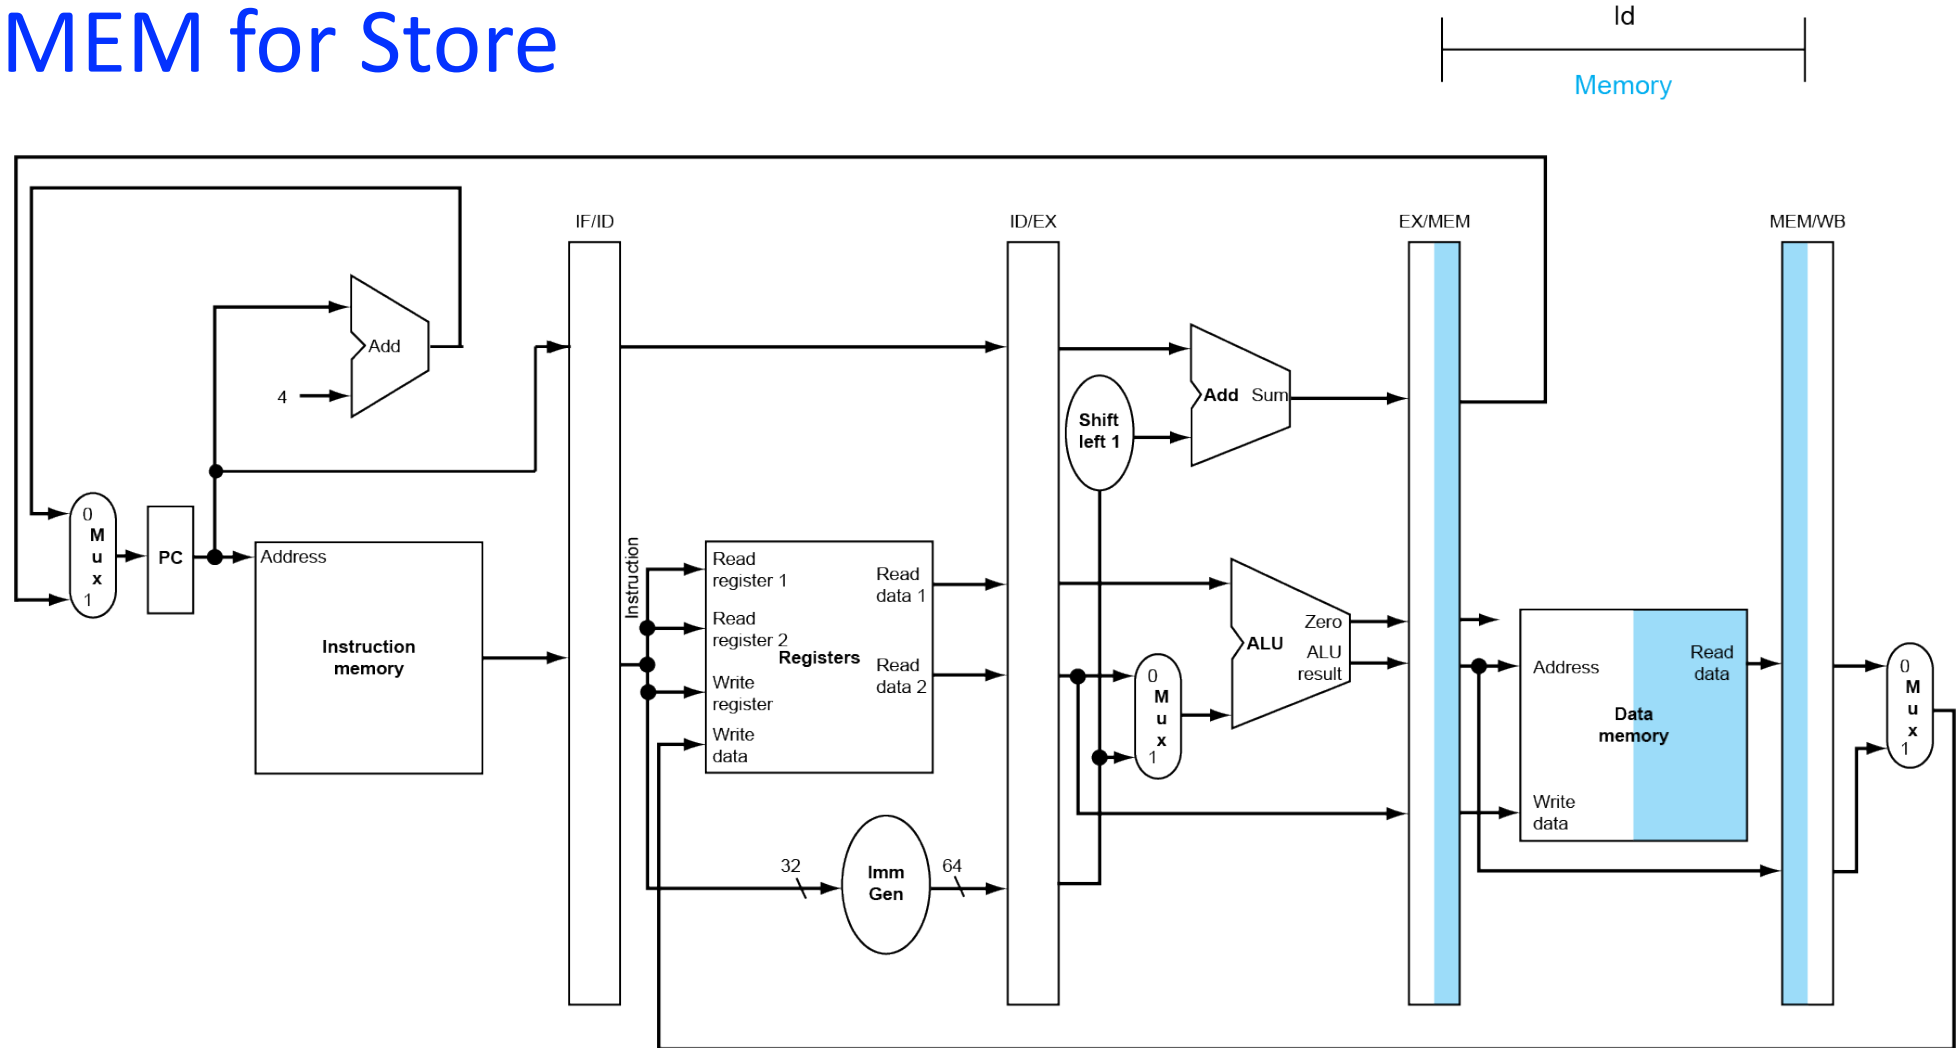
Corrected Datapath for Load



EX for Store

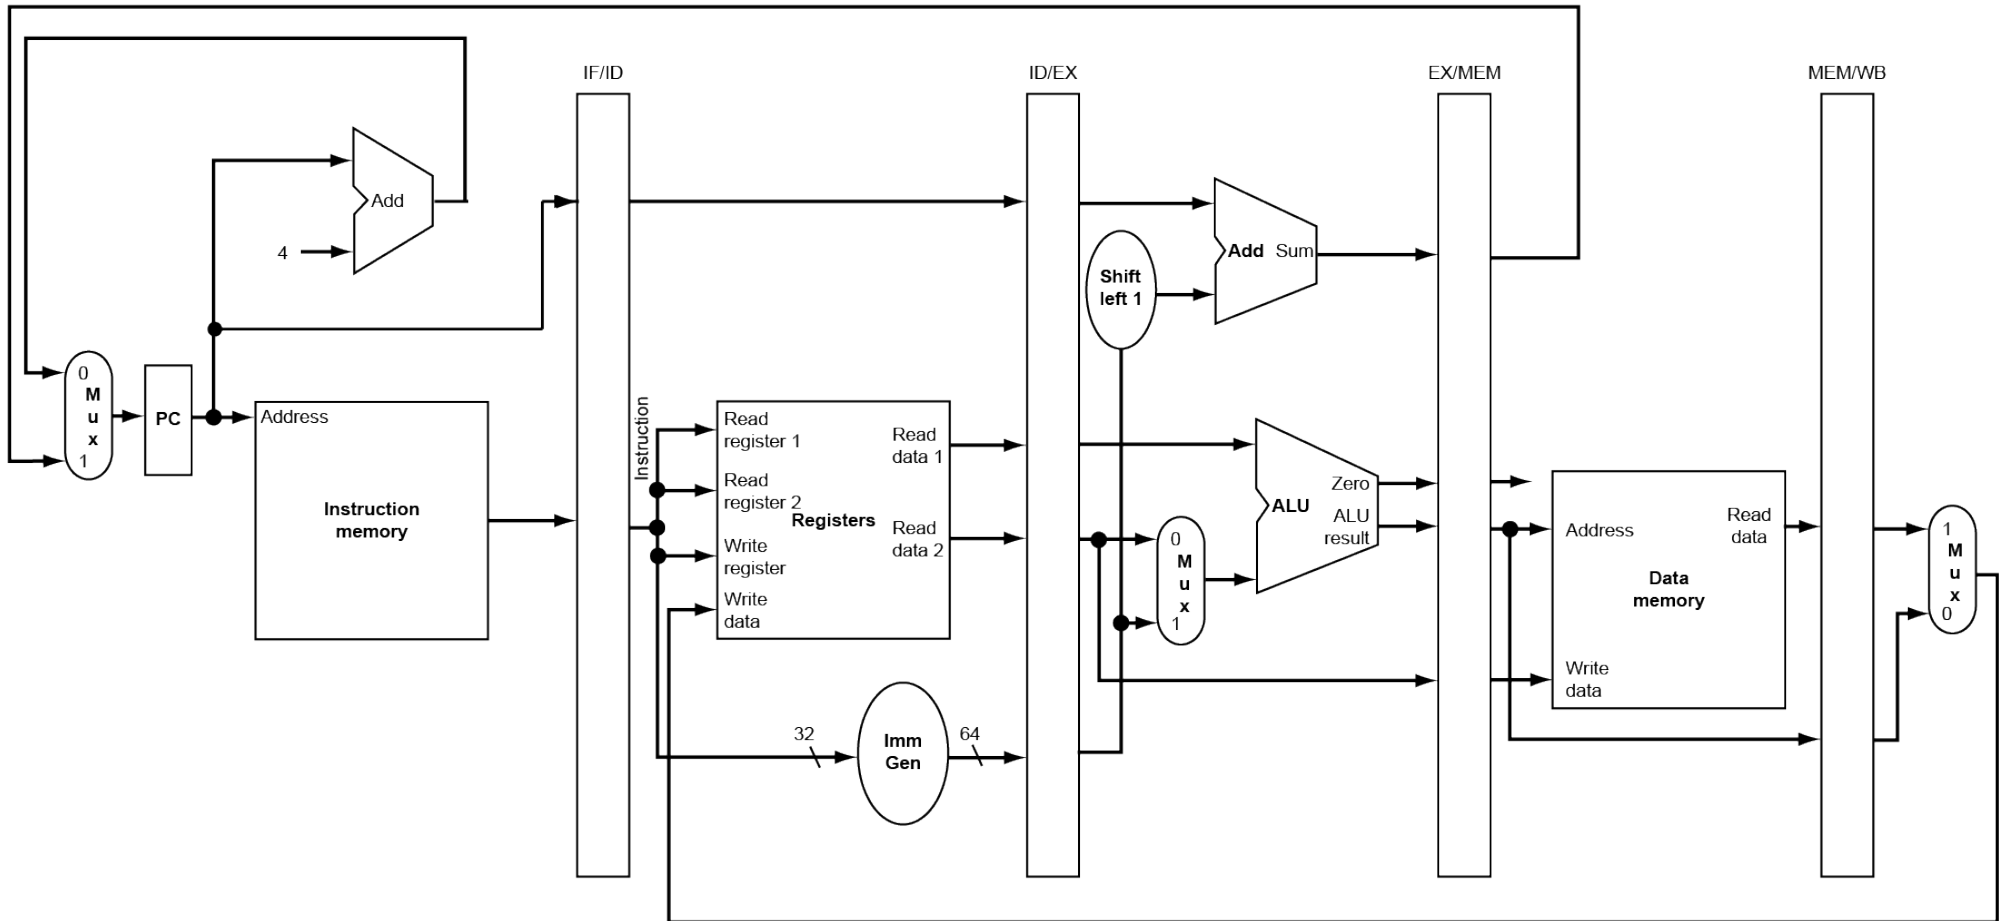


MEM for Store



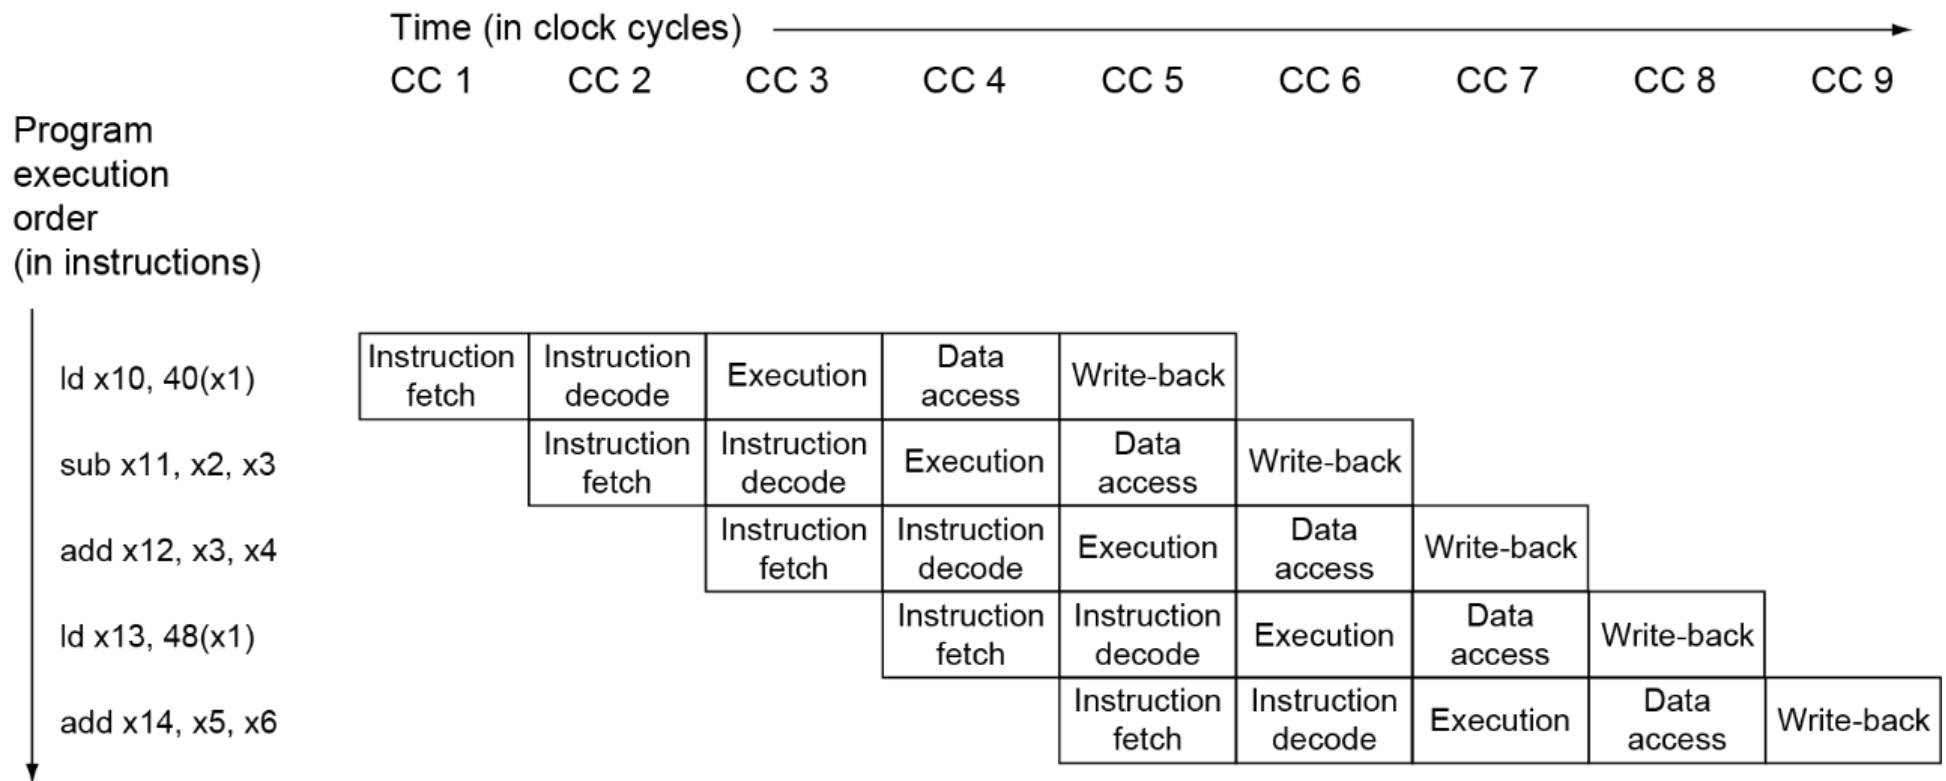
WB for Store

sd
Write-back



Multi-Cycle Pipeline Diagram

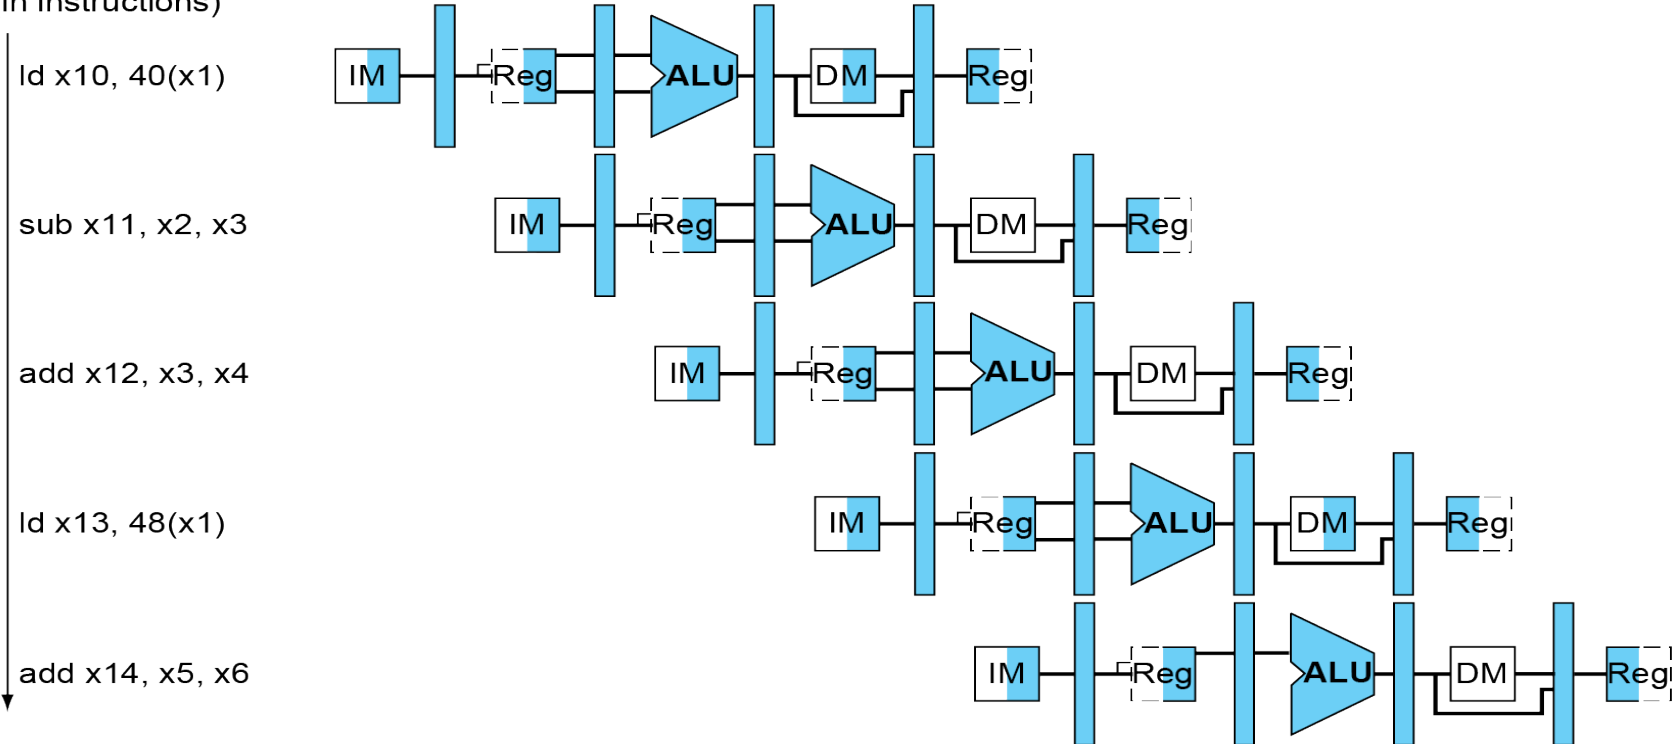
- Traditional form



Multi-Cycle Pipeline Diagram

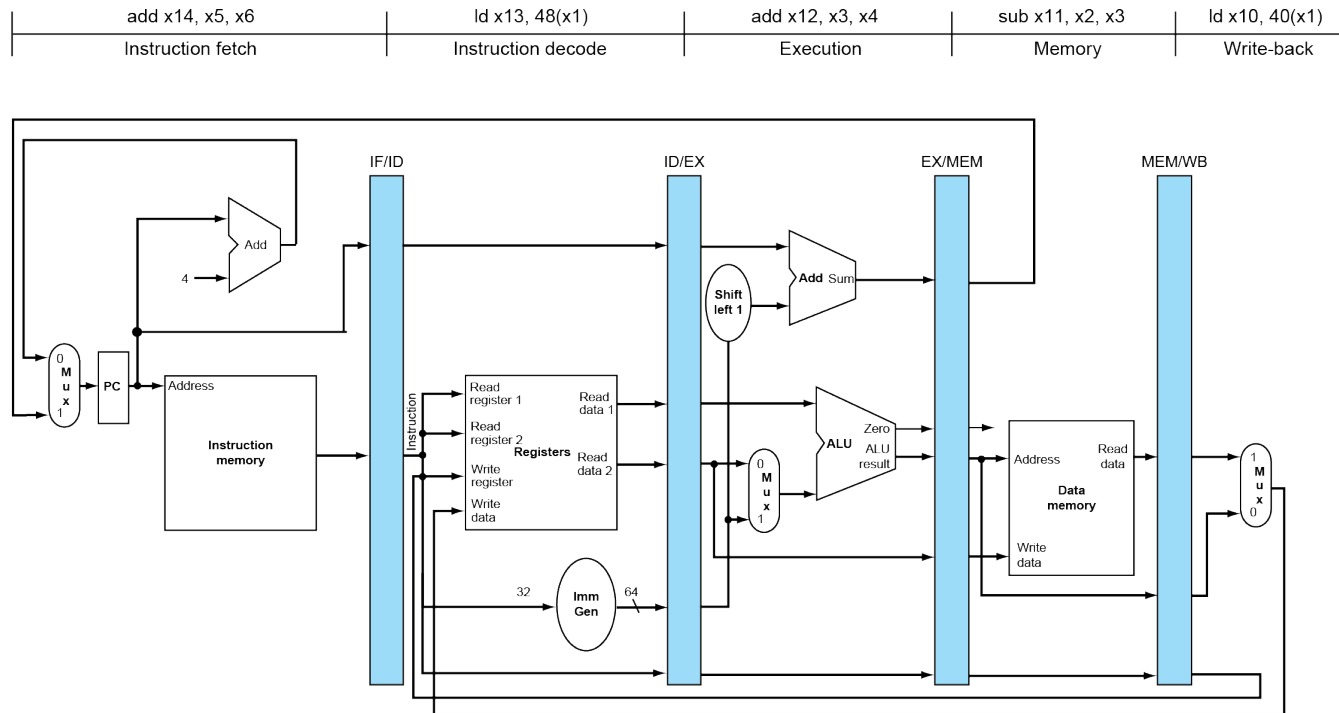
- Form showing resource usage

Program
execution
order
(in instructions)

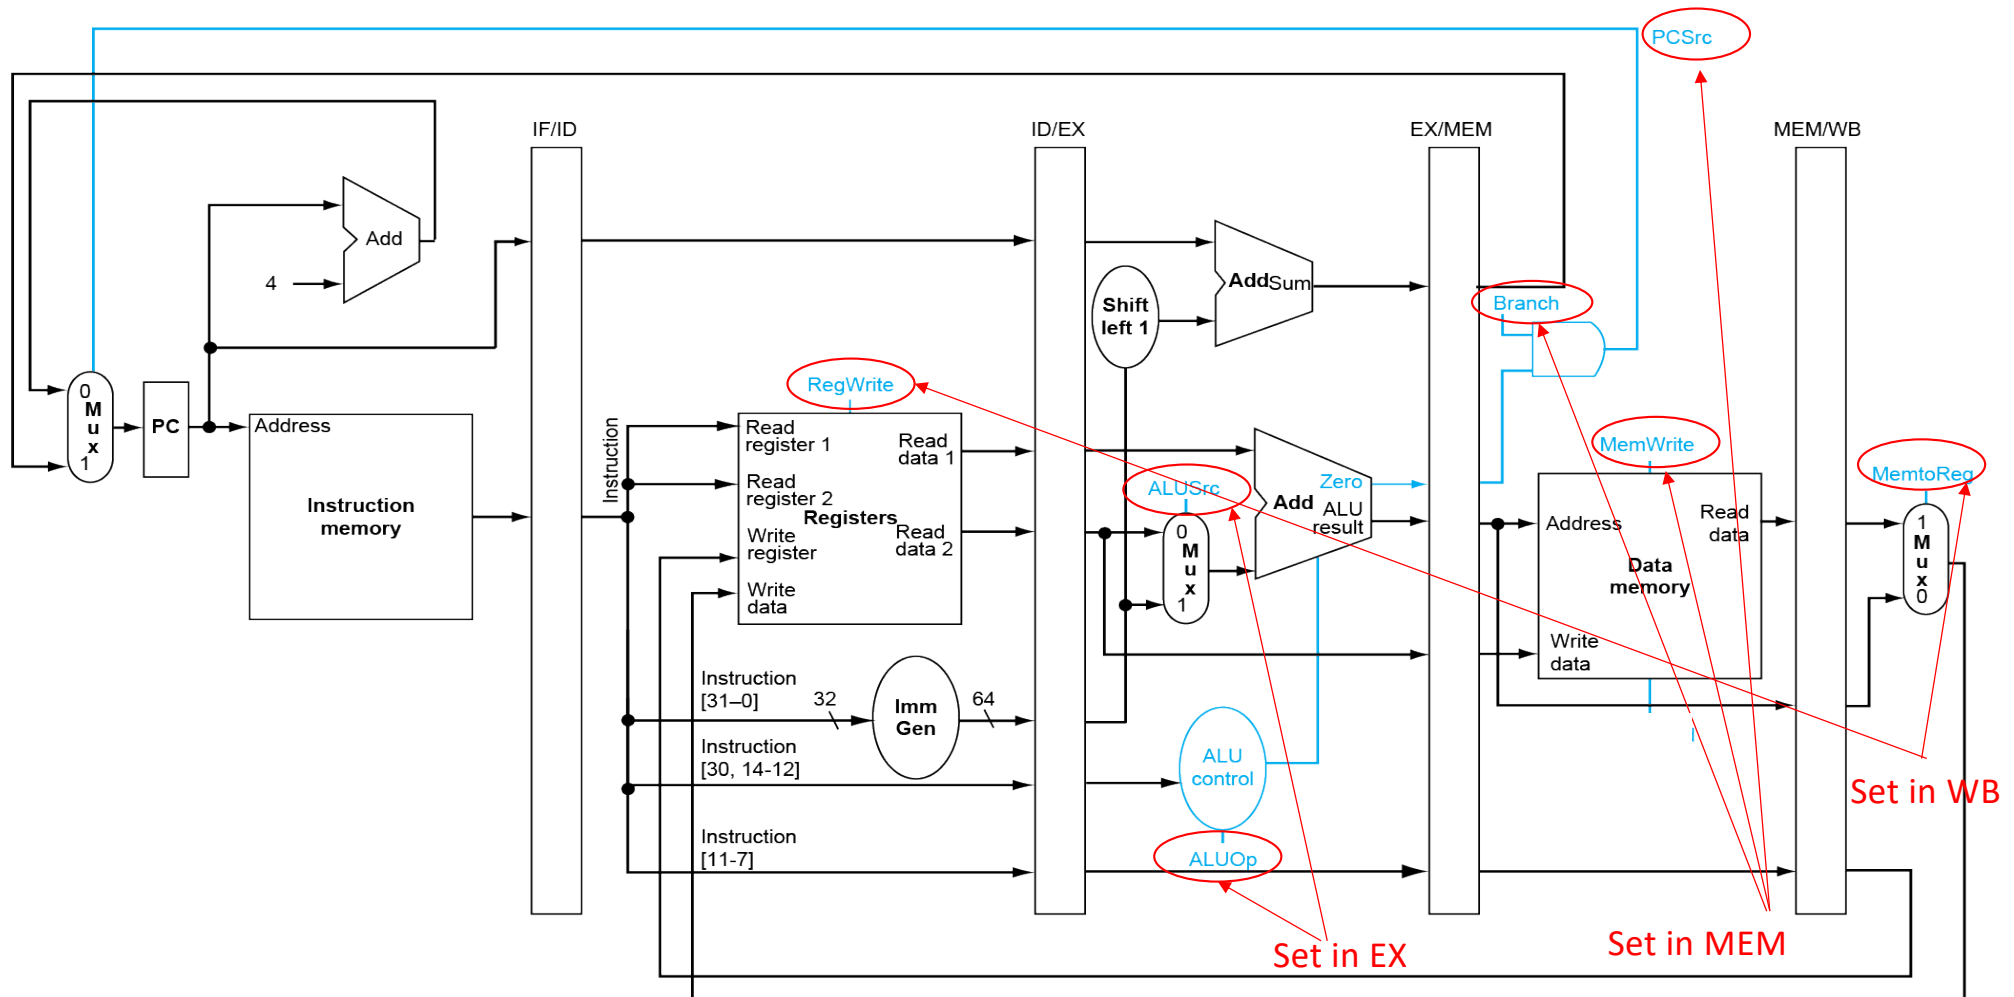


Single-Cycle Pipeline Diagram

- State of pipeline in a given cycle

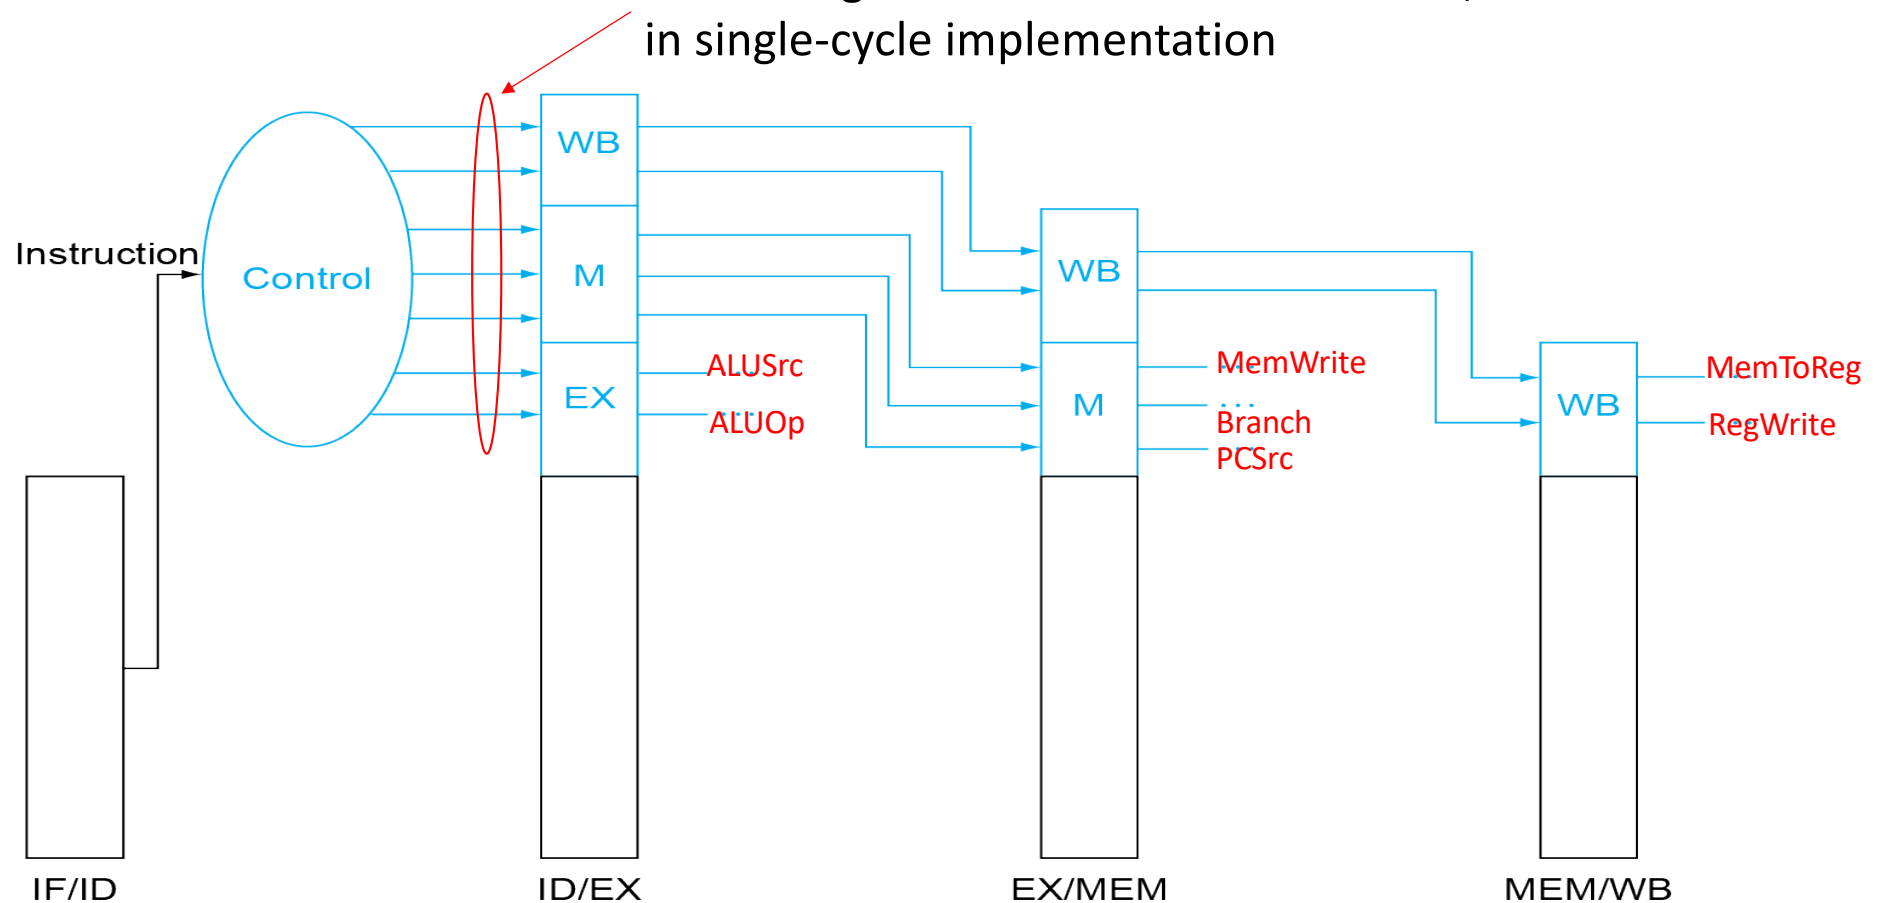


Pipelined Control (Simplified)



Pipelined Control

Control signals derived from instruction, same as in single-cycle implementation



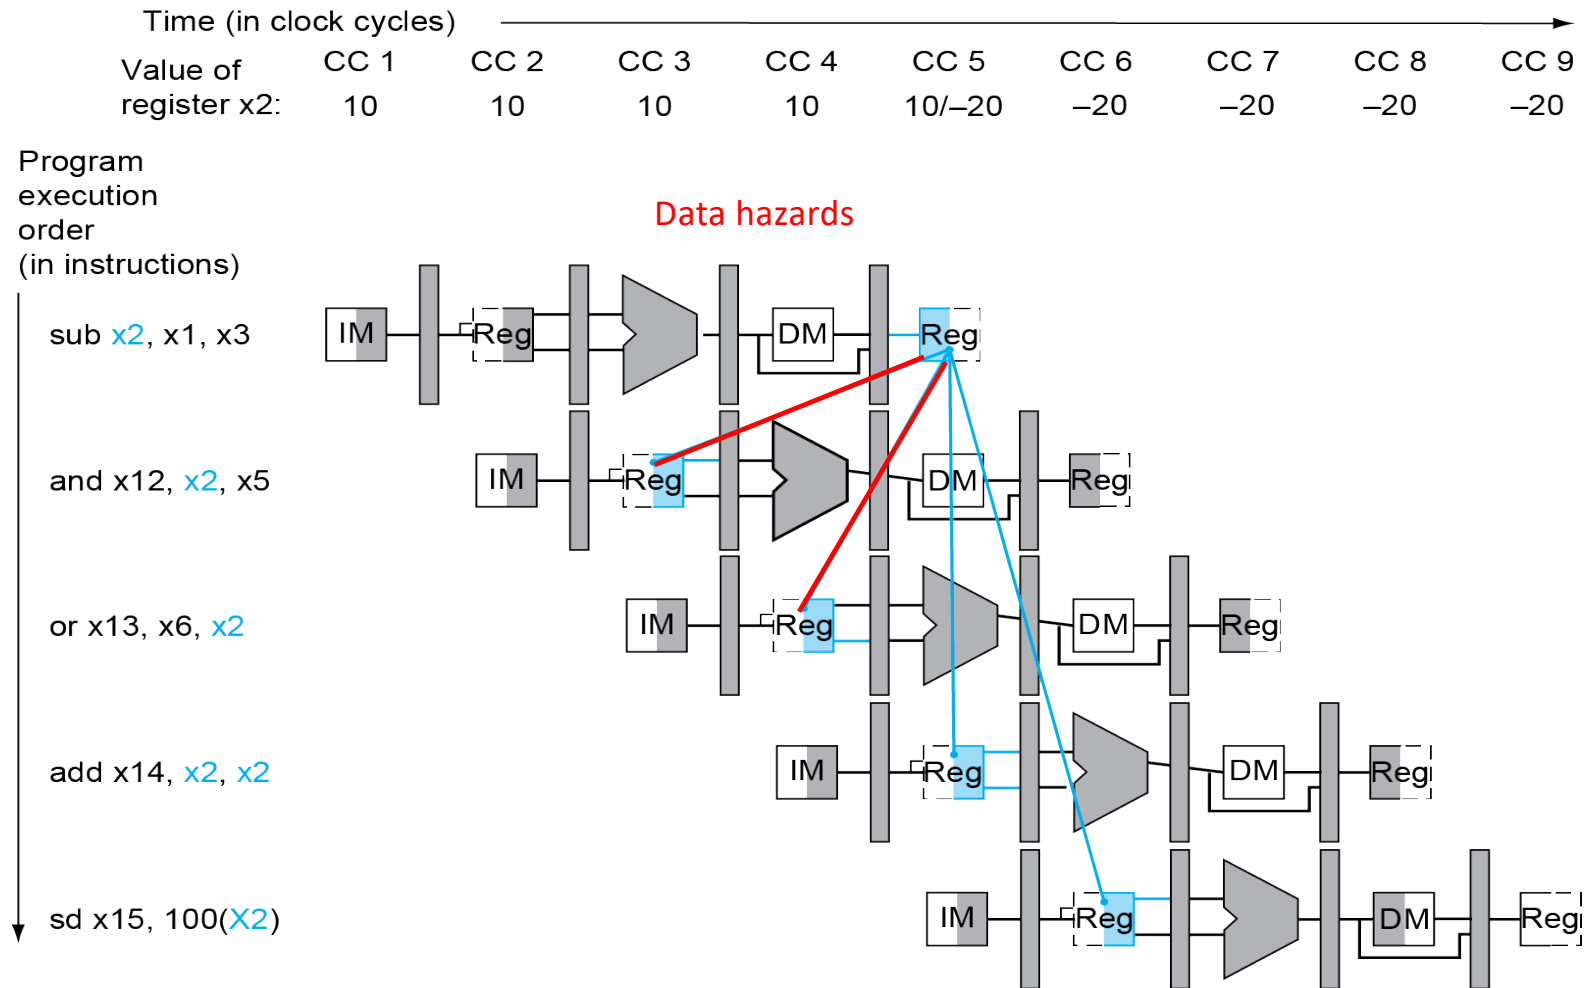
Data Hazards in ALU Instructions

- Example instruction sequence:

```
sub  x2, x1, x3
and  x12, x2, x5
or   x13, x6, x2
add  x14, x2, x2
sd   x15, 100(x2)
```

- Solution: forwarding (aka bypassing)
 - Use result when it's computed; don't wait for it to be stored in register

Dependencies & Forwarding



Detecting the Need to Forward

- Pass register numbers along pipeline
 - e.g., ID/EX.RegisterRs1: register number for Rs1 in ID/EX pipeline register
- ALU operand register numbers (in EX stage) are:
 - ID/EX.RegisterRs1
 - ID/EX.RegisterRs2
- Data hazards when
 - 1a. EX/MEM.RegisterRd = ID/EX.RegisterRs1
 - 1b. EX/MEM.RegisterRd = ID/EX.RegisterRs2
 - 2a. MEM/WB.RegisterRd = ID/EX.RegisterRs1
 - 2b. MEM/WB.RegisterRd = ID/EX.RegisterRs2

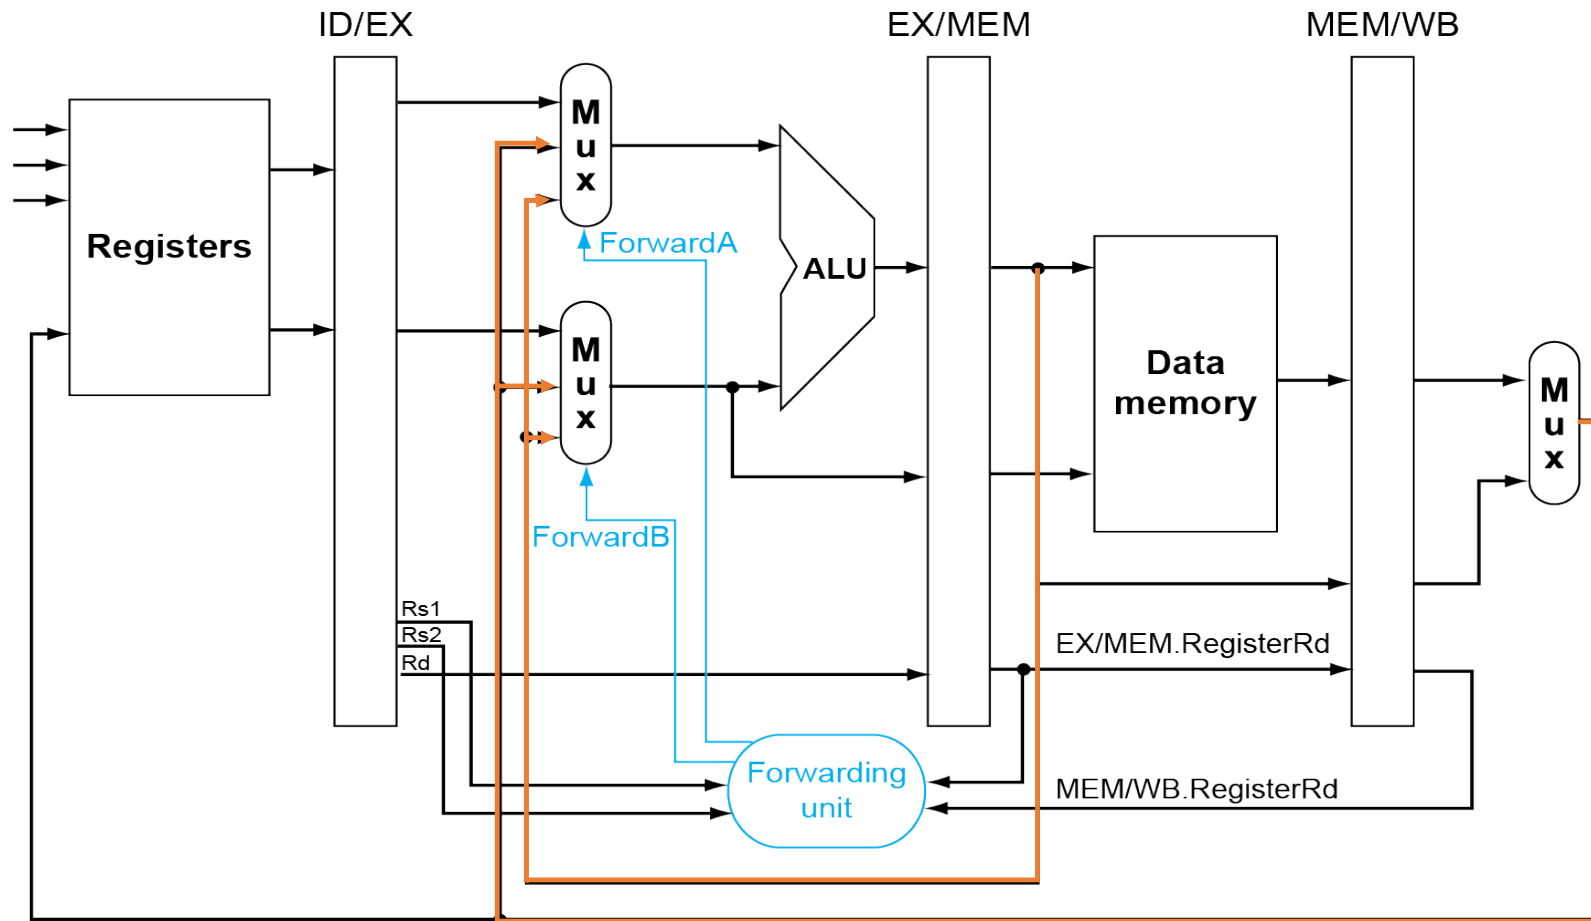
EX-hazard:
Fwd from EX/MEM
pipeline reg

Mem-hazard: Fwd from
MEM/WB
pipeline reg

Detecting the Need to Forward


- But only if forwarding instruction will write to a register!
 - EX/MEM.RegWrite, MEM/WB.RegWrite
- And only if Rd for that instruction is not x0
 - EX/MEM.RegisterRd \neq 0,
MEM/WB.RegisterRd \neq 0

Forwarding Paths




Double Data Hazard

- Consider the sequence:
 - add **x1**, x1, x2
 - add **x1**, **x1**, x3
 - add x1, **x1**, x4
- Both hazards occur
 - Want to use the most recent
- Revise MEM hazard condition
 - Only fwd if EX hazard condition isn't true



Load-Use Hazard Detection

- Load-use hazard cannot be resolved using forwarding alone

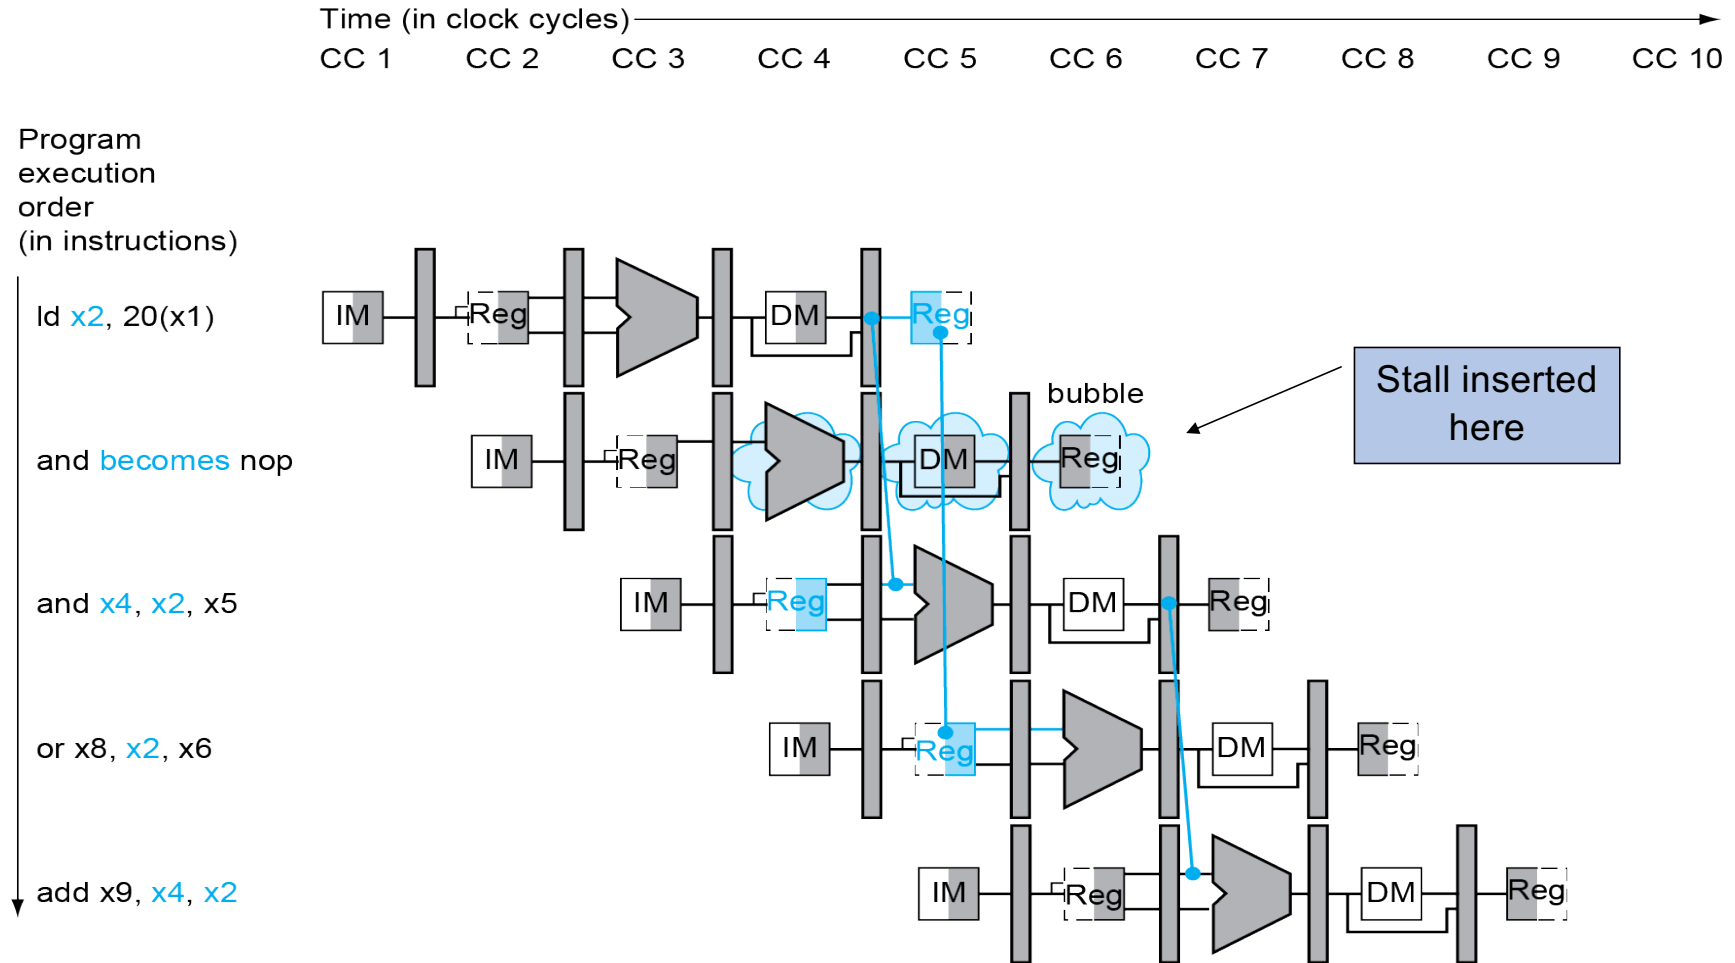


- Check load-use hazard using condition
 - $ID/EX.MemRead$ and $((ID/EX.RegisterRd = IF/ID.RegisterRs1) \text{ or } (ID/EX.RegisterRd = IF/ID.RegisterRs1))$
- If detected, stall and insert bubble

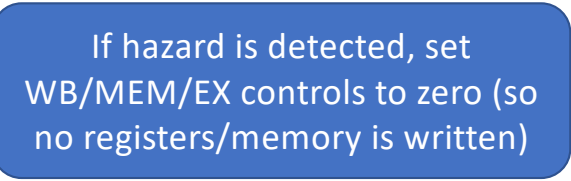
How to Stall the Pipeline

- Force control values in ID/EX register to 0
 - EX, MEM and WB do nop (no-operation)
- Prevent update of PC and IF/ID register
 - Using instruction is decoded again
 - Following instruction is fetched again
 - 1-cycle stall allows MEM to read data for 1 d
 - Can subsequently forward to EX stage

Load-Use Data Hazard



Datapath with Hazard Detection



If hazard is detected, set WB/MEM/EX controls to zero (so no registers/memory is written)

Summary

- Pipeline increases throughput by overlapping execution of multiple instructions
- Pipeline hazard
 - Structure (solution: add resources)
 - Data (solution: forwarding)
 - Control (next class)
- Pipeline stalls