

Pipelined CPU

Jinyang Li

Slides are based on Patterson and Hennessy

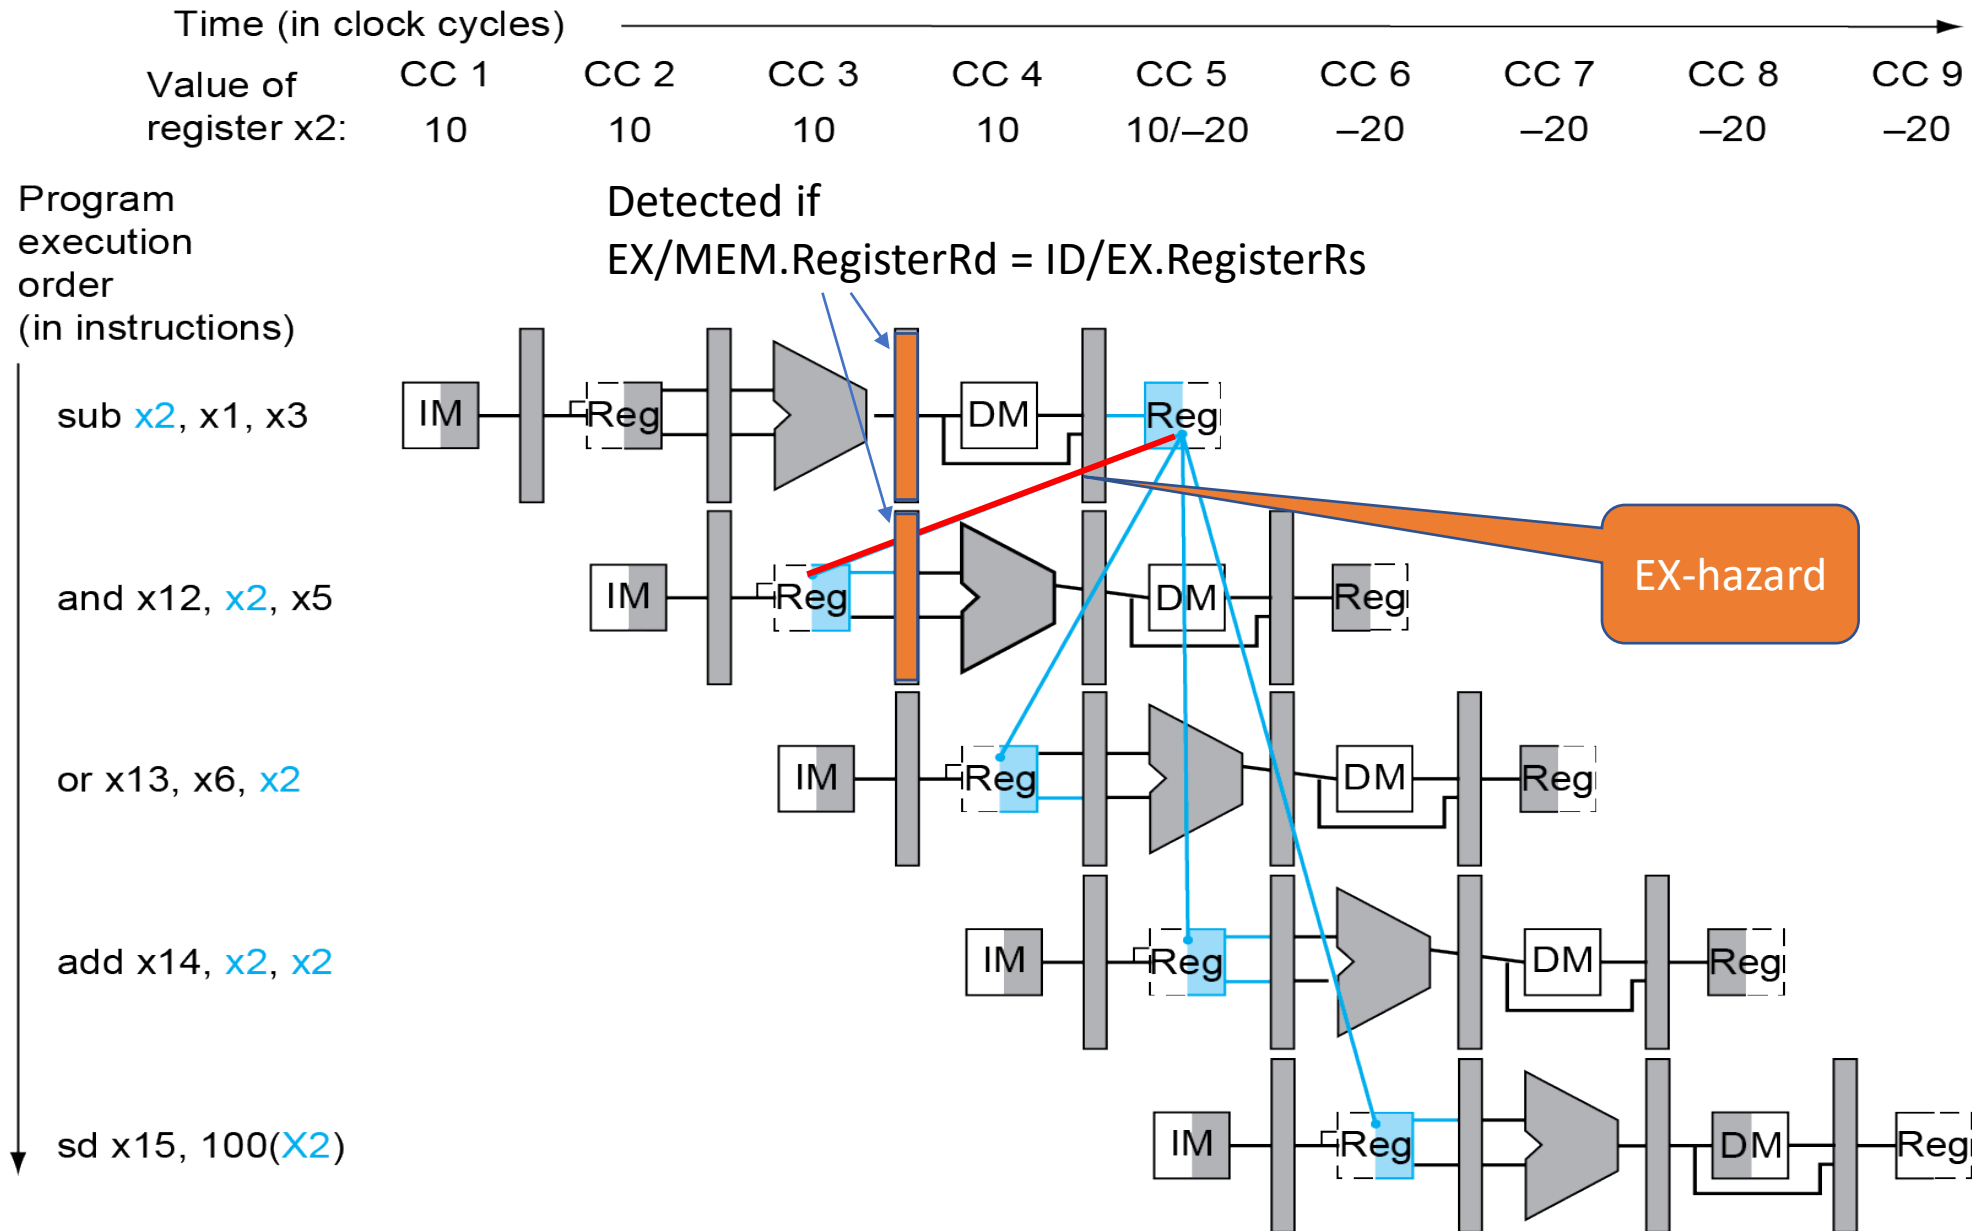
What we've learnt so far

- Single cycle RISC-V CPU design
- 5 stage pipelined RISC-V CPU
 - Why pipelining? Faster (ideal throughput speedup: #-of-stages)
- Pipelining challenges: hazards
 - Must stall (bubble) to ensure correctness
- 3 types of hazards:
 - Structure (To mitigate, add resources)
 - Data (To mitigate, ??)
 - Control (To mitigate, ??)

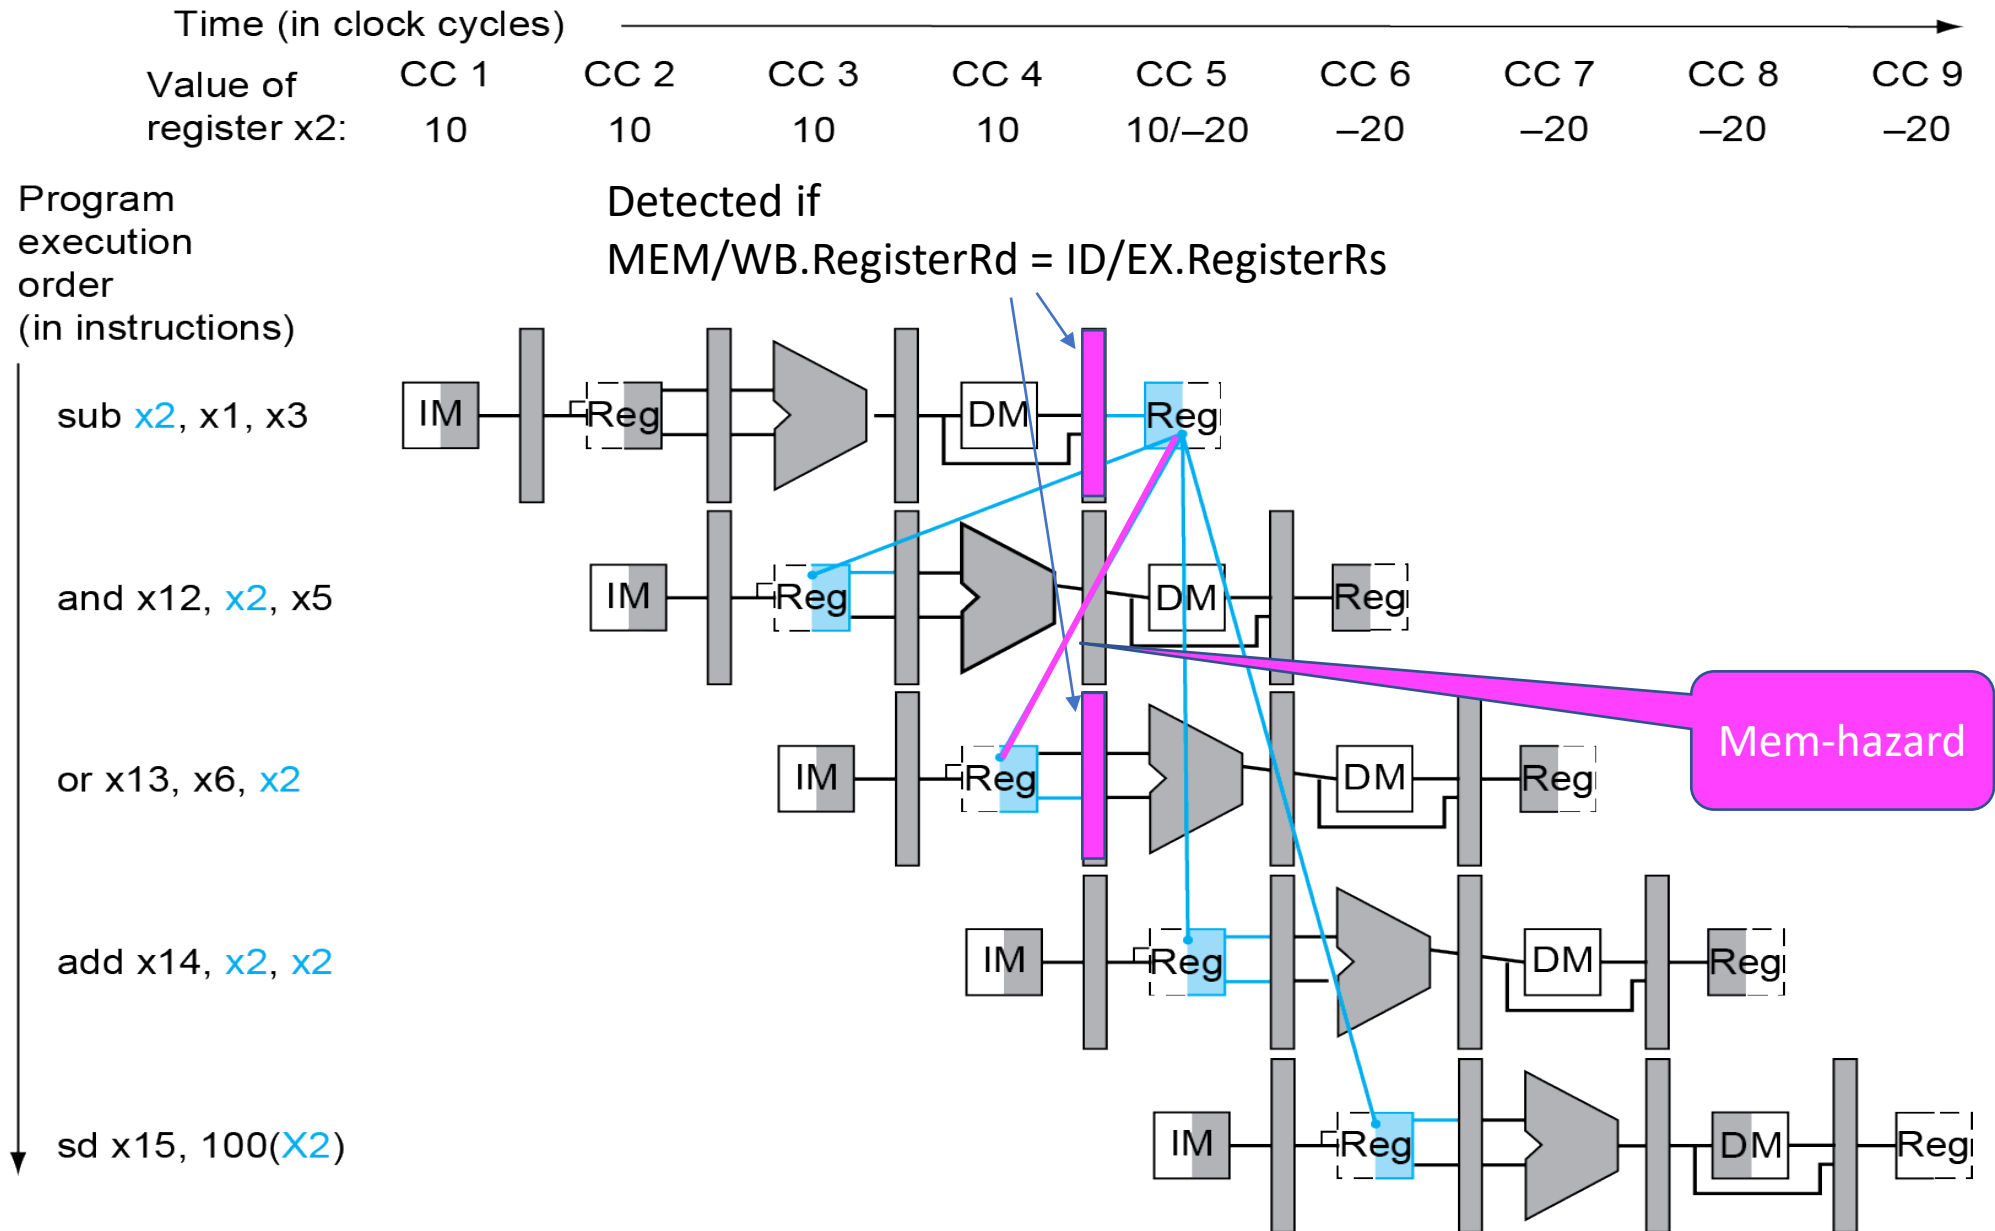
Today's lesson

- Handling data hazard
- Handling control hazard
 - Branch hazard
 - exceptions
- Briefly, more advanced topics:
 - Multiple-issue
 - Subword parallelism (SIMD)

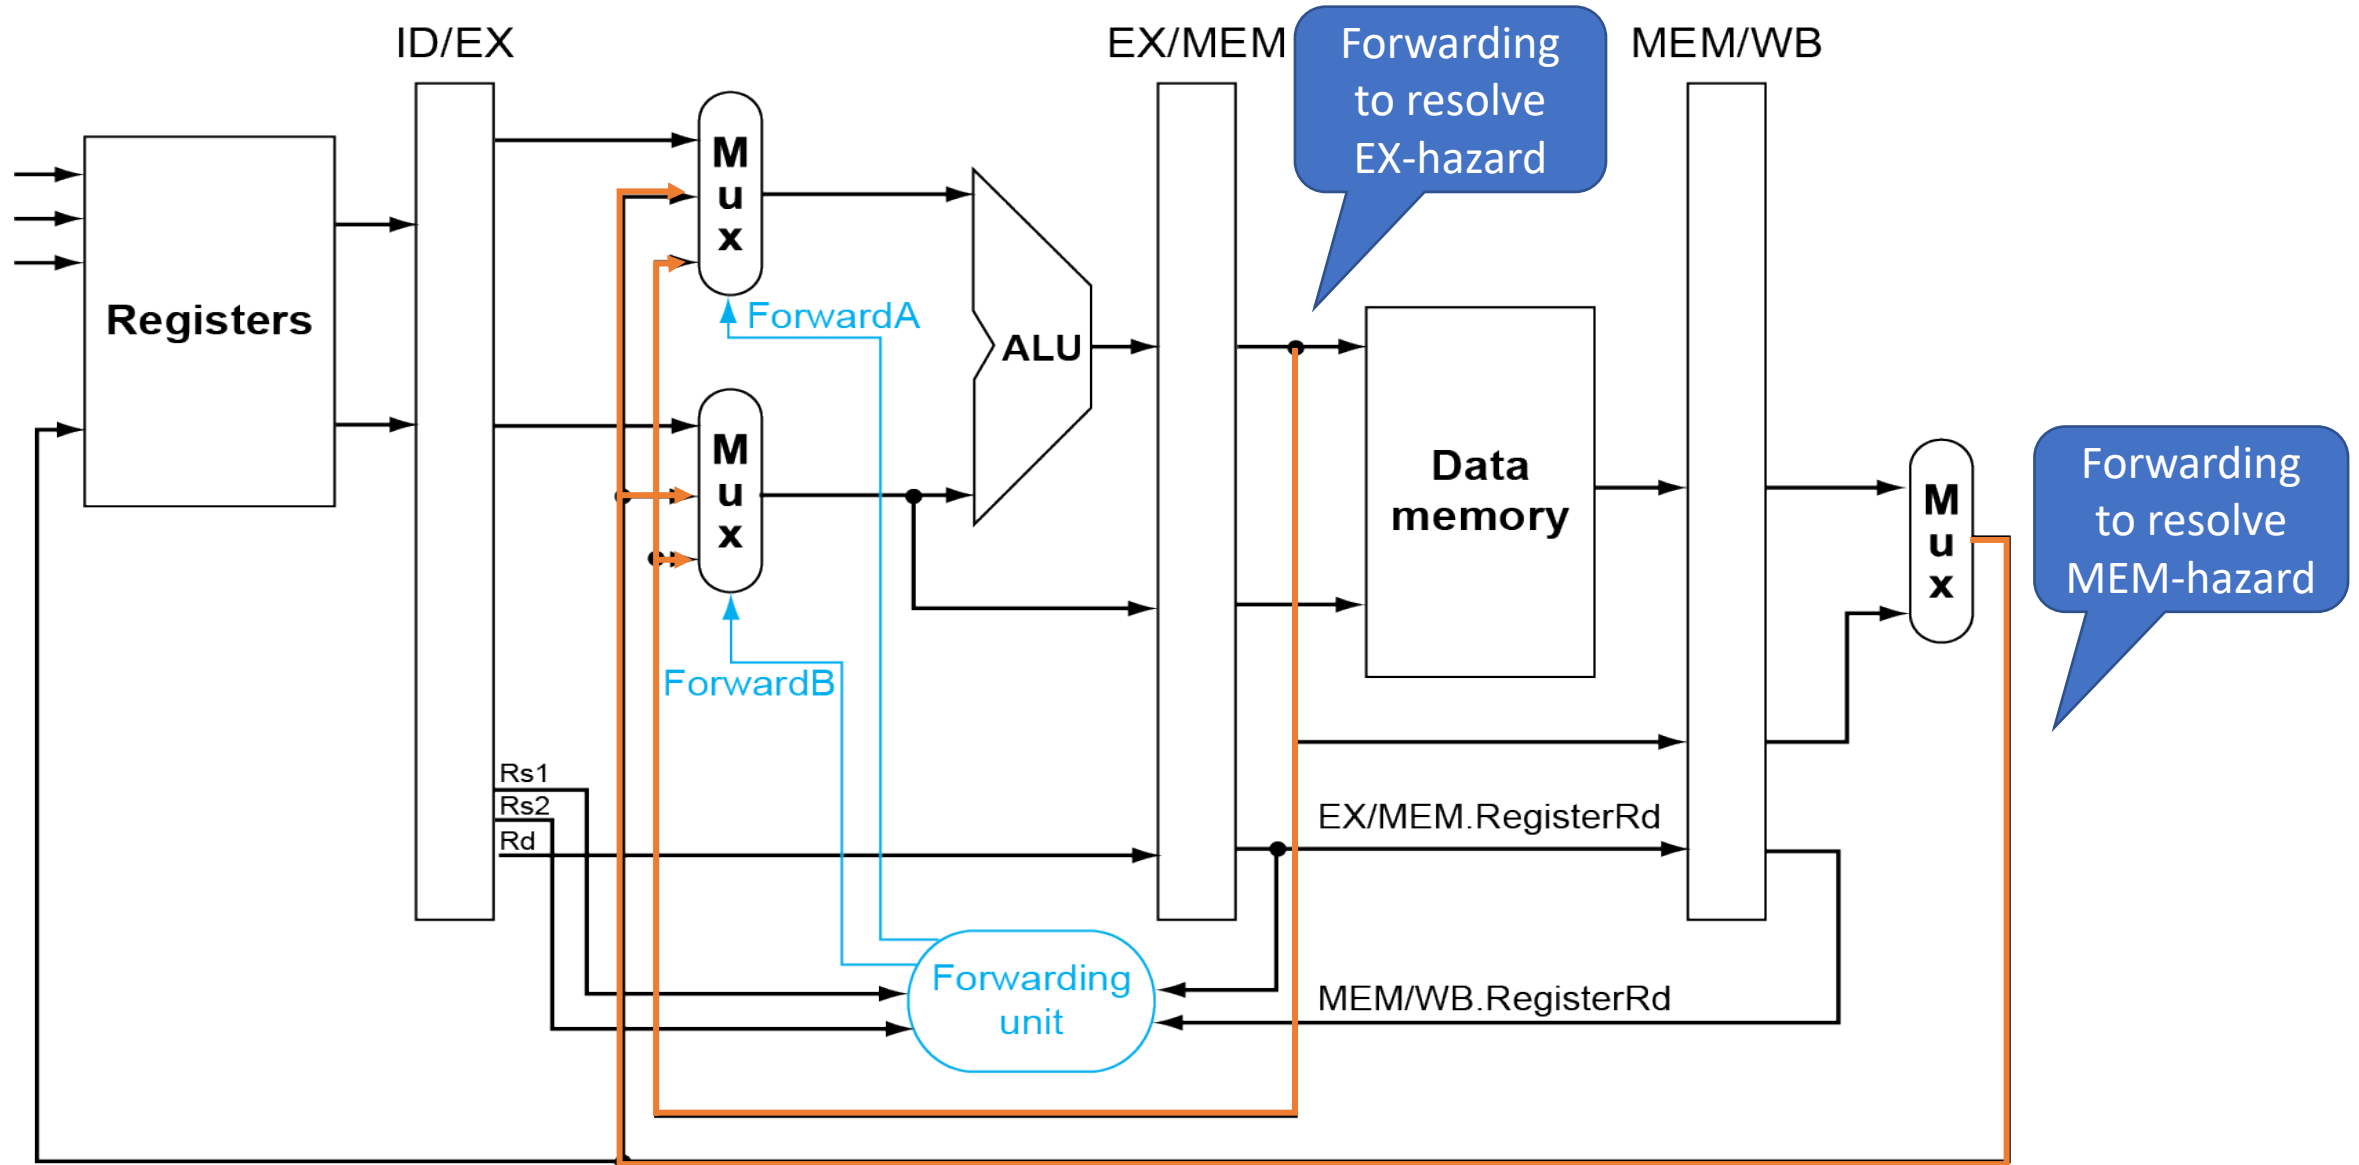
Data hazard



Data hazard

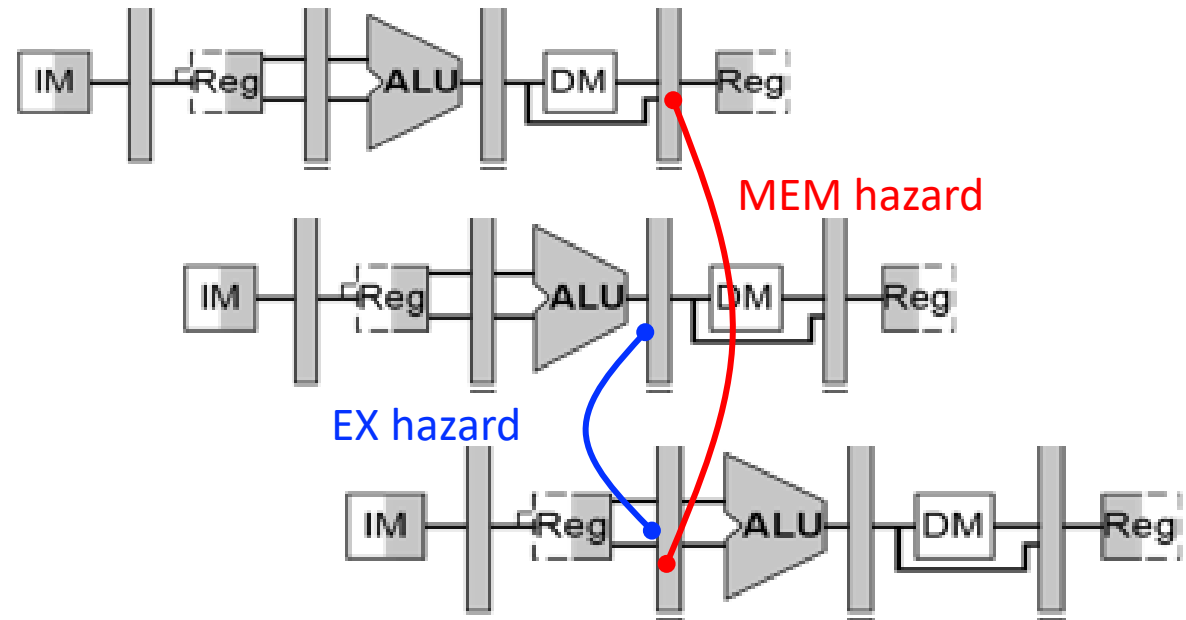


Forwarding Paths




Double Data Hazard

- Consider the sequence:
 add **x1**, x1, x2
 add **x1**, **x1**, x3
 add x1, **x1**, x4
- Both hazards occur
 - Want to use the most recent
- Revise MEM hazard condition
 - Only fwd if EX hazard condition isn't true



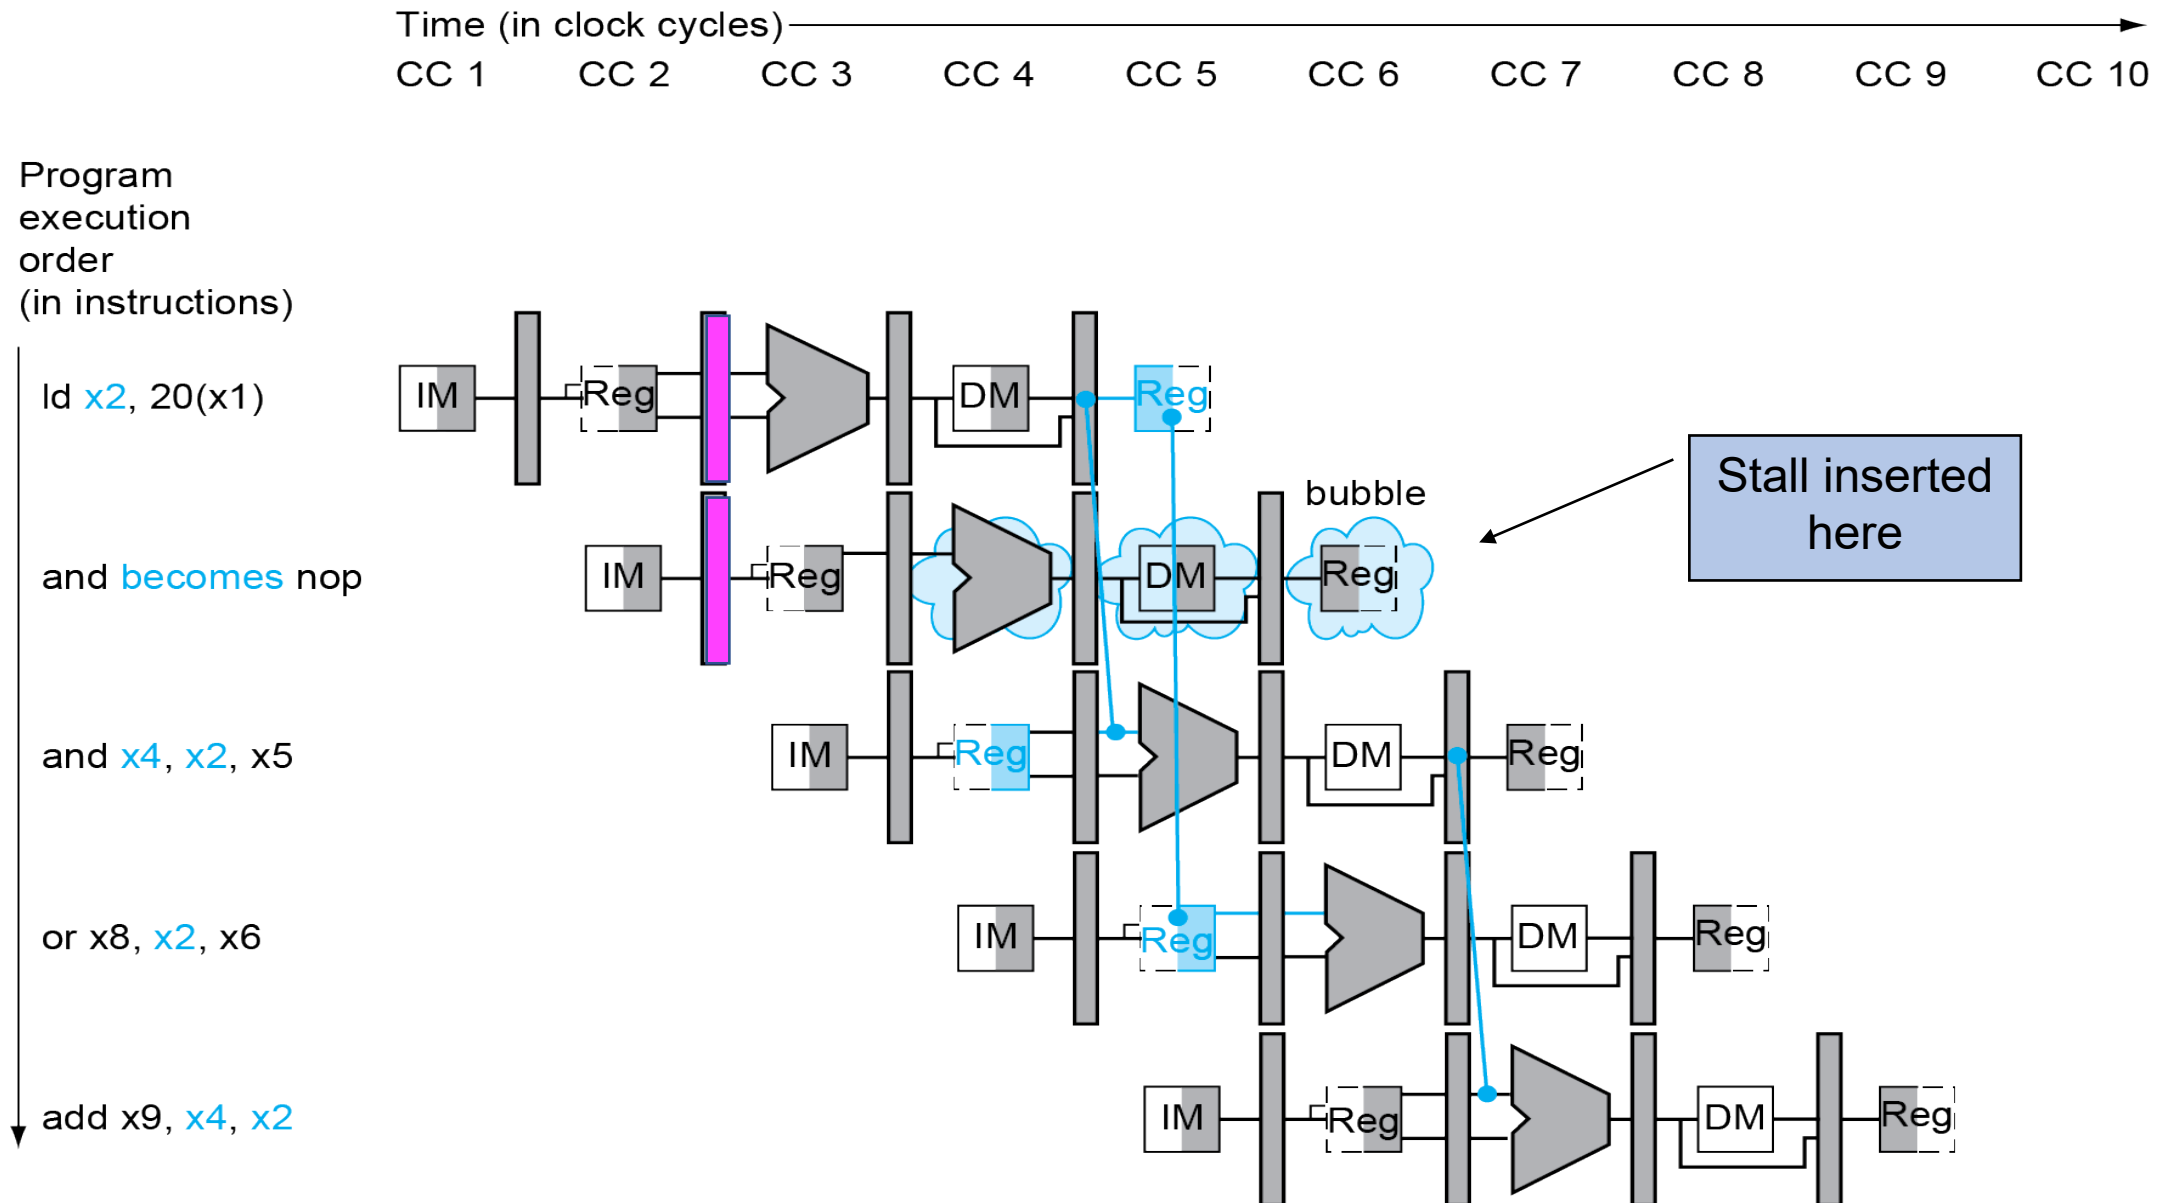
Load-Use Hazard cannot be forwarded away

- Load-use hazard cannot be resolved using forwarding alone

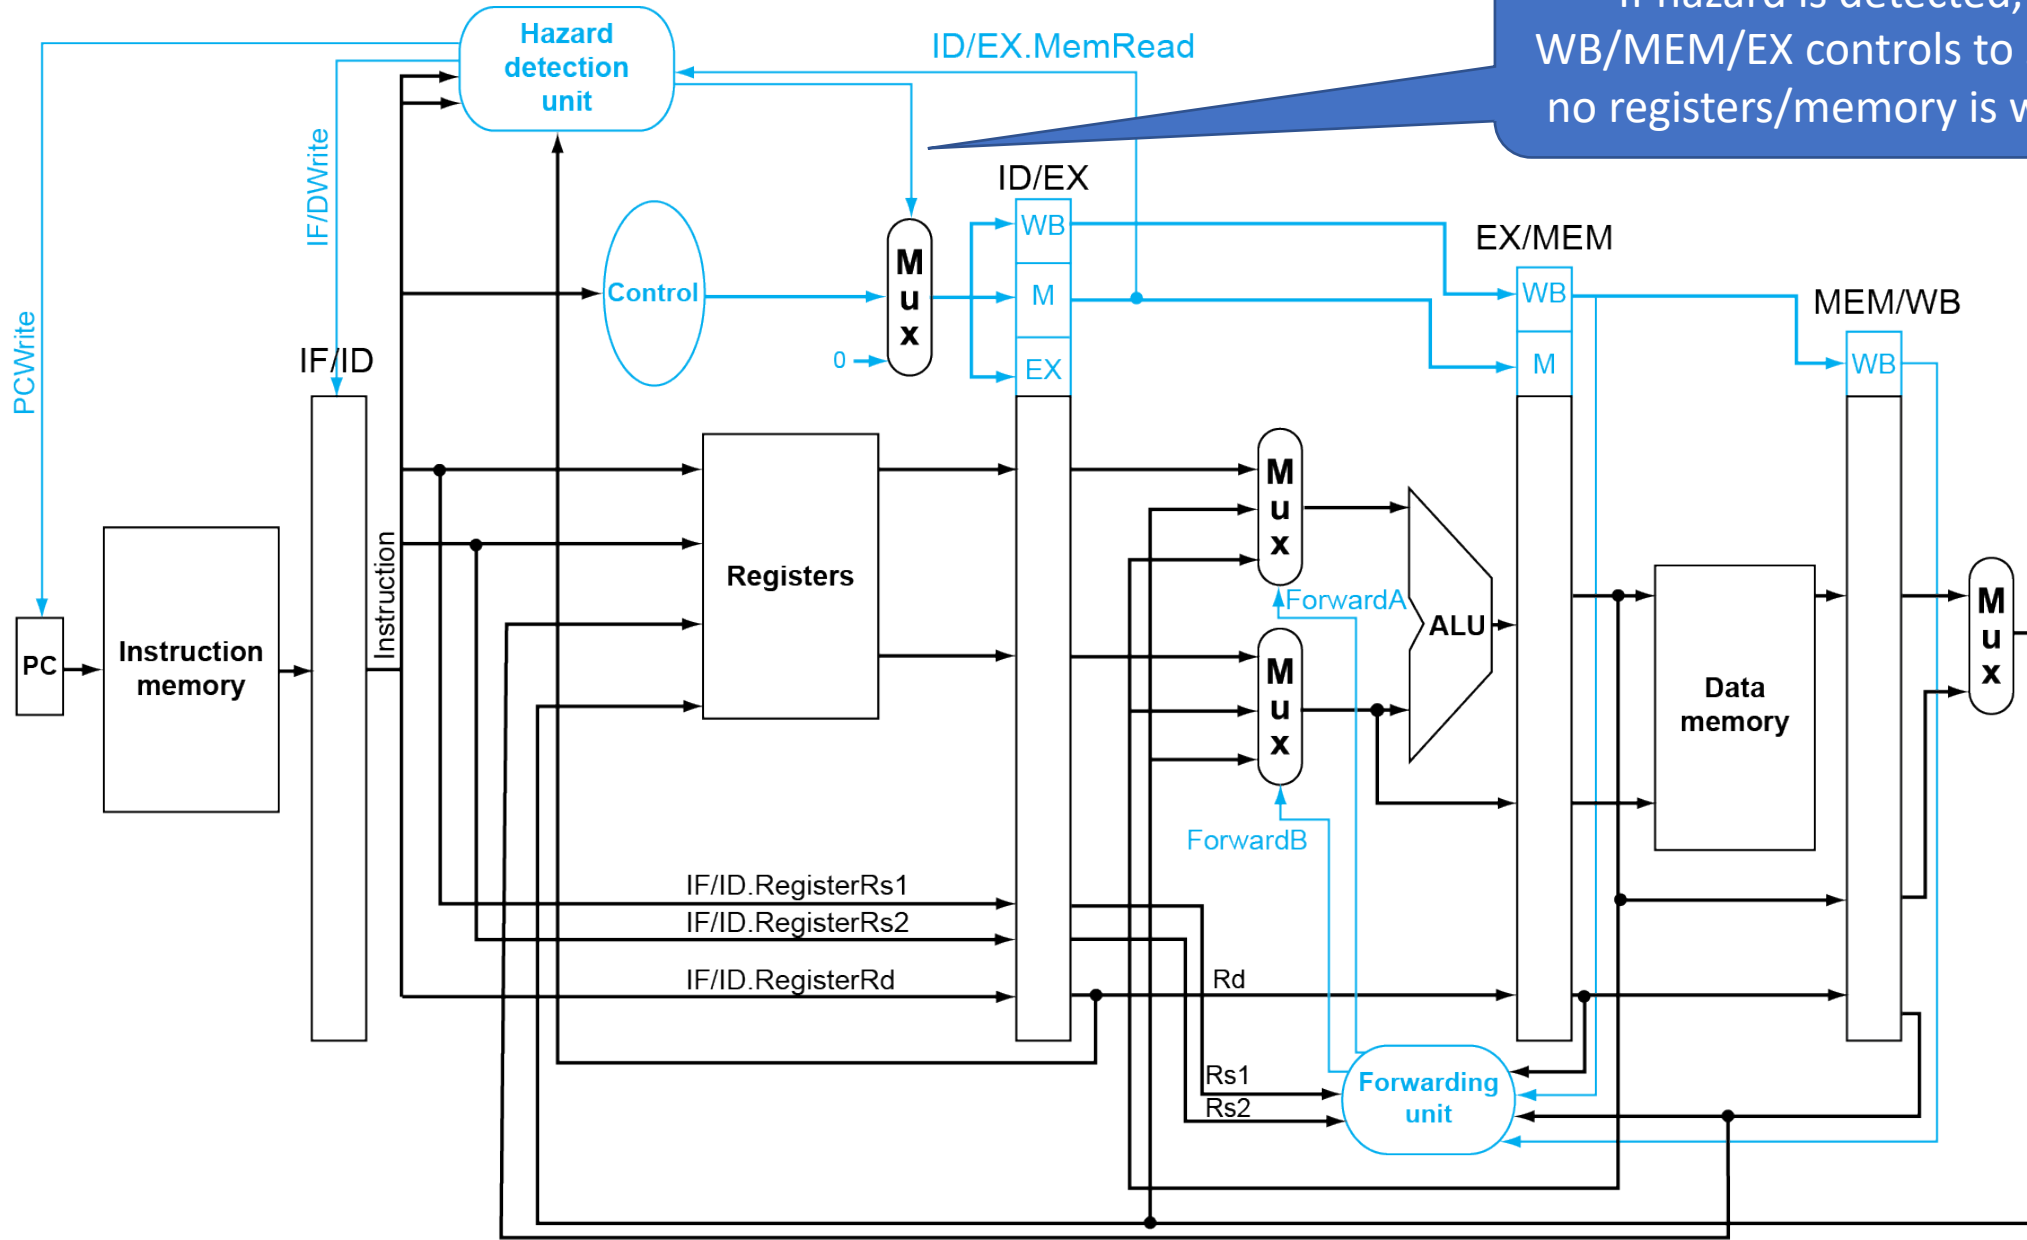


- Check load-use hazard using condition
 - $ID/EX.MemRead$ and $((ID/EX.RegisterRd = IF/ID.RegisterRs1) \text{ or } (ID/EX.RegisterRd = IF/ID.RegisterRs2))$
- If detected, stall and insert bubble

Stalling upon detection of load-use data hazard

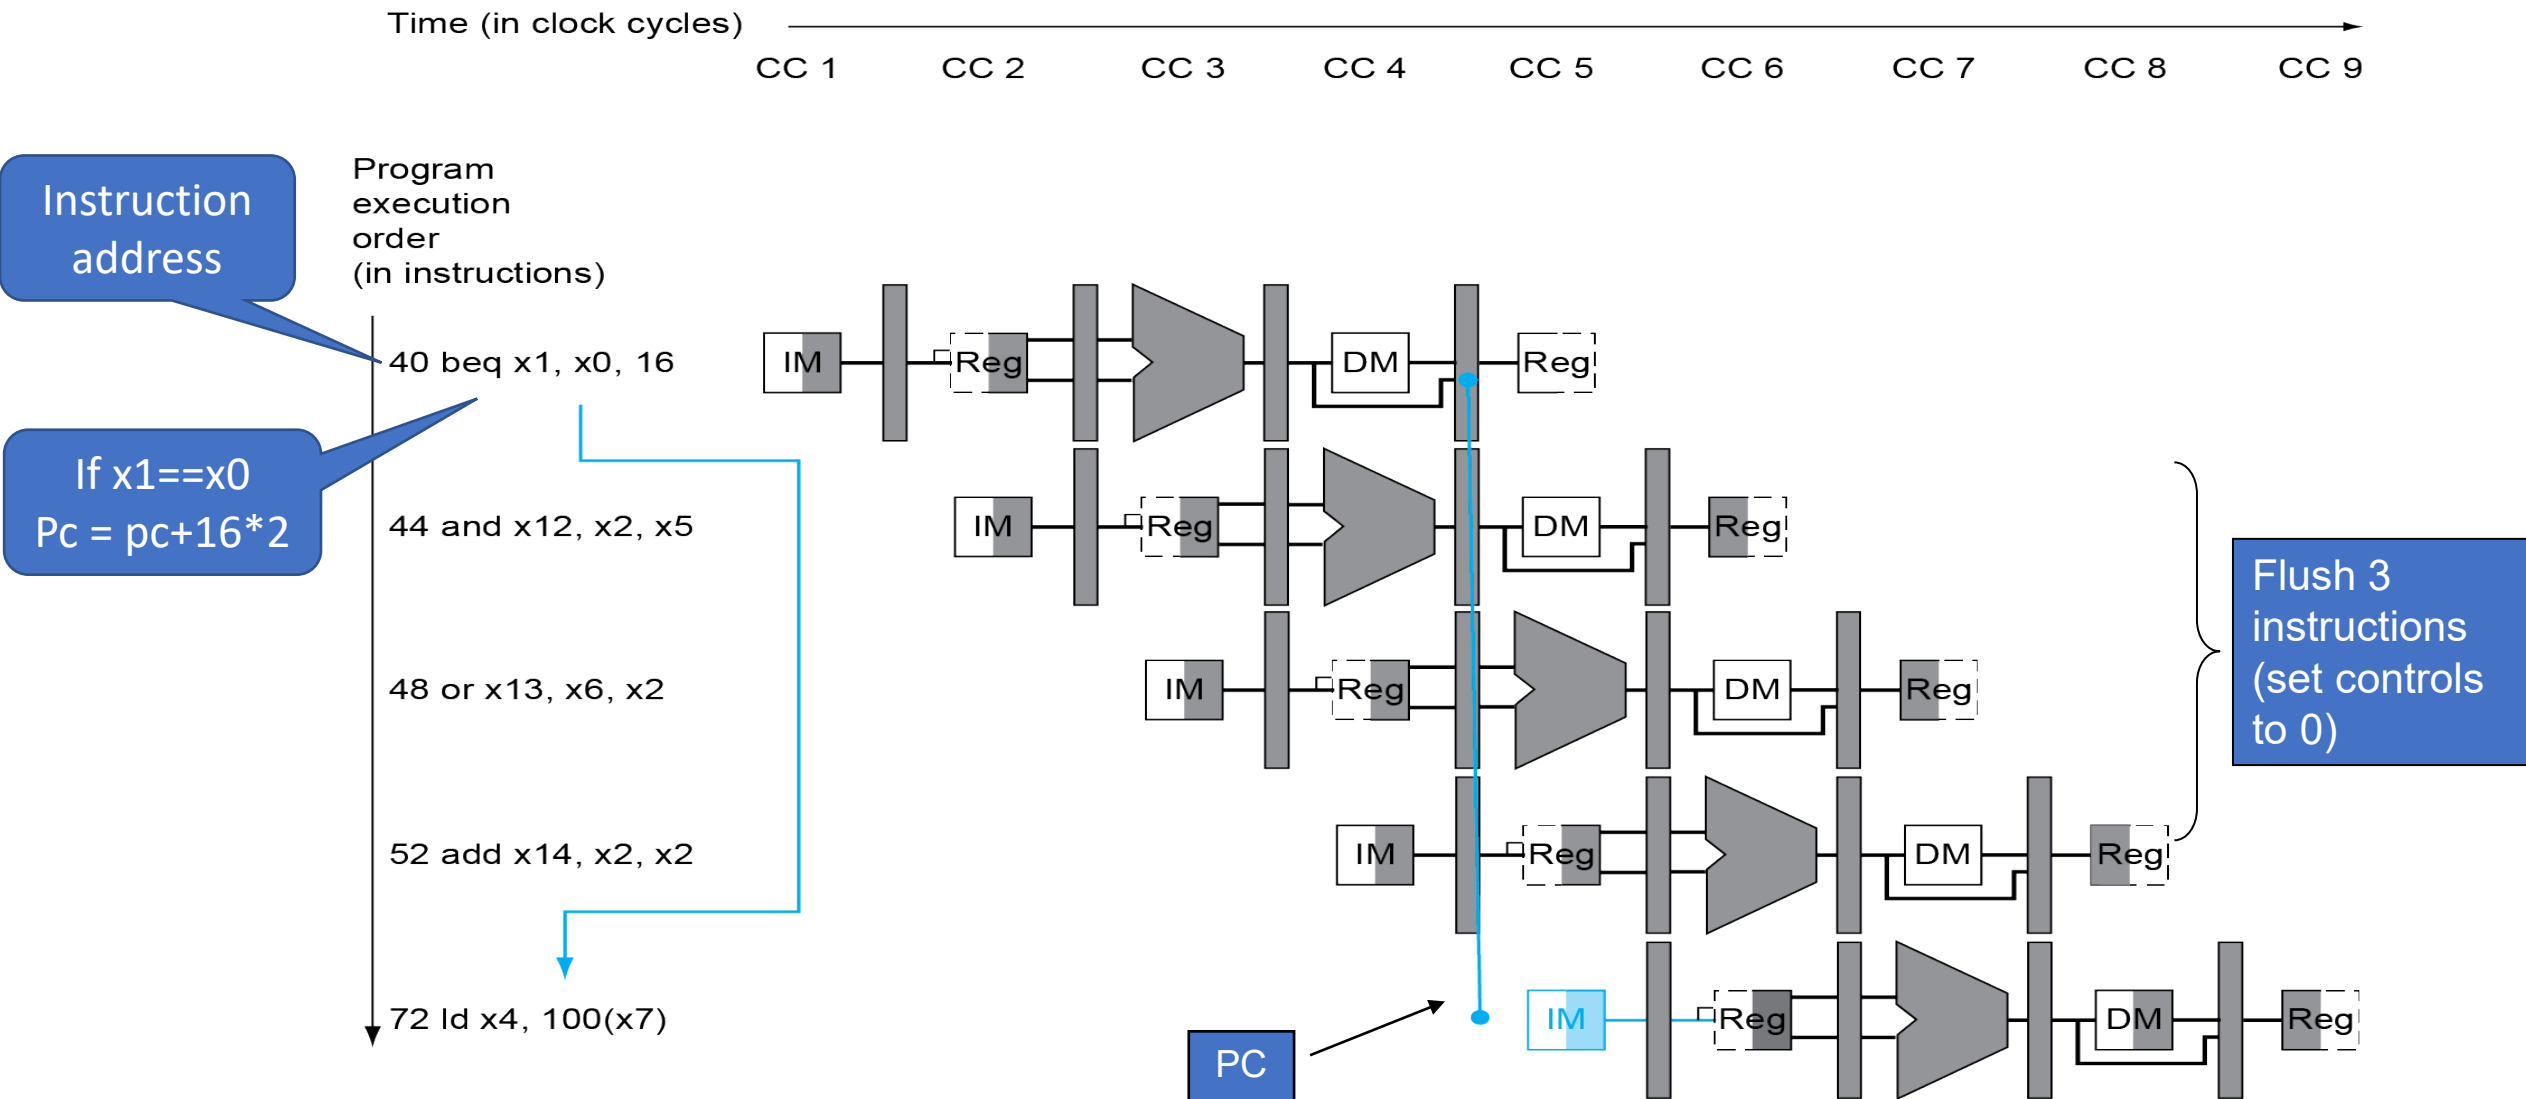


Datapath with Hazard Detection



If hazard is detected, set WB/MEM/EX controls to zero (so no registers/memory is written)

Branch Hazard: if outcome is determined in MEM...

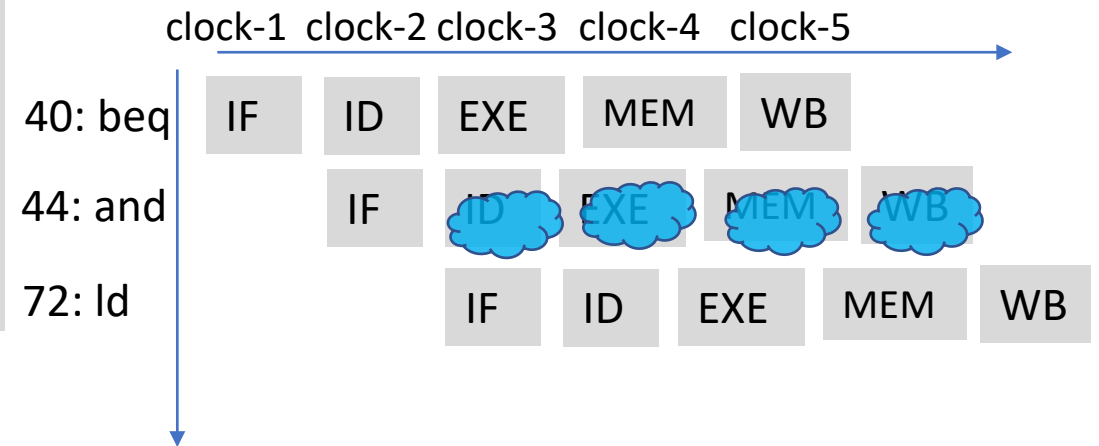


Reduce Branch Delay


- Add hardware to determine branch outcome earlier (e.g. in ID instead of MEM) → fewer instructions to flush

```
36:  sub  x10, x4, x8
40:  beq  x1,  x3, 16    // PC-relative branch
                        // to 40+16*2=72
44:  and  x12, x2, x5
48:  orr  x13, x2, x6
52:  add  x14, x4, x2
56:  sub  x15, x6, x7
...
72:  ld   x4, 50(x7)
```

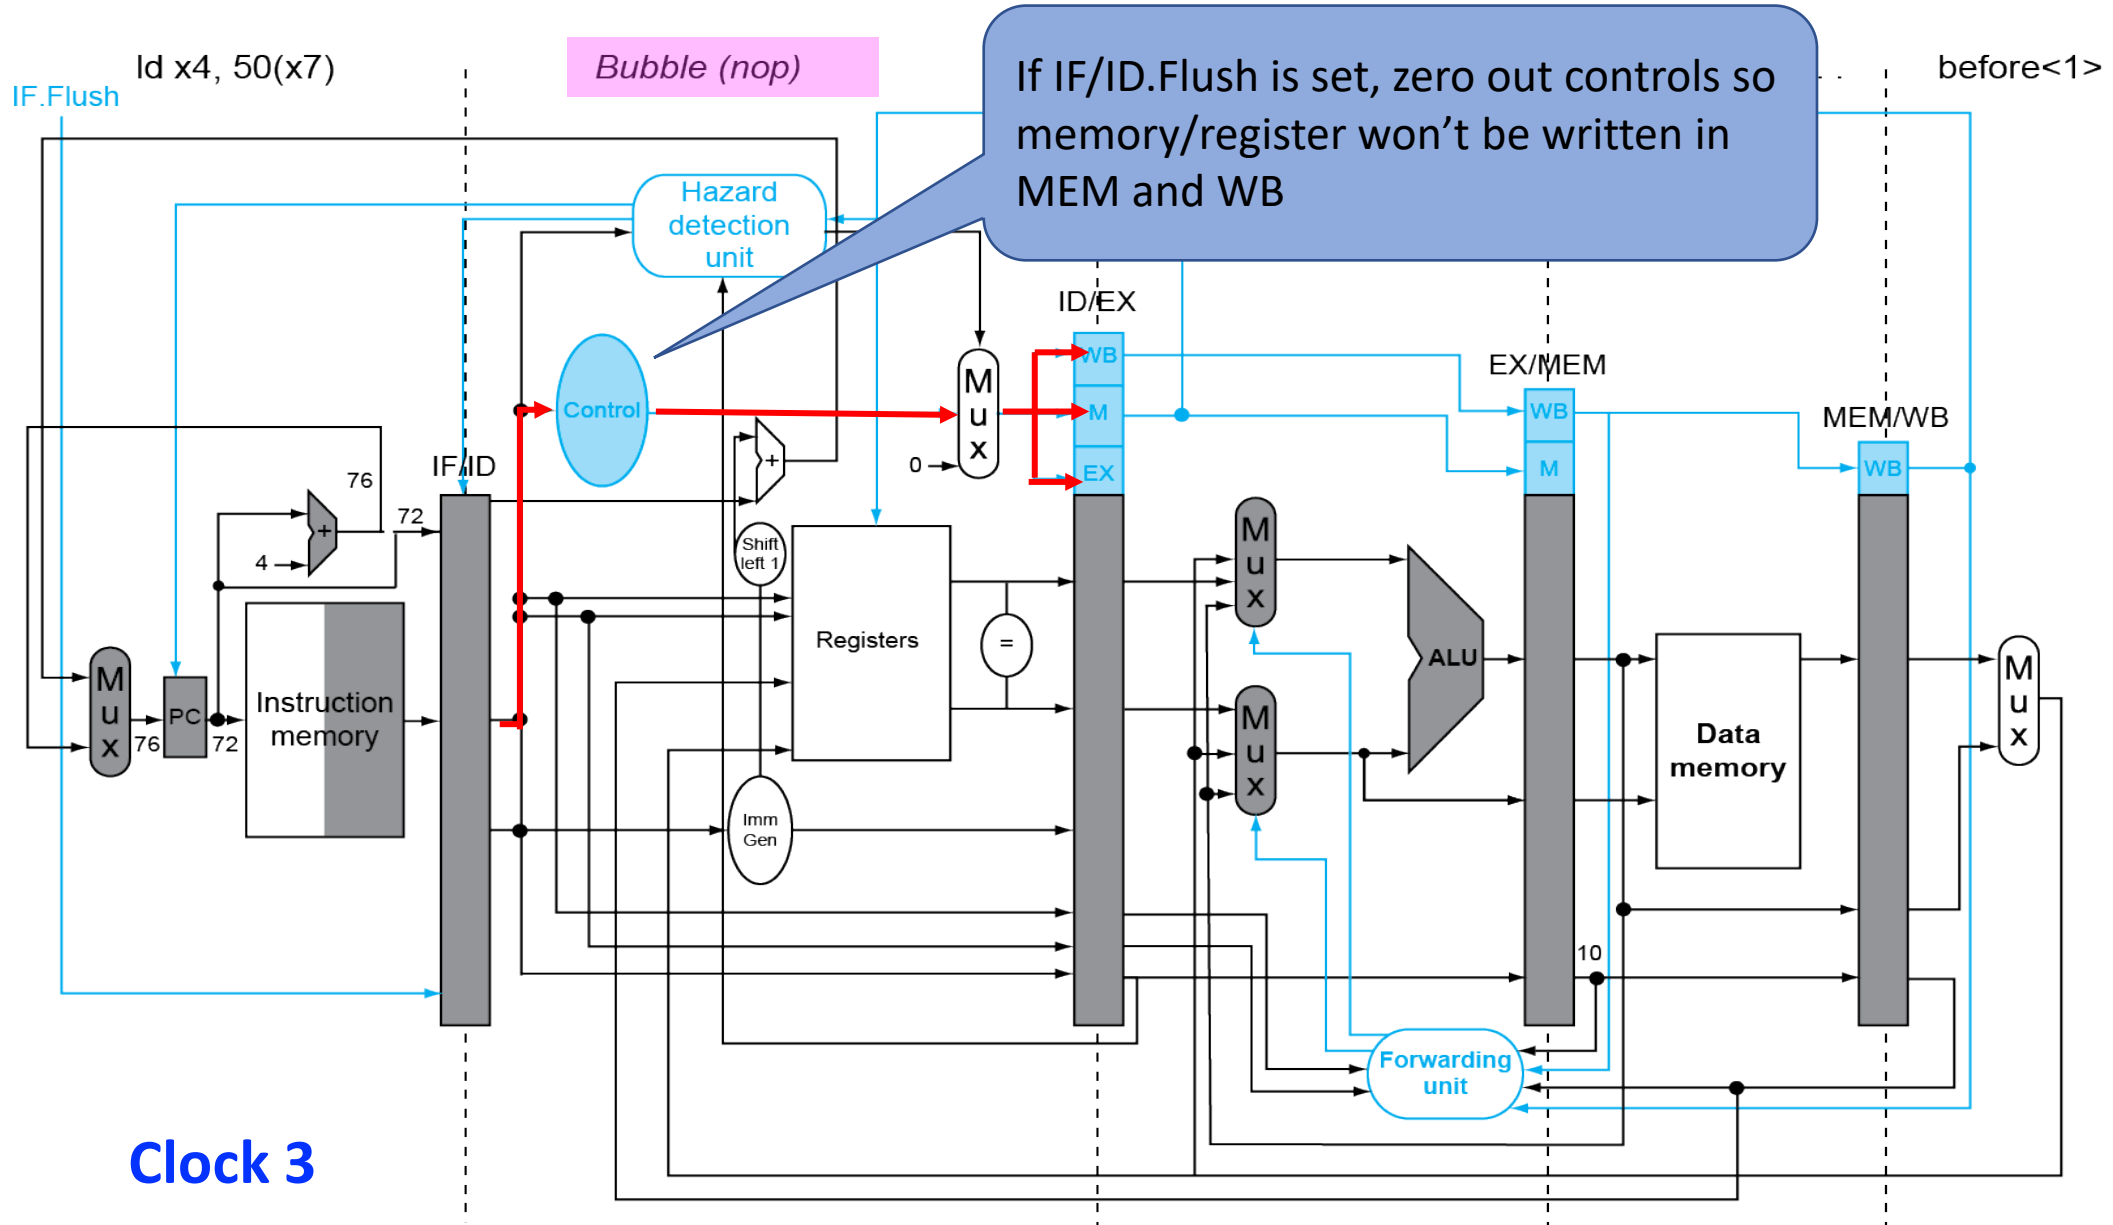
How many instructions to flush
if branch outcome is known in ID?



Branch determined in ID and is taken



Branch determined in ID and is taken




Dynamic Branch Prediction

- Our simple 5-stage pipeline's branch penalty is 1 bubble, but
 - In deeper pipelines, branch penalty is more significant
- Solution: dynamic prediction
 - Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
 - To execute a branch
 - Check table, expect the same outcome
 - Start fetching from fall-through or target
 - If wrong, flush pipeline and flip prediction

1-Bit Predictor: Shortcoming


- Inner loop branches mispredicted twice!



- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around

2-Bit Predictor

- Only change prediction on two successive mispredictions



Calculating Branch Target (needed if branch is predicted taken)

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch
- Branch target buffer
 - Cache of target addresses
 - Indexed by PC when instruction fetched
 - If hit and instruction is branch predicted taken, can fetch target immediately

Exceptions and Interrupts

- “Unexpected” events requiring change in flow of control
- Exception
 - Arises within the CPU
 - e.g., undefined opcode, syscall, memory permission error, divide by zero, ...
- Interrupt
 - From an external source like I/O controllers
 - E.g. keyboard stroke, packet arrival
- Dealing with them without sacrificing performance is hard

Handling Exceptions

- Save PC of offending (or interrupted) instruction
 - In RISC-V: Supervisor Exception Program Counter (SEPC)
- Save indication of the problem
 - In RISC-V: Supervisor Exception Cause Register (SCAUSE)
- Jump to handler (aka code in OS kernel)
 - Assume at 0000 0000 1C09 0000_{hex}
- OS's handler does either of the following:
 - Take corrective action; restart offending instruction
 - Terminate program; report cause

An Alternate Exception Handling Mechanism

- Vectored Exception/Interrupts (x86)
- Vector index (interrupt descriptor) specifies cause of exception
 - E.g. in x86, 0 (divide by zero), 6 (invalid opcode), ...
- OS sets up the Interrupt Descriptor Table
 - IDT[i] contains handler address, where i is vector index.

Exceptions in a Pipeline

- Another form of control hazard
- Consider malfunction on add in EX stage
 `add x1, x2, x1`
- What must be done?
 - Prevent x1 from being clobbered
 - Complete previous instructions
 - Flush `add` and subsequent instructions
 - Save PC (of offending instruction) in SEPC and set SCAUSE with cause
 - Transfer control to handler
- Similar to mispredicted branch
 - Use much of the same hardware

Exception Example

- Exception on **add** in

```
40      sub    x11, x2, x4
44      and    x12, x2, x5
48      orr    x13, x2, x6
4c      add    x1,  x2, x1
50      sub    x15, x6, x7
54      ldr    x16, 100(x7)
```

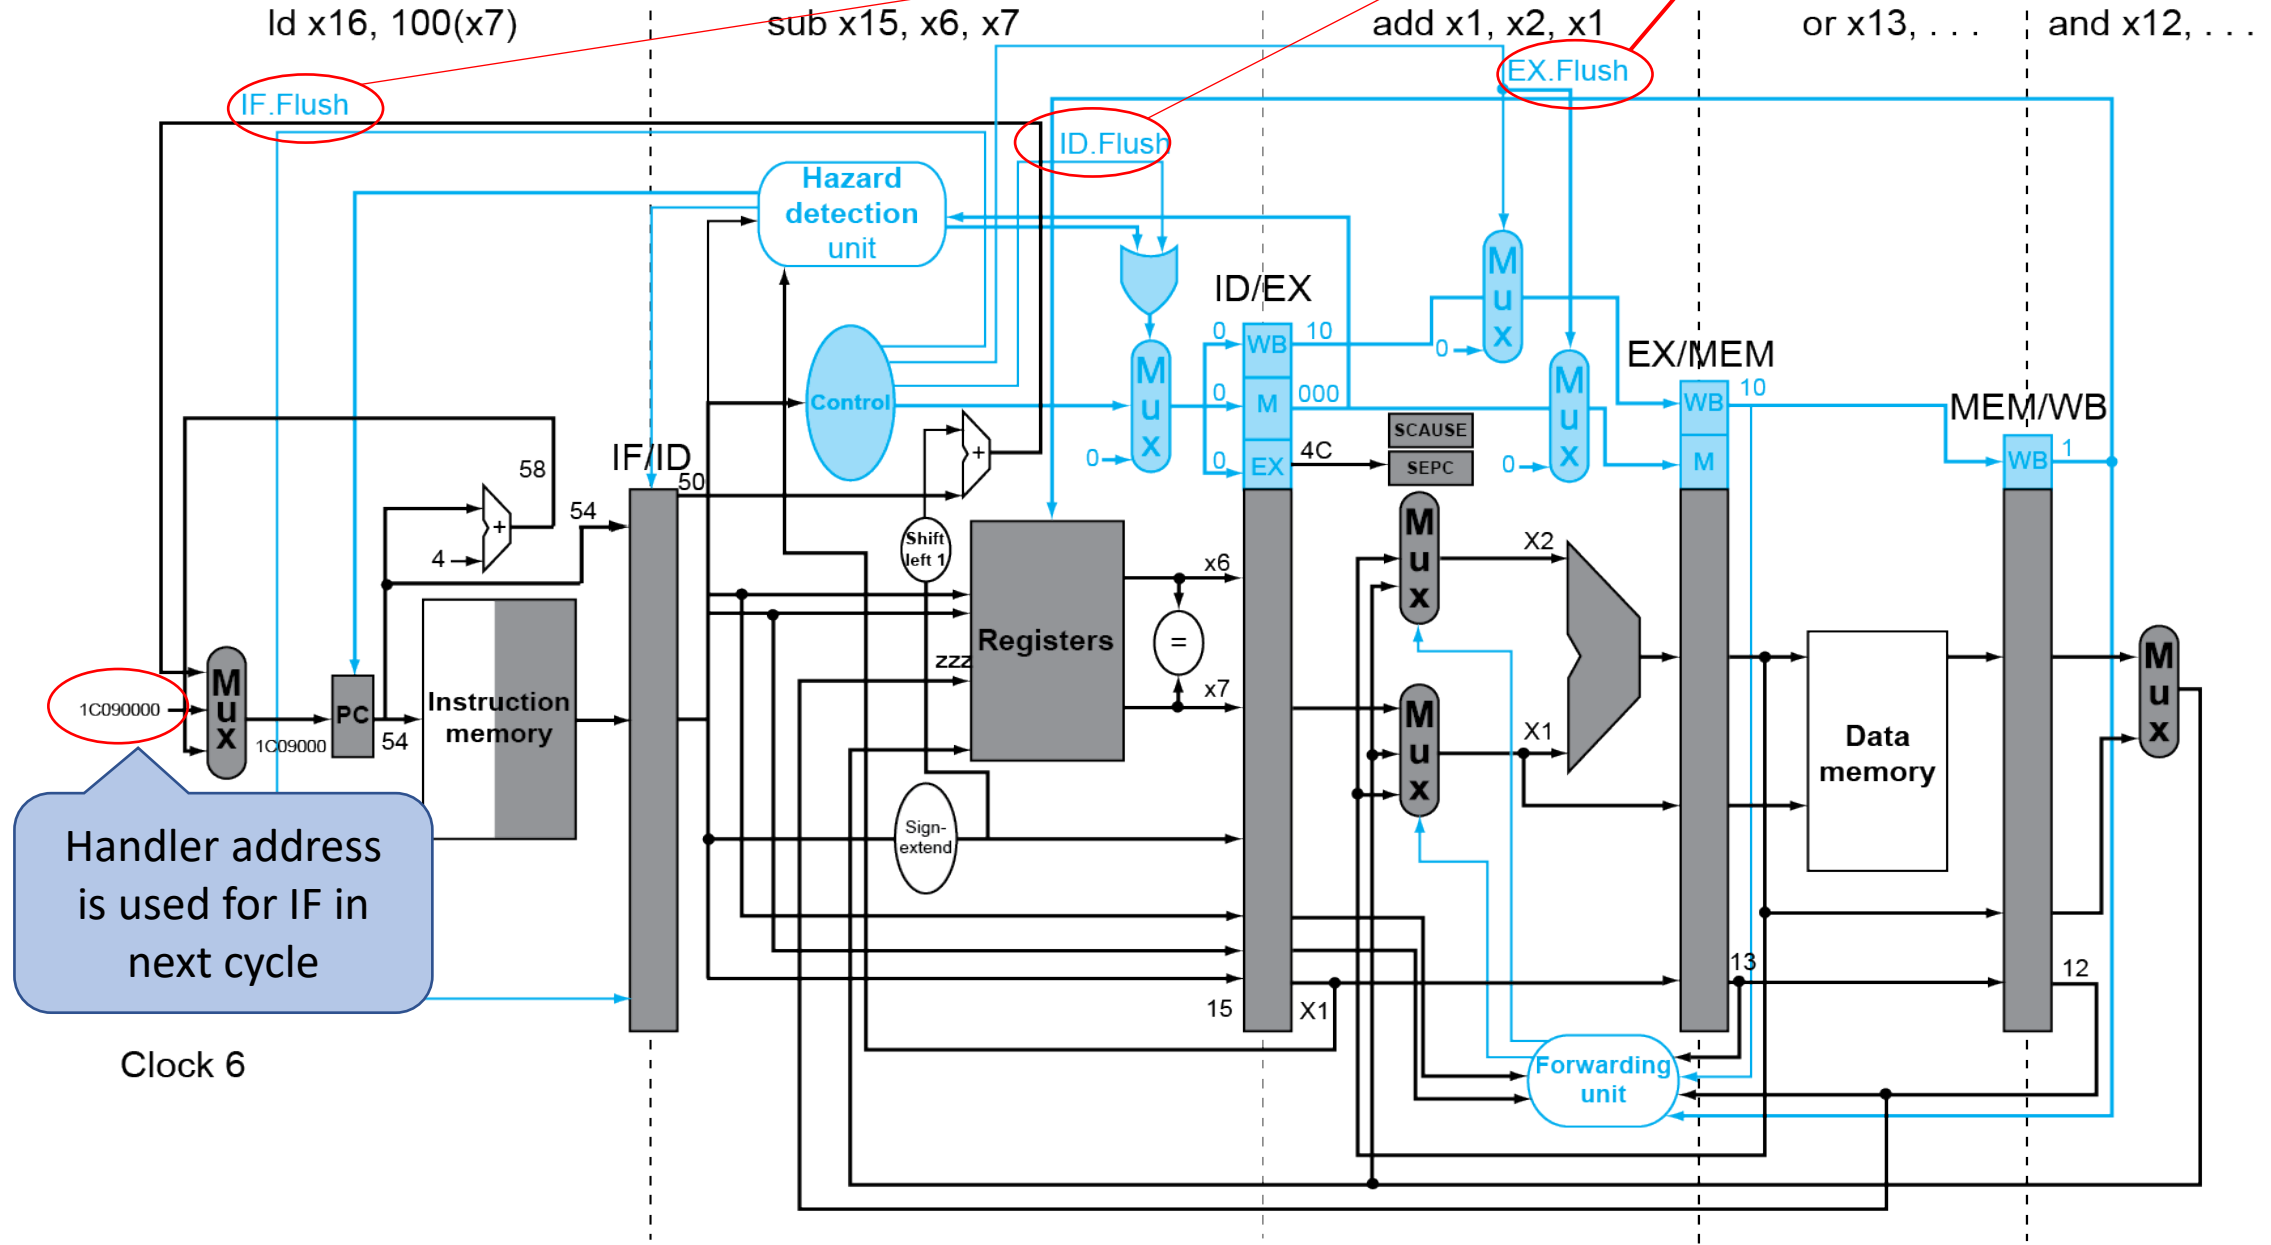
...

- Handler

```
1c090000    sd    x26, 1000(x10)
1c090004    sd    x27, 1008(x10)
```

...

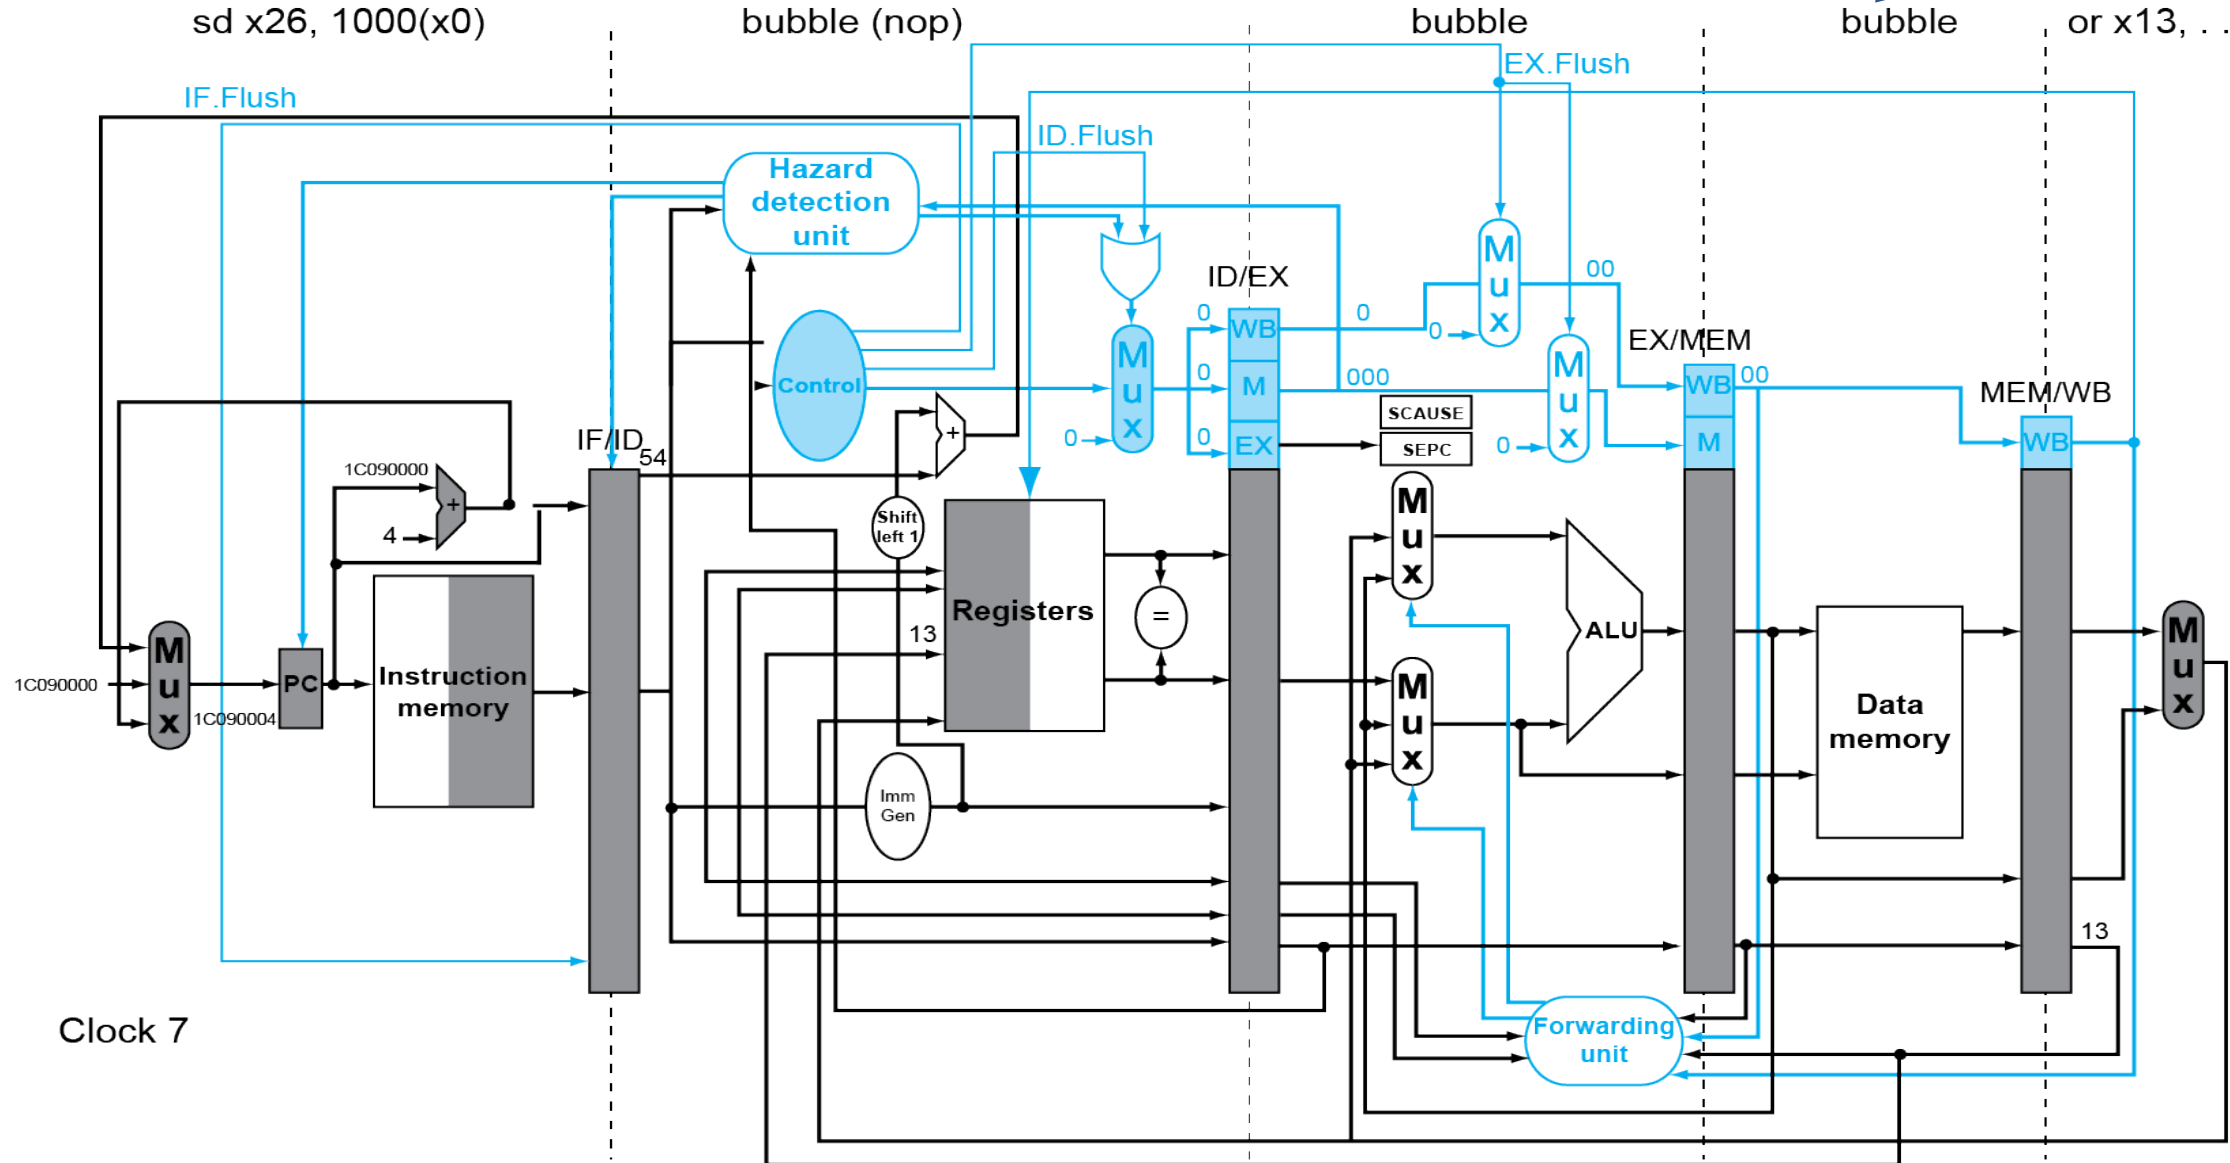
Exception Example



Exception Example

Instructions in exception handler

Offending instruction (add x1, x2, x1) becomes a bubble (so it won't clobber registers)



Today's lesson

- Handling data hazard
- Handling control hazard
 - Branch hazard
 - exceptions
- Briefly, more advanced topics:
 - Multiple-issue
 - Subword parallelism (SIMD)

Instruction-Level Parallelism (ILP)

- Pipelining: execute multiple instructions in parallel
- To increase ILP
 - Deeper pipeline
 - Less work per stage \Rightarrow shorter clock cycle
 - Multiple issue
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - Finish multiple Instructions Per Cycle ($IPC > 1$)
 - E.g., 4GHz 4-way multiple-issue
 - 16 billion instructions/sec, peak $IPC = 4$ ($CPI = 1/IPC = 0.25$)
 - Challenges: dependencies among multi-issued instructions
 - reduce peak IPC

Multiple Issue

- Static multiple issue (Very Long Instruction Word)
 - Compiler groups instructions to be issued together
 - Compiler detects and avoids hazards
- Dynamic multiple issue (superscalar)
 - CPU examines instruction stream and chooses instructions to issue each cycle
 - Compiler can help by reordering instructions
 - CPU resolves hazards using advanced techniques at runtime

Dynamic Pipeline Scheduling

- Allow the CPU to execute instructions out of order to avoid stalls
 - But commit result to registers in order
- Example

```
ld    x31, 20(x21)
add   x1, x31, x2
sub   x23, x23, x3
andi  x5, x23, 20
```

 - Can start sub while add is waiting for ld


Why Do Dynamic Scheduling?

- Why not just let the compiler schedule code?
- Not all stalls are predicable
 - e.g., cache misses
- Can't always schedule around branches
 - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

Complexities require power: ARM Cortex A53 and Intel i7

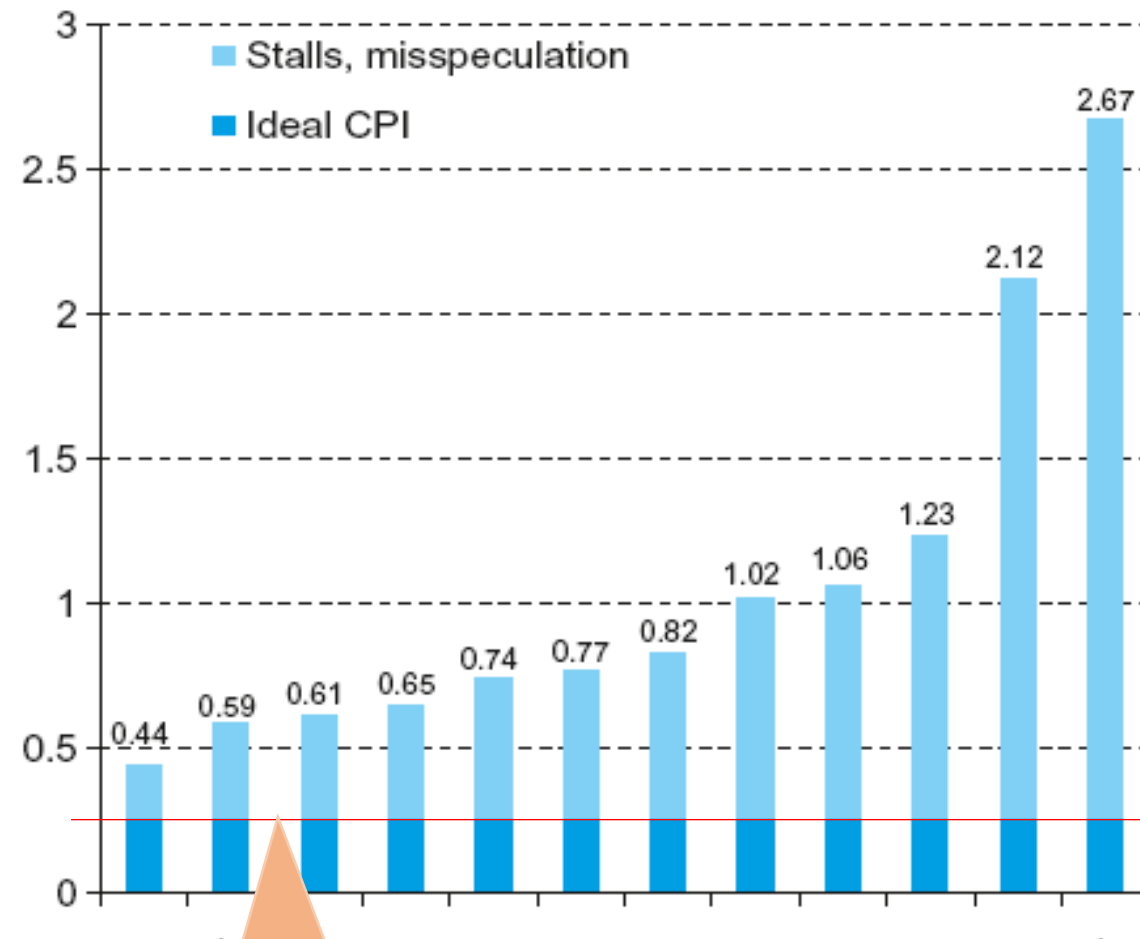
Processor	ARM A53	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	100 milliWatts (1 core @ 1 GHz)	130 Watts
Clock rate	1.5 GHz	2.66 GHz
Cores/Chip	4 (configurable)	4
Floating point?	Yes	Yes
Multiple issue?	Dynamic	Dynamic
Peak instructions/clock cycle	2	4
Pipeline stages	8	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	Hybrid	2-level
1 st level caches/core	16-64 KiB I, 16-64 KiB D	32 KiB I, 32 KiB D
2 nd level caches/core	128-2048 KiB	256 KiB (per core)
3 rd level caches (shared)	(platform dependent)	2-8 MB

Real world processor performance



ARM Cortex-A53

Ideal CPI=0.5
because of 2
issue



Ideal CPI=0.25
because of 4
issue


Intel core i7

Ways to improve performance: subword parallelism

- ML/graphics applications perform same operations on vectors
 - Partition a 128-bit adder into:
 - Sixteen 8-bit adds
 - Eight 16-bit adds
 - Four 32-bit adds
- Also called data-level parallelism, vector parallelism, or Single Instruction, Multiple Data (SIMD)
- Intel's AVX introduces 256-bit floating point registers
 - 8 single-precision ops
 - 4 double-precision ops


Example: unoptimized matrix multiplication

```
1. void dgemm (int n, double* A, double* B, double* C)
2. {
3.   for (int i = 0; i < n; ++i)
4.     for (int j = 0; j < n; ++j)
5.     {
6.       double cij = C[i+j*n];
7.       for(int k = 0; k < n; k++ )
8.         cij += A[i+k*n] * B[k+j*n];
9.       C[i+j*n] = cij;
10.    }
11. }
```




Matrix Multiply: using AVX SIMD

```
1. #include <x86intrin.h>
2. void dgemm (int n, double* A, double* B, double* C)
3. {
4.   for ( int i = 0; i < n; i+=4 )
5.     for ( int j = 0; j < n; j++ ) {
6.       __m256d c0 = _mm256_load_pd(C+i+j*n); /* c0 = C[i][j] */
7.       for( int k = 0; k < n; k++ )
8.         c0 = _mm256_add_pd(c0, /* c0 += A[i][k]*B[k][j] */
9.                           _mm256_mul_pd(_mm256_load_pd(A+i+k*n),
10.                                          _mm256_broadcast_sd(B+k+j*n)));
11.       _mm256_store_pd(C+i+j*n, c0); /* C[i][j] = c0 */
12.     }
13. }
```



Performance Impact



Summary on CPU design

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
 - Dependencies limit achievable parallelism
 - Complexity leads to the power wall