

Pipelined CPU

Jinyang Li

Slides are based on Patterson and Hennessy

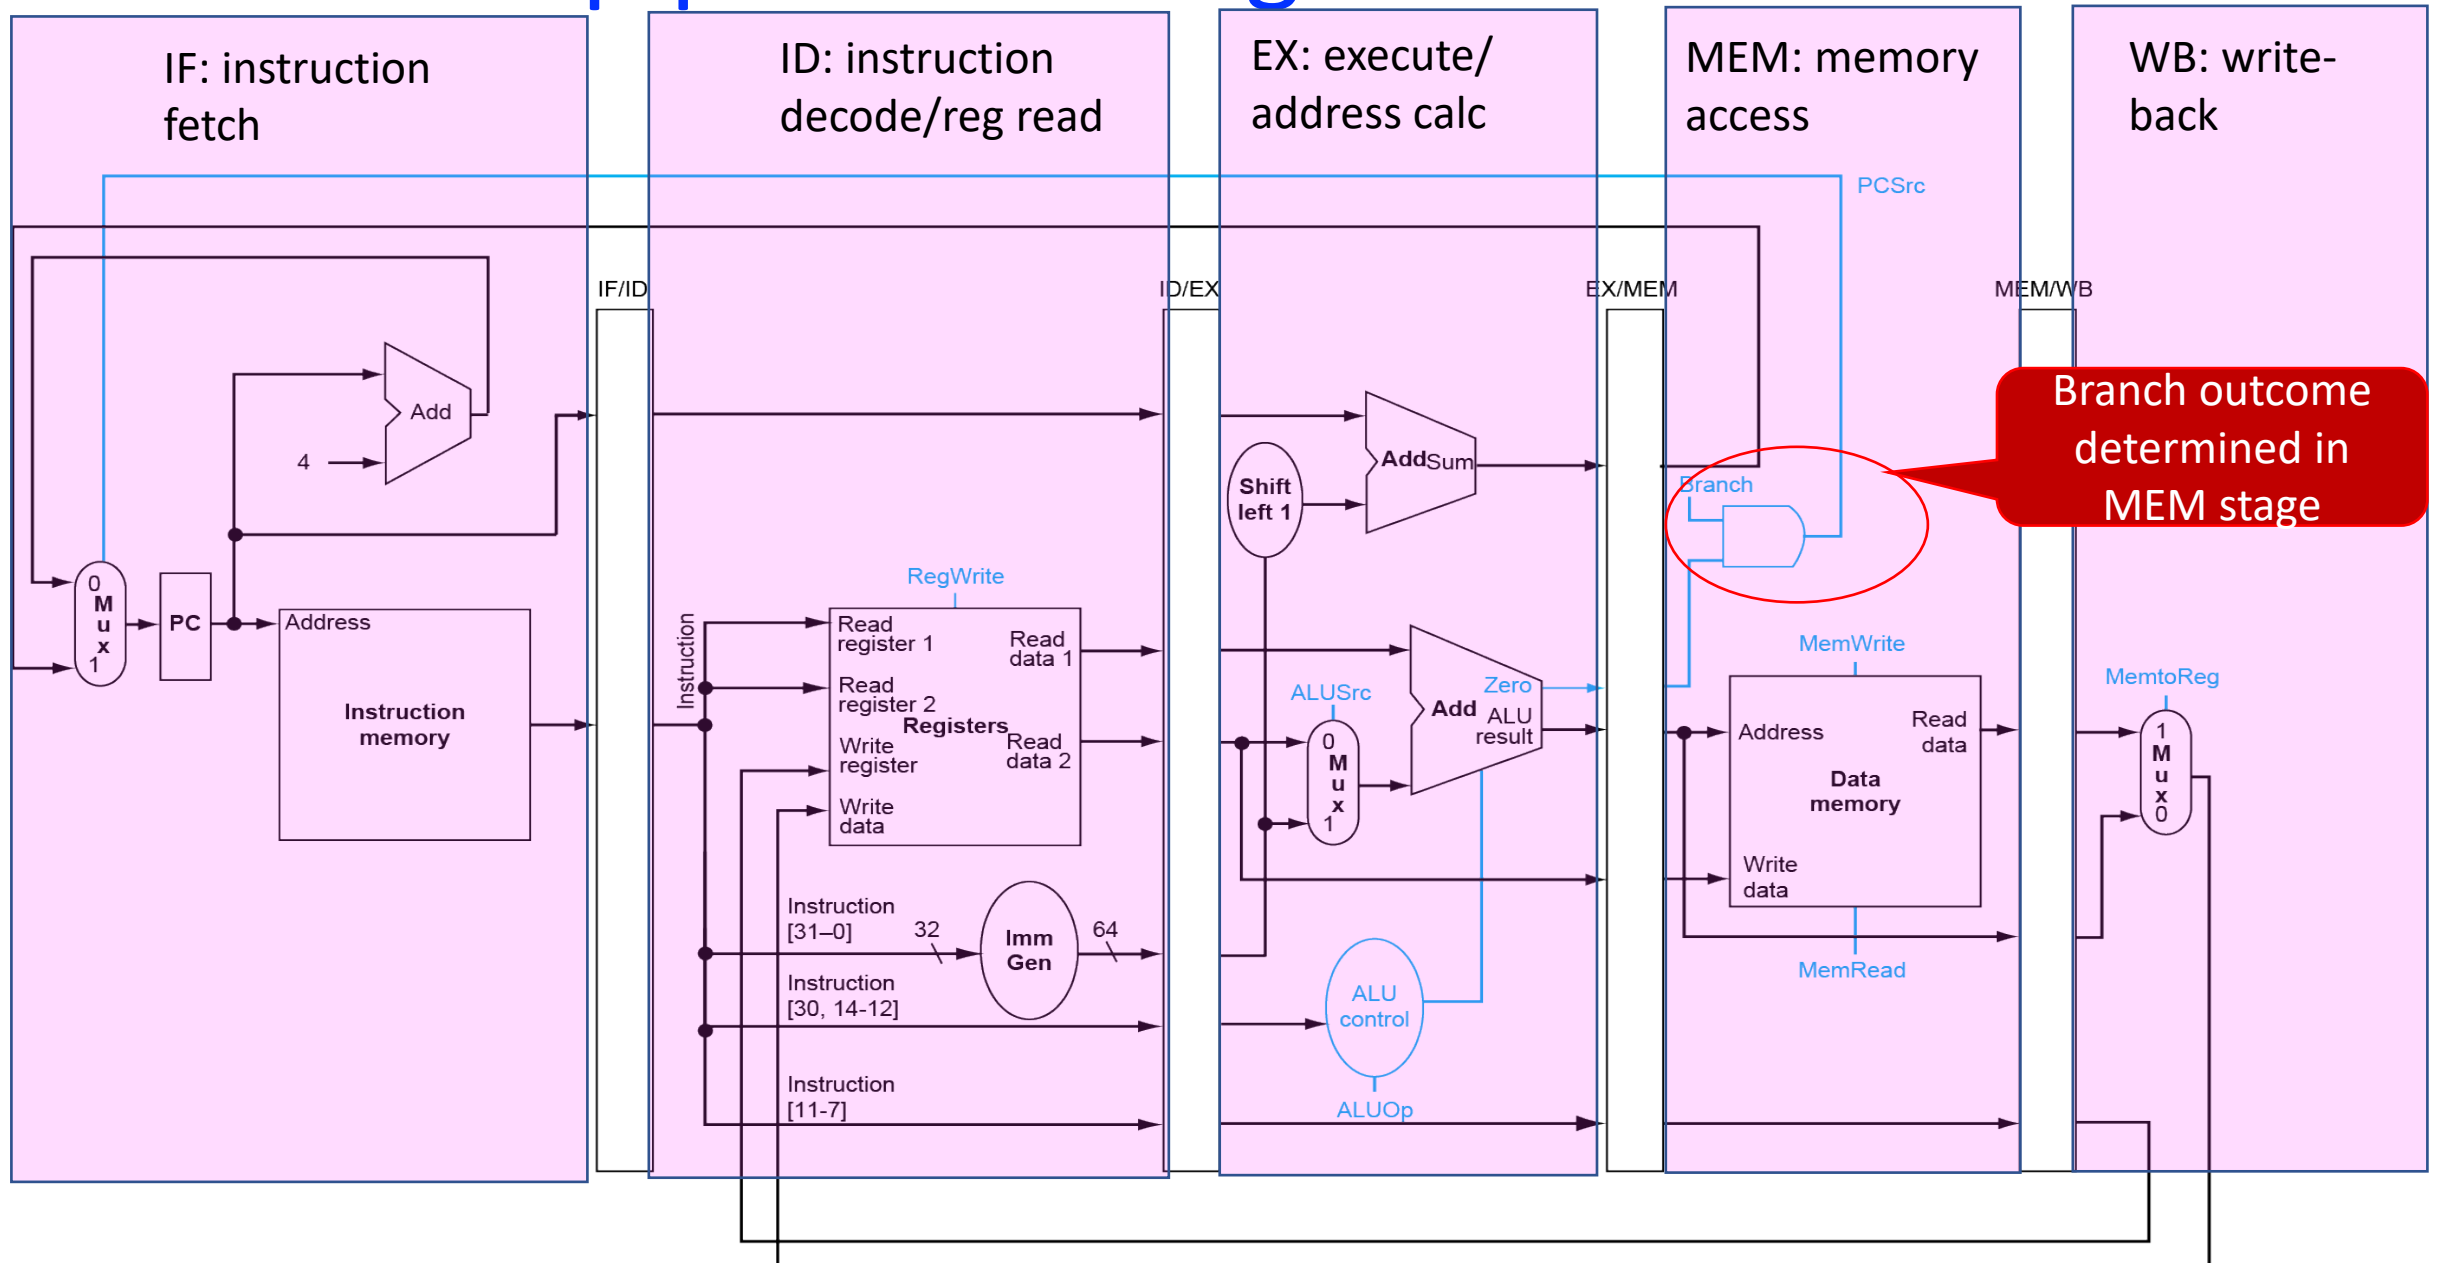
What we've learnt so far

- Single cycle RISC-V CPU design
- 5 stage pipelined RISC-V CPU
- Challenges: hazards:
 - Structure (To mitigate, add resources)
 - Data (To mitigate, do forwarding/bypassing)
 - Load-use must stall
 - Control (??)

Today's lesson

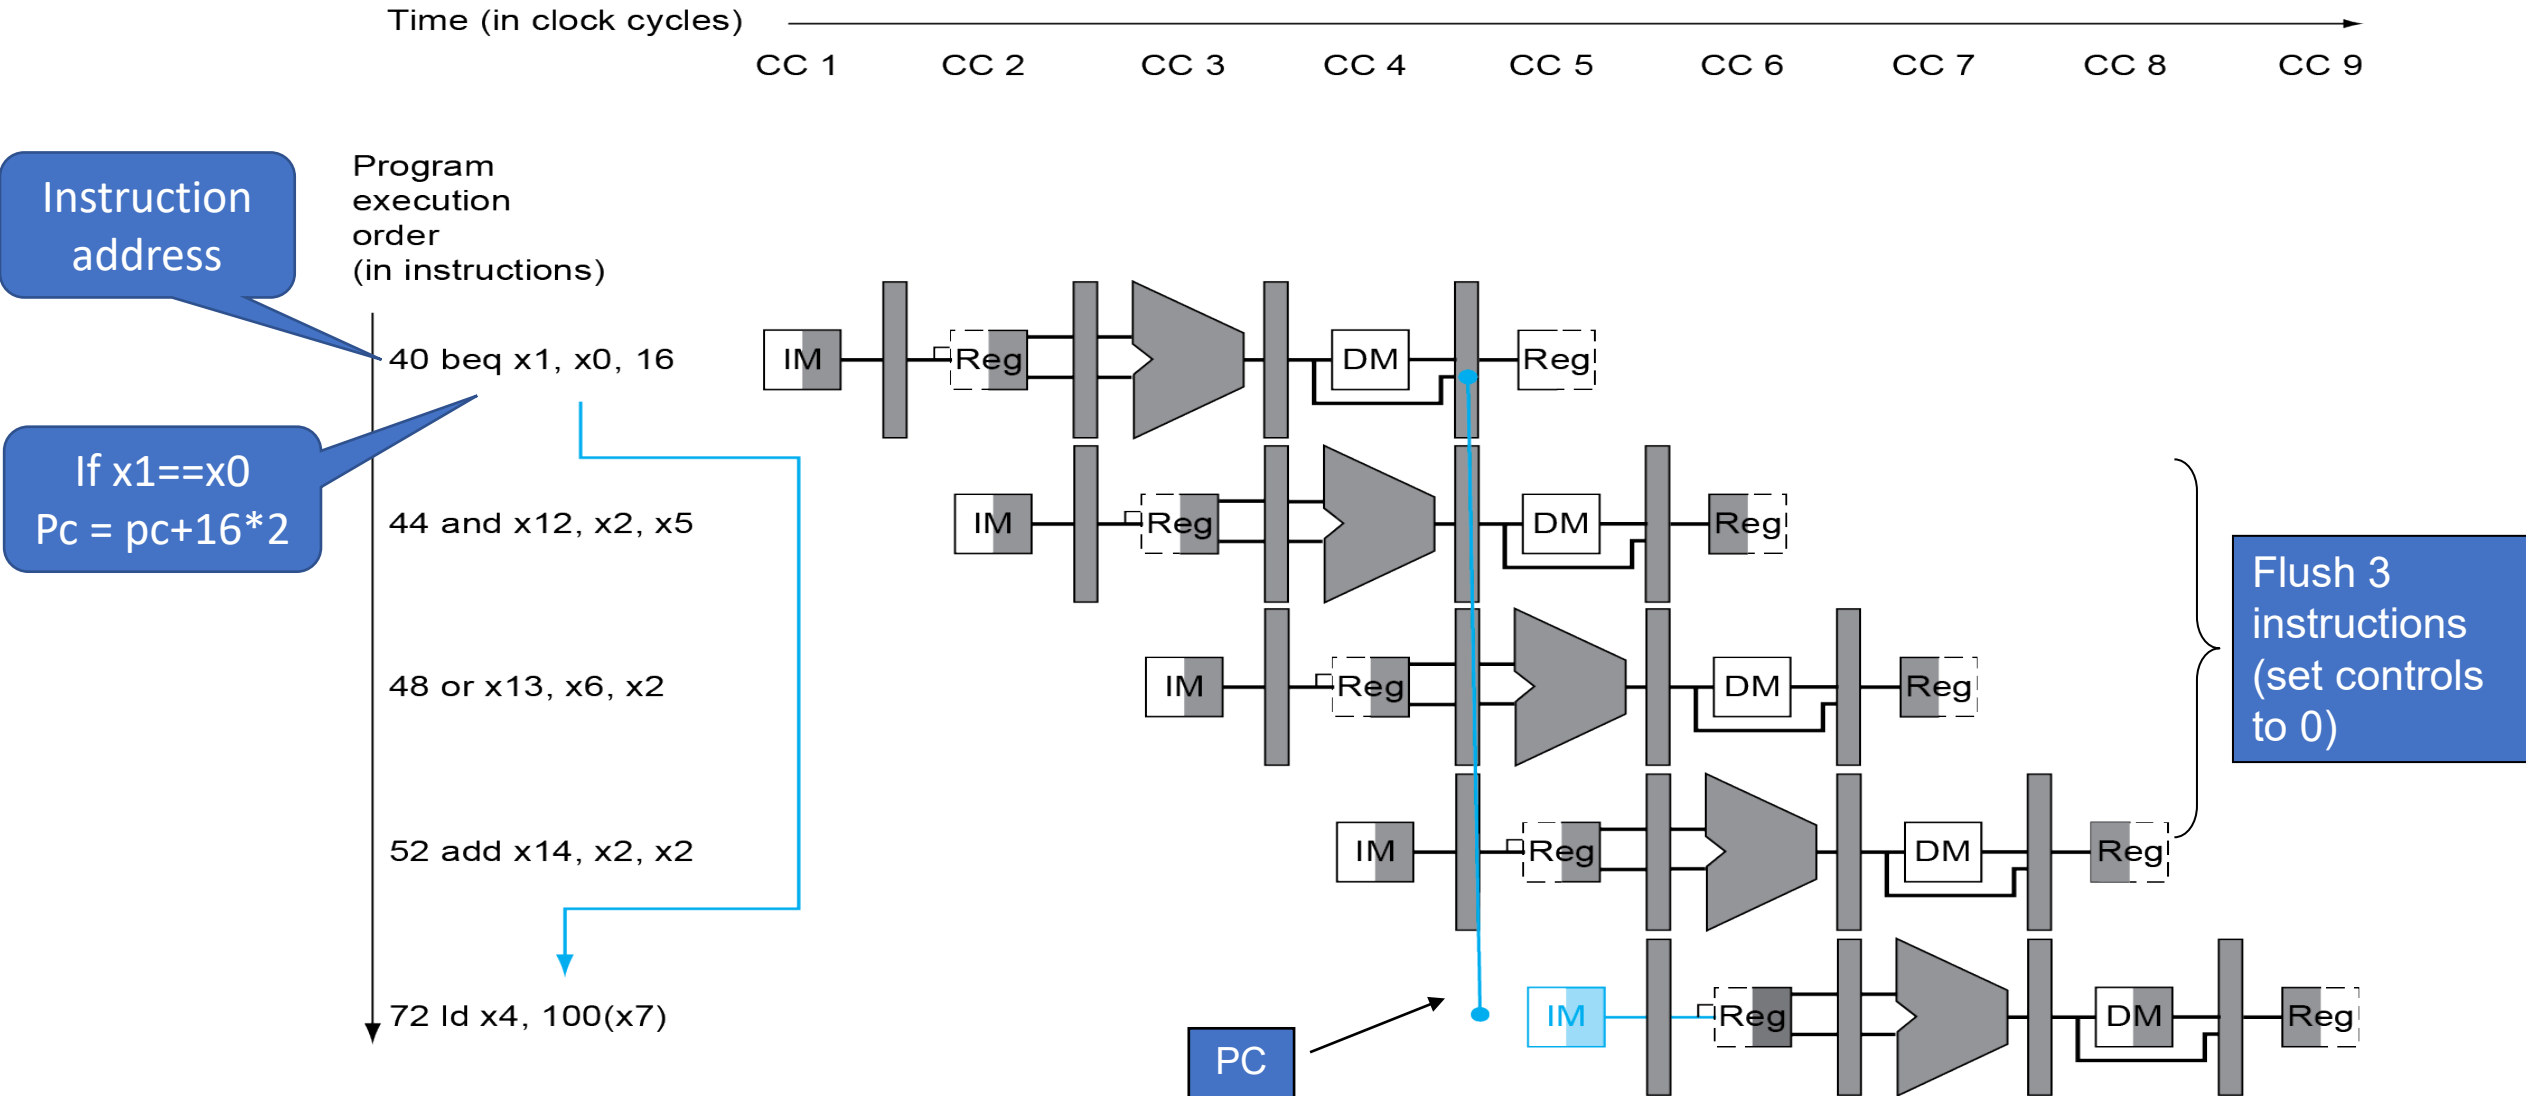
- Handling control hazard
 - Branch hazard
 - exceptions
- Briefly, more advanced topics:
 - Multiple-issue
 - Subword parallelism (SIMD)

Recall our pipelined design so far



Branch Hazard

- If branch outcome is determined in MEM

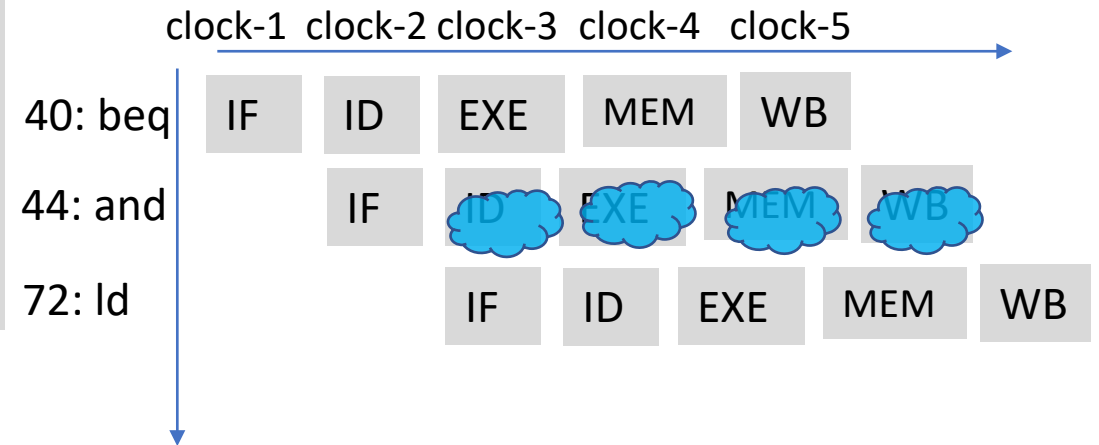


Reducing Branch Delay

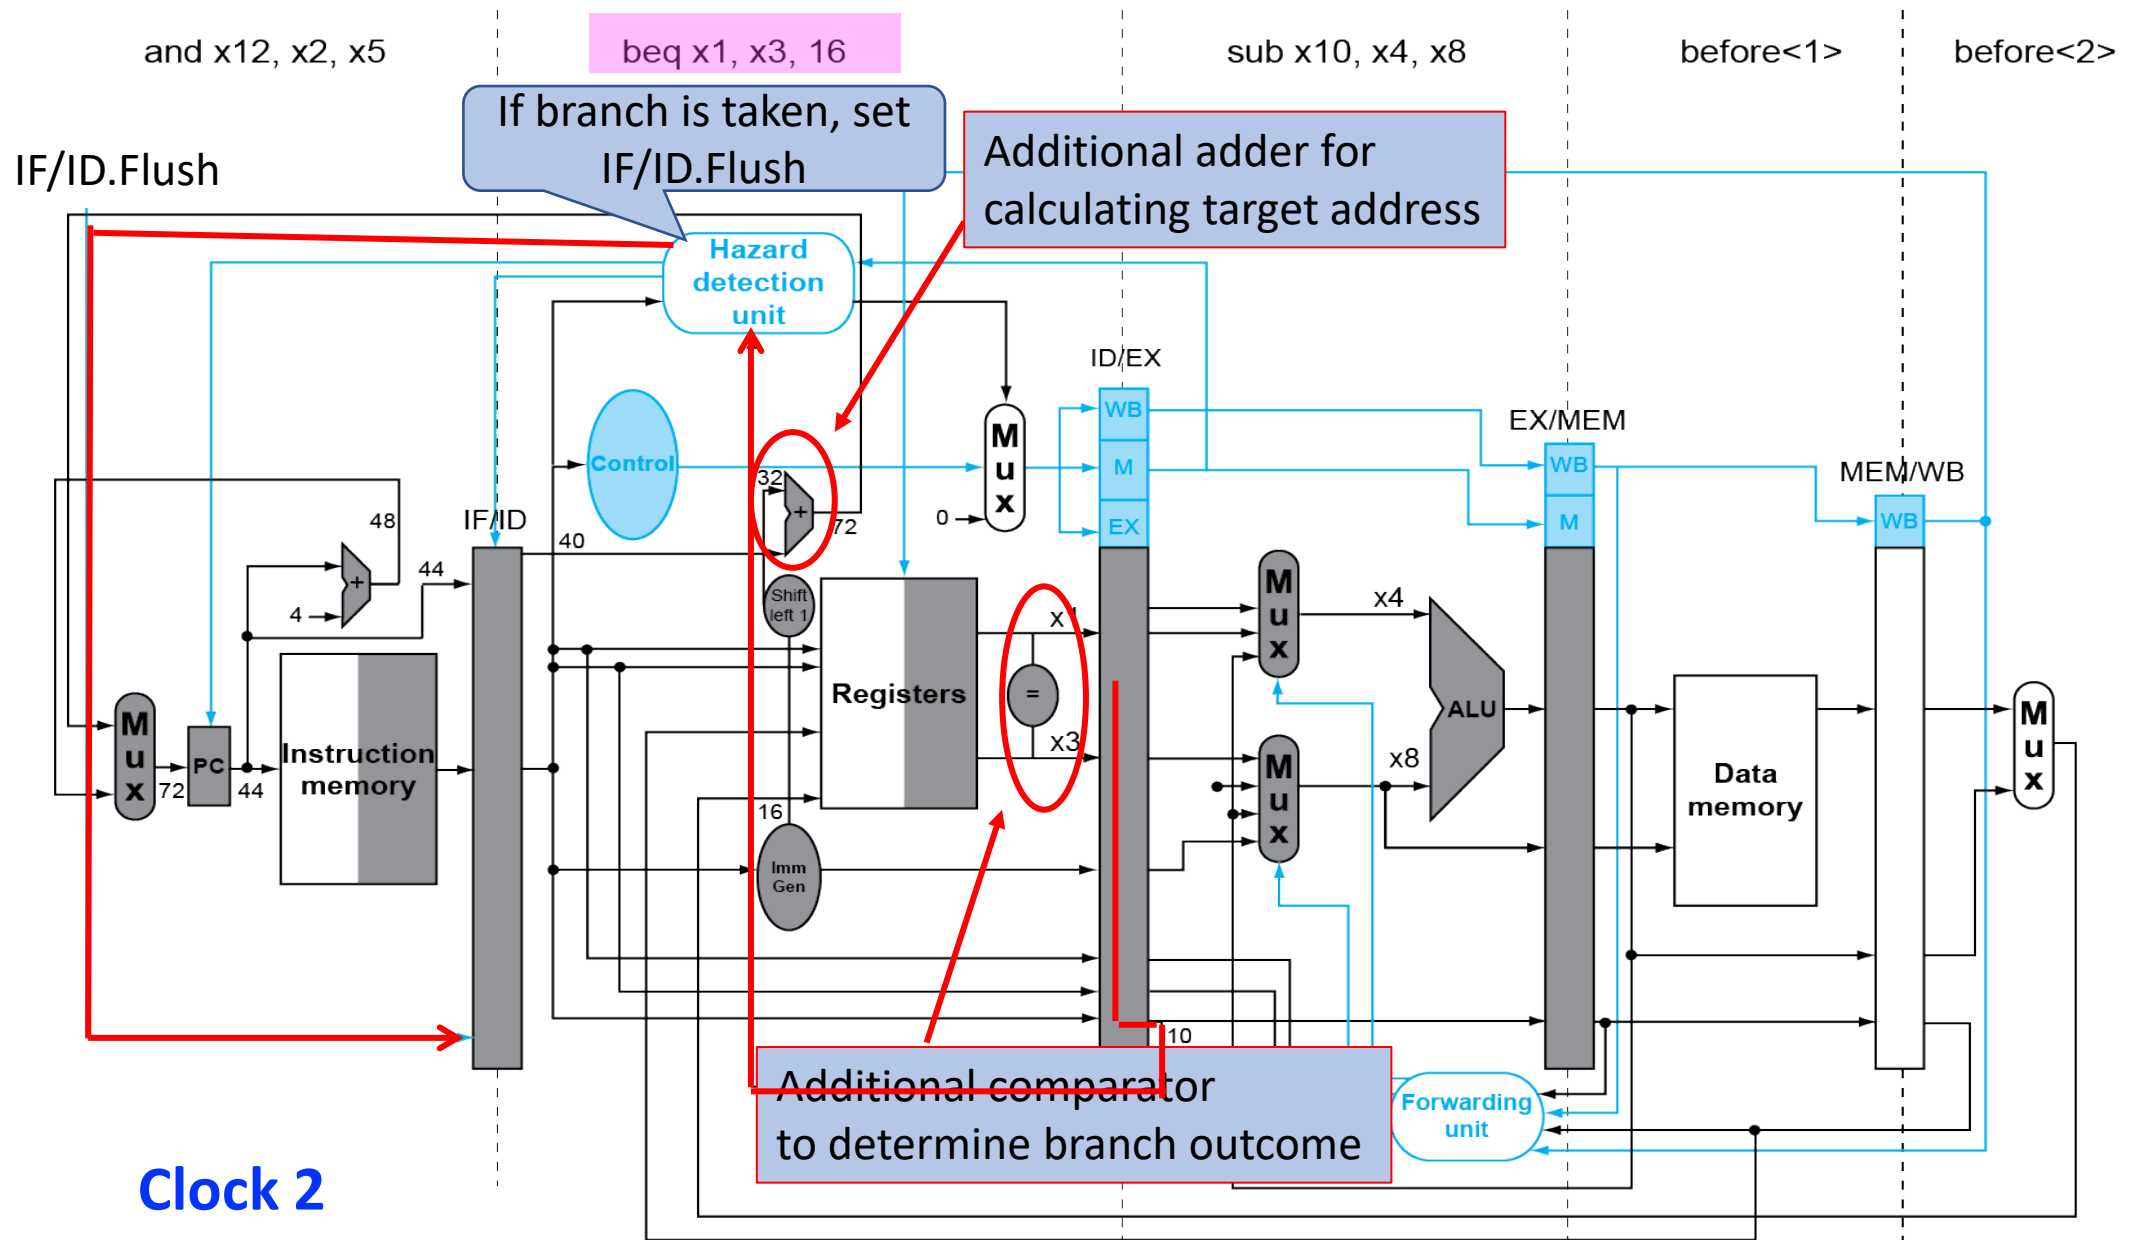
- Add hardware to determine branch outcome earlier (e.g. ID instead of MEM) → fewer instructions to flush

```
36:  sub    x10, x4, x8
40:  beq    x1,  x3, 16    // PC-relative branch
                           // to 40+16*2=72
44:  and    x12, x2, x5
48:  orr    x13, x2, x6
52:  add    x14, x4, x2
56:  sub    x15, x6, x7
   ...
72:  ld     x4,  50(x7)
```


How many instructions to flush
if branch outcome is known in ID?



Branch determined in ID and is taken



Branch determined in ID and is taken




Dynamic Branch Prediction

- Our simple 5-stage pipeline's branch penalty is 1 bubble, but
 - In deeper pipelines, branch penalty is more significant
- Solution: dynamic prediction
 - Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken) } 1-bit predictor
 - To execute a branch
 - Check table, fetch fall-through or target based on predicted outcome.
 - If wrong, flush pipeline and flip prediction

1-bit Predictor's shortcoming

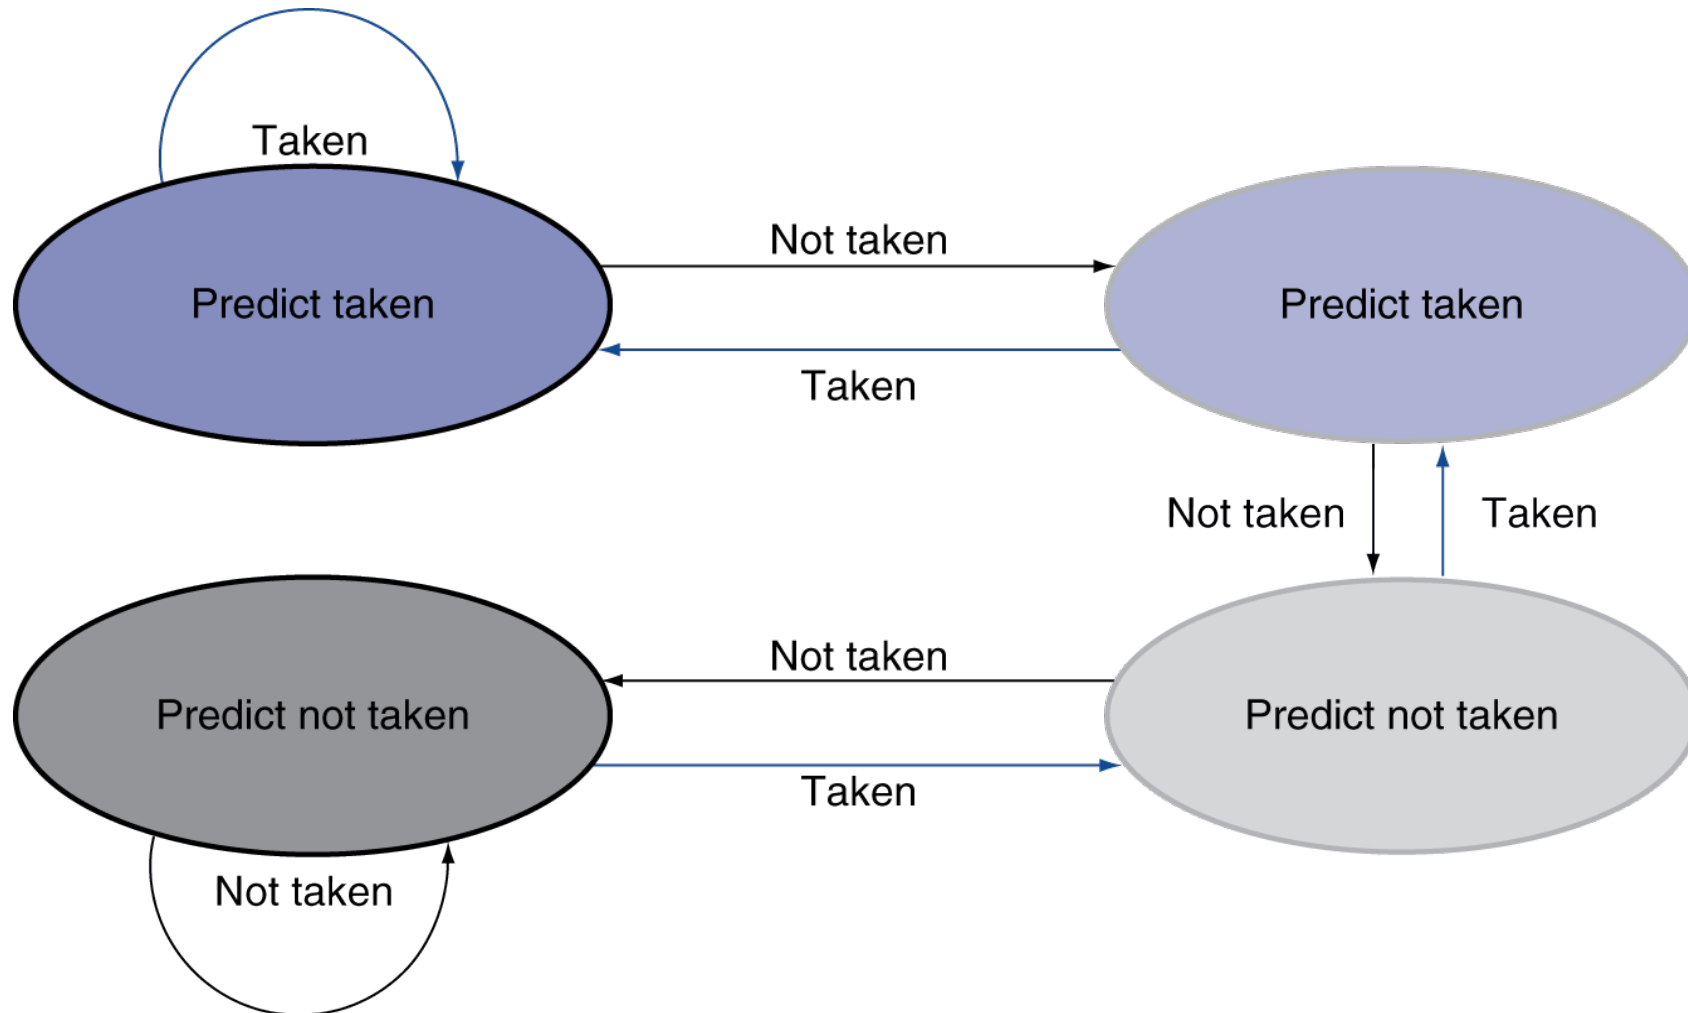
- Inner loop branches mispredicted twice!



- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around

2-Bit Predictor

- Only change prediction on two successive mispredictions



Calculating Branch Target (needed if branch is predicted taken)

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch
- Branch target buffer
 - Cache of target addresses
 - Indexed by PC when instruction fetched
 - If hit and instruction is branch predicted taken, fetch target immediately

Beyond basic pipelining

- ILP: execute multiple instructions in parallel
- To increase ILP
 - Deeper pipeline
 - Less work per stage \Rightarrow shorter clock cycle
 - Multiple issue
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - Finish multiple Instructions Per Cycle ($IPC > 1$)
 - E.g., 4GHz 4-way multiple-issue
 - 16 billion instructions/sec, peak $IPC = 4$ ($CPI = 1/IPC = 0.25$)
 - Challenges: dependencies among multi-issued instructions
 - reduce peak IPC

Dynamic Multiple Issue

- “Superscalar” processors
- CPU decides whether to issue 0, 1, 2, ... each cycle
 - Avoiding structural and data hazards
- Dynamic pipeline scheduling:
 - Execute instructions out of order to avoid stalls
 - But commit results to registers in order

```
ld    x31, 20(x21)
add   x1, x31, x2
sub   x23, x23, x3
```

Can start sub while add is waiting for ld

Does Multiple Issue Work?

- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
 - e.g., pointer aliasing
- Memory delays and limited bandwidth
 - Hard to keep pipelines full
- Speculation can help if done well

Other ways to improve performance: subword parallelism

- ML/graphics applications perform same operations on vectors
 - Partition a 128-bit adder into:
 - Sixteen 8-bit adds
 - Eight 16-bit adds
 - Four 32-bit adds
- Also called data-level parallelism, vector parallelism, or Single Instruction, Multiple Data (SIMD)
- Intel's AVX introduces 256-bit floating point registers
 - 8 single-precision ops
 - 4 double-precision ops

Other ways to improve performance: loop unrolling (software)

- Replicate loop body to expose more parallelism
 - Reduces loop-control overhead
 - Avoid loop-carried dependency
 - Store followed by a load of the same register

Other ways to improve performance: loop unrolling


- Loop unrolling: replicate loop body
- Unrolling exposes more parallelism for multi-issue pipelined CPU
 - Reduces loop-control overhead
 - Avoid loop-carried dependency
 - Store followed by a load of the same register

```
long x[100];  
...  
for (int i = 0; i < 100; i++) {  
    x[i] *= 2;  
}
```

```
//suppose x1=&x[0]  
//x2 = &x[100]  
Loop:  
    ld  x31, 0(x1)  
    add x31, x31, x30  
    sd  x30, 0(x1)  
    addi x1, x1, 8  
    blt x1, x2, Loop
```

```
//suppose x20=&x[0]  
//x22 = &x[100]  
Loop:  
    ld  x31, 0(x1)  
    add x31, x31, x30  
    sd  x30, 0(x1)  
    ld  x29, 8(x1)  
    add x29, x29, x28  
    sd  x28, 8(x1)  
    addi x20, x20, 16  
    blt x20, x22, Loop
```

Performance Impact



Summary on CPU design

- ISA influences design of datapath and control
- Pipelining improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
 - Dependencies limit achievable parallelism
- SIMD achieves data-level parallelism