

Full Name:_____

Univ ID:_____

Quiz II, Fall 2018 Date: Nov 1, 2018

Instructions:

- This exam takes 70 minutes. Read through all the problems and complete the easy ones first.
- This exam is CLOSED BOOK. You may use a double-sided A4 cheatsheet that you've prepared yourself. Any electronic devices or other paper materials are forbidden.

1 (40)	2 (25)	3 (25)	4 (10)	Bonus (15)	Total (100+15)

Notice:

Unless otherwise noted, answer all questions in the quiz assuming a little-endian 64-bit x86 machine. There are two appendices at the end of the exam booklet: 1) a cheatsheet of x86 instructions 2) ASCII table.

1 (40 points) Multiple choices. Circle *all* correct or applicable answers:

A. Facebook boasts 3 billion users now. If you are to accurately keep track of the number of Facebook users in a C variable called `counter`, what type can `counter` be?

1. `short`
2. `unsigned short`
3. `int`
4. `unsigned int`
5. `long`
6. `unsigned long`

B. Which of the following statements are true?

1. Accessing data stored in memory is as fast as accessing data stored in CPU registers.
2. Accessing data stored in memory is much slower than accessing data in CPU registers.
3. A C program is compiled into x86 instructions which are directly executed by the CPU.
4. A Java program is compiled into x86 instructions which are directly executed by the CPU.
5. All x86 instructions are 8 bytes long.

C. Normally, the `%rip` register is updated after executing the current instruction to point to the next instruction. Which of the following instructions *might* cause `%rip` to point to somewhere other than the next instruction?

1. `jle`
2. `mov`
3. `add`
4. `call`
5. `ret`
6. `sub`

D. Which of the following statements about the stack are true?

1. The callee passes back the function's return value to the caller by storing it on the stack.
2. The caller passes the function's arguments to the callee by storing all of them on the stack.
3. The instruction `subq $0x18, %rsp` grows the stack by 24 bytes.
4. The instruction `subq $0x18, %rsp` shrinks the stack by 24 bytes.
5. Global variables are allocated on the stack.
6. Malloc-ed space are allocated on the stack.

E. When executing a *buggy* program, which of the following instructions *might* cause the running program to encounter a "segmentation fault" upon its execution?

1. `movq (%rax,%rbx), %rbx`
2. `movq %rcx, %rbx`
3. `ret`
4. `leaq (%rax,%rbx), %rcx`
5. `subq %rax, (%rbx)`

F. Let `a` be an array of `int` elements. Suppose `%rdi` currently store the address of `a[0]`, and `%rsi` stores some index `i` of type `long`. Which of the following sequence of x86 instructions result in `%eax` storing `a[i]`?

1. `leal (%rdi, %rsi, 4), %eax`
2. `movl (%rdi, %rsi, 4), %eax`
3. `movl (%rsi, %rdi, 4), %eax`
4. `leal (%rdi, %rsi, 8), %eax`
5. `movl (%rdi, %rsi, 8), %eax`
6. `movl (%rsi, %rdi, 8), %eax`
7. `salq $2, %rsi`
`addq %rdi, %rsi`
`movl (%rdi), %eax`
8. `salq $2, %rsi`
`movl (%rsi, %rdi), %eax`

G. Suppose register `%rax` stores C variable `long x`, which of the following instructions effectively implement `x *= 2`?

1. `movq (%rax,%rax), %rax`
2. `leaq (%rax,%rax), %rax`
3. `leaq (, %rax, 2), %rax`
4. `addq %rax, %rax`
5. `shlq $1, %rax`
6. `salq $1, %rax`
7. `shrq $1, %rax`
8. `sarq $1, %rax`

H. Which of the following `mov` instructions are **valid** x86 instructions?

1. `movq $0x400456, %rip`
2. `movq (%rax), (%rbx)`
3. `movb (%rax), %bl`
4. `movb %al, (%bl)`
5. `movq %al, (%bl)`

2 C and assembly: (25 points)

A (5 points). Please complete the following skeleton function `my_memcpy` to copy `n` bytes from buffer `src` to buffer `dst`. You must only add C code in the space denoted by `_____`.

```
void my_memcpy(void *dst, void *src, unsigned int n)
{
    char *s = (char *)src;
    char *d = (char *)dst;

    for (_____; _____; _____) {
        _____;
    }
}
```

B (5 points). The following `main` function is to copy `int` array `a1` to another `int` array `a2`. Please complete the line to properly invoke `my_memcpy` to perform the copying.

```
int main() {
    int a1[2] = {1, 2};
    int a2[2];

    my_memcpy(_____, _____, _____);
    printf("a2: %d %d %d\n", a2[0], a2[1]); //this should print out "a2: 1 2"
}
```

C (5 points). Given the following `main` function, what would its output be?

```
int main() {
    char c1[3] = {'a', 'b', 'c'};
    char *c2;
    c2 = &c1[1];
    my_memcpy(c2, c1, 2);
    printf("c2: %c %c\n", c2[0], c2[1]);
}
```

The following assembly corresponds to the `main` function in the previous question (C.).

```
main:
    subq    $24, %rsp
    movb    $97, (%rsp)
    movb    $98, 1(%rsp)
    movb    $99, 2(%rsp)

    _____, %rdi

    _____, %rsi
    movl    $2, %edx
    call    mymemcpy
    ...
    call    printf
    addq    $24, %rsp
    ret
```

D (5 points). Which one of the above assembly instructions initialize the first element of the `c1` array to contain ASCII character 'a'? (You may just circle the instruction.)

E (5 points). Complete the two missing x86 assembly instructions marked by _____.

3 Assembly Mystery: (25 points)

Ben Bitdiddle is trying to figure out what the following disassembled function `mysterycount` does:

```
mysterycount:
    movl    $0, %edx
    movl    $0, %eax
.L5:
    cmpl    %edi, %edx
    jge     .L3
    movslq  %edx, %rcx
    cmpl    $0, (%rsi,%rcx,4)
    jle     .L4
    addl    $1, %eax
.L4:
    addl    $1, %edx
    jmp     .L5
.L3:
    ret
```

Alyssa P. Hacker decides to give Ben some hints on what `mysterycount` does by providing him with a skeleton C function. She also told Ben that `mysterycount` tries to count the number of elements in a given array that match a specific criteria.

```
1: _____
2: mysterycount( _____, _____ )
3: {
4:     _____ count = 0;

5:     for (_____ i = 0; _____; i++) {

6:         if (_____ ) {
7:             count++;
8:         }
9:     }
10: return count;
11: }
```

A. (4 points) Where is the C variable `count` stored?

B. (4 points) What is the type of `count`? Write it down on line 4. What is the type of the return value of `mysterycount`? Write it down on line 1.

C. (4 points) Which function signature (omiting the return value) should `mysterycount` be on line 2?

1. `mystery(int *array, int n)`
2. `mystery(int n, int *array)`
3. `mystery(int *array, unsigned int n)`
4. `mystery(unsigned int n, int *array)`
5. `mystery(char *array, int n)`
6. `mystery(int n, char *array)`
7. `mystery(char *array, unsigned int n)`
8. `mystery(unsigned int n, char *array)`

D. (4 points) Where is the loop variable `i` stored ?

E. (4 points) What is the type of the loop variable `i`?

F. (5) Please complete the rest of the skeleton code by filling out

4 Buffer Overflow (10 points + 15 bonus points)

Ben Bitdiddle has thought of a strategy to defend against buffer overflow. The following C program gives a concrete example of Ben's strategy. After allocating a buffer on the stack, Ben stores the size of the buffer BUFSZ at the end of buf (line 4). After calling a dangerous function such as gets (line 5) which might have overflowed the buffer, Ben terminates the program if the stored buffer size has changed (lines 6-8).

Note gets(buf) reads from terminal into buf until a newline character '\n', which it replaces with '\0'.

```
1: #define BUFSZ 4
2: void echo() {
3:     char buf[BUFSZ+sizeof(int)];
4:     *(int *) (buf+BUFSZ) = BUFSZ; //write BUFSZ as an int at address value buf+BUFSZ
5:     gets(buf);
6:     if ((* (int *) (buf+BUFSZ)) != BUFSZ) {
7:         printf("stack overflow detected!\n");
8:         exit(1); // terminate the program
9:     }
10:    return;
11:}

void divulge_secret() {
    // omitted
}

int main() {
    echo();
}
```

The corresponding disassembled program is as follows:

```
00000000000006ca <echo>:
6ca: 48 83 ec 18          sub    $0x18,%rsp
6ce: c7 44 24 0c 04 00 00 movl    $0x4,0xc(%rsp)
6d5: 00
6d6: 48 8d 7c 24 08       lea     0x8(%rsp),%rdi
6db: b8 00 00 00 00       mov     $0x0,%eax
6e0: e8 ab fe ff ff       callq  <gets>
6e5: 83 7c 24 0c 04       cmpl    $0x4,0xc(%rsp)
6ea: 75 05                jne     6f1 <echo+0x27>
6ec: 48 83 c4 18          add     $0x18,%rsp
6f0: c3                  retq
6f1: 48 8d 3d ac 00 00 00 lea     0xac(%rip),%rdi #this and next instruction do printf
6f8: e8 83 fe ff ff       callq  <printf>
6fd: bf 01 00 00 00       mov     $0x1,%edi #this and the next instruction do exit(1)
702: e8 99 fe ff ff       callq  <exit>

0000000000000707 <main>:
70b: b8 00 00 00 00       mov     $0x0,%eax
710: e8 b5 ff ff ff       callq  6ca <echo>
715: b8 00 00 00 00       mov     $0x0,%eax
71e: c3                  retq

0000000000000f876 <divulge_secret>:
# ... omitted....
```


Suppose the value of `%rsp` is `0x00007fffffffcd4f8` upon entering function `echo` just before executing its first instruction `sub $0x18, %rsp`.

A. (5 points) What are the values of the 8 bytes starting at address `0x00007fffffffcd4f8`? (Please write the 8 byte values *one by one* from lower to higher addresses. Each byte value should be represented with 2 hex-digits and successive byte values should be separated with a space.)

B. (5 points) What is the starting address of the buffer `buf`? (Hint: `%rdi` stores the value of the first argument when calling a function.)

The remaining questions are **bonus** questions.

C. (5 points) Suppose an attacker enters a sequence of 25 input bytes 31 32 33 34 35 36 37 38 31 32 33 34 35 36 37 38 76 f8 00 00 00 00 00 00 0a (each byte is represented by 2 hex digits. Note the last byte 0a is the ASCII value of the newline character). What are the values of the 8 bytes starting at address 0x00007fffffffdd4f8 when `callq <gets>` returns? (Write the 8 byte values *one by one* from lower to higher addresses.)

D. (5 points) Following the scenario described in **C.**, will the running program return to a corrupted return address? Please explain.

E. (5 points) What is the sequence of bytes that the attacker should enter in order to cause the program to execute the `divulge_secret` function.

— END OF QUIZ II —

Appedix: X86 Cheatsheet

4.1 Registers

x86 registers are 8-bytes. Additionally, the lower order bytes of these registers can be independently accessed as 4-byte, 2-byte, or 1-byte register. The register names are:

8-byte register	Bytes 0-3	Bytes 0-1	Byte 0 (lowest order byte)
%rax	%eax	%ax	%al
%rbx	%ebx	%bx	%bl
%rcx	%ecx	%cx	%cl
%rsi	%esx	%si	%sil
%rdi	%edx	%di	%dil

...the rest is omitted...

4.2 Instructions

Instruction suffixes:

"byte" (b)	1-byte
"word" (w)	2-bytes
"doubleword" (l)	4-bytes
"quardword" (q)	8-bytes

Complete memory addressing mode: A memory operand of the form $D(Rb, Ri, S)$ accesses memory at address $D + \text{val}(Rb) + \text{val}(Ri) * S$, where $\text{val}(Rb)$ and $\text{val}(Ri)$ refer to the value of registers Rb and Ri respectively, D is a constant, and S is a constant of value 1, 2, 4, or 8.

Sign extension and zero extension:

$\text{movzq } S, D$	copy 4-byte-sized S to 8-byte-sized D and fill in the higher order 4 bytes of D with zero bytes.
$\text{movslq } S, D$	copy 4-byte-sized S to 8-byte-sized D and sign extend the higher order 4 bytes of D , i.e. fill with 0s if S 's sign bit is zero and fill with 1s if S 's sign bit is one.

Basic Arithmetic instructions that you might not remember:

$\text{sal / shl } k, D$	Left shift destination D by k bits
sar	Arithmetic right shift destination D by k bits
shr	Logical right shift destination D by k bits

Jump instructions:

Jump instruction following `cmp S, D`:

<code>jmp</code>	Unconditional jump
<code>je</code>	Jump if D is equal to S
<code>jne</code>	Jump if D is not equal to S
<code>jg</code>	Jump if D is greater than S (signed)
<code>jge</code>	Jump if D is greater or equal than S (signed)
<code>jl</code>	Jump if D is less than S (signed)
<code>jle</code>	Jump if D is less or equal than S (signed)
<code>ja</code>	Jump if D is above S (unsigned)
<code>jae</code>	Jump if D above or equal S(unsigned)
<code>jb</code>	Jump is D is below S (unsigned)
<code>jbe</code>	Jump if D is below or equal S (unsigned)

4.3 Calling convention

Argument Passing:

Which argument	Stored in register
1	<code>%rdi</code>
2	<code>%rsi</code>
3	<code>%rdx</code>
4	<code>%rcx</code>
5	<code>%r8</code>
6	<code>%r9</code>
7 and up	passed on stack

Return value (if any) is stored in `%rax`

Caller save registers: `%rax`, `%rcx`, `%rdx`, `%rdi`, `%rsi`, `%r8-11`

Callee save registers: `%rbx`, `%rbp`, `%r12-15`

Appendix: ASCII

The following table contains the 128 ASCII characters.

Oct	Dec	Hex	Char	Oct	Dec	Hex	Char
000	0	00	NUL '\0'	100	64	40	@
001	1	01	SOH (start of heading)	101	65	41	A
002	2	02	STX (start of text)	102	66	42	B
003	3	03	ETX (end of text)	103	67	43	C
004	4	04	EOT (end of transmission)	104	68	44	D
005	5	05	ENQ (enquiry)	105	69	45	E
006	6	06	ACK (acknowledge)	106	70	46	F
007	7	07	BEL '\a' (bell)	107	71	47	G
010	8	08	BS '\b' (backspace)	110	72	48	H
011	9	09	HT '\t' (horizontal tab)	111	73	49	I
012	10	0A	LF '\n' (new line)	112	74	4A	J
013	11	0B	VT '\v' (vertical tab)	113	75	4B	K
014	12	0C	FF '\f' (form feed)	114	76	4C	L
015	13	0D	CR '\r' (carriage ret)	115	77	4D	M
016	14	0E	SO (shift out)	116	78	4E	N
017	15	0F	SI (shift in)	117	79	4F	O
020	16	10	DLE (data link escape)	120	80	50	P
021	17	11	DC1 (device control 1)	121	81	51	Q
022	18	12	DC2 (device control 2)	122	82	52	R
023	19	13	DC3 (device control 3)	123	83	53	S
024	20	14	DC4 (device control 4)	124	84	54	T
025	21	15	NAK (negative ack.)	125	85	55	U
026	22	16	SYN (synchronous idle)	126	86	56	V
027	23	17	ETB (end of trans. blk)	127	87	57	W
030	24	18	CAN (cancel)	130	88	58	X
031	25	19	EM (end of medium)	131	89	59	Y
032	26	1A	SUB (substitute)	132	90	5A	Z
033	27	1B	ESC (escape)	133	91	5B	[
034	28	1C	FS (file separator)	134	92	5C	\ '\\'
035	29	1D	GS (group separator)	135	93	5D]
036	30	1E	RS (record separator)	136	94	5E	^
037	31	1F	US (unit separator)	137	95	5F	_
040	32	20	SPACE	140	96	60	`
041	33	21	!	141	97	61	a
042	34	22	"	142	98	62	b
043	35	23	#	143	99	63	c
044	36	24	\$	144	100	64	d
045	37	25	%	145	101	65	e
046	38	26	&	146	102	66	f
047	39	27		147	103	67	g
050	40	28	(150	104	68	h
051	41	29)	151	105	69	i
052	42	2A	*	152	106	6A	j
053	43	2B	+	153	107	6B	k
054	44	2C	,	154	108	6C	l
055	45	2D	-	155	109	6D	m
056	46	2E	.	156	110	6E	n
057	47	2F	/	157	111	6F	o
060	48	30	0	160	112	70	p
061	49	31	1	161	113	71	q
062	50	32	2	162	114	72	r
063	51	33	3	163	115	73	s
064	52	34	4	164	116	74	t
065	53	35	5	165	117	75	u
066	54	36	6	166	118	76	v
067	55	37	7	167	119	77	w
070	56	38	8	170	120	78	x
071	57	39	9	171	121	79	y
072	58	3A	:	172	122	7A	z
073	59	3B	;	173	123	7B	{
074	60	3C	<	174	124	7C	
075	61	3D	=	175	125	7D	}
076	62	3E	>	176	126	7E	~
077	63	3F	?	177	127	7F	DEL