

Full Name:_____

Mock Midterm Practice (Fall2022)

Instructions:

- Quiz II takes 70 minutes. Read through all the problems and complete the easy ones first.
- This exam is **closed book**, except that you may bring a single double-sided page of prepared note.

1 (xx/32)	2 (xx/20)	3 (xx/28)	4 (xx/10)	Total (xxx/90)

This exam assumes 64-bit x86 hardware (little Endian) unless otherwise mentioned.

1 Machine Representation (32 points, 4 points each):

Multiple choices. Circle *all* answers that apply.

A. Suppose register `%rbx` corresponds to some C variable `v`. Given instruction `addq $4, %rbx`, what are *all* the potential types of `v` and the corresponding C statement of the instruction?

1. type: `int` , C statement: `v +=4;`
2. type: `unsigned int` , C statement: `v +=4;`
3. type: `long` , C statement: `v +=4;`
4. type: `unsigned long`, C statement: `v +=4;`
5. type: `int *` , C statement: `v +=4;`
6. type: `int *` , C statement: `v++;`
7. type: `long *` , C statement: `v +=4;`
8. type: `long *` , C statement: `v ++;`
9. none of the above

B. Suppose `%rdi` and `%rsi` corresponds to C variables `x` and `y`, respectively. Given instruction `movq (%rdi, %rsi, 8), %rax`, what are the most likely types of `x` and `y`, respectively?

1. `long` and `long`
2. `long *` and `long`
3. `long *` and `long *`
4. `int` and `int`
5. `int *` and `int`
6. `int *` and `int *`
7. none of the above

C. Which of the following machine instructions change the value of register `%rsp`?

1. `ret`
2. `pushq %rax`
3. `popq %rbx`
4. `call`
5. `movq %rax, (%rsp)`
6. none of the above

D. Which of the following instructions read from or write to memory?

1. `ret`
2. `pushq %rax`
3. `popq %rbx`
4. `call`
5. `addq %rax, %rbx`
6. none of the above

E. Suppose `%rsi` corresponds to C variable `y` of some pointer type. Which of the following instructions dereference the pointer `y`?

1. `leaq (%rsi), %rax`
2. `movq (%rsi), %rax`
3. `movq %rsi, %rax`
4. `subq %rax, (%rsi)`
5. `subq %rax, %rsi`
6. none of the above

F. Consider the following recursive function:

```
void foo(int c) {  
    if (c <= 0)  
        return;  
    foo(c-1);  
}
```

What is the minimum stack size (in bytes) needed in order to execute `foo(128)` successfully¹?

1. 8096
2. 1024
3. 512
4. 128
5. any value

G. Which of the following statements on segmentation faults are true?

1. Performing any out-of-bounds array access will result in an immediate segmentation fault.
2. Segmentation faults only occur when an instruction tries to write to memory.
3. Dereferencing a null pointer will always result in a segmentation fault.
4. Performing pointer arithmetic will always result in a segmentation fault.
5. none of the above

H. Suppose *local variable* `a` is defined as `int a[16];` Which of the following statements are true?

1. `a` takes up space on the stack.
2. `a` takes up space on the heap.
3. `subq $64, %rsp` allocates space for `a` and `addq $64, %rsp` de-allocates space for `a`.
4. `addq $64, %rsp` allocates space for `a` and `subq $64, %rsp` de-allocates space for `a`.
5. none of the above

¹You should assume that the compiler does not perform tail call optimization to avoid allocating a stack frame per function call.

2 Basic C (20 points, 4 points each)

Answer the following multiple-choice questions. Circle *all* answers that apply. Each is 5 points.

A. In the following code, what's the most likely outcome of line 4?

```
1: int a = 5;
2: int *p;
3: p = (int *)a;
4: printf("%d\n", *p);
```

1. It will produce a compilation error.
2. It will cause a segmentation fault.
3. It will print 5.
4. It will print some memory address.
5. None of the above

B. In the following code, what should be at line 3 in order for character 'd' to be printed at line 4?

```
1: char *s = "abcdef";
2: char v;
3: v = _____;
4: printf("%c\n", v);
```

1. `v = s+3;`
2. `v = *(s+3);`
3. `v = s[3];`
4. `v = *(++s);`
5. `v = ++s;`
6. None of the above

C. What is the output of the following code snippet?

```
char a[5] = {0, 0, -1, -1, -1};  
int *p;  
p = (int *) (a + 1);  
printf("%d\n", *p);
```

1. a positive number
2. a negative number
3. 0
4. Segmentation fault
5. Compilation error

D. Suppose `x` is of type `unsigned int`. Which of the following statement computes 0 if and only if the i -th bit of `x` (starting from the left) is zero? (The 0-th bit corresponds to the most significant bit).

1. `x & 0x80000000`
2. `(x << i) >> i`
3. `x & (1 << (31-i))`
4. `x | (0x80000000 >> i)`
5. `x & (0x7fffffff >> i)`
6. None of the above

E. Suppose `x` is of type `unsigned int`. Which of the following statement sets the i -th bit of `x` (starting from the left) to be one? (The 0-th bit corresponds to the most significant bit).

1. `x &= (1 << i)`
2. `x |= (1 << i)`
3. `x |= (1 << (31-i))`
4. `x &= (1 << (31-i))`
5. `x &= ~(1 << (31-i))`
6. None of the above

3 C to assembly (28 points):

Ben Bitdiddle wrote the following `swap` function to swap two long integers. Its corresponding assembly code is also shown below.

```
1 void
2 swap(long *a, long *b)
3 {
4     long tmp;
5     tmp = *a;
6     *a = *b;
7     *b = tmp;
8 }
```

```
1 swap:
2 movq (%rdi), %rax
3 movq (%rsi), %rdx
4 movq %rdx, (%rdi)
5 movq %rax, (%rsi)
```

A. (2 points) Where are function arguments `a` and `b` stored at, respectively?

Answer: _____

B. (4 points) For each C statement, which is its corresponding assembly instruction (or set of assembly instructions)? Fill in the line number of the assembly instructions that correspond to each C statement in the table below.

C line number	Assembly line number
5	
6	
7	
6	

Ben Bitdiddle then implemented the `reverse_array` function to reverse the elements in an array of long integers. Note that `reverse_array` uses his previously implemented `swap` as a helper function. Ben's `reverse_array` C function and its corresponding assembly are shown below:

```
// sz is the number of elements in array a
1 void reverse_array(long *a, long sz)
2 {
3     long s = 0;
4     long e = sz-1;
5     while (s < e) {
6         swap((long *)a[s], (long *)a[e]);
7         s++;
8         e--;
9     }
10 }
```

```
1 reverse_str:
2     pushq    %r12
3     pushq    %rbp
4     pushq    %rbx
5     movq     %rdi, %r12
6     leaq     -1(%rsi), %rbx
7     movq     $0, %rbp
8 .L3:
9     cmpq     %rbx, %rbp
10    jge      .L6
11    movq     (%r12,%rbx), %rsi
12    movq     (%r12,%rbp), %rdi
13    call     swap
14    addq     $1, %rbp
15    subq     $1, %rbx
16    jmp      .L3
17 .L6:
18    popq     %rbx
19    popq     %rbp
20    popq     %r12
21    ret
```

C. (4 points) Where are the C variables `s` and `e` stored at, respectively?

Answer: _____

D. (2 points) Which line of assembly performs computes the condition that determines whether the loop in `reverse_array` should terminate or continue?

Answer: _____

E. (2 points) Which line(s) of assembly copy the values to be the arguments used by `swap` function?

Answer: _____

F. (2 points) Which C variable does `%r12` contain in the body of the loop?

1. `a`
2. `s`
3. `e`
4. `sz`
5. None of the above

G. (4 points) Suppose `%rsp` is `0x00007fffffffffff20` *before* executing the first instruction of `reverse_str` (aka line 2 of assembly). What is the value of `%rsp` *immediately before* and *immediately after* executing line 13 of assembly?

Value of `%rsp` immediately before: _____

Value of `%rsp` immediately after: _____

Finally, Ben Bitdiddle writes the following program to test the correctness of `reverse_array`.

```
int
main()
{
    long a[4] = {1, 2, 3, 4};
    reverse_array(a, 4);
    assert(a[0] == 4);
    assert(a[1] == 3);
    assert(a[2] == 2);
    assert(a[3] == 1);
    return 0;
}
```

H. (4 points) Ben found out that running his program results in segmentation fault, and the offending instruction is the first instruction of `swap`, aka `movq (%rdi), %rax`. If Ben is to print out the value of `%rdi` in GDB, what is the value? Answer: _____

I. (4 points) Please fix Ben's bug. You may directly write your corrections on the given C functions.

4 More C Programming (10 points):

Please complete the following code to implement a function `hex2int` that converts a hex string to its integer value (by interpreting the bit pattern represented by the hex string as 2's complement).

```
// return the integer value of the ASCII hex digit c
// you may assume hex digits are always given in lower case
char hex2digit(char c) {

}

int hex2int(char *s) {
    assert(strlen(s) == 8); //assume the string always contains exactly 8 hex "digits"

}

void main() {
    char *s;
    s = "ffffffff";
    int x = hex2int(s);
    assert(x == -1);
    s = "0000000f";
    x = hex2int(s);
    assert(x == 15);
}
```

1. (5 points) Please complete the helper function `hex2digit` that converts a hex character to an integer value in the range $[0, 15]$. You may assume that hex characters are always in lower case.
2. (5 points) Complete `hex2int` to convert a hex string representing an integer in 2's complement to the corresponding integer. The expected return value of this function is demonstrated using two example inputs in `main`. You may assume that the hex string is always exactly 8 characters long and in lower case.

—END of Quiz II—

5 Appendix: X86 Cheatsheet

5.1 Registers

x86 registers are 8-bytes. Additionally, the lower order bytes of these registers can be independently accessed as 4-byte, 2-byte, or 1-byte register. The register names are:

8-byte register	Bytes 0-3	Bytes 0-1	Byte 0 (lowest order byte)
%rax	%eax	%ax	%al
%rbx	%ebx	%bx	%bl
%rcx	%ecx	%cx	%cl
%rsi	%esx	%si	%sil
%rdi	%edi	%di	%dil

...the rest is omitted...

5.2 Instructions

Instruction suffixes:

"byte" (b)	1-byte
"word" (w)	2-bytes
"doubleword" (l)	4-bytes
"quadword" (q)	8-bytes

Complete memory addressing mode: A memory operand of the form $D(Rb, Ri, S)$ accesses memory at address $D + \text{val}(Rb) + \text{val}(Ri) * S$, where $\text{val}(Rb)$ and $\text{val}(Ri)$ refer to the value of registers Rb and Ri respectively, D is a constant, and S is a constant of value 1, 2, 4, or 8.

Sign extension and zero extension:

$\text{movzq } S, D$	copy 4-byte-sized S to 8-byte-sized D and fill in the higher order 4 bytes of D with zero bytes
$\text{movslq } S, D$	copy 4-byte-sized S to 8-byte-sized D and sign extend the higher order 4 bytes of D , i.e. fill with 0s if S 's sign bit is zero and fill with 1s if S 's sign bit is one.

Basic Arithmetic instructions that you might not remember:

$\text{sal / shl } k, D$	Left shift destination D by k bits
sar	Arithmetic right shift destination D by k bits
shr	Logical right shift destination D by k bits

Jump instructions:

Jump instruction following `cmp S, D`:

<code>jmp</code>	Unconditional jump
<code>je</code>	Jump if D is equal to S
<code>jne</code>	Jump if D is not equal to S
<code>jg</code>	Jump if D is greater than S (signed)
<code>jge</code>	Jump if D is greater or equal than S (signed)
<code>jl</code>	Jump if D is less than S (signed)
<code>jle</code>	Jump if D is less or equal than S (signed)
<code>ja</code>	Jump if D is above S (unsigned)
<code>jae</code>	Jump if D above or equal S(unsigned)
<code>jb</code>	Jump is D is below S (unsigned)
<code>jbe</code>	Jump if D is below or equal S (unsigned)

5.3 Calling convention

Argument Passing:

Which argument	Stored in register
1	<code>%rdi</code>
2	<code>%rsi</code>
3	<code>%rdx</code>
4	<code>%rcx</code>
5	<code>%r8</code>
6	<code>%r9</code>
7 and up	passed on stack

Return value (if any) is stored in `%rax`

Caller save registers: `%rax`, `%rcx`, `%rdx`, `%rdi`, `%rsi`, `%r8-11`

Callee save registers: `%rbx`, `%rbp`, `%r12-15`

6 Appendix: ASCII

The following table contains the 128 ASCII characters.

Oct	Dec	Hex	Char	Oct	Dec	Hex	Char
000	0	00	NUL '\0'	100	64	40	@
001	1	01	SOH (start of heading)	101	65	41	A
002	2	02	STX (start of text)	102	66	42	B
003	3	03	ETX (end of text)	103	67	43	C
004	4	04	EOT (end of transmission)	104	68	44	D
005	5	05	ENQ (enquiry)	105	69	45	E
006	6	06	ACK (acknowledge)	106	70	46	F
007	7	07	BEL '\a' (bell)	107	71	47	G
010	8	08	BS '\b' (backspace)	110	72	48	H
011	9	09	HT '\t' (horizontal tab)	111	73	49	I
012	10	0A	LF '\n' (new line)	112	74	4A	J
013	11	0B	VT '\v' (vertical tab)	113	75	4B	K
014	12	0C	FF '\f' (form feed)	114	76	4C	L
015	13	0D	CR '\r' (carriage ret)	115	77	4D	M
016	14	0E	SO (shift out)	116	78	4E	N
017	15	0F	SI (shift in)	117	79	4F	O
020	16	10	DLE (data link escape)	120	80	50	P
021	17	11	DC1 (device control 1)	121	81	51	Q
022	18	12	DC2 (device control 2)	122	82	52	R
023	19	13	DC3 (device control 3)	123	83	53	S
024	20	14	DC4 (device control 4)	124	84	54	T
025	21	15	NAK (negative ack.)	125	85	55	U
026	22	16	SYN (synchronous idle)	126	86	56	V
027	23	17	ETB (end of trans. blk)	127	87	57	W
030	24	18	CAN (cancel)	130	88	58	X
031	25	19	EM (end of medium)	131	89	59	Y
032	26	1A	SUB (substitute)	132	90	5A	Z
033	27	1B	ESC (escape)	133	91	5B	[
034	28	1C	FS (file separator)	134	92	5C	\ '\\'
035	29	1D	GS (group separator)	135	93	5D]
036	30	1E	RS (record separator)	136	94	5E	^
037	31	1F	US (unit separator)	137	95	5F	_
040	32	20	SPACE	140	96	60	`
041	33	21	!	141	97	61	a
042	34	22	"	142	98	62	b
043	35	23	#	143	99	63	c
044	36	24	\$	144	100	64	d
045	37	25	%	145	101	65	e
046	38	26	&	146	102	66	f
047	39	27	'	147	103	67	g
050	40	28	(150	104	68	h
051	41	29)	151	105	69	i
052	42	2A	*	152	106	6A	j
053	43	2B	+	153	107	6B	k
054	44	2C	,	154	108	6C	l
055	45	2D	-	155	109	6D	m
056	46	2E	.	156	110	6E	n
057	47	2F	/	157	111	6F	o
060	48	30	0	160	112	70	p
061	49	31	1	161	113	71	q
062	50	32	2	162	114	72	r
063	51	33	3	163	115	73	s
064	52	34	4	164	116	74	t
065	53	35	5	165	117	75	u
066	54	36	6	166	118	76	v
067	55	37	7	167	119	77	w
070	56	38	8	170	120	78	x
071	57	39	9	171	121	79	y
072	58	3A	:	172	122	7A	z
073	59	3B	;	173	123	7B	{
074	60	3C	<	174	124	7C	
075	61	3D	=	175	125	7D	}
076	62	3E	>	176	126	7E	~
077	63	3F	?	177	127	7F	DEL

7 Appendix III: strstr, strcpy, strncpy, strlen

STRSTR(3) Linux Programmer's Manual STRSTR(3)

NAME

strstr - locate a substring

SYNOPSIS

```
#include <string.h>
```

```
char *strstr(const char *haystack, const char *needle);
```

DESCRIPTION

The `strstr()` function finds the first occurrence of the substring `needle` in the string `haystack`. The terminating null bytes (`'\0'`) are not compared.

RETURN VALUE

The `strstr()` function returns a pointer to the beginning of the substring in `haystack`, or `NULL` if the substring is not found.

STRCPY(3) Linux Programmer's Manual STRCPY(3)

NAME

strcpy, strncpy - copy a string

SYNOPSIS

```
#include <string.h>
```

```
char *strcpy(char *dest, const char *src);
char *strncpy(char *dest, const char *src, size_t n);
```

DESCRIPTION

The `strcpy()` function copies the string pointed to by `src`, including the terminating null byte (`'\0'`), to the buffer pointed to by `dest`. The strings may not overlap, and the destination string `dest` must be large enough to receive the copy.

The `strncpy()` function is similar, except that at most `n` bytes of `src` are copied. Warning: If there is no null byte among the first `n` bytes of `src`, the string placed in `dest` will not be null-terminated.

If the length of `src` is less than `n`, `strncpy()` writes additional null bytes to `dest` to ensure that a total of `n` bytes are written.

RETURN VALUE

The `strcpy()` and `strncpy()` functions return a pointer to the destination string `dest`.

STRLEN(3) Linux Programmer's Manual STRLEN(3)

NAME

strlen - calculate the length of a string

SYNOPSIS

```
#include <string.h>
```

```
size_t strlen(const char *s);
```

DESCRIPTION

The `strlen()` function calculates the length of the string pointed to by `s`, excluding the terminating null byte (`'\0'`).

RETURN VALUE

The `strlen()` function returns the number of characters in the string pointed to by `s`.

Q1

A 3 4 6
B 2
C 1 2 3 4
D 1 2 3 4
E 2 4
F 2
G 3
H 1 3

Q2

A 2
B 2,3
C 2
D 3
E 3

Q3

(a) %rdi, %rsi
(b) 5 --> 2
6 --> 3, 4
7 --> 5

(c) s is %rbp
e is %rbx

(d) line 9

(e) line 11 and 12

(f) 1

(g) before: 0x00007fffffffffff08
after: 0x00007fffffffffff00

(h) 0x0000000000000001

(i) swap(&a[s], &a[e]);

Q4

```
char hex2digit(char c)
{
    if (c >= '0' && c <= '9') {
        return c - '0';
    } else if (c >= 'a' && c <= 'f') {
        return c - 'a' + 10;
    }
    assert(0);
}
```

Page 15 of 11

```
int hex2int(char *s) {
    unsigned result = 0;
    for (int i = 0; i < 8; i++) {
        result = result << 4 + hex2digit(s[i]);
    }
    return (int)result;
}
```