

CSO-Recitation 10

CSCI-UA 0201-007

R10: Assessment 08 & Buffer overflow & Compiler optimization

Today's Topics

- Assessment 08
- Breakout exercises
- Buffer overflow
- Compiler optimization

Assessment 08

Q1 Set_five

Given the following C function from Lab 1,

```
void set_five(int *p)
{
    *p = 5;
}

void test()
{
    int p = 0;
    set_five(&p);
}
```

The assembly for `set_five` function is:

```
0x00000000000005fa <+0>: ???
0x0000000000000600 <+6>: retq
```

The assembly for `test` function is:

```
0x0000000000000601 <+0>:    sub    $0x10,%rsp
0x0000000000000605 <+4>:    movl   $0x0,0xc(%rsp)
0x000000000000060d <+12>:    lea    0xc(%rsp),%rdi
0x0000000000000612 <+17>:    callq  0x5fa <set_five>
0x0000000000000617 <+22>:    add    $0x10,%rsp
0x000000000000061b <+26>:    retq
```

Q1.1 %rsp

Under normal program execution, suppose the value of %rsp is 0x7fff856001d8 **just prior to** executing the first instruction of `test`. What is the value of %rsp just prior to executing the first instruction of `set_five`?

- A. 0x7fff856001d8
- B. 0x7fff856001c8
- C. 0x7fff856001e8
- D. 0x7fff856001c0**
- E. 0x7fff856001d0
- F. 0x7fff856001c4
- G. 0x7fff856001cc
- H. None of the above

- `val(%rsp)= 0x7fff856001d8`
- `sub $0x10, %rsp`
 - `val(%rsp)= 0x7fff856001c8`
- `callq 0x5fa <set_five>`
- `callq label`
 - DECREASES %rsp by 8
 - THEN stores the return address at the memory location given by the new %rsp
 - THEN jumps to the operand
- `callq 0x5fa <set_five>`
 - `val(%rsp)= 0x7fff856001c0`

Given the following C function from Lab 1,

```
void set_five(int *p)
{
    *p = 5;
}

void test()
{
    int p = 0;
    set_five(&p);
}
```

The assembly for `set_five` function is:

```
0x00000000000005fa <+0>: ???
0x0000000000000600 <+6>: retq
```

The assembly for `test` function is:

```
0x0000000000000601 <+0>: sub    $0x10,%rsp
0x0000000000000605 <+4>: movl   $0x0,0xc(%rsp)
0x000000000000060d <+12>: lea    0xc(%rsp),%rdi
0x0000000000000612 <+17>: callq  0x5fa <set_five>
0x0000000000000617 <+22>: add    $0x10,%rsp
0x000000000000061b <+26>: retq
```

Q1.2

Under normal program execution, what is the 8-byte value stored under the address specified by `%rsp` **just prior to** executing the first instruction of `set_five`?

- A. `0x7fff856001d8`
- B. `0x7fff856001c0`
- C. `0x0000000000000060d`
- D. `0x00000000000000612`
- E. `0x00000000000000617`
- F. It could be any arbitrary 8-byte value

- `callq label`
 - DECREASES `%rsp` by 8
 - THEN stores the **return address** at the memory location given by the new `%rsp`
 - THEN jumps to the operand
- `callq 0x5fa <set_five>`
 - `val((%rsp))`
=`0x00000000000000617`

Q1.3

After executing instruction 0x00000000000000600 <+6>: retq in set_five, what's new %rip value?

A. 0x0000000000000060d

B. 0x00000000000000612

C. 0x00000000000000617

D. 0x00000000000000604

E. 0x00000000000000608

- `retq`
 - Jumps to the location given by the value in memory located at `%rsp`
 - `%rip = mem[%rsp]`
 - `val((%rsp))`
`=0x00000000000000617`
 - THEN INCREASES `%rsp` by 8
 - `%rsp = %rsp + 8`

Q1.4 p's location

Where is the local variable p stored?

- A. some register
- B. memory (data segment)
- C. memory (stack)**
- D. memory (heap)

Given the following C function from Lab 1,

```
void set_five(int *p)
{
    *p = 5;
}

void test()
{
    int p = 0;
    set_five(&p);
}
```

The assembly for `set_five` function is:

```
0x00000000000005fa <+0>: ???
0x0000000000000600 <+6>: retq
```

The assembly for `test` function is:

```
0x0000000000000601 <+0>: sub    $0x10,%rsp
0x0000000000000605 <+4>: movl   $0x0,0xc(%rsp)
0x000000000000060d <+12>: lea    0xc(%rsp),%rdi
0x0000000000000612 <+17>: callq  0x5fa <set_five>
0x0000000000000617 <+22>: add    $0x10,%rsp
0x000000000000061b <+26>: retq
```

- Local variables -> stack
- If C's primitive data type and pointer
 - registers
 - stack
- Here, 0xc(%rsp) -> some place in memory
- -> stack

Q1.5 p's location

If your answer of 1.4 is memory, where in memory (aka what address) is p stored (assuming the value of %rsp is 0x7fff856001d8 just prior to executing the first instruction of test)?

A. 0x7fff856001d8

B. 0x7fff856001c8

C. 0x7fff856001d4

D. 0x7fff856001d0

E. 0x7fff856001cc

- `val(%rsp) = 0x7fff856001d8`
- `sub $0x10, %rsp`
 - `val(%rsp) = 0x7fff856001c8`
- where in memory is p stored
 - `0xc(%rsp)`
 - `0x7fff856001d4`

Given the following C function from Lab 1,

```
void set_five(int *p)
{
    *p = 5;
}

void test()
{
    int p = 0;
    set_five(&p);
}
```

The assembly for `set_five` function is:

```
0x00000000000005fa <+0>: ???
0x0000000000000600 <+6>: retq
```

The assembly for `test` function is:

```
0x0000000000000601 <+0>: sub    $0x10,%rsp
0x0000000000000605 <+4>: movl   $0x0,0xc(%rsp)
0x000000000000060d <+12>: lea    0xc(%rsp),%rdi
0x0000000000000612 <+17>: callq  0x5fa <set_five>
0x0000000000000617 <+22>: add    $0x10,%rsp
0x000000000000061b <+26>: retq
```

Q1.6 set_five

What's the missing first instruction of set_five (aka the instruction corresponding to ???)

A. ``movl $0x5, (%rdi)``

B. ``movq $0x5, (%rdi)``

C. ``movl $0x5, (%rsi)``

D. ``movq $0x5, (%rdi)``

E. ``movl $0x5, %edi``

F. ``movq $0x5, %rdi``

G. ``movl $0x5, %esi``

H. ``movq $0x5, %rsi``

- `*p=5;`
- `p -> the first argument`
 - store in `%rdi`
- `mov $0x5, (%rdi)`
- `p` is `int *` type
 - `movl $0x5, (%rdi)`

Given the following C function from Lab 1,

```
void set_five(int *p)
{
    *p = 5;
}

void test()
{
    int p = 0;
    set_five(&p);
}
```

The assembly for `set_five` function is:

```
0x00000000000005fa <+0>: ???
0x0000000000000600 <+6>: retq
```

The assembly for `test` function is:

```
0x0000000000000601 <+0>: sub    $0x10,%rsp
0x0000000000000605 <+4>: movl   $0x0,0xc(%rsp)
0x000000000000060d <+12>: lea    0xc(%rsp),%rdi
0x0000000000000612 <+17>: callq  0x5fa <set_five>
0x0000000000000617 <+22>: add    $0x10,%rsp
0x000000000000061b <+26>: retq
```

Q2 array → Q2.1

What is the value of *c[1] after executing line 11?

A. 1

B. 2

C. 3

D. 4

E. 5

F. None of the above

- c's type is int **
- c[i] == *(c+i)
 - c+i -> int ** (pointer arithmetic)
 - c[1]==*(c+1) -> int *
 - c[1] == b == &b[0]
- c[1]++;
 - pointer arithmetic
 - c[1] = c[1]+1 == &b[1]
- *c[1]==b[1]==4

Given the following C code

```
1: void foo(int **arr, long i)
2: {
3:     arr[i]++;
4: }
5:
6: void test()
7: {
8:     int a[2] = {1, 2};
9:     int b[2] = {3, 4};
10:    int *c[2] = {a, b};
11:    foo(c, 1);
12:}
```

Q2.2 arr[i]++

If Line 3 is realized using one instruction, what's that instruction?

- A. ``addl $0x8, (%rdi,%rsi,8)``
- B. ``addl $0x4, (%rdi,%rsi,8)``
- C. ``addl $0x8, (%rdi,%rsi,4)``
- D. ``addl $0x4, (%rdi,%rsi,4)``
- E. ``addq $0x8, (%rdi,%rsi,8)``
- F. ``addq $0x4, (%rdi,%rsi,8)``
- G. ``addq $0x8, (%rdi,%rsi,4)``
- H. ``addq $0x4, (%rdi,%rsi,4)``

- `arr[i]++;`
- `arr[i]`
 - `(%rdi, %rsi, size)`
 - `arr -> int **, arr[i] -> int * -> 8-byte`
 - `(%rdi, %rsi, 8)`
- `arr[i] = arr[i]+1;`
- `add: dest = dest + src`
 - `dest == arr[i] -> int * -> 8-byte`
 - `addq`
- `src`
 - `arr[i]=arr[i]+1` is a pointer arithmetic
 - `" +1 " -> 1*size(element)`
 - `arr[i] -> int *, arr[i][j] -> int -> 4-byte`
 - `1*4 == 0x4`
- `addq $0x4, (%rdi,%rsi,8)`

Q3 → Q3.1 location of p

Where is the local variable t in test stored?

- A. some register
- B. memory (data segment)
- C. memory (stack)
- D. memory (heap)

Given the following C code

```
1: typedef struct kv_pair {
2:     long key;
3:     char* val;
4: } kv_pair;
5:
6: void init_pair(kv_pair *p)
7: {
8:     p->key = -1;
9:     p->val = NULL;
10:}
11:
12: void test()
13: {
14:     kv_pair t;
15:     init_pair(&t);
16:}
```

Q3.2 p->val

If Line 9 is realized using one instruction, what is that instruction?

- A. ``movl $0x0,0x4(%rdi)``
- B. ``movq $0x0,0x4(%rdi)``
- C. ``movl $0x0,0x8(%rdi)``
- D. ``movq $0x0,0x8(%rdi)``
- E. ``movl $0x0,0x4(%rsi)``
- F. ``movq $0x0,0x4(%rsi)``
- G. ``movl $0x0,0x8(%rsi)``
- H. ``movq $0x0,0x8(%rsi)``

- `p->val = NULL;`
- `p` is `kv_pair *` type
 - `val(p)` is stored in memory (stack)
 - `(%rdi)`
- the start address of `p->val`
 - `0x8(%rdi)`
 - because first 8-byte stores “key”, next 8-byte stores “val”
- “val” has type `char *`
 - 8-byte
 - `movq`
- `movq $0x0, 0x8(%rdi)`

Exercises

Discuss the answer

Ex1

1. Consider the following scenario:

%rsp is 0x7fffffff448

The following instructions execute:

...

pushq %rbp

pushq %rax

...

What is the new value for %rsp?

- `pushq src`
 - DECREASES %rsp by 8
 - THEN stores the operand at the memory location given by the new %rsp
- %rsp -> 0x7fffffff448
 - 0x7fffffff440
 - 0x7fffffff438
- 0x7fffffff438

Ex2

2. Consider the following scenario:

%rsp is 0x7fffffff448

The following instructions execute:

``

callq my_cool_function

``

While `my_cool_function` is executing, what is the value for %rsp?

When `my_cool_function` executes `retq`, what will be the value for %rsp?

- `callq label`
 - DECREASES %rsp by 8
 - THEN stores the **return address** at the memory location given by the new %rsp
 - THEN jumps to the operand
- %rsp -> 0x7fffffff448
 - 0x7fffffff440
- `retq`
 - Jumps to the location given by the value in memory located at %rsp
 - THEN INCREASES %rsp by 8
- 0x7fffffff448

Ex3

3. Consider the following scenario:

%rsi is 5

%rdi is 8

The following instructions execute:

...

leaq 40(%rdi, %rsi, 8), %rax

...

What is the value of %rax?

Memory looks like this:

0x00:	10
0x08:	24
0x10:	32
0x18:	59
0x20:	23
0x28:	1
0x30:	66
0x38:	10000000
0x40:	2607
0x48:	2019
0x50:	111
0x58:	17
0x60:	32

- “lea” is no memory access, only do arithmetic calculation
- leaq 40(%rdi, %rsi, 8), %rax
- $\text{val}(\%rax) = \text{val}(\%rdi) + 8 * \text{val}(\%rsi) + 40$
 - $= 8 + 8 * 5 + 40 = 88 = 0x58$

Ex4

4. Consider the following scenario:

%rsi is 5

%rdi is 8

The following instructions execute:

...

```
movq 40(%rdi, %rsi, 8), %rax
```

...

What is the value of %rax?

Memory looks like this:

0x00:	10
0x08:	24
0x10:	32
0x18:	59
0x20:	23
0x28:	1
0x30:	66
0x38:	10000000
0x40:	2607
0x48:	2019
0x50:	111
0x58:	17
0x60:	32

- `movq 40(%rdi, %rsi, 8), %rax`
- `val(%rax) = mem[val(%rdi)+8*val(%rsi)+40]`
 - = `mem[0x58]`
 - = **17**

Buffer Overflow

Not all buggy memory references access “illegal” memory

Buffer Overflow

- Have learnt about the memory layout
- If an instruction tries to access some invalid memory
 - Segmentation fault occurs
- But not all buggy memory references access “illegal” memory
 - Buffer overflow exploits

Buffer Overflow

- Buffer overflow on the stack
- Buffer overflow overwrites the return address
 - attackers may carefully chosen return address, executes malicious code
 - code injection buffer overflow attacks

Defenses

- Write correct code to avoid overflow vulnerability
 - Use safe APIs to limit buffer lengths
 - Use a memory-safe language
- Mitigate attacks despite buggy code
 - will be an always on-going project, attack and defense itself are alternately developed
 - Security research domain
 - One idea to prevent control flow hijacking: catch over-written return address before invocation
 - place special value (“canary”) just beyond buffer
 - GCC implementation: -fstack-protector

Compiler optimization

Tries to minimize or maximize some attributes of an executable computer program

Optimizing compiler

- Goal: generate efficient, correct machine code
 - generally implemented using a sequence of *optimizing transformations*
 - algorithms which take a program and transform it to produce a semantically equivalent output program that uses fewer resources and/or executes faster
- The compiler performs optimization based on the knowledge it has of the program
 - Compiling multiple files at once to a single output file mode allows the compiler to use information gained from all of the files when compiling each of them
- Common optimization
 - code motion
 - use simpler instructions
 - reuse common subexpressions

Optimization -- GCC

- Without any optimization option, the compiler's goal is to reduce the cost of compilation and to make debugging produce the expected results
- Turning on optimization flags makes the compiler
 - attempt to improve the performance and/or code size at the expense of compilation time and possibly the ability to debug the program
- When debugging your code, it may help to disable optimizations
 - Add `-O0` to the Makefile CFLAGS
- Depending on the target and how GCC was configured, a slightly different set of optimizations may be enabled at each -O level
 - You can invoke GCC with `-Q --help=optimizers` to find out the exact set of optimizations that are enabled at each level

Optimization -- GCC

- gcc's optimization levels: -O, -O2, -O3, -Og, -O0, -Os, -Ofast
- With -O, the compiler tries to reduce code size and execution time
 - without performing any optimizations that take a great deal of compilation time
- -O2 optimize even more
 - turns on all optimization flags specified by -O, and it also turns on some other optimization flags: e.g. -finline-small-functions...
- -Og optimize debugging experience
 - offering a reasonable level of optimization while maintaining fast compilation and a good debugging experience
 - enables all -O1 optimization flags except for those that may interfere with debugging