

CSO-Recitation 10

CSCI-UA 0201-007

R10: Assessment 08 & Buffer overflow & Compiler optimization

Today's Topics

- Assessment 08
- Breakout exercises
- Buffer overflow
- Compiler optimization

Exercises

Discuss the answer

Ex1

1. Consider the following scenario:

%rsp is 0x7fffffff448

The following instructions execute:

...

pushq %rbp

pushq %rax

...

What is the new value for %rsp?

- `pushq src`
 - DECREASES %rsp by 8
 - THEN stores the operand at the memory location given by the new %rsp
- %rsp -> 0x7fffffff448
 - 0x7fffffff440
 - 0x7fffffff438
- 0x7fffffff438

Ex2

2. Consider the following scenario:

%rsp is 0x7fffffff448

The following instructions execute:

``

callq my_cool_function

``

While `my_cool_function` is executing, what is the value for %rsp?

When `my_cool_function` executes `retq`, what will be the value for %rsp?

- callq label
 - DECREASES %rsp by 8
 - THEN stores the **return address** at the memory location given by the new %rsp
 - THEN jumps to the operand
- %rsp -> 0x7fffffff448
 - 0x7fffffff440
- retq
 - Jumps to the location given by the value in memory located at %rsp
 - THEN INCREASES %rsp by 8
- 0x7fffffff448

Ex3

3. Consider the following scenario:

%rsi is 5

%rdi is 8

The following instructions execute:

...

leaq 40(%rdi, %rsi, 8), %rax

...

What is the value of %rax?

Memory looks like this:

0x00:	10
0x08:	24
0x10:	32
0x18:	59
0x20:	23
0x28:	1
0x30:	66
0x38:	10000000
0x40:	2607
0x48:	2019
0x50:	111
0x58:	17
0x60:	32

- “lea” is no memory access, only do arithmetic calculation
- leaq 40(%rdi, %rsi, 8), %rax
- $\text{val}(\%rax) = \text{val}(\%rdi) + 8 * \text{val}(\%rsi) + 40$
 - $= 8 + 8 * 5 + 40 = 88 = 0x58$

Ex4

4. Consider the following scenario:

%rsi is 5

%rdi is 8

The following instructions execute:

...

```
movq 40(%rdi, %rsi, 8), %rax
```

...

What is the value of %rax?

Memory looks like this:

0x00:	10
0x08:	24
0x10:	32
0x18:	59
0x20:	23
0x28:	1
0x30:	66
0x38:	10000000
0x40:	2607
0x48:	2019
0x50:	111
0x58:	17
0x60:	32

- `movq 40(%rdi, %rsi, 8), %rax`
- `val(%rax) = mem[val(%rdi)+8*val(%rsi)+40]`
 - = `mem[0x58]`
 - = **17**

Buffer Overflow

Not all buggy memory references access “illegal” memory

Buffer Overflow

- Have learnt about the memory layout
- If an instruction tries to access some invalid memory
 - Segmentation fault occurs
- But not all buggy memory references access “illegal” memory
 - Buffer overflow exploits

Buffer Overflow

- Buffer overflow on the stack
- Buffer overflow overwrites the return address
 - attackers may carefully chosen return address, executes malicious code
 - code injection buffer overflow attacks

Defenses

- Write correct code to avoid overflow vulnerability
 - Use safe APIs to limit buffer lengths
 - Use a memory-safe language
- Mitigate attacks despite buggy code
 - will be an always on-going project, attack and defense itself are alternately developed
 - Security research domain
 - One idea to prevent control flow hijacking: catch over-written return address before invocation
 - place special value (“canary”) just beyond buffer
 - GCC implementation: -fstack-protector

Compiler optimization

Tries to minimize or maximize some attributes of an executable computer program

Optimizing compiler

- Goal: generate efficient, correct machine code
 - generally implemented using a sequence of *optimizing transformations*
 - algorithms which take a program and transform it to produce a semantically equivalent output program that uses fewer resources and/or executes faster
- The compiler performs optimization based on the knowledge it has of the program
 - Compiling multiple files at once to a single output file mode allows the compiler to use information gained from all of the files when compiling each of them
- Common optimization
 - code motion
 - use simpler instructions
 - reuse common subexpressions

Optimization -- GCC

- Without any optimization option, the compiler's goal is to reduce the cost of compilation and to make debugging produce the expected results
- Turning on optimization flags makes the compiler
 - attempt to improve the performance and/or code size at the expense of compilation time and possibly the ability to debug the program
- When debugging your code, it may help to disable optimizations
 - Add `-O0` to the Makefile CFLAGS
- Depending on the target and how GCC was configured, a slightly different set of optimizations may be enabled at each -O level
 - You can invoke GCC with `-Q --help=optimizers` to find out the exact set of optimizations that are enabled at each level

Optimization -- GCC

- gcc's optimization levels: -O, -O2, -O3, -Og, -O0, -Os, -Ofast
- With -O, the compiler tries to reduce code size and execution time
 - without performing any optimizations that take a great deal of compilation time
- -O2 optimize even more
 - turns on all optimization flags specified by -O, and it also turns on some other optimization flags: e.g. -finline-small-functions...
- -Og optimize debugging experience
 - offering a reasonable level of optimization while maintaining fast compilation and a good debugging experience
 - enables all -O1 optimization flags except for those that may interfere with debugging