

# CSO-Recitation 11

CSCI-UA 0201-007

R11: Assessment 09 & Dynamic memory allocation

# Today's Topics

- Assessment 09
- Dynamic memory allocation
  - implement your malloc & free

# Assessment 09

# Q1 seg fault

Which of the following statements on segmentation faults are true?

- A. Performing any out-of-bounds array access will result in an immediate segmentation fault.
- B. Segmentation faults only occur when an instruction tries to write to memory.
- C. Dereferencing a null pointer will always result in a segmentation fault.
- D. Performing pointer arithmetic will always result in a segmentation fault.
- E. none of the above

## Q2 local variable

Suppose *local variable* `a` is defined as `int a[16]`; Which of the following statements are true?

A. `a` takes up space on the stack.

B. `a` takes up space on the heap.

C. `subq $64, %rsp` allocates space for `a` and `addq $64, %rsp` de-allocates space for `a`.  
stack goes downwards

D. `addq $64, %rsp` allocates space for `a` and `subq $64, %rsp` de-allocates space for `a`.

E. none of the above

# Q3 str\_concat

The following C function str\_concat appends the src string to the dst string, overwriting the terminating null byte at the end of dst, and then adds a terminating null byte.

- **Q3.1** line 5
- Please fill in the code at line 6 (must be a one liner). *To facilitate automatic grading, please do not have any spaces in your C statement, and make sure to include the end of the statement semicolon.*
- `dst[len+i]=src[i];`

```
1: void str_concat(char *dst, char *src)
2: {
3:     int len = strlen(dst);
4:     int i;
5:     for (i = 0; src[i]!='\0'; i++) {
6:         ???
7:     }
8:     ???
9: }
```

## Q3.2 line 7

Which of the following C statement can be used at line 8 correctly, without compilation nor runtime error?

A. `dst[len] = '\0';`

B. `dst[len] = NULL;`

C. `dst[len] = 0;`

D. `dst[len+i] = '\0';`

E. `dst[len+i] = NULL;`

F. `dst[len+i] = 0;`

NULL: means a null pointer

Q4. Given the following C program (which invokes str\_concat defined in Q3)

```
void dangerous()
{
    char buf1[8] = "hello";
    char buf2[8] = "world";
    str_concat(buf1, buf2);
}

int main()
{
    dangerous();
    printf("i wonder if this program is correct\n");
}
```

Suppose the following assembly is generated for the above C program (including str\_concat). Please assume that the addresses to the left of each instruction shown below are the actual addresses where the instructions are stored at during runtime.

```
000000000550068a <str_concat>:
    68a:    48 83 ec 28          subq    $0x18,%rsp
    ...omitted....
    711:    c3                  ret

00000000055006ba <dangerous>:
    6ba:    48 83 ec 10          subq    $0x10,%rsp
    6be:    48 b8 68 65 6c 6c 6f movq    $0xf6c6c6568,%rax
    6c8:    48 89 44 24 08        movq    %rax,0x8(%rsp)
    6cd:    48 b8 77 6f 72 6c 64 movq    $0x646c726f77,%rax
    6d7:    48 89 04 24           movq    %rax,(%rsp)
    6db:    48 89 e6              movq    %rsp,%rsi
    6de:    48 8d 7c 24 08        leaq    0x8(%rsp),%rdi
    6e3:    e8 82 ff ff ff        callq   66a <str_concat>
    6e8:    48 83 c4 10           addq    $0x10,%rsp
    6ec:    c3                  ret

00000000055006ed <main>:
    6ed:    48 83 ec 08          subq    $0x8,%rsp
    6f1:    b8 00 00 00 00        movq    $0x0,%eax
    6f6:    e8 bf ff ff ff        callq   6ba <dangerous>
    6fb:    48 8d 35 a2 00 00 00 leaq    0xa2(%rip),%rsi
    702:    bf 01 00 00 00        movq    $0x1,%edi
    707:    b8 00 00 00 00        movq    $0x0,%eax
    70c:    e8 2f fe ff ff        callq   540 <__printf_chk@plt>
    711:    b8 00 00 00 00        movq    $0x0,%eax
    716:    48 83 c4 08           addq    $0x8,%rsp
    71a:    c3                  ret
```



# Q4

```
000000000550068a <str_concat>:
 68a: 48 83 ec 28      subq    $0x18,%rsp
    ...omitted....
 711: c3              ret

00000000055006ba <dangerous>:
 6ba: 48 83 ec 10      subq    $0x10,%rsp
 6be: 48 b8 68 65 6c 6c 6f movq    $0x6f6c6c6c6568,%rax
 6c8: 48 89 44 24 08    movq    %rax,0x8(%rsp)
 6cd: 48 b8 77 6f 72 6c 64 movq    $0x646c726f77,%rax
 6d7: 48 89 04 24      movq    %rax,(%rsp)
 6db: 48 89 e6         movq    %rsp,%rsi
 6de: 48 8d 7c 24 08    leaq    0x8(%rsp),%rdi
 6e3: e8 82 ff ff ff   callq   66a <str_concat>
 6e8: 48 83 c4 10      addq    $0x10,%rsp
 6ec: c3              ret

00000000055006ed <main>:
 6ed: 48 83 ec 08      subq    $0x8,%rsp
 6f1: b8 00 00 00 00    movq    $0x0,%eax
 6f6: e8 bf ff ff ff   callq   6ba <dangerous>
 6fb: 48 8d 35 a2 00 00 00 leaq    0xa2(%rip),%rsi
 702: bf 01 00 00 00    movq    $0x1,%edi
 707: b8 00 00 00 00    movq    $0x0,%eax
 70c: e8 2f fe ff ff   callq   540 <__printf_chk@plt>
 711: b8 00 00 00 00    movq    $0x0,%eax
 716: 48 83 c4 08      addq    $0x8,%rsp
 71a: c3              ret
```

-0x8

← rsp



```

000000000550068a <str_concat>:
    68a:      48 83 ec 28          subq    $0x18,%rsp
    ...omitted....
    711:      c3                 ret

00000000055006ba <dangerous>:
    6ba:      48 83 ec 10          subq    $0x10,%rsp
    6be:      48 b8 68 65 6c 6c 6f movq    $0x6f6c6c6568,%rax
    6c8:      48 89 44 24 08          movq    %rax,0x8(%rsp)
    6cd:      48 b8 77 6f 72 6c 64 movq    $0x646c726f77,%rax
    6d7:      48 89 04 24          movq    %rax,(%rsp)
    6db:      48 89 e6            movq    %rsp,%rsi
    6de:      48 8d 7c 24 08          leaq    0x8(%rsp),%rdi
    6e3:      e8 82 ff ff ff          callq   66a <str_concat>
    6e8:      48 83 c4 10          addq    $0x10,%rsp
    6ec:      c3                 ret

00000000055006ed <main>:
    6ed:      48 83 ec 08          subq    $0x8,%rsp
    6f1:      b8 00 00 00 00          movq    $0x0,%eax
    6f6:      e8 bf ff ff ff          callq   6ba <dangerous>
    6fb:      48 8d 35 a2 00 00 00 leaq     0xa2(%rip),%rsi
    702:      bf 01 00 00 00          movq    $0x1,%edi
    707:      b8 00 00 00 00          movq    $0x0,%eax
    70c:      e8 2f fe ff ff          callq   540 <__printf_chk@plt>
    711:      b8 00 00 00 00          movq    $0x0,%eax
    716:      48 83 c4 08          addq    $0x8,%rsp
    71a:      c3                 ret

```

0x7ffd6a5e3010



to executing the first instruction of dangerous

# Q4

```

000000000550068a <str_concat>:
    68a:    48 83 ec 28             subq    $0x18,%rsp
    ...omitted....
    711:    c3                     ret

00000000055006ba <dangerous>:
    6ba:    48 83 ec 10             subq    $0x10,%rsp
    6be:    48 b8 68 65 6c 6c 6f    movq    $0x6f6c6c6568,%rax
    6c8:    48 89 44 24 08           movq    %rax,0x8(%rsp)
    6cd:    48 b8 77 6f 72 6c 64    movq    $0x646c726f77,%rax
    6d7:    48 89 04 24             movq    %rax,(%rsp)
    6db:    48 89 e6               movq    %rsp,%rsi
    6de:    48 8d 7c 24 08          leaq    0x8(%rsp),%rdi
    6e3:    e8 82 ff ff ff          callq   66a <str_concat>
    6e8:    48 83 c4 10             addq    $0x10,%rsp
    6ec:    c3                     ret

00000000055006ed <main>:
    6ed:    48 83 ec 08             subq    $0x8,%rsp
    6f1:    b8 00 00 00 00          movq    $0x0,%eax
    6f6:    e8 bf ff ff ff          callq   6ba <dangerous>
    6fb:    48 8d 35 a2 00 00 00    leaq    0xa2(%rip),%rsi
    702:    bf 01 00 00 00          movq    $0x1,%edi
    707:    b8 00 00 00 00          movq    $0x0,%eax
    70c:    e8 2f fe ff ff          callq   540 <__printf_chk@plt>
    711:    b8 00 00 00 00          movq    $0x0,%eax
    716:    48 83 c4 08             addq    $0x8,%rsp
    71a:    c3                     ret

```

0x7ffd6a5e3030

0x7ffd6a5e3028

0x7ffd6a5e3020

0x7ffd6a5e3018

0x7ffd6a5e3010

0x7ffd6a5e3008



← rsp

Suppose the value of %rsp is 0x7ffd6a5e3020 just prior to executing the first instruction of dangerous

# Q4

```

000000000550068a <str_concat>:
 68a: 48 83 ec 28      subq    $0x18,%rsp
...omitted....
 711: c3              ret

00000000055006ba <dangerous>:
 6ba: 48 83 ec 10      subq    $0x10,%rsp
 6be: 48 b8 68 65 6c 6c 6f  movq    $0x6f6c6c6568,%rax
 6c8: 48 89 44 24 08    movq    %rax,0x8(%rsp)
 6cd: 48 b8 77 6f 72 6c 64  movq    $0x646c726f77,%rax
 6d7: 48 89 04 24      movq    %rax,(%rsp)
 6db: 48 89 e6         movq    %rsp,%rsi
 6de: 48 8d 7c 24 08    leaq    0x8(%rsp),%rdi
 6e3: e8 82 ff ff ff    callq   66a <str_concat>
 6e8: 48 83 c4 10      addq    $0x10,%rsp
 6ec: c3              ret

00000000055006ed <main>:
 6ed: 48 83 ec 08      subq    $0x8,%rsp
 6f1: b8 00 00 00 00    movq    $0x0,%eax
 6f6: e8 bf ff ff ff    callq   6ba <dangerous>
 6fb: 48 8d 35 a2 00 00 00  leaq    0xa2(%rip),%rsi
 702: bf 01 00 00 00    movq    $0x1,%edi
 707: b8 00 00 00 00    movq    $0x0,%eax
 70c: e8 2f fe ff ff    callq   540 <__printf_chk@plt>
 711: b8 00 00 00 00    movq    $0x0,%eax
 716: 48 83 c4 08      addq    $0x8,%rsp
 71a: c3              ret

```

Suppose the value of %rsp is 0x7ffd6a5e3020 just prior to executing the first instruction of dangerous

0x7ffd6a5e3030


0x7ffd6a5e3028

0x7ffd6a5e3020

0x7ffd6a5e3018

0x7ffd6a5e3010

0x7ffd6a5e3008



← rsp

```

1: void str_concat(char *dst, char *src)
2: {
3:     int len = strlen(dst);
4:     int i;
5:     for (i = 0; src[i]!='\0'; i++) {
6:         ???
7:     }
8:     ???
9: }

```

# Q4

```

000000000550068a <str_concat>:
    68a:    48 83 ec 28      subq    $0x18,%rsp
    ...omitted....
    711:    c3              ret

00000000055006ba <dangerous>:
    6ba:    48 83 ec 10      subq    $0x10,%rsp
    6be:    48 b8 68 65 6c 6c 6f  movq    $0x6f6c6c6c6568,%rax
    6c8:    48 89 44 24 08      movq    %rax,0x8(%rsp)
    6cd:    48 b8 77 6f 72 6c 64  movq    $0x646c726f77,%rax
    6d7:    48 89 04 24          movq    %rax,(%rsp)
    6db:    48 89 e6           movq    %rsp,%rsi
    6de:    48 8d 7c 24 08      leaq    0x8(%rsp),%rdi
    6e3:    e8 82 ff ff ff      callq   66a <str_concat>
    6e8:    48 83 c4 10        addq    $0x10,%rsp
    6ec:    c3              ret

00000000055006ed <main>:
    6ed:    48 83 ec 08      subq    $0x8,%rsp
    6f1:    b8 00 00 00 00      movq    $0x0,%eax
    6f6:    e8 bf ff ff ff      callq   6ba <dangerous>
    6fb:    48 8d 35 a2 00 00 00  leaq    0xa2(%rip),%rsi
    702:    bf 01 00 00 00      movq    $0x1,%edi
    707:    b8 00 00 00 00      movq    $0x0,%eax
    70c:    e8 2f fe ff ff      callq   540 <__printf_chk@plt>
    711:    b8 00 00 00 00      movq    $0x0,%eax
    716:    48 83 c4 08        addq    $0x8,%rsp
    71a:    c3              ret

```

Suppose the value of %rsp is 0x7ffd6a5e3020 just prior to executing the first instruction of dangerous

0x7ffd6a5e3030


0x7ffd6a5e3028

0x7ffd6a5e3020

0x7ffd6a5e3018

0x7ffd6a5e3010

0x7ffd6a5e3008



← rsp

```

1: void str_concat(char *dst, char *src)
2: {
3:     int len = strlen(dst);
4:     int i;
5:     for (i = 0; src[i]!='\0'; i++) {
6:         ???
7:     }
8:     ???
9: }

```

# Q4

```

000000000550068a <str_concat>:
 68a: 48 83 ec 28      subq    $0x18,%rsp
...omitted....
 711: c3              ret

00000000055006ba <dangerous>:
 6ba: 48 83 ec 10      subq    $0x10,%rsp
 6be: 48 b8 68 65 6c 6c 6f  movq    $0x6f6c6c6c6568,%rax
 6c8: 48 89 44 24 08    movq    %rax,0x8(%rsp)
 6cd: 48 b8 77 6f 72 6c 64  movq    $0x646c726f77,%rax
 6d7: 48 89 04 24      movq    %rax,(%rsp)
 6db: 48 89 e6         movq    %rsp,%rsi
 6de: 48 8d 7c 24 08    leaq    0x8(%rsp),%rdi
 6e3: e8 82 ff ff ff    callq   66a <str_concat>
 6e8: 48 83 c4 10      addq    $0x10,%rsp
 6ec: c3              ret

00000000055006ed <main>:
 6ed: 48 83 ec 08      subq    $0x8,%rsp
 6f1: b8 00 00 00 00    movq    $0x0,%eax
 6f6: e8 bf ff ff ff    callq   6ba <dangerous>
 6fb: 48 8d 35 a2 00 00 00  leaq    0xa2(%rip),%rsi
 702: bf 01 00 00 00    movq    $0x1,%edi
 707: b8 00 00 00 00    movq    $0x0,%eax
 70c: e8 2f fe ff ff    callq   540 <__printf_chk@plt>
 711: b8 00 00 00 00    movq    $0x0,%eax
 716: 48 83 c4 08      addq    $0x8,%rsp
 71a: c3              ret

```

0x7ffd6a5e3030

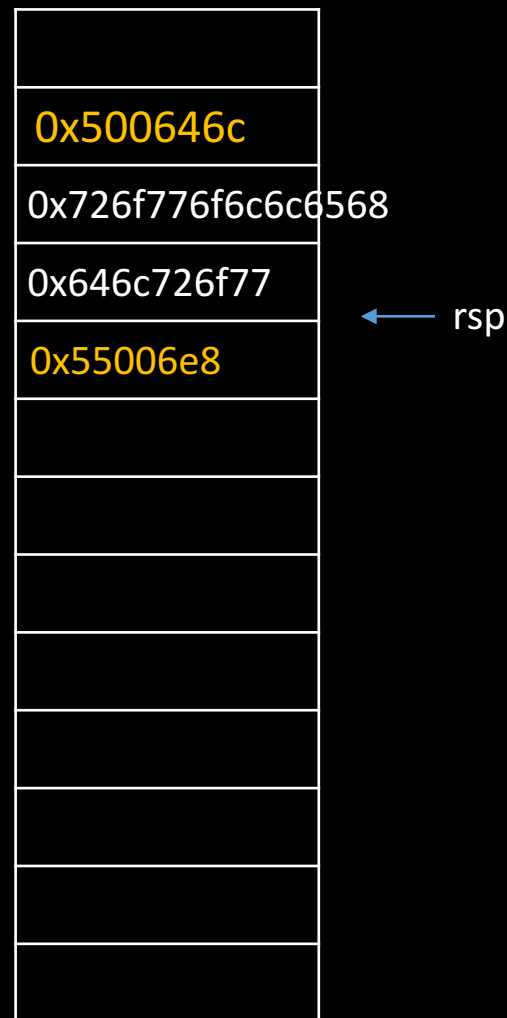
0x7ffd6a5e3028

0x7ffd6a5e3020

0x7ffd6a5e3018

0x7ffd6a5e3010

0x7ffd6a5e3008



Suppose the value of %rsp is 0x7ffd6a5e3020 just prior to executing the first instruction of dangerous

```

000000000550068a <str_concat>:
    68a:      48 83 ec 28                subq    $0x18,%rsp
    ...omitted....
    711:      c3                       ret

00000000055006ba <dangerous>:
    6ba:      48 83 ec 10                subq    $0x10,%rsp
    6be:      48 b8 68 65 6c 6c 6f      movq    $0x6f6c6c6568,%rax
    6c8:      48 89 44 24 08              movq    %rax,0x8(%rsp)
    6cd:      48 b8 77 6f 72 6c 64      movq    $0x646c726f77,%rax
    6d7:      48 89 04 24                movq    %rax,%rsp
    6db:      48 89 e6                   movq    %rsp,%rsi
    6de:      48 8d 7c 24 08              leaq    0x8(%rsp),%rdi
    6e3:      e8 82 ff ff ff              callq   66a <str_concat>
    6e8:      48 83 c4 10                addq    $0x10,%rsp
    6ec:      c3                       ret

00000000055006ed <main>:
    6ed:      48 83 ec 08                subq    $0x8,%rsp
    6f1:      b8 00 00 00 00              movq    $0x0,%eax
    6f6:      e8 bf ff ff ff              callq   6ba <dangerous>
    6fb:      48 8d 35 a2 00 00 00        leaq    0xa2(%rip),%rsi
    702:      bf 01 00 00 00              movq    $0x1,%edi
    707:      b8 00 00 00 00              movq    $0x0,%eax
    70c:      e8 2f fe ff ff              callq   540 <__printf_chk@plt>
    711:      b8 00 00 00 00              movq    $0x0,%eax
    716:      48 83 c4 08                addq    $0x8,%rsp
    71a:      c3                       ret

```

[illegible]

```
%rip -> 0x500646c
```

# Q4

- **Q4.1 &buf1[0]**
- Suppose the value of %rsp is 0x7ffd6a5e3020 just prior to executing the first instruction of dangerous, what is the address of the first element of buf1 (aka &buf1[0])?
- 0x7ffd6a5e3018
- **Q4.2 &buf2[0]**
- Using the same premise of Q4.1 earlier, what is the address of the first element of buf2 (aka &buf2[0])?
- 0x7ffd6a5e3010



# Q4

- **Q4.3**

- Using the same premise of Q4.1 earlier, what are the 8 bytes stored in the memory address 0x7ffd6a5e3020 (which is the value of %rsp just prior to executing the first instruction of dangerous)?
- 0x55006fb

# Q4


• **Q4.4** Which of the following statements are true?

- A. This program has no buffer overflow bugs and will execute correctly.
- B. This program has a buffer overflow bug, but it will nevertheless execute without a problem because the compiler has protected the stack using a canary.
- C. This program has a buffer overflow bug, but it will nevertheless execute without a problem because the compiler has allocated extra space on the stack that cushions the overflow.
- D. This program has a buffer overflow bug which is likely to manifest as a segmentation fault.
- E. buf1 is overflown during execution.
- F. buf2 is overflown during execution.

# Q4

- **Q4.5** last instruction

If running this program results in a segmentation fault. What is the last instruction executed before the segmentation fault occurs?

- A. The ret instruction in main function
- B. The ret instruction in dangerous function
- C. The ret instruction in str\_concat function
- D. The instruction to deallocate stack in dangerous, i.e. addq \$0x10,%rsp.

# Q4

- **Q4.6 Bonus question**
- If running this program results in a segmentation fault, what is the memory address that corresponds to the illegal memory access? You should assume the same premise as Q4.1.
- 0x500646c

# Dynamic Memory Allocation

For when static memory isn't enough

# Why Dynamic Memory?

- You don't always know how much memory you will need for your program
- What if you want to write a program that finds the average value in a column?
- If you did write such a program, how do you handle a user giving you a really big file, bigger than you expected?
- Even if you made sure you specified a really big global variable as a static buffer, people might still give you bigger files
  - And why go through that trouble anyway instead of just having dynamic memory?

# Dynamic memory and the stack

- Does the stack give us dynamic memory?
  - In a sense, yes
  - However, it isn't always suitable, because the memory gets reused after we return from a function call
  - By default the stack is also only a few megabytes in size

# Dynamic memory on the heap

- We can use the `sbrk` syscall to ask the operating system to give us more heap space
- We can also use it to give back to the operating system
- However, in the real world, programmers don't often do this themselves
  - Why?
- Instead, we usually use a library that handles things for us
  - API: `malloc` and `free`
  - Dynamic memory allocator



# Malloc and Free

- **Malloc** allocates us a contiguous section of memory
  - It returns a `void*`, which is just “pointer to anything”
  - So you cast the result of malloc to what you want, e.g. `int *`
  - Malloc can return NULL if there was an error
- **Free** gives the memory back to the allocator
  - DO NOT call free twice on the same section of memory
    - This is undefined behavior
  - What you call free on must be the result of malloc
    - (or calloc or realloc, but I won't discuss those)

# Allocators


- Allocators can't move data around
- How do you track what parts of the heap are freed or malloced? How do you track their sizes?
  - The trick is to **store metadata along with the data** in the heap to create a “linked list”
    - implicit list, explicit list
  - You store the **status** of the chunk (free or allocated), and the **size of the data** (which effectively points to the next chunk)
- When someone asks for memory, what do you give them?
  - There are a number of different strategies
- How to give back the memory?

# Malloc using Implicit list

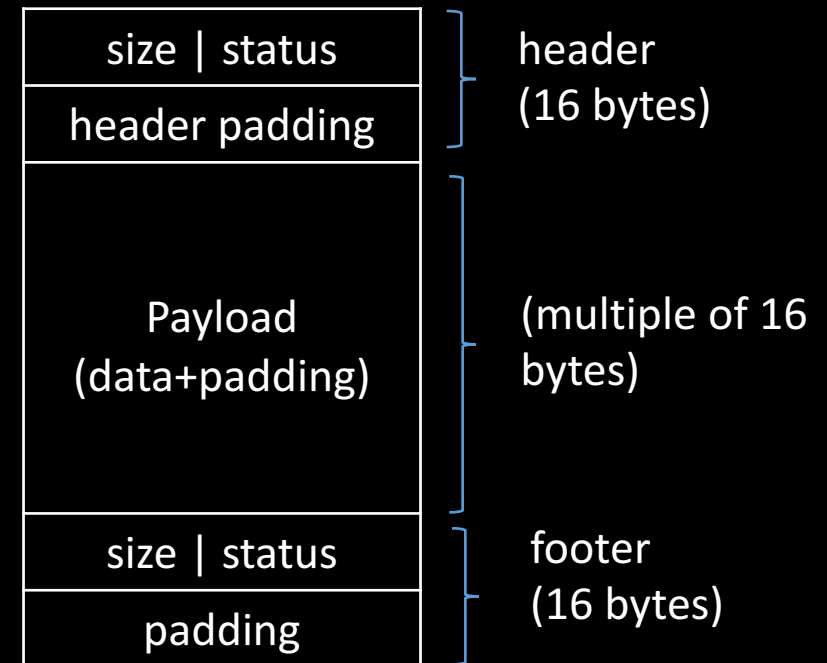
1. Structure of implicit list
2. Where to place an allocation?
3. Splitting a free block
4. Coalescing a free block

# Malloc using Implicit list

- Structure of implicit list
  - Implicit list means that it does not use pointers explicitly, but it can find the next node just like a linked list.



e.g. `p=malloc(20);`



# Malloc using Implicit list

- Where to place an allocation?
- Different algorithms:
  - First fit → easy & fast; cause fragmentation at beginning of the heap
  - Best fit → good for utilization; slower
  - Next fit → faster than first fit; even worse fragmentation

# Malloc using Implicit list


- Splitting a free block
  - Happens when we do memory allocation, e.g. `p10=malloc(16)`
  - (I have found a suitable free block)
    - find the next chunk (according to the size you want)
    - set size & status
- Coalescing a free block
  - Happens when we do free memory, e.g. `free(p10)`
  - After free, merge this free block with its next(& prev) free neighbor
    - find the next and prev chunk, and check its status
    - set size & status
      - (don't forget the footer)

# Implement Malloc using Implicit list

- Malloc <malloc(size)>
  - get the size of the chunk to allocate
  - find a free chunk
  - ask the OS for chunk
    - Split chunk if necessary
  - set this chunk status to be allocated
  - return pointer to the payload
- Free <free(p)>
  - go to the header from payload
  - set this chunk status to be free
  - coalesce

# Malloc using Explicit free list

- Based on the implicit list, because the implicit list is too slow
- Structure of explicit free list
  - Maintain a linked list by adding 2 pointers: next & prev
    - points to the next/previous **free** chunk
    - only the free chunk needs to be recorded, and the allocated ones do not need





# Implement Malloc using Explicit free list

- Malloc <malloc(size)>
  - get the size of the chunk to allocate
  - find a free chunk (in your linked list – linked list traverse)
    - delete this chunk from the linked list
  - ask the OS for chunk
    - Split chunk if necessary
    - insert the new free chunk to the linked list
  - set this chunk status to be allocated
  - return pointer to the payload
- Free <free(p)>
  - go to the header from payload
  - free the chunk
    - set this chunk status to be free
    - initial the next & prev pointer
  - coalesce
    - delete some chunk(s) from the linked list
  - insert this new free block into the linked list