

CSO-Recitation 11

CSCI-UA 0201-007

R11: Assessment 09 & Dynamic memory allocation

Today's Topics

- Assessment 09
- Dynamic memory allocation
 - implement your malloc & free

Dynamic Memory Allocation

For when static memory isn't enough

Why Dynamic Memory?

- You don't always know how much memory you will need for your program
- What if you want to write a program that finds the average value in a column?
- If you did write such a program, how do you handle a user giving you a really big file, bigger than you expected?
- Even if you made sure you specified a really big global variable as a static buffer, people might still give you bigger files
 - And why go through that trouble anyway instead of just having dynamic memory?

Dynamic memory and the stack

- Does the stack give us dynamic memory?
 - In a sense, yes
 - However, it isn't always suitable, because the memory gets reused after we return from a function call
 - By default the stack is also only a few megabytes in size

Dynamic memory on the heap

- We can use the `sbrk` syscall to ask the operating system to give us more heap space
- We can also use it to give back to the operating system
- However, in the real world, programmers don't often do this themselves
 - Why?
- Instead, we usually use a library that handles things for us
 - API: `malloc` and `free`
 - Dynamic memory allocator

Malloc and Free

- **Malloc** allocates us a contiguous section of memory
 - It returns a `void*`, which is just “pointer to anything”
 - So you cast the result of malloc to what you want, e.g. `int *`
 - Malloc can return NULL if there was an error
- **Free** gives the memory back to the allocator
 - DO NOT call free twice on the same section of memory
 - This is undefined behavior
 - What you call free on must be the result of malloc
 - (or calloc or realloc, but I won't discuss those)

Allocators


- Allocators can't move data around
- How do you track what parts of the heap are freed or malloced? How do you track their sizes?
 - The trick is to **store metadata along with the data** in the heap to create a “linked list”
 - implicit list, explicit list
 - You store the **status** of the chunk (free or allocated), and the **size of the data** (which effectively points to the next chunk)
- When someone asks for memory, what do you give them?
 - There are a number of different strategies
- How to give back the memory?

Malloc using Implicit list

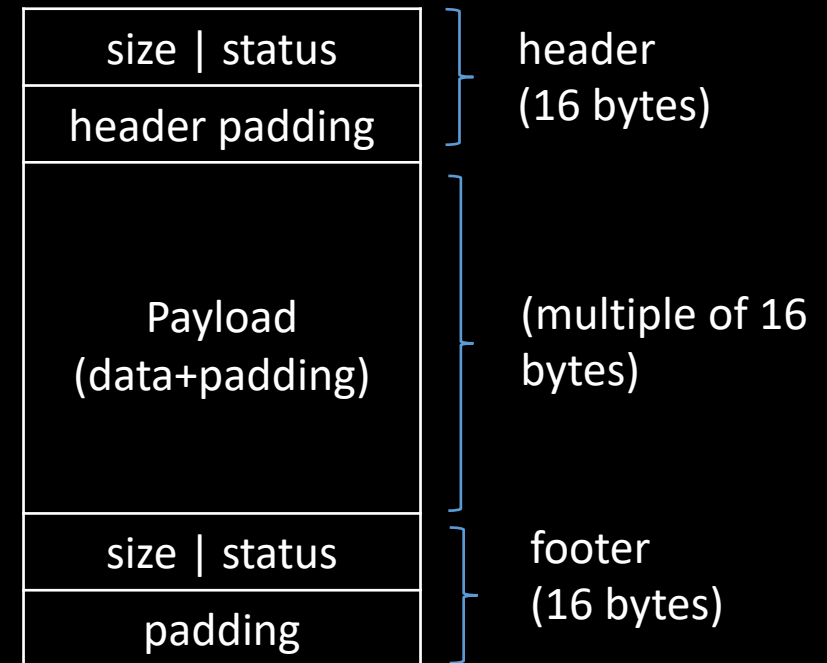
1. Structure of implicit list
2. Where to place an allocation?
3. Splitting a free block
4. Coalescing a free block

Malloc using Implicit list

- Structure of implicit list
 - Implicit list means that it does not use pointers explicitly, but it can find the next node just like a linked list.



e.g. `p=malloc(20);`



Malloc using Implicit list

- Where to place an allocation?
- Different algorithms:
 - First fit → easy & fast; cause fragmentation at beginning of the heap
 - Best fit → good for utilization; slower
 - Next fit → faster than first fit; even worse fragmentation

Malloc using Implicit list


- Splitting a free block
 - Happens when we do memory allocation, e.g. `p10=malloc(16)`
 - (I have found a suitable free block)
 - find the next chunk (according to the size you want)
 - set size & status
- Coalescing a free block
 - Happens when we do free memory, e.g. `free(p10)`
 - After free, merge this free block with its next(& prev) free neighbor
 - find the next and prev chunk, and check its status
 - set size & status
 - (don't forget the footer)

Implement Malloc using Implicit list

- Malloc <malloc(size)>
 - get the size of the chunk to allocate
 - find a free chunk
 - ask the OS for chunk
 - Split chunk if necessary
 - set this chunk status to be allocated
 - return pointer to the payload
- Free <free(p)>
 - go to the header from payload
 - set this chunk status to be free
 - coalesce

Malloc using Explicit free list

- Based on the implicit list, because the implicit list is too slow
- Structure of explicit free list
 - Maintain a linked list by adding 2 pointers: next & prev
 - points to the next/previous **free** chunk
 - only the free chunk needs to be recorded, and the allocated ones do not need



Implement Malloc using Explicit free list

- Malloc <malloc(size)>
 - get the size of the chunk to allocate
 - find a free chunk (in your linked list – linked list traverse)
 - delete this chunk from the linked list
 - ask the OS for chunk
 - Split chunk if necessary
 - insert the new free chunk to the linked list
 - set this chunk status to be allocated
 - return pointer to the payload
- Free <free(p)>
 - go to the header from payload
 - free the chunk
 - set this chunk status to be free
 - initial the next & prev pointer
 - coalesce
 - delete some chunk(s) from the linked list
 - insert this new free block into the linked list